

Stuff

Michael Langbein

April 7, 2021

Contents

1 Math	15
1.1 Logic	16
1.1.1 Algebra	16
1.1.2 Quantifiers	16
1.1.3 Inference	16
1.1.4 Consistency and well-definedness	16
1.1.5 Experimental and protocol design	16
1.2 Set theory	17
1.2.1 Relations	17
1.2.2 Orders	17
1.2.3 Partitions	18
1.2.4 Recursion	18
1.3 Combinatorics	19
1.3.1 Sums and asymptotics	19
1.3.2 Products of sums and convolutions	19
1.3.3 Cardinality	19
1.3.4 Counting	19
1.3.5 Generating functions	20
1.4 General algebra	21
1.4.1 Vector spaces	21
1.4.2 Inner product spaces	22
1.4.3 Excursion: More on orthogonality and preview of function decompositions	24
1.5 Linear algebra	26
1.5.1 Spaces	26
1.5.2 Change of basis	28
1.5.3 Linear transformations	28
1.5.4 Linear independence	30
1.5.5 Determinant	30
1.5.6 Rank nullity Theorem	31
1.5.7 Special matrices	32
1.5.8 Eigenvalues and eigenvectors	32
1.5.9 Singular value decomposition	32
1.5.10 Inverse	33
1.5.11 Applications	34
1.5.11.1 Systems of linear equations	34
1.5.11.2 Matrix factorisation	36
1.6 Fourier and Laplace	37
1.6.1 Discrete Fourier Analysis	37
1.6.1.1 The Fourier Base	37
1.6.1.2 Obtaining the amplitudes	38
1.6.2 As a matrix operation	38
1.6.3 Fast Fourier Transform	39
1.6.3.1 Backtransformation	39
1.6.4 Fourier for musical frequencies	39

1.6.4.1	Goertzel	40
1.6.4.2	PCI	40
1.6.5	Continuous Fourier Analysis	40
1.6.5.1	Comparing Fourier to other important series	40
1.6.5.2	Proving that Fourier functions form a basis	41
1.6.5.3	Fourier transform on vector valued functions	41
1.6.6	Fourier Transforms and Convolutions	41
1.7	Probability	44
1.7.1	Basics	44
1.7.1.1	Probability space	44
1.7.1.2	Uncountable sets and infinitessimals	45
1.7.1.3	Probabilistic fallacies	45
1.7.2	Probability distributions	45
1.8	Statistics	46
1.8.1	Correlation	46
1.8.2	Linear regression	46
1.8.3	Principal component analysis	46
1.8.4	Tests	46
1.9	Graph Theory	47
1.9.1	General properties	47
1.9.1.1	Degrees	47
1.9.2	Walks and paths	47
1.9.2.1	Relations and adjacency matrices	48
1.9.3	Planar graphs	48
1.9.3.1	Trees	49
1.9.3.2	Triangulation	50
1.9.4	Stable marriage and Gayle-Shapely	50
1.10	Calculus	52
1.10.1	Hyperreals	52
1.10.2	Limits	53
1.10.3	Sequences and series	53
1.10.3.1	Tailor	53
1.10.3.2	Fourier	53
1.10.3.3	Laplace	53
1.10.4	Eulers formula	53
1.10.5	Integration	54
1.10.5.1	Integration strategies	54
1.10.6	Vector calculus	54
1.11	Number theory and cryptography	55
1.11.1	Divisability	55
1.11.2	Hashing	56
1.11.3	Encryption	56
1.11.3.1	Symmetric encryption	56
1.11.3.2	Asymmetric encryption	56
1.12	Computation	58
1.12.1	Finite state machines	58
1.12.1.1	Definition	58
1.12.1.2	Getting the language of a given machine M	58
1.12.1.3	Getting the simplest machine for a given language	59
1.12.1.4	Buildung machines from simpler machines	59
1.12.1.5	Induction using the invariant principle	59
1.12.2	Turing machines	59
1.13	Dynamic Programming	60
1.13.1	Overlapping subproblems (always) and optimal substructure (most)	60
1.13.2	Worked example: possible ways to sum coins	60
1.13.3	A strategy for finding the recursion	61

1.13.4	Markow chains	61
1.13.5	Markov decision process with known state-transition-function	61
1.13.6	Reinforcement learning: MDP without explicitly known state-transition-function	61
1.14	Distributed systems	62
1.14.1	Model	62
1.14.2	Axioms	62
1.14.3	Theorems	62
1.14.3.1	Coverings	62
1.14.3.2	Byzantine agreement	63
1.14.3.3	Weak agreement	63
1.14.3.4	Byzantine fireing squad	63
1.14.3.5	Approximate agreement	63
1.14.3.6	Clock syncing	63
1.14.3.7	When timing does and does not matter	63
1.14.3.8	fip theorem	64
1.14.3.9	CAP-Theorem	64
1.14.4	Important algorithms	64
1.14.4.1	one-, two- and tree-phase commit	64
1.14.4.2	paxos algorithm	64
1.14.4.3	raft algorithm	64
1.14.4.4	map reduce algorithm	64
1.14.5	Databases	64
1.14.5.1	ACID	64
1.14.5.2	NoSQL Database types	65
1.14.5.3	Normalisation	65
1.14.5.4	Indices	67
1.14.5.5	Spatial indices	67
1.14.6	Distributed cache: redis	67
1.15	Machine learning	68
1.15.1	Neural networks	68
1.15.1.1	Backpropagation	68
1.15.1.2	Universal approximation	69
1.15.1.3	Some data in n dimensions requires more than n neurons in a layer	69
1.15.1.4	Some functions can be better approximated with a deep net than with a shallow one	69
1.15.1.5	Convolutional networks	70
1.15.1.6	Self organizing maps	71
1.15.2	Computer vision	71
1.15.3	Feature extraction and dimensionality reduction	71
1.15.4	Symbolic AI	71
2	Algorithms and data-structures	75
2.1	General algorithm theory	76
2.1.1	Invariant principle	76
2.1.2	Well ordering	76
2.1.3	Program verification according to Floyd	76
2.1.3.1	Proving 'partial correctness' with invariant principle	76
2.1.3.2	Proving 'termination' with well ordering principle	76
2.1.4	Notation	76
2.1.5	Order of basic operations	77
2.1.6	Optimization	77
2.1.7	Memoization: Remembering the state in loops	77
2.1.8	Using sub problems: Dynamic programming	77
2.1.9	Avoiding the need for dynamic programming: Solving recurrences to closed form	78
2.1.9.1	Solving linear recurrences	78
2.2	Important algorithms	79
2.2.1	Merge-sort	79

2.2.2	Matrix inverse	79
2.2.3	Polynomials	79
2.2.3.1	Coefficient-representation of polynomials	79
2.2.3.2	Point-value-representation of polynomials	80
2.2.3.3	Naive transformation from coefficient to point-value and vice-versa	81
2.2.3.4	Fast transformation	81
2.2.4	Stream algorithms	82
2.3	Data-structures	83
2.3.1	Stack	83
2.3.2	Linked lists	84
2.3.3	Binary trees	84
2.3.4	Hash-tables	85
3	Practical computer knowledge	87
3.1	Hardware	88
3.1.1	Memory	88
3.1.1.1	How data is stored	88
3.1.2	Processors	88
3.1.3	Cables and Busses	88
3.1.4	Harddrives	88
3.2	Networking	89
3.2.1	Protocols	89
3.2.1.1	Layer 1: Physical	89
3.2.1.2	Layer 2 : Data link	89
3.2.1.3	Layer 3 : Network	89
3.2.1.4	Layer 4 : Transport	90
3.2.1.5	Layer 5: Application layer	91
3.2.1.6	Layer 6 and higher	92
3.2.2	Email	92
3.2.3	Security	92
3.2.3.1	Encryption	92
3.2.4	Making things talk: Remote data transmission	92
3.2.4.1	ISDN (wired)	93
3.2.5	DSL (wired)	93
3.2.6	WiFi (really only those few meters up to your wired router)	94
3.2.7	GSM, EDGE, 3G=HSDPA, 4G=LTE=HSDPA(+) (radio)	94
3.2.8	GPS (satellite)	95
3.3	Internet - practical aspects	96
3.3.1	Nomenclature	96
3.3.2	Session basics	96
3.3.3	Authentication, SSO, and authorization	97
3.3.4	Security threats	100
3.3.5	Status codes	103
3.4	Linux render-stack	104
3.5	Linux sound-stack	105
3.5.1	Audio Standards	105
3.5.2	Hardware	105
3.5.3	Software	105
3.5.4	JACK	105
3.6	Legal and opensource	106
3.6.1	Popular licenses	106
3.6.2	Adding 3rd party licenses to your apache2-javascript product	106
3.6.3	Deutsches/Europäisches Recht	106
3.6.3.1	Haftung	106
3.6.3.2	Exportkontrolle	106

4 Programming languages	107
4.1 C	108
4.1.1 General theory	108
4.1.1.1 Memory allocation	108
4.1.1.2 Threading	108
4.1.1.3 Getting data from C to another programm	109
4.1.2 Quirks and features you need to know	109
4.1.2.1 * syntax	109
4.1.2.2 Pointer arithmetic	109
4.1.2.3 Array decay: Functions can't accept arrays	110
4.1.2.4 Array decay: Functions don't return arrays, either	110
4.1.2.5 Array syntax	111
4.1.2.6 Struct syntax	111
4.1.2.7 Passing functions as variables	111
4.1.2.8 Strings end with a \0	112
4.1.2.9 char[] does not equal char*	112
4.1.2.10 Header files	113
4.1.2.11 Building and Makefiles	113
4.1.2.12 Valgrind	114
4.1.2.13 OpenMP	114
4.1.3 CMake and Autotools	114
4.1.4 Best practice	115
4.2 C++	116
4.2.1 Some weird syntax and new concepts	116
4.2.1.1 References	116
4.2.1.2 Const keyword	116
4.2.2 Object orientation	117
4.2.2.1 Operator overloading	118
4.2.2.2 Rule of three	119
4.2.2.3 Smart pointers	120
4.2.2.4 Ownership: unique, shared and weak pointers	121
4.2.3 Templates	123
4.2.3.1 Function template	123
4.2.3.2 Class template	123
4.2.4 Threading	124
4.2.5 CMake	124
4.2.5.1 Variables	125
4.2.5.2 Logic	125
4.2.5.3 Command arguments	125
4.2.5.4 Important predefined variables	126
4.3 Java	127
4.3.1 General	127
4.3.1.1 Basic theory	127
4.3.1.2 Environment variables	127
4.3.1.3 Versions	127
4.3.2 Maven	128
4.3.3 Threading	128
4.3.4 Functional programming	131
4.3.5 Annotations	132
4.3.6 Transforming textformats: marshaling	133
4.3.6.1 CSV to POJO and back: Bindy	133
4.3.6.2 JSON to POJO and back: Jackson	133
4.3.6.3 XML to POJO and back: JAXB	133
4.3.7 Reading from documents	133
4.3.7.1 Extracting from xml: XPath	133
4.3.7.2 Extracting from json: JsonPath	133

4.3.8	Databases	133
4.3.8.1	First level db-access: JDBC	133
4.3.8.2	Second level db-access: JPA and Hibernate	134
4.3.9	Swing	136
4.3.9.1	Lists	136
4.3.9.2	Trees	137
4.3.9.3	Writing custom components	138
4.3.10	Servlets and JSP	138
4.3.10.1	Building with maven	139
4.3.10.2	Jsp and forms	140
4.3.10.3	Servlet context listener	140
4.3.10.4	Servlet lifecycle	141
4.3.11	Webservices	141
4.3.11.1	Rest-container: Jersey	141
4.3.11.2	Soap-container: Apache CXF	142
4.3.11.3	JSP and JSTL	142
4.3.12	Unit testing	143
4.3.13	Logging	143
4.4	Python	145
4.4.1	List comprehensions	145
4.4.2	Functional programming	145
4.4.3	Decorators	145
4.4.4	Metaprogramming	146
4.4.4.1	Changing the way an object handles operations	146
4.4.4.2	Creating new operators	146
4.4.4.3	Extending existing datastructures	147
4.4.4.4	Named tuples	148
4.4.5	Plotting	148
4.4.6	SymPy	149
4.4.7	Numpy and Scipy	150
4.4.8	Gradual typing	151
4.4.9	Piping and currying	151
4.4.10	Generators	151
4.4.11	Oslash: advanced functional programming	151
4.4.12	Multimedia: Pygame	153
4.4.13	Regex	153
4.4.14	Modules and Packages	153
4.4.15	Miniconda	154
4.4.16	IPythonNotebook aka. Jupyter	154
4.4.17	Pandas	154
4.5	Racket	155
4.5.1	Basics	155
4.5.2	Macros	156
4.5.3	Nomenclature	156
4.5.4	Datastructures	156
4.6	Javascript	157
4.6.1	The js-runtime	158
4.6.2	Asynchronous programming	158
4.6.3	Module system	161
4.6.3.1	Building with webpack	161
4.6.4	Promises, async-await and rxjs	164
4.6.5	Webworkers: multithreading	164
4.6.6	WASM	164
4.7	CSS	165
4.7.1	Position	165
4.7.2	Display	165

4.7.2.1	Flexbox	165
4.8	xml	166
4.8.1	Structure	166
4.8.2	Namespaces: xmlns	166
4.8.3	Schema definitions: xsd	166
4.8.4	Marshalling	167
4.9	Design patterns	168
4.9.1	UML	168
4.9.1.1	Classes	168
4.9.1.2	Sequences	168
4.9.1.3	UseCases	168
4.9.2	Dependencies versus microservices	168
4.9.3	Strategy pattern	169
4.9.4	State pattern	169
4.9.5	Observer pattern	169
4.9.6	Orchestration: Workflow pattern	170
4.9.7	Choreography: Actor pattern	170
4.9.8	Monads	170
4.9.8.1	Maybe-monad	170
5	Stacks, platforms and libraries	173
5.1	Git	174
5.1.1	Lingo	174
5.1.2	Changing origin	174
5.1.3	Merge vs. rebase	175
5.1.4	Creating your own repository	176
5.2	Docker	177
5.2.1	Command overview	177
5.2.2	Dockerfile syntax	178
5.2.3	Composefile syntax	178
5.2.4	Layers and continuing failed builds	178
5.3	Spring	180
5.3.1	Dependency injection	180
5.3.2	Aspects	181
5.3.3	Templates	183
5.3.4	Bean lifecycle	183
5.3.5	Wireing	184
5.3.6	Mocking	184
5.3.7	Spring MVC	184
5.3.7.1	General structure	184
5.3.7.2	Configuration	185
5.3.7.3	Application: a minimal setup	185
5.3.7.4	Application: what goes into the root-config?	186
5.3.8	Adding further servlets to spring mvc	189
5.3.9	Spring REST services	189
5.3.10	Spring Database access	189
5.4	Spring Boot	190
5.4.1	Example	190
5.4.2	Configuration	190
5.4.3	Database access	190
5.4.4	Websites	190
5.4.5	Testing	191
5.4.6	Camel	192
5.5	Android	193
5.5.1	ADB	193
5.5.2	Basic concepts	193

5.5.3	Userinput	194
5.5.4	Activities	194
5.5.5	Threading	194
5.5.6	Drawing	195
5.6	C++ on android	198
5.7	VR on android	199
5.7.1	Basic app	199
5.8	Tensorflow	200
5.8.1	Low-level API	200
5.8.2	Even lower level, and understanding TF internals	200
5.8.3	High-level API	201
5.8.4	Topics	201
5.8.4.1	Convolutional nets	201
5.8.4.2	LSTM	202
5.8.4.3	Deep dream	202
5.9	Angular	204
5.9.1	Elements	204
5.9.2	Life cycle	204
5.9.3	Change detection	204
5.9.4	Service availability	204
5.9.5	Building blocks	204
6	Sciences	207
6.1	Remote sensing	208
6.1.1	Electromagnetic radiation	208
6.1.2	Orbital periods and acquisition	208
6.1.3	Important satellites and service providers	208
6.1.4	Datatypes, protocols and libraries	209
6.1.5	Projections	210
6.1.6	NDVI	210
6.1.7	Image (co-)registration	210
6.2	Economics	212
6.3	Projektwesen	213
7	Art	215
7.1	Design	216
7.1.1	Layout	216
7.1.2	UI	216
7.2	Math	217
7.2.1	Points	217
7.2.2	Geometric objects in matrix notation	217
7.2.2.1	Plane	217
7.2.2.2	Line	217
7.2.2.3	Intersection Line/Plane	217
7.2.3	Projections	218
7.2.3.1	Finding a perpendicular	218
7.2.3.2	Projecting one vector onto another	218
7.2.3.3	Flattening: expressing the points on a 3d-plane in a 2d-system	218
7.2.4	Rotation	219
7.2.4.1	In 2d	219
7.2.4.2	In 3d	219
7.2.4.3	Why is there one rotation axis in 2d, but 3 in 3d?	219
7.2.5	Implicit versus parameterized representation of bodies	219
7.2.6	Maps	220
7.2.6.1	Projections	220
7.2.6.2	Zoom to width/height	220
7.2.6.3	Apparent height	220

7.3	Graphics Programming	221
7.3.1	Rendering in the browser	221
7.3.2	SVG	221
7.3.3	D3	221
7.3.4	WASM	222
7.3.5	WebGl	223
7.3.5.1	Shader gallery	236
7.3.6	CSS	236
7.3.6.1	Positioning	236
7.3.6.2	Events	236
7.3.6.3	Animations	236
8	Applications	237
8.1	Where should we meet?	238
8.2	Effective wind on a ship	239
8.3	Making a touch-instrument	240
8.3.1	Music	240
8.3.1.1	Notes	240
8.3.1.2	Harmony	240
8.3.1.3	Songs	240
8.3.2	Deconstructing $w(t)$ into f_x and a_x : the FFT	240
8.4	Drawing a room onto a cylinder	242
8.4.1	Projecting a line in space onto a cylinder	242
8.4.1.1	Projecting a point onto the cylinder	242
8.4.1.2	Projecting a plane onto the cylinder	242
8.4.2	Rolling out the cylinder	242
8.5	Playing with projections and transformations	245
8.6	Fourier on sound and geometric objects	246
8.6.1	Basics: Fourier in one dimension	246
8.6.1.1	Representing a function under a Fourier base	246
8.6.1.2	Bringing in arbitrary intervals	246
8.6.1.3	Implementation and verification in python	247
8.6.1.4	Our first frequency domain operations: adding overtones	247
8.6.1.5	Windowing	248
8.6.2	Fourier in more than one dimensions	251
8.6.2.1	Multiple input parameters	251
8.6.2.2	Multiple output parameters: Vector valued functions	251
8.7	Estimating the time needed for a task	254
8.8	Scheduling	259
8.8.1	Running time	259
8.8.2	Partial correctness	259
8.8.3	Termination	259
8.9	Webcamgiant	260
8.9.1	Webcam to RPy	260
8.9.2	RPy to computer	260
8.9.3	Computer to Oculus	260
8.10	Labeling radar-movies with likely summer-thunderstorms	261
8.11	Backpropagation for hydrological parameter estimation	262
8.11.1	Hydrological model	262
8.11.2	Optimisation using tensorflow	263
8.12	Hydrological models as Hidden Markov Models	264

Preface

I think that the world is in the process of becoming something amazing. Things that I saw in animated series as a kid are becoming reality. I don't just want to be a passive spectator of this process, but rather have my own active part in it, be it ever so small.

Chapter 1

Math

1.1 Logic

1.1.1 Algebra

\wedge and \vee are distributive:

$$(A \wedge B) \vee C \equiv (A \vee C) \wedge (B \vee C)$$

$$(A \vee B) \wedge C \equiv (A \wedge C) \vee (B \wedge C)$$

This is easiest seen by drawing a Venn-diagram.

1.1.2 Quantifiers

$$\exists!x : Q(x) \equiv (\exists x : Q(x)) \wedge (\forall x, y : Q(x) \wedge Q(y) \rightarrow y = x)$$

1.1.3 Inference

Simple if statements:

$$A \rightarrow B \equiv \neg A \vee B$$

$$\neg(A \rightarrow B) \equiv A \wedge \neg B$$

Applying this to function-statements yields:

$$\neg(\forall x : A(x) \rightarrow B(x)) \equiv \exists x : A(x) \wedge \neg B(x)$$

If-and-only-if statements:

$$A \leftrightarrow B \equiv (A \rightarrow B) \wedge (B \rightarrow A) \equiv (A \wedge B) \vee (\neg A \wedge \neg B)$$

$$\neg(A \leftrightarrow B) \equiv (A \wedge \neg B) \vee (B \wedge \neg A) \equiv (\neg A \rightarrow B) \wedge (A \rightarrow \neg B) \equiv A \leftrightarrow \neg B$$

Applying this to function-statements yields:

$$\neg(\forall x : A(x) \leftrightarrow B(x)) \equiv \exists x : (\neg A(x) \rightarrow B(x)) \wedge (A(x) \rightarrow \neg B(x))$$

1.1.4 Consistency and well-definedness

Whenever you receive a set of axioms (like *imagine there was a complex number i*), you should only accept working with them when they don't lead to any logical inconsistencies. However, it is hard to prove that a set of axioms doesn't lead to inconsistencies. Usually, you start with a really basic set (ZFC axioms) and prove that you can construct structures (like for example Dedekind cuts) that have as properties the new axioms you want to establish. Then your axioms are consistent as long as ZFC is consistent.

A concept is well defined when it is a function, that is, when for every input-data there is always at most one output-data.

1.1.5 Experimental and protocol design

This section deals with logic-puzzles.

Multistep processes At every step of your experiment, utilize the full spectrum of your measuring device. This way, each measurement contains the maximal information content. For example, in the "twelve coins" problem, at every one of your three steps the scale should be able to tip left, right or not at all.

1.2 Set theory

1.2.1 Relations

A relation is injective if any y belongs only to one x - if at all.

Definition 1. *Injective (1-1):* $\forall x_1, x_2 \in X : x_1 R y_0 \wedge x_2 R y_0 \rightarrow x_1 = x_2$

A relation is surjective if any y belongs to at least one x .

Definition 2. *Surjective (onto):* $\forall y \in Y : \exists x \in X : x R y$

Note that none of the above definitions means that there is a y for any x . Under both definitions, there can be X 's that have none or more than one y .

While these definitions are useful, it is sometimes easier to use the following definitions based on the number of in- or outgoing arrows:

Table 1.1: Types of binary relations

	out	in
≤ 1	function	injective (1-1)
≥ 1	total	surjective (onto)

- function: # out ≤ 1
- injective: # in ≤ 1
- total: # out ≥ 1
- surjective: # in ≥ 1
- bijective: # out = # in = 1

Theorem 1. Let R be a relation on $A \times B$. Then: $R : \text{injective} \wedge R : \text{function} \rightarrow |A| \leq |B|$

Proof. Suppose that R :injective and R :function. Proof that $|A| \leq |B|$

R :injective, thus: # edges $\leq |B|$
 R :function, thus: $|A| \leq$ # edges
 Thus $|A| \leq$ # edges $\leq |B|$

Thus $|A| \leq |B|$

□

Theorem 2. Let R be a relation on $A \times B$. Then: $R : \text{surjective} \wedge R : \text{total} \rightarrow |A| \geq |B|$

Theorem 3. Let R be a relation on $A \times B$. Then: $R : \text{bijective} \rightarrow |A| = |B|$

The combination of function and surjectivity is sometimes written as A surj B ; the combination of totality and injectivity as A inj B .

This begs a question: $\neg A \text{surj } B \leftrightarrow A \text{inj } B$? No, this doesn't hold. A counterexample would be ... But we can proof that $A \text{surj } B \leftrightarrow B \text{inj } A$:

1.2.2 Orders

Definition 3. *Partial order:* A relation R on a set S is called a partial order if it is reflexive, antisymmetric and transitive.

- *reflexive:* $\forall x \in S : x R x$
- *antisymmetric:* $\forall x, y \in S : x R y \wedge y R x \rightarrow x = y$

- *transitive:* $\forall x, y, z \in S : xRy \wedge yRz \rightarrow xRz$

Definition 4. *Total order:* a relation R on a set S is called a total order if it is a partial order and also comparable

- *comparable, a.k.a. total:* $\forall a, b \in S : aRb \vee bRa$

Definition 5. *Topological order*

Definition 6. *Closure:*

A set S and a binary operator $*$ are said to exhibit closure if applying the binary operator to two elements S returns a value which is itself a member of S .

The closure of a set A is the smallest closed set containing A . Closed sets are closed under arbitrary intersection, so it is also the intersection of all closed sets containing A . Typically, it is just A with all of its accumulation points.

1.2.3 Partitions

1.2.4 Recursion

1.3 Combinatorics

1.3.1 Sums and asymptotics

We begin with a few trivial results. Closed form on the right. The methods we develop for sums will also work for products, since any product can be converted into a sum by taking its logarithm.

$$\sum_{n=0}^N n = \frac{N(N+1)}{2}$$

$$\sum_{n=0}^N x^n = \frac{1-x^{N+1}}{1-x}$$

We might get to these results with the perturbation method:

- $1 + x + x^2 \dots + x^N = S$
- $-x - x^2 \dots - x^{N+1} = -xS$
- $1 - x^{N+1} = S - xS$

It is very important to note that this method only works for calculating *finite* partial sums. Nothing guarantees us that such a method would work as well for infinite sums. In fact, we can disprove this.

Theorem 4. A example where using the perturbation method on infinite series leads to a contradiction would be ...

However, we are always allowed to calculate a partial sum by some variant of the perturbation method and then taking the limit of $n \rightarrow \infty$.

1.3.2 Products of sums and convolutions

Let $S_A = \sum_n a_n$ and $S_B = \sum_m b_n$. To calculate $S_A S_B$ we sum all elements in the following table:

Table 1.2: Multiplication of sums

	a_0	a_1	a_2	a_3	
b_0	$a_0 b_0$	$a_1 b_0$	$a_2 b_0$	\dots	
b_1	$a_0 b_1$	$a_1 b_1$	\dots		
b_2	$a_0 b_2$	\dots			
b_3	\dots				

Note that the diagonal from ll to ur in this table is a convolution.

1.3.3 Cardinality

1.3.4 Counting

Permutations are the number of rearrangements where order matters. Safe-combinations should really be called safe permutations. Permutations of unique items are relatively straightforward, but things get a little involved when we deal with permutations of grouped items.

- with repetition (*4-slot lock*): n^4
- without repetition (*first 3 places in race*): $\frac{n!}{(n-3)!}$

The number of unique permutations of a string consisting of n_a a's, n_b b's, n_c c's, etc. is

$$\frac{(\sum n_i)!}{\prod (n_i)!}$$

Here, $(\sum n_i)!$ is the number of permutations if every letter in the string was unique, that is, if we could distinguish between two a's. We then reduce this number by the number of duplicates by dividing through $n_a!$, the number of rearrangements of a's.

Note how the binomial coefficient is a special case of our string-permutation: it is the number of unique permutations in a string of length n with only two letters in it. In fact, the number of string-permutations is called the *multinomial* coefficient. The formula is easily proven by induction on the number of groups i . In the base-case with two groups we trivially get the binomial coefficient. Then do the induction step.

Combinations are sets where the order does *not* matter: $(2, 4, 3) = (2, 3, 4) = (3, 2, 4) = \dots$

- without repetition (*lottery of 6*): $\frac{n!}{(n-6)!6!} = \binom{n}{6}$, aka. the *binomial coefficient*, aka. n choose k . Note the term $6!$ in the denominator. n choose k is exactly a permutation without repetition (that is, where order matters), reduced by the different ways to order things ($6!$)).
- with repetition (*draw 5 coins out of your pocket*): $\binom{n-1+5}{5}$, where n is the number of different coin-denominations. Its pretty cool how this formula is obtained: write all coin-denominations in a list. Under that list, write a program, consisting of arrows for moving right and x 'es for picking once. This is the same as picking 5 slots for x 'es out of $(n - 1 + 5)$.

1.3.5 Generating functions

In the previous section we have found a bunch of formulas for getting the possible selections in different simple situations. But what if we need to combine a few simple selection-scenarios into one complicated scenario?

Imagine you need to buy n items. There are two kinds of items: apples and bananas. Apples always come in packs of six, and bananas have two different sub-kinds. How many ways are there then to get n items?

This is where we use generating functions.

Algebraically, what's happening is that taking an ordinary generating function is a bijection between the vector space of sequences and the ring of formal power series. While in a vector space we have vector addition and scalar multiplication, in a ring we have element addition and element multiplication (and even division and differentiation!). So working in the powerseries-ring allows for other operations than working in the series-vectorspace. This is very simmilar to what we do when doing the Fourier-transform: we dive into a different domain where some operations are easier to do. However, Fourier is merely a change of basis, going from one vector space to another. Here, we move from a vector space to a ring - a whole different universe.

Let $A = (a_0, a_1, \dots)$ be a sequence where a_n represents the number of ways to select from group \mathcal{A} . Then the generating function G_A would be defined as:

$$G_A(x) = a_0 + a_1x + a_2x^2 + \dots = \sum_n a_n x^n$$

For any sequence we can define a generating function. Now, to get the desired result, we first transform several sequences into their generating functions, find their closed forms, multiply them, and transform the result of the multiplication back into a sequence.

In detail:

- From (a_0, a_1, \dots) create $G_A(x)$ and find its closed form.
- From (b_0, b_1, \dots) create $G_B(x)$ and find its closed form.
- Get $G_C = G_A G_B$
- Get (c_0, c_1, \dots) from G_C

This works because of the following theorem:

Theorem 5 (Convolutions with generating functions). *Let $A = (a_0, a_1, \dots)$ be a sequence where a_n represents the number of ways to select from group \mathcal{A} . Let $B = (b_0, b_1, \dots)$ be a sequence where b_n represents the number of ways to select from group \mathcal{B} . Let $\mathcal{A} \cap \mathcal{B} = \emptyset$. Then c_n is the number of ways to select n items from $\mathcal{A} \cup \mathcal{B}$. It can be obtained by getting the n -th coefficient of $G_C = G_A G_B$*

1.4 General algebra

In the following sections, we will take care to differentiate between the definition of a concept (like of an inner product) and its implementation. The definition of an inner product is based on properties that a thing has to fulfill, whereas the implementation begins with a definition and is then followed by a proof that the properties hold under that definition.

We will mention the following implementations: the vectorspace of coordinate free oriented lengths, the vectorspace \mathbb{R}^n , which is the same as oriented length but put inside of a coordinate system(euclidian or angular), the vectorspace of matrices, and $C_{[a,b]}$, the vectorspace of continuous functions.

1.4.1 Vector spaces

Definition 7 (Vector space). *A vector space V is a set closed over two operations: scalar multiplication and vector addition. These two operations must fullfill the following properties:*

- V is closed under scalar multiplication and vector-addition: $a\vec{v} \in V, \vec{v} + \vec{w} \in V$
- vector-addition is commutative: $\vec{v} + \vec{w} = \vec{w} + \vec{v}$
- vector-addition is associative: $(\vec{u} + \vec{v}) + \vec{w} = \vec{u} + (\vec{v} + \vec{w})$
- $\vec{0}$ is the additive identity: $\vec{v} + \vec{0} = \vec{v}$
- $a(b\vec{v}) = (ab)\vec{v}$
- Vector-addition and scalar-multiplication are distributive (part 1): $a(\vec{v} + \vec{w}) = a\vec{v} + a\vec{w}$
- Vector-addition and scalar-multiplication are distributive (part 2): $(a + b)\vec{v} = a\vec{v} + b\vec{v}$

The most common vectorspaces are certainly \mathbb{R}^n and functionspaces. The implementations of the above defined scalar product and vector addition are trivial.

Subspaces of vectorspaces A set U is a subspace of a vectorspace V , iff $U \subset V$ and U is a vectorspace.

linear independence The elements in A , a subset of a vectorspace, are linearly independent iff $\forall \alpha_1, \dots, \alpha_n : [(\sum \alpha_i \vec{a}_i = 0) \leftrightarrow (\alpha_1 = \alpha_2 = \dots = 0)]$. Note that this reduces automatically to $\forall \alpha_1, \dots, \alpha_n : [\sum_i^n \alpha_i b_i = 0 \rightarrow (\alpha_1 = \dots = \alpha_n = 0)]$, because the \leftarrow case is always true.

Consequently, linear dependence is defined as $B : ld \equiv \exists \alpha_1, \dots, \alpha_2 : [(\alpha_1 \neq 0 \vee \dots \vee \alpha_n \neq 0) \wedge (\sum_i^n \alpha_i b_i = 0)]$.

Proof. Prove that iff $B : ld$, then one of the elements of B is a linear combination of the others. \square

Consider the span of two integers (like in the die-hard water-jug problem). They are always linearly dependent. For example a base $B = [2, 3]$ is linearly dependent because $\frac{-3}{2}2 + 1 \cdot 3 = 0$.

Bases A set B is a base to a vectorspace V iff $B \subseteq V \wedge \forall v \in V : \exists! \alpha_1, \dots, \alpha_n : v = \sum_i \alpha_i b_i$. It is easy to prove that this means that $B : baseV \leftrightarrow (B : li \wedge B : spanV)$.

Definition 8 (The dimension of a vectorspace S). *is defined as the size of its base B_S : $dim_S = |B_S|$.*

That means to get a base for \mathbb{R}^2 , we never need more than two vectors. We'll prove this for the example of $S = \mathbb{R}^2$:

Proof. For any three vectors chosen from \mathbb{R}^2 , at least one must be a linear combination of the others. $\vec{a}, \vec{b}, \vec{c} \in \mathbb{R}^2$ Proof that $\vec{a} : ld(\vec{b}, \vec{c}) \vee \vec{b} : ld(\vec{a}, \vec{c}) \vee \vec{c} : ld(\vec{a}, \vec{b})$

Without loss of generality, assume $\vec{a} : li(\vec{b}, \vec{c})$ and $\vec{b} : li(\vec{a}, \vec{c})$ Proof that $\vec{c} : ld(\vec{a}, \vec{b})$

\vec{a} and \vec{b} form a base for \mathbb{R}^2 . That means any $\vec{x} \in \mathbb{R}^2$ is a linear combination of these two ... including \vec{c} .

Thus $\vec{c} : ld(\vec{a}, \vec{b})$

Thus $\vec{a} : ld(\vec{b}, \vec{c}) \vee \vec{b} : ld(\vec{a}, \vec{c}) \vee \vec{c} : ld(\vec{a}, \vec{b})$

□

Theorem 6 (Every vectorspace has a basis).

Theorem 7 (All bases of a vectorspace S have the same size).

1.4.2 Inner product spaces

Vector spaces don't define anything about lengths, angles or projections. This lack is alleviated by adding a inner product.

Definition 9 (Inner product space). *The inner product is defined as any operation that has the following properties:*

- $(a\vec{u}) \cdot \vec{v} = a(\vec{u} \cdot \vec{v})$
- $(\vec{u} + \vec{v}) \cdot \vec{w} = \vec{u} \cdot \vec{w} + \vec{v} \cdot \vec{w}$
- $\vec{u} \cdot \vec{v} = \vec{v} \cdot \vec{u}$
- $\vec{v} \neq \vec{0} \rightarrow \vec{v} \cdot \vec{v} > 0$

In inner product spaces, we can define norms, orthogonality, and angles.

As for norms:

$$|\vec{v}|^2 = \vec{v} \cdot \vec{v}$$

And orthogonality:

$$\vec{v} \perp \vec{u} \leftrightarrow \vec{v} \cdot \vec{u} = 0$$

And finally angles:

$$\cos\theta = \frac{\vec{v} \cdot \vec{u}}{|\vec{v}||\vec{u}|}$$

As one nice little application, consider this statement.

Proof. Suppose $|\vec{u}| = |\vec{v}|$. Proof that $\vec{u} + \vec{v} \perp \vec{u} - \vec{v}$

This is to prove that $(\vec{u} + \vec{v}) \cdot (\vec{u} - \vec{v}) = 0$

The above can be rewritten to $\vec{u} \cdot \vec{u} - \vec{u} \cdot \vec{v} + \vec{v} \cdot \vec{u} - \vec{v} \cdot \vec{v}$

The two middle terms cancel out, and the two outer terms equal $|\vec{u}|^2$ and $|\vec{v}|^2$, respectively.

Using the given, these terms are equal.

Thus $\vec{u} + \vec{v} \perp \vec{u} - \vec{v}$

□

Other properties of the norm are also easily proved:

- $|a\vec{v}| = |a||\vec{v}|$
- $|\vec{v} + \vec{w}| = |\vec{v}| + |\vec{w}|$
- $|\vec{v}| \geq 0$
- $|\vec{v}| = 0 \leftrightarrow \vec{v} = \vec{0}$

Two more important statements that can be proven for the general inner product spaces are the pythagorean theorem and the Cauchy-Schwartz inequality.

Proof. Pythagorean theorem.

Suppose $\vec{u} \perp \vec{v}$. Proof that $|\vec{u}|^2 + |\vec{v}|^2 = |\vec{u} + \vec{v}|^2$

|
Thus $|\vec{u}|^2 + |\vec{v}|^2 = |\vec{u} + \vec{v}|^2$

□

Proof. Cauchy-Schwartz inequality.

Proof that $|\vec{u}||\vec{v}| \geq |\vec{u} \cdot \vec{v}|$

|
Thus $|\vec{u}||\vec{v}| \geq |\vec{u} \cdot \vec{v}|$

□

An implementation of this inner product in dimension-free oriented length-space would be:

$$\vec{v} \cdot \vec{w} = |\vec{v}||\vec{w}| \cos \theta$$

Based on this definition, the projection of \vec{v} onto \vec{u} is defined as:

$$P_{\vec{u}}(\vec{v}) = |\vec{v}| \cos \theta \frac{\vec{u}}{|\vec{u}|} = \frac{\vec{u} \cdot \vec{v}}{|\vec{u}|^2} \vec{u}$$

The direct equivalent of this inner product from oriented length space to \mathbb{R}^n would be:

$$\vec{v} \cdot \vec{w} = \sum_n v_n w_n$$

The following is a proof that the two implementations of inner product are equivalent.

Proof. Suppose $\vec{u} = u_1 \vec{e}_1 + u_2 \vec{e}_2 + u_3 \vec{e}_3$ and $\vec{v} = v_1 \vec{e}_1 + v_2 \vec{e}_2 + v_3 \vec{e}_3$. Proof that $\vec{u} \cdot \vec{v} = \sum_n v_n u_n$

$$\vec{u} \cdot \vec{v} = (u_1 \vec{e}_1 + u_2 \vec{e}_2 + u_3 \vec{e}_3) \cdot (v_1 \vec{e}_1 + v_2 \vec{e}_2 + v_3 \vec{e}_3)$$

This is written out as $u_1 v_1 (\vec{e}_1 \cdot \vec{e}_1) + u_1 v_2 (\vec{e}_1 \cdot \vec{e}_2) + \dots$

Of this, almost all terms cancel out, leaving $u_1 v_1 + u_2 v_2 + u_3 v_3$

Thus $\vec{u} \cdot \vec{v} = \sum_n v_n u_n$

□

Note that if we were to chose a *nonorthonormal* basis, the inner product would not be reduced so nicely.

Table 1.3: Important implemantations of vector spaces, inner product spaces and algebras

Vector space scalar prod	Inner product space inner prod	Algebra norm
addition	norm	outer prod
\mathbb{R}^n	$\sum_n u_i v_i$	$\sqrt{\sum_i v_i^2}$
\mathbb{R}^n in polar		
$C_{[a,b]}$	$\int_a^b u(t)v(t) dt$	
X on Ω	$E[XY]$	$E[X]$
matrices		$(\sum_x \sum_y a_x b_y)_{i,j}$ (aka.linear algebra)
G^3	$\vec{v} \cdot \vec{u}$	$\vec{v} \wedge \vec{u}$ (aka.geometric algebra)
\mathbb{R}		(aka. ordinary algebra)
Booleans		(aka. boolean algebra)

A few comments to the different spaces shown here. X on Ω is a inner product space very similar to $C_{[a,b]}$, but note that Ω itself is not necessarily even a vectorsace.

Here is one more example of an inner product. Consider the vector-space $L^2(\mathbb{R}^2 \rightarrow \mathbb{R}^3)$, that is, the space of square-integrable functions that take two input arguments \vec{x} and return a 3-vector \vec{y} . How would you define an inner product to turn this into an inner-product-space? Well, the choice is up to you, but most often one choses:

$$\langle f, g \rangle := \int_{\mathbb{R}^2} \langle f(\vec{x}), g(\vec{x}) \rangle_i d\vec{x}$$

, where $\langle \vec{a}, \vec{b} \rangle_i$ is defined as $\sum_{i=0}^3 a_i b_i$. You can prove for yourself that this is indeed an inner product!

1.4.3 Excursion: More on orthogonality and preview of function decompositions

Orthogonality turns out to be an important concept for statistics and signal analysis, so we'll look at it in a little more detail here. Why is orthogonality so important though? Linear independence allows us to take a complex signal and decompose it into simpler, independent parts. Orthogonality ensures that these parts are easy to calculate.

Although conceptually similar, orthogonality is a stricter concept than linear independence. It requires an inner product space instead of just a vector space. Also, two vectors may be linearly independent, but not orthogonal (although we can use Gram-Schmidt orthogonalisation to make any li vectors orthogonal).

Proof. If a set of vectors is orthogonal, then it is linearly independent: $B : \text{orth} \rightarrow B : \text{li}$.

Suppose that $\forall b_i, b_j \in B : b_i \cdot b_j = 0$ Proof that $\forall \alpha_1, \dots, \alpha_n : \sum \alpha_i b_i = 0 \rightarrow \alpha_1 = \dots = \alpha_n = 0$

Let $\alpha_1^0, \dots, \alpha_n^0$ be chosen. Suppose $\sum \alpha_i^0 b_i = 0$ Proof that $\alpha_1 = \dots = \alpha_n = 0$

Proof that $\alpha_j^0 = 0$ for any $j \in [1, n]$

$$\sum \alpha_i^0 b_i = 0$$

Multipled by b_j :

$$\sum \alpha_i^0 b_i \cdot b_j = 0 \cdot b_j$$

With $B : \text{orth}$:

$$\alpha_j^0 = 0$$

Thus $\alpha_j^0 = 0$ for any $j \in [1, n]$

Thus $\alpha_1 = \dots = \alpha_n = 0$

Thus $\forall \alpha_1, \dots, \alpha_n : \sum \alpha_i b_i = 0 \rightarrow \alpha_1 = \dots = \alpha_n = 0$

□

This is profound. For example, the cos-sin-Fourier-basis is hard to prove to be linearly independent. But we can use the *stricter* property of orthogonality to prove that it must also be linearly independent.

It is good to know that although orthogonality helps us to prove linear independence, it doesn't help us to prove that a set is a base, because for that we also need the set to span the whole space.

Proof. In the infinite dimensional case, a orthogonal set $B \subseteq V$ does not have to be a base of V . A good example would be a set of linear functions in $V = C_{[a,b]}$ - they can never span quadratic functions. □

Fourier decomposition In every vectorspace a vector can be expressed as a sum of the basevectors like this: $v = \alpha_1 b_1 + \alpha_2 b_2 + \dots$ If the base is orthonormal, we additionally get the benefit that the coefficients α are very easy to calculate: $\alpha_i = v \cdot b_i$. This way of calculating the coefficients is called the Fourier decomposition.

Proof. Let B be an orthonormal base of V . Then for any $\vec{v} \in V$ the n th coefficient α_n can be easily calculated as $\langle \vec{v}, \vec{b}_n \rangle$

For any vectorspace it holds that $\forall \vec{v} \in V : \exists \alpha_1, \dots, \alpha_N : \sum \alpha_k \vec{b}_k = \vec{v}$.

Suppose B to be orthonormal. Proof that $\alpha_n = \langle \vec{v}, \vec{b}_n \rangle$

$$\langle \vec{v}, \vec{b}_n \rangle = \langle \sum \alpha_k \vec{b}_k, \vec{b}_n \rangle$$

$$= \alpha_0 \langle \vec{b}_0, \vec{b}_n \rangle + \alpha_1 \langle \vec{b}_1, \vec{b}_n \rangle + \dots + \alpha_n \langle \vec{b}_n, \vec{b}_n \rangle + \dots + \alpha_N \langle \vec{b}_N, \vec{b}_n \rangle$$

$$= \alpha_0 0 + \alpha_1 0 + \dots + \alpha_n 1 + \dots + \alpha_N 0$$

Thus $\alpha_n = \langle \vec{v}, \vec{b}_n \rangle$

□

This is much easier than the case where the base is *not* orthonormal. If that is the case, we have to calculate the coefficients α_n by using the projections:

$$\vec{v} = \sum \alpha_k \vec{b}_k, \text{ with } \alpha_k = \gamma_k P_{\vec{b}_k}(\vec{v}) = \gamma_k \frac{\vec{b}_n \cdot \vec{v}}{|\vec{b}_n|^2} \vec{b}_n, \text{ with } \gamma_k \text{ to be determined.}$$

An alternative, but equally expensive method would be to use linear algebra:

$$\vec{v} = [\vec{b}_1, \vec{b}_2, \dots, \vec{b}_N] \vec{\alpha}$$

Calling the matrix $[\vec{b}_1, \vec{b}_2, \dots, \vec{b}_N]$ the basematrix B , we get:

$$\vec{v} = B \vec{\alpha}$$

$$B^{-1} \vec{v} = \vec{\alpha}$$

This looks simple enough, but unfortunately, inverting a matrix is a $\Theta(N^3)$ operation, and matrix multiplication is still a $\Theta(N^{>2.3})$ operation. Contrary to that, the evaluation of a polynomial is a $\Theta(N)$ operation when using Horner's method.

Gram-Schmidt orthogonalisation For every set of n vectors we can find a set of orthogonal vectors like this:

- $b_1 = v_1$
- $b_2 = v_2 - \text{prj}(v_2, b_1) = v_2 - \frac{v_2 \cdot b_1}{b_1 \cdot b_1} b_1$
- $b_3 = v_3 - \text{prj}(v_3, b_2) - \text{prj}(v_3, b_1)$
- ...

Orthogonal complement

1.5 Linear algebra

In the previous section on general algebra we dealt with vector spaces (subspaces, linear independence, bases) and inner product spaces (norms, orthogonality). Here we'll mostly deal with the inner-product space of vectors and occasionally with the vector-space of matrices (but really, almost exclusively with the former ...).

1.5.1 Spaces

We'll work a lot with a few vector-spaces and transformations from and to those spaces.

Definition 10. Let W be a space. Then $V \subseteq W$ is a subspace, if:

- $\forall v_1, v_2 \in V : v_1 + v_2 \in W$
- $\forall v \in V : \forall r \in R : rv \in V$

Definition 11 (Nullspace). \mathcal{N}_A is the nullspace of A . It is defined as $\mathcal{N}_A = \{x | Ax = 0\}$

Definition 12 (Columnspace). \mathcal{C}_A is the columnspace of A . It is defined as $\mathcal{C}_A = \{y | Ax = y\}$

Definition 13 (Rowspace). \mathcal{R}_A is the rowspace of A . It is defined as $\mathcal{R}_A = \{y | A^T x = y\}$

Definition 14 (Rank). r_A is the rank of A . It is defined as the dimension of \mathcal{C}_A , $\dim_{\mathcal{C}_A}$

Definition 15 (Nullity). n_A is the nullity of A . It is defined as the dimension of \mathcal{N}_A , $\dim_{\mathcal{N}_A}$

rops and cops as matrix multiplication ...

$$\text{rops}(A) = \text{rops}(I)A = RA$$

$$\text{cops}(A) = A \text{cops}(I) = AC$$

rops don't change \mathcal{N}_A Row- and column-operations (rops and cops) seem somewhat trivial at first and not worth any proof writing efforts. However, many theorems of linear algebra are much easier proved if we first reduce the matrices to their RRLE (reduced row linear echelon) form.

Theorem 8. Row-operations on A don't change \mathcal{N}_A .

Proof. Let $A' = \text{rops}(A) = RA$ Proof that $\mathcal{N}_A = \mathcal{N}_{A'}$

Let $x_0 : Ax_0 = 0$ Proof that $\exists y_0 : A'y_0 = 0$

Try $y_0 = x_0$ Proof that $A'y_0 = 0$

$$RAy_0 = 0$$

Thus $A'y_0 = 0$

Thus $\exists y_0 : A'y_0 = 0$

Let $y_0 : A'y_0 = 0$ Proof that $\exists x_0 : Ax_0 = 0$

This is equivalent to stating that $\exists x_0 : R^{-1}A'x_0 = 0$

Try $x_0 = y_0$ Proof that $R^{-1}A'x_0 = 0$

$$R^{-1}A'x_0 = 0$$

Thus $R^{-1}A'x_0 = 0$

Thus $\exists x_0 : Ax_0 = 0$

Thus $\mathcal{N}_A = \mathcal{N}_{A'}$

□

Therefore, when searching for the special solutions to a problem $Ab = 0$, we can use Gauss-Elimination and RREF without any problems.

cops don't change $C(A)$...

rops don't change $C(A)$ if A is invertible ...

We can now print an overview of the different spaces that are associated with a matrix A of dimension $m \times n$ and rank r .



The rowspace of A can be visualized using the line-intersection view of matrix-equations: it contains all the points that lie in the intersection of all the lines that make up the matrix. The columnspace of A can be visualized using the vector-image of A : it contains all the points that are spanned by A . Notice how we included the previous theorem: any combination x of a particular solution x_r and any vector in the nullspace x_n is also a solution.

We should look in more detail at this graphic. Note, for example, that \mathcal{N}_A and \mathcal{C}_{A^T} seem to be orthogonal. Indeed:

Theorem 9 ($\forall v \in \mathcal{N}_A, w \in \mathcal{C}_{A^T} : v \perp w$).

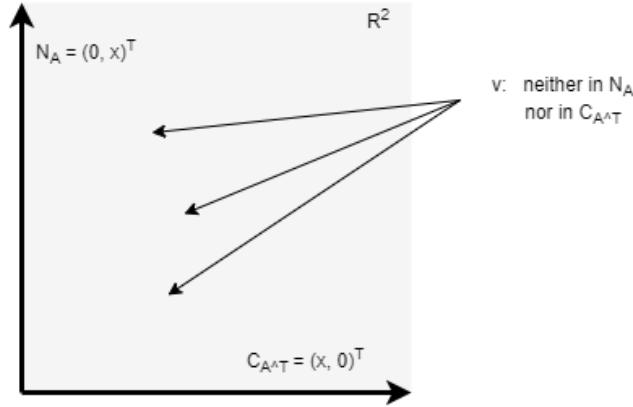
Proof.

$$\begin{aligned} w^T v = 0 & \quad \text{with } \exists u : \mathbf{A}^T u = w \\ u^T \mathbf{A} v = 0 & \quad \text{with } \mathbf{A} v = 0 \\ u^T 0 = 0 & \quad \text{which is trivially true.} \end{aligned} \tag{1.1}$$

Thus, \mathcal{N}_A and \mathcal{C}_{A^T} only intersect in $\vec{0}$. □

However, note that $\mathbb{R}^n \geq \mathcal{N}_A \cup \mathcal{C}_{A^T}$. As an example, consider $\mathbf{A} = \begin{bmatrix} 1 & 0 \\ 0 & 0 \end{bmatrix}$. We get $\mathcal{N}_A = \begin{bmatrix} 0 \\ x \end{bmatrix}$ and $\mathcal{C}_{A^T} = \begin{bmatrix} x \\ 0 \end{bmatrix}$. Now consider $\vec{v} = \begin{bmatrix} 1 \\ 1 \end{bmatrix}$, which is neither in \mathcal{N}_A nor in \mathcal{C}_{A^T} . As illustration, look at fig. 1.5.1.

Figure 1.1: When A is not full rank, there are infinitely many vectors v that are neither in \mathcal{R}_A nor in \mathcal{N}_A . When it comes to spatial dimensions, $1 + 1 \neq 2$!



1.5.2 Change of basis

Let V be a vector space. Let 0 be the canonical basis for that vector space. Let $A = \{\vec{a}_1, \dots, \vec{a}_N\}$ and $B = \{\vec{b}_1, \dots, \vec{b}_N\}$ be two other basis for that vectorspace. Let \mathbf{A} be the matrix $[\vec{a}_1 \dots \vec{a}_N]$ and $\mathbf{B} = [\vec{b}_1 \dots \vec{b}_N]$

Every vector \vec{v} can be expressed as a linear sum of the basisvectors in A , that is $\vec{v} = \sum_n \alpha_n \vec{a}_n$. That same thing in matrix notation: $\vec{v} = \mathbf{A}(\vec{v})_A$, where $(\vec{v})_A$ is the coordinates α of \vec{v} with respect to the basis A . Correspondingly, for B we have $\vec{v} = \sum_n \beta_n \vec{b}_n = \mathbf{B}(\vec{v})_B$.

Note that within A and B , we express the basevectors with respect to the canonical basis 0 , that is, we should really write $\mathbf{A} = [(\vec{a}_1)_0 \dots (\vec{a}_N)_0]$. Note also that $[(\vec{a}_1)_A \dots (\vec{a}_N)_A] = \mathbf{I}$, the identity matrix.

We can use this to obtain a simple formula for the change of basis.

$$\vec{v} = \mathbf{A}(\vec{v})_A$$

$$\vec{v} = \mathbf{B}(\vec{v})_B$$

$$(\vec{v})_B = \mathbf{B}^{-1} \mathbf{A}(\vec{v})_A$$

But inverses are notoriously hard to calculate. Fortunately, there is another approach. Call $\mathbf{T}_{BA} = [(\vec{a}_1)_B \dots (\vec{a}_N)_B]$ the *transition matrix*. We can prove that $\mathbf{B}^{-1} \mathbf{A} = \mathbf{T}_{BA}$:

$$\begin{aligned}
 \mathbf{B}^{-1} \mathbf{A}(\vec{v})_A &= \mathbf{T}_{BA}(\vec{v})_A \\
 &= \sum_n (\vec{v}_n)_A (\vec{a}_n)_B \\
 &= \sum_n (\vec{v}_n)_A \mathbf{B}^{-1} \mathbf{A}(\vec{a}_n)_A \\
 &= \mathbf{B}^{-1} \mathbf{A} \sum_n (\vec{v}_n)_A (\vec{a}_n)_A \\
 &= \mathbf{B}^{-1} \mathbf{A} \mathbf{I}(\vec{v})_A \\
 (\vec{v})_A &= (\vec{v})_A
 \end{aligned} \tag{1.2}$$

Using $\mathbf{T}_{BA} = \mathbf{B}^{-1} \mathbf{A}$, a lot of statements are trivial to prove:

- $\mathbf{T}_{BA} = \mathbf{T}_{AB}^{-1}$
- $\mathbf{T}_{CA} = \mathbf{T}_{CB} \mathbf{T}_{BA}$

1.5.3 Linear transformations

Definition 16. Let U and V be two vector spaces and $f : U \rightarrow V$. Then f is a linear transform if

- f preserves scalar multiplication: $f(\alpha \vec{u}) = \alpha f(\vec{u})$

- f preserves vector addition: $f(\vec{u}_1 + \vec{u}_2) = f(\vec{u}_1) + f(\vec{u}_2)$

There are a bunch of properties to linear transformations that can be useful to us.

There is a unique linear transform from the basis of U to any set of vectors in V that we want. In other words, any linear transform f is completely determined by the matrix $[f(\vec{b}_1) \dots f(\vec{b}_N)] = [\vec{v}_1 \dots \vec{v}_N]$. That means if we don't know the transform, but we do know the results of the transform on a basis, then we can reconstruct the transform with certainty.

Let $B = \{\vec{b}_1, \dots, \vec{b}_N\}$ be a basis for U . Let $\{\vec{v}_1, \dots, \vec{v}_N\}$ be any vectors in V that we may choose. Proof that there is a unique function $f : U \rightarrow V$ such that $f(\vec{b}_i) = \vec{v}_i$

Try $f(\vec{x}) = \mathbf{V}(\vec{x})_B$ Proof that $f(\vec{b}_i) = \vec{v}_i$ and f is a linear transform.

Part 1: Proof that $f(\vec{b}_i) = \vec{v}_i$

$$f(\vec{b}_i) = \mathbf{V}(\vec{b}_i)_B = \mathbf{V}\vec{e}_i = \vec{v}_i$$

Thus $f(\vec{b}_i) = \vec{v}_i$

Part 2: Proof that f is unique

We could not have obtained any other form of f than $f(\vec{x}) = \mathbf{V}(\vec{x})_B$. This is because for *any* linear transform from $U \rightarrow V$ we have:

$$f(\vec{x}) = f(\mathbf{B}(\vec{x})_B) = f(\sum_n (x_n)_B \vec{b}_n) = \sum_n (x_n)_B f(\vec{b}_n)$$

Using the result from part 1, this cannot be any other function than:

$$\sum_n (x_n)_B f(\vec{b}_n) = \sum_n (x_n)_B \vec{v}_n = \mathbf{V}(\vec{x})_B$$

Thus f is unique

Part 3: Proof that f is a linear transform

|

Thus f is a linear transform

Thus $f(\vec{b}_i) = \vec{v}_i$ and f is a linear transform.

Thus there is a unique function $f : U \rightarrow V$ such that $f(\vec{b}_i) = \vec{v}_i$

□

A whole bunch of other properties are now easily proved. Let f and g be linear transforms from U to V . The following are also linear transforms:

- αf
- $f + g$
- f^{-1} (if it exists)
- fg (here $g : V \rightarrow W$)

Let $f : U \rightarrow V$ be a linear transform. Then the following are equivalent:

- If $f(\vec{u}) = \vec{0}$, then $\vec{u} = \vec{0}$
- f is one-to-one
- f maps linearly independent vectors to linearly independent vectors.

Prove that a transform can be split up into multiple transforms on the basis vectors. As an example, consider the case of a rotation. A diagonal rotation can be reproduced by a rotation first around one, then around another axis.

A linear transformation can also be a change of basis when it is on a vectorspace and invertible.

1.5.4 Linear independence

Definition 17. B is linearly independent if $\forall b \in B : b \neq \sum r_n b_n$, with $b_n \in B/b$.

Theorem 10. A better definition could be stated as such: B is linearly independent, if the only solution to $Bx = 0$ is the zero-vector.

Proof. Suppose $\forall \vec{b} \in \mathbf{B} : \vec{b} \neq \sum_{B/b} \alpha_i \vec{b}_i$ Proof that $\mathbf{B}\vec{x} = \vec{0} \rightarrow \vec{x} = \vec{0}$

Suppose $\mathbf{B}\vec{x} = \vec{0}$ Proof that $\vec{x} = \vec{0}$

By contradiction. Assume $\vec{x} \neq \vec{0}$ Proof that there is a contradiction.

We have $\mathbf{B}\vec{x} = \vec{0}$

Consider the vector \vec{b}_3 .

Thus: $\sum_{B/b_3} x_i \vec{b}_i = -x_3 \vec{b}_3$

Or: $\sum_{B/b_3} \frac{-x_i}{x_3} \vec{b}_i = -x_3 \vec{b}_3$

This contradicts the statement that $\forall \vec{b} \in \mathbf{B} : \vec{b} \neq \sum_{B/b} \alpha_i \vec{b}_i$

Thus there is a contradiction.

Thus $\vec{x} = \vec{0}$

Thus $\mathbf{B}\vec{x} = \vec{0} \rightarrow \vec{x} = \vec{0}$

Suppose $\mathbf{B}\vec{x} = \vec{0} \rightarrow \vec{x} = \vec{0}$ Proof that $\forall \vec{b} \in \mathbf{B} : \vec{b} \neq \sum_{B/b} \alpha_i \vec{b}_i$

Let $\vec{b}_0 \in \mathbf{B}$ Proof that $\vec{b}_0 \neq \sum_{B/b_0} \alpha_i \vec{b}_i$

By contradiction. Assume $\vec{b}_0 = \sum_{B/b_0} \alpha_i \vec{b}_i$ Proof that there is a contradiction.

Since $\vec{b}_0 = \sum_{B/b_0} \alpha_i \vec{b}_i$

we get $\vec{0} = \sum_B \alpha_i \vec{b}_i$, where $\alpha_0 = -1$

In matrix notation: $\vec{0} = \mathbf{B}\vec{\alpha}$, with a non-zero $\vec{\alpha}$.

This contradicts our assumption that $\mathbf{B}\vec{x} = \vec{0} \rightarrow \vec{x} = \vec{0}$

Thus there is a contradiction.

Thus $\vec{b}_0 \neq \sum_{B/b_0} \alpha_i \vec{b}_i$

Thus $\forall \vec{b} \in \mathbf{B} : \vec{b} \neq \sum_{B/b} \alpha_i \vec{b}_i$

□

Definition 18. Let V be a subspace. $B \subseteq V$ is a basis for V if B is linearly independent and $\forall v \in V : v = \sum r_n b_n$, with $b_n \in B$

1.5.5 Determinant



Definition Consider a 2×2 matrix \mathbf{A}

$$\mathbf{A} = \begin{bmatrix} a_{11} & a_{12} \\ a_{21} & a_{22} \end{bmatrix}$$

$\det \mathbf{A}$ is defined as $a_{11}a_{22} - a_{12}a_{21}$. The definitions for higher order matrix-determinants are similar, but so extremely tedious that we just ignore them here.

Proof. Let $\mathbf{A}\vec{x} = 0$ and $\vec{x} \neq \vec{0}$. Then it must hold that $\det \mathbf{A} = 0$. In fact, this holds in two directions:

$$\exists \vec{x} \neq \vec{0} : \mathbf{A}\vec{x} = 0 \leftrightarrow \det \mathbf{A} = 0$$

We can prove this making use of theorem 10. From the given it follows that \mathbf{A} is linearly dependent. Thus $\exists \vec{a} \in \mathbf{A} : \vec{a} = \sum_{A/a} \alpha_i \vec{a}_i$. In the case of a 2×2 matrix this expression becomes very simple: $\vec{a}_1 = \alpha \vec{a}_2$. Breaking \vec{a}_1 into its components

- $a_{11} = \alpha a_{12}$
- $a_{21} = \alpha a_{22}$

Putting this into the definition of a determinant we get:

$$\begin{aligned} \det \mathbf{A} &= a_{11}a_{22} - a_{12}a_{21} \\ &= \alpha a_{12}a_{22} - \alpha a_{12}a_{22} \\ &= 0 \end{aligned} \tag{1.3}$$

For higher order matrices, the proof follows from induction or something like that. \square

Interpretation The size of the determinant can be seen as the scaling-factor of the transformation described by the matrix A . It is also a measure of how much linearly independent the rows/cols of A are. The size of the determinant equals the size of the (hyper-)parallelogram spanned by the columns. If two vectors are almost linearly dependent, they will be almost parallel, leading to a very small area of the parallelogram. So if you have a small determinant, your columns are almost dependent. If you have a large one, your columns are very orthogonal.


1.5.6 Rank nullity Theorem

Theorem 11. $\mathcal{N}_A = \{0\} \rightarrow \det A \neq 0$

Proof. This is almost self-proving. $\mathcal{N}_A = \{0\}$ is equivalent to writing $\mathbf{A}\vec{x} = \vec{0} \rightarrow \vec{x} = \vec{0}$. By theorem 10 this means that \mathbf{A} is linearly independent. And by theorem 1.5.5 this means that $\det A \neq 0$. \square

Theorem 12. Let $A : X \rightarrow Y$. Then:

$$\mathcal{N}_A = \{0\} \rightarrow \mathcal{C}_A = Y$$



Theorem 13. Rank nullity: this is a fundamental theorem of linear algebra. Let A be of dimension $r \times c$ and full rank. Then:

$$r_A + n_A = c$$

Theorem 14.

$$r_A = \dim \mathcal{C}_A = \dim \mathcal{R}_A$$

Theorem 15.

$$A : \text{basis}_{\mathcal{C}_A} \rightarrow r_A = n$$

1.5.7 Special matrices

Symmetric matrices are simple but very useful.

Definition 19 (A matrix \mathbf{A} is symmetric). iff $\mathbf{A} = \mathbf{A}^T$

Positive (semi-)definite matrix are a special case of symmetric matrices.

Definition 20 ($\mathbf{A} : \text{PSD} \leftrightarrow \mathbf{A} : \text{sym} \wedge \forall z : z^T \mathbf{A} z \geq 0$). From the first condition, $\mathbf{A} : \text{sym}$, it follows that the eigenvectors are orthogonal (see eq. 19), from the second one, $\forall z : z^T \mathbf{A} z \geq 0$, it follows that the eigenvalues are all greater or equal than 0 (see eq. ...).

Gram matrices appear often in important theorems. They turn out to be PSD.

Definition 21 (A Gram matrix). is a matrix $\mathbf{A} = \mathbf{B}^T \mathbf{B}$

Theorem 16 ($\text{rank}_A = \text{rank}_{A^T A}$).

Theorem 17 ($A^T A : \text{PSD}$).

Orthogonal matrices make tons of calculations simpler.

Definition 22 (A matrix \mathbf{A} is orthogonal). iff $\forall x_1, x_2 \in \mathbf{A}, x_1 \neq x_2 : x_1 \perp x_2$

Theorem 18 (If A is orthogonal, then $A^{-1} = A^T$). This is trivially proven. By the definition of orthogonality we have $\mathbf{A}^T \mathbf{A} = \mathbf{I}$

1.5.8 Eigenvalues and eigenvectors

$$\begin{aligned} Ax &= \lambda x \\ (A - \lambda I)x &= 0 \\ \det A - \lambda I &= 0 \text{ using eq. 1.5.5} \end{aligned} \tag{1.4}$$

Any square matrix can be eigenvector (aka. spectrally) decomposed: $\mathbf{A} = \mathbf{X} \boldsymbol{\Lambda} \mathbf{X}^{-1}$.

Theorem 19. $[\mathbf{A} : \text{sym} \rightarrow \mathbf{X} : \perp]$

$$\begin{aligned} < Ax, y > &= < x, A^T y > \\ \text{Since } A : \text{sym} : < Ax, y > &= < x, Ay > \\ \text{Since } x, y \text{ are eigenvalues of } A : < \lambda_x x, y > &= < x, \lambda_y y > \\ \lambda_x < x, y > &= \lambda_y < x, y > \end{aligned} \tag{1.5}$$

Thus, either $\lambda_x = \lambda_y$ or $x \perp y$

Theorem 20 (As a corollary, the eigenvalue decomposition can be simplified). : $\mathbf{A} : \text{sym} \rightarrow \mathbf{A} = \mathbf{X} \boldsymbol{\Lambda} \mathbf{X}^{-1} = \mathbf{X} \boldsymbol{\Lambda} \mathbf{X}^T$

1.5.9 Singular value decomposition

Square matrices may be decomposed to $\mathbf{A} = \mathbf{X} \boldsymbol{\Lambda} \mathbf{X}^{-1}$. More generally, any matrix may be decomposed to $\mathbf{A} = \mathbf{U} \boldsymbol{\Sigma} \mathbf{V}^T$. Before proving this, lets first interpret that statement from a matrix point of view.

Consider some basis $B_{C_{A^T}}$. Through the transformation $\mathbf{A} : \mathbb{R}^n \rightarrow \mathbb{R}^m$ that basis $B_{C_{A^T}}$ is transformed into C_A . So far so easy - but we have an additional requirement: we want that transformed basis to be a basis itself: $\mathbf{A} B_{C_{A^T}} = B_{C_A}$.

There may be many bases $B_{C_{A^T}}$, but with that additional requirement, it must truely be a special basis. Let's call it \mathbf{V} , and the resulting new basis $\mathbf{AV} = \mathbf{U}$. Also, we want \mathbf{U} not only to be a basis, but also normal; so let's normalize it to $\mathbf{U}\boldsymbol{\Sigma}$. With all that we get: $\mathbf{AV} = \mathbf{U}\boldsymbol{\Sigma}$.

So these are our requirements. Let's jump into a proof.

Consider $\mathbf{A}^T \mathbf{A}$. It is a Gram-matrix, and as such PSD. We have:

$$\mathbf{A}^T \mathbf{A} v_i = \lambda_i v_i$$

by eigenvalue decomposition. Define $\sigma_i = \sqrt{\lambda_i}$. This exists, because $\mathbf{A}^T \mathbf{A} : \text{PSD}$.

$$\mathbf{A}^T \mathbf{A} v_i = \sigma_i^2 v_i$$

Case 1 : Assume that \mathbf{A} is full rank; i.e. $\forall i : \sigma_i > 0$. Let $u_i = \frac{1}{\sigma_i} \mathbf{A} v_i$. Expressed in matrix-form this gives us

$$\begin{aligned} \mathbf{U} &= \mathbf{A} \mathbf{V} \Sigma^{-1} \\ \mathbf{U} \Sigma &= \mathbf{A} \mathbf{V} \\ \mathbf{U} \Sigma \mathbf{V}^T &= \mathbf{A} \end{aligned} \tag{1.6}$$

Note that u_i is, by construction, an eigenvector of $\mathbf{A} \mathbf{A}^T$:

$$\begin{aligned} \mathbf{A} \mathbf{A}^T u_i &= \mathbf{A} \mathbf{A}^T \mathbf{A} v_i \frac{1}{\sigma_i} \\ &= A \sigma_i^2 v_i \frac{1}{\sigma_i} \\ &= \sigma_i^2 A v_i \frac{1}{\sigma_i} \\ &= \sigma_i^2 u_i \end{aligned} \tag{1.7}$$

Rephrased It is really worth it to rephrase that proof again, this time fully matrix-oriented.

Our goal is to find a matrix V such that $V : B_{C_{A^T}}$ and also $AV = U' : B_{C_A}$. Normalize U' to $U' = U\Sigma$, where Σ is a diagonal matrix of scalings. Thus our requirements can be written as $AV = U\Sigma$ or, since V is a basis and as such orthogonal, $A = U\Sigma V^T$.

Let's find those very special V and U . To find V consider $A^T A$ (note that $A^T A : \text{PSD}$):

$$\begin{aligned} A^T A &= (U\Sigma V^T)^T U\Sigma V^T \\ &= V\Sigma^T U^T U\Sigma V^T \\ &= V\Sigma U^T U\Sigma V^T && \text{because } \Sigma \text{ is diagonal} \\ &= V\Sigma\Sigma V^T && \text{because } U \text{ is orthogonal, so that } U^{-1} = U^T \end{aligned} \tag{1.8}$$

Thus V are just eigenvectors of $A^T A$ for eigenvalues σ_i^2 . Analogously we find U for $\mathbf{A} \mathbf{A}^T$.

Example for a shear-matrix : SVD means that any matrix-transformation can be rewritten as a rotation, scale and another roation. Even a shear-matrix!

```

M = [
    [1, .5],
    [0, 1]
]

U, S, VT = np.linalg.svd(M)

U @ (np.eye(2) * S) @ VT

```

SVD is often used for image compression, too. Just leave out some lower σ_i 's, then you can also leave out equally many columns / rows of U and V^T , respectively.

1.5.10 Inverse

If A has dimensions $n \times n$, then the inverse A^{-1} such that

$$AA^{-1} = A^{-1}A = I$$

exists iff $\det A \neq 0$.

Nonsquare matrices do not have an inverse, but they might have a right- or left-inverse. Consider $C = AA^T$. This is a square matrix, so it might have an inverse:

$$AA^T(AA^T)^{-1} = I$$

Calling $A^T(AA^T)^{-1} = A^{RI}$ the right inverse:

$$AA^{RI} = I$$

For C^{-1} to exist, we require C to be full-rank, which means that A must be full row rank. This requires $r \leq c$, in other words, A being a broad matrix.

Proof. If A^{RI} exists, then A^{LI} does not □

Proof. A^{RI} is not unique. Let A be of dimension $r \times c$ and full rank, with $r < c$. By rank nullity we have

$$n_A = c - r > 0$$

That is, the nullspace is nonempty. Thus $\exists x : Ax = 0$ Now try $B' = B + x$ We get

$$AB' = AB + Ax = I + 0$$

□

1.5.11 Applications

1.5.11.1 Systems of linear equations

If there are more variables than equations, the system is underdetermined. If there are more equations than variables, the system is overdetermined. A potentially solveable system is one where there are equally many variables as equations. But even then we must distinguish two cases.

Solving well determined systems There are two cases: a system is either consistent or inconsistent. The following statements are all equivalent, meaning that any one of them is related to any other one in an if-and-only-if way.

- the system is consistent
- the matrix is invertible
- the determinant is nonzero
- there is exactly one solution

We proof a few of those equivalences just for the hell of it.

Proof. There is exactly one solution if and only if the determinant is nonzero.

Proof that $|A| = 0 \leftrightarrow \exists! \vec{x} : A\vec{x} = \vec{b}$

|

Thus $|A| = 0 \leftrightarrow \exists! \vec{x} : A\vec{x} = \vec{b}$ □

On the other hand, in the inconsistent case, the following statements are equivalent:

- the system is inconsistent
- the matrix is singular (aka. noninvertible)
- the determinant equals zero
- one (or more) row (or column) is linearly dependent of the others

Solving overdetermined systems: least squares

Solving underdetermined systems: geometric bodies I like to think of underdetermined systems as (linear) geometric bodies, written in their parameterized form. A line in \mathbb{R}^3 is described by a 3×1 matrix (or rather, it's column space), a plane in \mathbb{R}^3 by a 3×2 matrix. However, it is important to note one distinction: geometric objects don't need to go through the origin, a matrix system however does. A line that does not go through the origin needs a base vector, like so:

$$\vec{x} = \begin{bmatrix} d \\ \alpha \end{bmatrix} + \vec{b}$$

A plane that does not go through the origin also needs a base vector \vec{b} , like so:

$$\vec{x} = \begin{bmatrix} \vec{p}_1 & \vec{p}_2 \end{bmatrix} \begin{bmatrix} \alpha \\ \beta \end{bmatrix} + \vec{b}$$

Here is a problem that bothered me for a while: a line needs one parameter, a plane two. An ellipsoid, too, needs two parameters. Are there any linear geometric objects in \mathbb{R}^3 that require more than three parameters? The answer is: no. Here is the proof.

Proof. For any 3×4 matrix, there is a 3×3 matrix that has the same column space, that is, that describes the same geometrical body.

Proof that $\forall A(3 \times 4) \exists A'(3 \times 3) : \mathcal{C}_A = \mathcal{C}_{A'}$

Let $A_0(3 \times 4)$. Proof that $\exists A' : \mathcal{C}_{A_0} = \mathcal{C}_{A'}$

We know from 1.4.1 that $\exists \vec{a}_0 \in A_0 : \vec{a}_0 : ld$. So Try $A' = A_0 / \vec{a}_0$.

Indeed, now A_0 and A' both form a base of the same space. So they must have the same column space.

Thus $\exists A' : \mathcal{C}_{A_0} = \mathcal{C}_{A'}$

Thus $\forall A(3 \times 4) \exists A'(3 \times 3) : \mathcal{C}_A = \mathcal{C}_{A'}$

□

However, there are *nonlinear* objects in \mathbb{R}^3 that require more than three parameters! Many curves in 3d require many parameters. *But* those curves don't form a vector-space, while lines and planes do (as long as they go through the origin).

Table 1.4: Influence of rank on solutions

	$m < n$	$m = n$	$m > n$
$r < m$	n - m free variables m - r conditions on $b \in \mathcal{C}_A$		
$r = m$ (full row rank)	1 pivot per row n - r = n - m free variables 0 conditions on $b \in \mathcal{C}_A$	X	
$r < n$			m - r conditions on $b \in \mathcal{C}_A$ n - r free variables
$r = n$ (full column rank)	X		1 pivot per column m - r = m - n conditions on $b \in \mathcal{C}_A$ 0 free variables, thus $\mathcal{N}_A = \{0\}$

Summary

$\mathbf{Ax} = \mathbf{b}$ reduces to $\mathbf{A}'\mathbf{x}' = \mathbf{0}$

Theorem 21. A problem of the form $\mathbf{Ax} = \mathbf{b}$ can be re-expressed as $\mathbf{A}'\mathbf{x}' = \mathbf{0}$, where $\mathbf{A}' = [\mathbf{A}, -\mathbf{b}]$

Proof. Proof that $\mathbf{Ax} = \mathbf{b}$ can be re-expressed as $\mathbf{A}'\mathbf{x}' = \mathbf{0}$

$$\begin{aligned}
 Ax &= b \\
 Ax - b &= 0 \\
 A_1x_1 + A_2x_2 + \dots + b &= 0 \\
 \text{Let } A' = [A, b] \text{ and } x_{n+1} = -1. \text{ Then:} \\
 A'x' &= 0
 \end{aligned}$$

Now we can use the nullspace of A' to find the solutionspace of A .

$$\begin{aligned}
 \mathcal{N}_{[Ab]} &= \{x' | [Ab]x' = 0\} \\
 &= \{x' | Ax'_{1:m} = -bx'_{m+1}\} \\
 \text{A subset of that nullspace equals the solutionset for } Ax = b: \\
 \mathcal{N}_{[Ab]} \text{ where } [x'_{m+1} = -1] &= \{x' | Ax'_{1:m} = b\}
 \end{aligned}$$

Thus $Ax = b$ can be re-expressed as $A'x' = 0$

□

Solving $\mathbf{Ax} = \mathbf{b}$

Theorem 22. If we can only find any one particular solution x_p such that $Ax_p = b$, then we get the whole solutionspace as $\mathcal{N}_A + x_p$.

Proof. Let $x_p : Ax_p = b$. Proof that $\mathcal{S}_{Ax=b} = \mathcal{N}_A + x_p$

We'll make use of the fact that $\mathcal{N}_A + x_p = \{x + x_p | Ax = 0\} = \{x | Ax = Ax_p\}$

Let $x_0 \in \mathcal{N}_A + x_p$. Proof that $x_0 \in \mathcal{S}_{Ax=b}$, i.o.w. $Ax_0 = b$

$$\begin{aligned}
 x_0 &\in \mathcal{N}_A + x_p \\
 x_0 &\in \{x + x_p | Ax = 0\} \\
 x_0 &\in \{x | A(x - x_p) = 0\} \\
 \text{Thus } Ax_0 &= Ax_p \\
 \text{Since } Ax_p &= b, \text{ it must be that } Ax_0 = b.
 \end{aligned}$$

Thus $x_0 \in \mathcal{S}_{Ax=b}$, i.o.w. $Ax_0 = b$

Let $x_0 \in \mathcal{S}_{Ax=b}$. Proof that $x_0 \in \mathcal{N}_A + x_p$, i.o.w. $Ax_0 = Ax_p$

$$\begin{aligned}
 \text{Because } x_0 &\in \mathcal{S}_{Ax=b}, \text{ we have } Ax_0 = b. \\
 \text{Also, it was given that } Ax_p &= b.
 \end{aligned}$$

Thus $x_0 \in \mathcal{N}_A + x_p$, i.o.w. $Ax_0 = Ax_p$

Thus $\mathcal{S}_{Ax=b} = \mathcal{N}_A + x_p$

□

1.5.11.2 Matrix factorisation

Eigenvalue decomposition

$$A = V\Lambda V^{-1}$$

Singular value decomposition is eigenvalue decomposition, generalized to non-square matrices.

$$A = U\Sigma V^T$$


Nonnegative matrix factorisation : Consider a dataset A , mapping people (rows) to properties (columns). You are looking for some hidden, small set of features, that groups of people have in common. In neural networks we can reconstruct images from a minimal amount of hidden features by funnelling the image through a very small hidden layer out to a large output layer. We can do the very same thing here!

$$A = UV$$

Where A has dimension $r \times c$, U associates people with their hidden features/groups $r \times f$ and V associates features/groups with the properties $f \times c$.

1.6 Fourier and Laplace

Figure 1.2: Time Base versus Fourier Base



1.6.1 Discrete Fourier Analysis

We'll begin with the discrete case. For one, this case is the basis of the FFT. Also, the math is a lot easier in the discrete case. Let N be the samplesize and δ be the samplefrequency (often equal to 44100 Hz or 48000 Hz). Let V be the samplespace of signals of size N . This is a vectorspace, equipped with an inner product of the form $\langle \vec{a}, \vec{b} \rangle = \frac{1}{N} \sum_{k=0}^{N-1} a_k(b_k)^*$, where $(*)^*$ is the complex complement.

1.6.1.1 The Fourier Base

We'll use the base $B = \{\vec{w}_N^n | n \in [0, N - 1]\}$. Here, \vec{w}_N^n is the vector we get by evaluating $e^{i2\pi \frac{n}{N} k}$ at the points $k = 0, k = 1, \dots, k = N - 1$. Thus if we call $w_N = e^{i2\pi \frac{1}{N}}$, then the vector \vec{w}_N^n consists of the elements w_N^{nk} . Throughout this chapter, k will be the index of time/the signalvector and n will be the index of frequency/the basevectors.

In the case of $N = 4$ we obtain the matrix:

$$\begin{bmatrix} e^{0.25i\pi} & e^{0.5i\pi} & e^{0.75i\pi} & e^{1.0i\pi} \\ e^{0.5i\pi} & e^{1.0i\pi} & e^{1.5i\pi} & e^{2.0i\pi} \\ e^{0.75i\pi} & e^{1.5i\pi} & e^{2.25i\pi} & e^{3.0i\pi} \\ e^{1.0i\pi} & e^{2.0i\pi} & e^{3.0i\pi} & e^{4.0i\pi} \end{bmatrix}$$

Note that fourier bases are symmetric matrices.

The Fourier-base is a very particular choice: each element of each base-vector turns out to be one of the N complex N th roots of one. We could have chosen any base-vectors, but these complex roots will turn out to have a few useful properties that we will exploit to speed up the Fourier transformation. These properties are:

1. $w_N^x = (w_N)^x$
2. $w_N^{2x} = w_{N/2}^x$
3. $w_N^{2x+N} = w_N^{2x}$
4. $w_N^{x+N/2} = -w_N^x$

It is also easy to prove that this base is orthogonal.

Proof. The vectors \vec{w}_N^n are orthogonal.

Suppose $n \neq m$. Proof that $\langle \vec{w}_N^n, \vec{w}_N^m \rangle = 0$

$$\langle \vec{w}_N^n, \vec{w}_N^m \rangle = \frac{1}{N} \sum_{k=0}^{N-1} e^{i2\pi \frac{n-m}{N} k}$$

Using the result $\sum_{n=0}^{N-1} e^{xn} = \frac{1-e^{xN}}{1-e^x}$, we find:

$$\frac{1}{N} \sum_{k=0}^{N-1} e^{i2\pi \frac{n-m}{N} k} = \frac{1}{N} \frac{1-e^{i2\pi(n-m)}}{1-e^{i2\pi \frac{n-m}{N}}}$$

It is easy to see that $e^{i2\pi(n-m)} = 1$. So the above term must equal 0.

Thus $\langle \vec{w}_N^n, \vec{w}_N^m \rangle = 0$

Suppose $n = m$. Proof that $\langle \vec{w}_N^n, \vec{w}_N^m \rangle = 1$

We use the same line of reasoning as above to obtain:

$$\langle \vec{w}_N^n, \vec{w}_N^m \rangle = \frac{1}{N} \frac{1-e^{i2\pi(0)}}{1-e^{i2\pi \frac{0}{N}}}$$

In the limit, this equals 1.

Thus $\langle \vec{w}_N^n, \vec{w}_N^m \rangle = 1$

□

Proof of spanning

1.6.1.2 Obtaining the amplitudes

After having proven that B is an orthonormal base for V , we can get to the core of the Fourier analysis: given a signal \vec{v} , how do we obtain the amplitudes in $\vec{v} = \sum_{n=0}^{N-1} \alpha_n \vec{w}_n$? Well, from paragraph 1.4.3 we know that

$$\alpha_n = \langle \vec{v}, \vec{w}_N^n \rangle = \frac{1}{N} \sum_{k=0}^{N-1} v_k e^{i2\pi \frac{n}{N} k}$$

Using Horners method, this is a $\Theta(n)$ operation. Executing it on all n coefficients, the whole process becomes a $\Theta(n^2)$ operation. A naive implementation might look like this:

```
import numpy as np

def amplitudes(signal):
    N = len(signal)
    amps = []
    for n in range(N):
        sm = 0
        for k in range(N):
            wnk = np.exp(-1j * 2 * np.pi * n * k / N)
            sm += signal[k] * wnk
        amps.append(sm/N)
    return amps

sig = [2, 1, 1, 4]
print amplitudes(sig)
```

Notice that this is a matrix operation.

1.6.2 As a matrix operation

Proof. The Fourier transform of $\vec{a} + \vec{b}$ equals the transform of \vec{a} plus the transform of \vec{b}

Note that the Fourier transform can be rewritten in matrix form.

$$\langle \vec{a}, \vec{f}_m \rangle (m) = \begin{bmatrix} f_{1,1} & \dots & f_{F,T} \\ \dots & \dots & \dots \\ f_{F,1} & \dots & f_{1,T} \end{bmatrix} \begin{bmatrix} a_1 \\ a_2 \\ \dots \\ a_T \end{bmatrix} = \begin{bmatrix} \langle \vec{a}, \vec{f}_1 \rangle \\ \langle \vec{a}, \vec{f}_2 \rangle \\ \dots \\ \langle \vec{a}, \vec{f}_F \rangle \end{bmatrix}$$

This means that the rules for matrix multiplication apply to the Fourier transformation, amongst which that matrix multiplication is distributive:

$$A(\vec{a} + \vec{b}) = A\vec{a} + A\vec{b}$$

Remember that a matrix-multiplication is always a change of basis: the vector \vec{v}_I (as displayed in the euclidian system \mathbf{I}) can also be expressed relative to the Fourier-base \mathbf{F} as \vec{v}_F .

$$\mathbf{F}^{-1}\vec{v}_I = \vec{v}_F$$

where (according to 1.5.2):

$$\mathbf{F}^{-1} = [\vec{x}_F \quad \vec{y}_F \quad \vec{z}_F]$$

□

1.6.3 Fast Fourier Transform

In the previous section we have evaluated the polynomial $\sum_{k=0}^{N-1} v_k e^{i2\pi \frac{n}{N} k}$ like this:

```
for k in range(N):
    wnk = np.exp(1j * 2 * np.pi * n * k / N)
    sm += signal[k] * wnk
```

Really, this is just the same as evaluating the polynomial $A(x) = \sum_{k=0}^{N-1} v_k x^k$, where $x = e^{i2\pi \frac{n}{N} 1} = w_N^n$ (using property 1). In other words, we calculated $A(w_N^n)$.

However, it turns out that we can simplify this process. Before we detail how exactly this section is to be simplified, we need to realize that for any polynomial $A(x)$, it holds that $A(x) = A_{even}(x^2) + x A_{odd}(x^2)$

Knowing that, we can split the evaluation in half:

$$A(w_N^n) = A_{even}(w_N^{2n}) + w_N^n A_{odd}(w_N^{2n})$$

And using the properties 2-4 of w_N^n , we can calculate:

$$\begin{aligned} A(w_N^n) &= A_{even}(w_N^{2n}) + w_N^n A_{odd}(w_N^{2n}) \\ &= A_{even}(w_{N/2}^n) + w_N^n A_{odd}(w_{N/2}^n) \end{aligned} \tag{1.9}$$

$$\begin{aligned} A(w_N^{n+N/2}) &= A_{even}(w_N^{2n+N}) + w_N^{n+N/2} A_{odd}(w_N^{2n+N}) \\ &= A_{even}(w_N^{2n}) - w_N^n A_{odd}(w_N^{2n}) \\ &= A_{even}(w_{N/2}^n) - w_N^n A_{odd}(w_{N/2}^n) \end{aligned} \tag{1.10}$$

This way, we obtain:

```
def fft(signal):
    N = len(signal)

    if N == 1:
        return signal

    sigE = []
    sigU = []
    for k in range(N):
        if k%2 == 0:
            sigE.append(signal[k])
        else:
            sigU.append(signal[k])

    ampsE = fft(sigE)
    ampsU = fft(sigU)

    amps = []
    for n in range(N/2):
        wn = np.exp(-1j * 2 * np.pi * n / N)
        an = ampsE[n] + wn * ampsU[n]
        an2N = ampsE[n] - wn * ampsU[n]
        amps.insert(n, an)
        amps.insert(n + N/2, an2N)

    return amps
```

1.6.3.1 Backtransformation

1.6.4 Fourier for musical frequencies

Musical notes are special. Here, we don't want to get the full spectrum that FFT delivers, but only a few selected frequencies.

1.6.4.1 Goertzel

Goertzel continues to work with the Fourier base. It differs from FFT in that only a few of the elements that FFT returns are actually evaluated.

1.6.4.2 PCI

PCI (pre-calculated inverse algorithm) discards of the Fourier base and instead uses the musical frequencies in the base. This results in a non-orthogonal base, but that won't hinder us.

$$\begin{aligned}\vec{s} &= \sum_F a_F e^{-2\pi i f t}, \text{ with } F = \{440, 465, \dots\} \\ &= \begin{bmatrix} \dots & e^{-2\pi i f t} & \dots \\ \dots & & \dots \end{bmatrix} \vec{a} \\ \vec{a} &= \begin{bmatrix} \dots & e^{-2\pi i f t} & \dots \\ \dots & & \dots \end{bmatrix}^{-1} \vec{s}\end{aligned}\tag{1.11}$$

1.6.5 Continuous Fourier Analysis

In Fourier analysis, we deal with the innerproduct space $C_{[a,b]}^2$: the space of square-integrable functions that are continuous from a to b . An orthogonal base for this vectorspace would be the Fourier base $\{\sin(\frac{j2\pi}{b-a}t), \cos(\frac{j2\pi}{b-a}t) | j \in \mathbb{N}\}$. Often, making use of Eulers formula, we instead write this base as $\{e^{i\frac{n2\pi}{b-a}t} | j \in \mathbb{N}\}$.

1.6.5.1 Comparing Fourier to other important series

Fourier versus Taylor Another famous expansion of functions is the Taylor-expansion. It is important to note that the Taylor expansion is a completely different beast from the Fourier expansion.

- Taylor works on locally differentiable functions, whereas Fourier works on globally integrable functions. You cannot recover the full function from its Taylor-expansion.
- The Taylor-base is non-orthogonal¹

Table 1.5: Fourier versus Taylor

	coefficients	basevectors
Fourier	$f(t) \cdot e^{i\frac{n2\pi}{b-a}t}$	$e^{i\frac{n2\pi}{b-a}t}$
Taylor	$f^{(n)}(t) t_0$	$\frac{(x-x_0)^n}{n!}$

Fourier versus PCA PCA comes a lot closer to Fourier in that in PCA we represent our data as a linear combination of orthogonal base-vectors. There are two differences though. In PCA the data is usually a matrix instead of a vector. And in PCA we don't know in advance what the set of base-vectors is going to be, but much rather chose the most fitting one for our datamatrix. It turns out that we can use the set of eigenvectors of the datamatrix².

Proof. There is a set of eigenvectors E of a matrix M that is an orthogonal basis for \mathcal{S}_M Proof that

Thus

□

¹proof required

²Strictly speaking, we don't work with the raw datamatrix, but rather its correaltion matrix. This is because for the eigenvector thing to work , we need the matrix to be symetric and centered around the origin, which the raw data matrix usually isn't.

1.6.5.2 Proving that Fourier functions form a basis

We can use Fourier to describe any functionspace - if we can prove that the Fourier functions do indeed form a basis for that space. This takes two steps: proving orthogonality and proving span.

Orthogonality This is very straightforward:

Span This requires more work. We're first going to have to prove The Weierstrass-Theorem. This theorem states that the set $B = \{x^j | j \in \mathbb{N}\}$ forms a basis for $C_{[a,b]}$. <https://psychedai.wordpress.com/2016/11/15/the-intuition-behind-bernsteins-proof-of-the-weierstrass-approximation-theorem/>

1.6.5.3 Fourier transform on vector valued functions

As long as a vector-valued function can be decomposed over a set of basis-vectors (never mind what basis vectors exactly, this works with any basis), we can still use the normal, scalar Fourier transform on them.

$$\begin{aligned} g : A^N &\rightarrow A^M \\ g(\vec{x}) &= \sum_m^M g_m(\vec{x}) \vec{e}_m \dots \text{decomposing the function over the output basis} \\ \mathcal{F}_g(\vec{u}) &= \sum_m^M \vec{e}_m \mathcal{F}_{g_m}(\vec{u}) \dots \text{Fourier transform over the input basis} \end{aligned} \tag{1.12}$$

Consider the example of a rgb-image.

$$\begin{aligned} c(x, y) &= \vec{e}_1 r(x, y) + \vec{e}_2 g(x, y) + \vec{e}_3 b(x, y) \\ \mathcal{F}_c \begin{bmatrix} f_x \\ f_y \end{bmatrix} &= \int_X \int_Y c(x, y) e^{-2\pi i(f_x x + f_y y)} dy dx \\ &= \int_X \int_Y \vec{e}_1 r(x, y) e^{-2\pi i(f_x x + f_y y)} dy dx + \dots \\ &= \vec{e}_1 \mathcal{F}_r \begin{bmatrix} f_x \\ f_y \end{bmatrix} + \vec{e}_2 \mathcal{F}_g \begin{bmatrix} f_x \\ f_y \end{bmatrix} + \vec{e}_3 \mathcal{F}_b \begin{bmatrix} f_x \\ f_y \end{bmatrix} \end{aligned} \tag{1.13}$$

Note that the resulting fourier spectrum consist of complex 3d-vectors.

Sometimes, however, you either have an input function that contains *multivector* elements, or you want to apply a filter to the fourier spectrum based on multivectors. In that case, you can find an introduction to geometric-algebra Fourier transforms here: <https://pdfs.semanticscholar.org/41ce/67428ee60748a4142dee0eea28ed997855e6.pdf>

1.6.6 Fourier Transforms and Convolutions

One major reason that Fourier transforms are so important in image processing is the convolution theorem which states that

If $f(x)$ and $g(x)$ are two functions with Fourier transforms $F(u)$ and $G(u)$, then the Fourier transform of the convolution $f(x) * g(x)$ is simply the product of the Fourier transforms of the two functions, $F(u)G(u)$.

Thus in principle we can undo a convolution. e.g. to compensate for a less than ideal image capture system:

- Take the Fourier transform of the imperfect image,
- Take the Fourier transform of the function describing the effect of the system,
- Divide the former by the latter to obtain the Fourier transform of the ideal image.
- Inverse Fourier transform to recover the ideal image.

This process is sometimes referred to as deconvolution.

		1D	2D
continuous	Transform	$\mathcal{F}_f(u) = \int_{-\infty}^{\infty} f(t)e^{-2\pi i tu} dt$	$\mathcal{F}_f(u, v) = \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} f(x, y)e^{-2\pi i(ux+vy)} dx dy$
	Convolution	$(f * g)(t) = \int_{-\infty}^{\infty} f(u)g(t-u) du$	$(f * g)(x, y) = \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} f(u, v)g(x-u, y-v) du dv$
discrete	Transform	$\mathcal{F}_f(u) = \sum_{-\infty}^{\infty} f(t)e^{-2\pi i tu}$	$\mathcal{F}_f(u, v) = \sum_{-\infty}^{\infty} \sum_{-\infty}^{\infty} f(x, y)e^{-2\pi i(ux+vy)}$
	Convolution	$(f * g)(t) = \sum_{-\infty}^{\infty} f(u)g(t-u)$	$(f * g)(x, y) = \sum_{-\infty}^{\infty} \sum_{-\infty}^{\infty} f(u, v)g(x-u, y-v)$

Motion blur, for example, can be simulated by convoluting an original image x with a simple line-shaped kernel g , yielding the blurred image y . Thus we have:

$$\begin{aligned} y &= x * g \\ Y &= XG \\ X &= Y/G \end{aligned} \tag{1.14}$$

Note that the division in the last line is pointwise (i.e. not a matrix division, which would require being able to compute an inverse).

We can implement this in python.

```

def motionBlurKernel(degrees, n):
    # rotated line
    c = int((n-1)/2)
    kernel = np.zeros((n, n))
    kernel[c, :] = np.ones(n)
    kernel = scn.rotate(kernel, degrees)

    # crop
    N, _ = kernel.shape
    delta = int((N - n) / 2)
    if delta > 0:
        kernel = kernel[delta:-delta, delta:-delta]

    # normalize
    kernel = kernel / np.sum(kernel)
    return kernel

def centerPad(kernel, n):
    l, _ = kernel.shape
    m = np.zeros((n, n))
    delta = int((n - l) / 2)
    m[delta:delta+l, delta:delta+l] = kernel
    return m

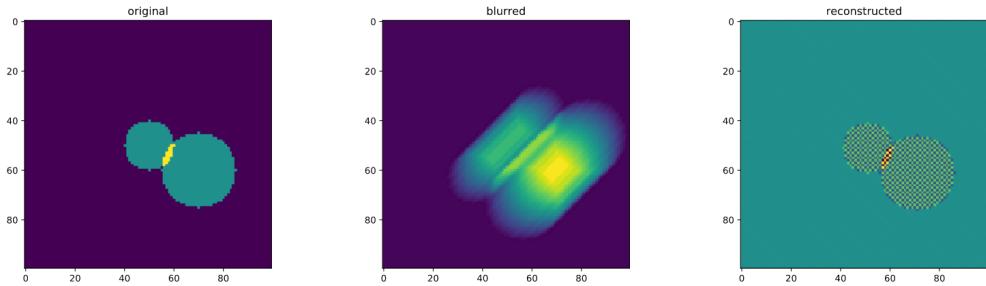
def filterMotionBlur(image, angle, width):
    L, _ = image.shape
    kernel = motionBlurKernel(angle, width)
    kernel_p = centerPad(kernel, L)
    F_image = scf.fftshift(np.fft.fft2(image))
    F_kernel = scf.fftshift(np.fft.fft2(kernel_p))
    F_reconstr = F_image / F_kernel
    reconstr = scf.ifftshift(np.fft.ifft2(F_reconstr))
    return reconstr

# image
point = pointMtrx(100, 100, 50, 50, 10)
point += pointMtrx(100, 100, 60, 70, 15)
noiseK = motionBlurKernel(45, 40)
image = scs.convolve2d(point, noiseK, 'same')

# reconstruction
pointReconstr = filterMotionBlur(image, 45, 40)

# plotting
fig, axes = plt.subplots(1, 3, figsize=(20, 5))
axes[0].imshow(point)
axes[0].set_title('original')
axes[1].imshow(image)
axes[1].set_title('blurred')
axes[2].imshow(np.real(pointReconstr))
axes[2].set_title('reconstructed')
```

Figure 1.3: Reconstructing an original image



1.7 Probability

1.7.1 Basics

1.7.1.1 Probability space

Probability works on some basic entities:

- Ω is a nonempty set called the samplespace.
- $\omega \in \Omega$ is called an outcome
- $E \subseteq \Omega$ is called an event

Definition 23 (Probability). *Probability is a measure on Ω . It is a total function $\Pr : \Omega \rightarrow \mathbb{R}$ such that:*

- $\forall \omega \in \Omega : \Pr[\omega] \geq 0$
- $\sum_{\omega \in \Omega} \Pr[\omega] = 1$

A probability measure together with a samplespace is called a probability space.

We define the probability of an event as:

$$\Pr[E] = \sum_{\omega \in E} \Pr[\omega]$$

Definition 24 (Random variable). *A random variable is a function mapping a ω from Ω to the reals.*

$$X(\omega) : \Omega \rightarrow \mathbb{R}$$

Note that a random variable strictly takes a single ω as argument, not a set of outcomes.

We then calculate the probability that a random variable X has a certain value x as such:

$$\Pr[X = x] = \sum_{X^{-1}(x)} \Pr[\omega]$$

Definition 25 (Expectation). *The expectation of a random variable is defined as*

$$E[X] = \sum_{\Omega} X(\omega) P(\omega)$$

Definition 26 (Conditional Probability).

$$\Pr[A|B] = \frac{\Pr[A \cap B]}{\Pr[B]}$$

As a nice little exercise, we prove the formula for the conditional probability of the complement of B .

Proof.

$$\Pr[A|\overline{B}] =$$

□

As yet another exercise, here is the formula of the probability of a union of arbitrary events:

Proof.

$$\Pr(\cup A_i) = \sum_i \Pr(A_i) - \sum_i \sum_{j>i} \Pr(A_i \cap A_j) + \sum_i \sum_{j>i} \sum_{k>j} \Pr(A_i \cap A_j \cap A_k) - \dots$$

This is proven by induction.

Base case: Proof that $\Pr(A_1 \cup A_2) = \Pr(A_1) + \Pr(A_2) - \Pr(A_1 \cap A_2)$

| This is trivially true when looking at a Venn Diagramm.

Thus $\Pr(A_1 \cup A_2) = \Pr(A_1) + \Pr(A_2) - \Pr(A_1 \cap A_2)$

Induction step. Suppose that... Proof that

|
Thus

□

1.7.1.2 Uncountable sets and infinitessimals

Conventional probability restricts Ω 's to only include sigma-algebras, that is, sets of which every subset is guaranteed to be measureable. This is necessary because it turns out that some uncountable sets do indeed have subsets that you cannot assign a probability to. However, we can sidestep this difficulty by restricting ourselves only to countable sets. If we need continuity, we can include infinitessimals in that set.

1.7.1.3 Probabilistic fallacies

- T-Test interpretation: If $\Pr[A|B] = x$, then this does *not* mean that $\Pr[A|\bar{B}] = 1 - x$.
- Prosecutors fallacy aka. inverse fallacy: $P(A|B) \neq P(B|A)$

1.7.2 Probability distributions

A probability distribution is a function from the domain of a random variable to its probability - in other words, a probability distribution yields the probability that a random variable will take on a certain value.

There is an abundance of ready made probability distributions to chose from, covering virtually all important situations. But care must be taken when deciding which distribution to apply to a certain problem.

The Bernoulli family based on modelling a series of coin-tosses.

- Bernoulli: heads or tails?
- Binomial: k heads in n trials
- Poisson: k heads in ∞ trials.

A remarkable feature of the Poisson-distribution is that it has only a parameter for the mean, but always the same variance.

The geometric family based on repeating an experiment until it succeeds.

1.8 Statistics

1.8.1 Correlation

Assume an inner product space.

1.8.2 Linear regression

Assume that reality can be modelled by a model like this one:

$$\vec{y} = \mathbf{X}\vec{w} + \vec{e}$$

Not knowing \vec{e} , our best guess at the outcome would be

$$\hat{\vec{y}} = \mathbf{X}\vec{w}$$

We can find the gradient of w by:

$$\frac{dE}{d\vec{w}} = -(\vec{y} - \hat{\vec{y}})\mathbf{X}$$

1.8.3 Principal component analysis

Consider the case where you have a three-dimensional euclidean coordinate system, but all your observation-points form a plane inside this space. Try to find a new orthogonal base that minimizes the shadow of the datapoints. It turns out that the best base you can find for this purpose is the set of eigenvectors of the datamatrix.

1.8.4 Tests

Often, statistical tests work *the wrong way round*. You have your hypothesis, like "these groups are different". Then you state the opposite, the *null-hypothesis*. Then you determine how likely your data is under the null-hypothesis - the *p-value*. And than you hope that the p-value is very low.

Students T-Test compares two averages (means) and tells you if they are different from each other. The null hypothesis is that all samples come from the same distribution. That means that under H_0 , the two means you obtained should be similar. So the interpretation goes like this:

- low p-value: the chance that your result would have happened under H_0 is low. It is very unlikely that the two means you have obtained actually come from the same distribution.
- high p-value: It is quite possible that both your groups come from the same distribution.

Note that there are two experimental setups that can use t-tests:

- unpaired: you have two groups, with no person in both groups. Example: Test one group with medication and another with placebo.
- paired: every person is first in the first and then in the second group. Example: test patients before and after treatment.

1.9 Graph Theory

1.9.1 General properties

1.9.1.1 Degrees

Theorem 23. In a directed graph $\sum_{V_G} \text{indeg}(v_n) = |E_G|$

Proof. By induction. Let $P(i)$ be $\sum_{V_G} \text{indeg}(v_n) = |E_G|$

Proof that $P(1)$ and $P(i) \rightarrow P(i+1)$ hold true.

Base case: $\sum_{V_G} \text{indeg}(v_n) = |E_G|$ where $E_G = \{(v_1, v_2)\}$
Proof that $P(1)$ holds true.

| This is trivially true.

Thus $P(1)$ holds true.

Induction step:

- Let $\sum_{V_G} \text{indeg}(v_n) = |E_G|$
- Let $V_{G'} = V_G \cup \{v_{N+1}\}$
- Let $E_{G'} = E_G \cup \{(x, y) \in V_G \times \{v_{N+1}\} \cup \{v_{N+1}\} \times V_G\}$

Proof that $P(i) \rightarrow P(i+1)$ holds true.

For the left side it holds that:

$$\sum_{V_{G'}} \text{indeg}(v_n) = \sum_{V_G} \text{indeg}(v_n) + \text{outdeg}(v_{N+1}) + \text{indeg}(v_{N+1})$$

On the other hand on the right side it holds that:

$$|E_{G'}| = |E_G| + \text{outdeg}(v_{N+1}) + \text{indeg}(v_{N+1})$$

| So the equation reduces to $0 = 0$, which is true.

Thus $P(i) \rightarrow P(i+1)$ holds true.

Thus $P(1)$ and $P(i) \rightarrow P(i+1)$ hold true. □

Lemma 1. $\sum_{V_G} \text{indeg}(v_n) = \sum_{V_G} \text{outdeg}(v_n)$

Lemma 2. In an undirected graph we have $\sum_{V_G} \deg(v_n) = 2|E_G|$

1.9.2 Walks and paths

Definition 27. A walk is any sequence of vertices that are connected by an edge.

Definition 28. A path is a walk where no vertex appears twice.

Theorem 24. The shortest walk between any two vertices in a graph is a path.

Proof. By contradiction. Suppose that the vertex x appears twice in the shortest walk. Proof that this is a contradiction.

| Since x appears twice, the proposed walk must be of the form $f - x - g - x - h$, where g is a sequence of 0 or more connected vertices. Then this walk can be shortened to $f - x - h$.

Thus this is a contradiction. □

Theorem 25. The longest path in a graph has length $|V_G| - 1$.

This is too trivial for a full-fledged proof: there are no repetitions in a path, so it can only have as many steps as it has vertices.

1.9.2.1 Relations and adjacency matrices

Let $R \circ Q$ be a composition of two relations. The number of paths of length exactly n between two vertices x and y , written as $\text{paths}_n(x, z)$, can be expressed as

$$\text{paths}_n(x, z) = |\{y | xRy \wedge yQz\}| = \sum_{y \in Y} xRy \cdot yQz$$

If R_{AM} and Q_{AM} are the adjacency matrix representations of the above relations, then:

$$\begin{aligned} R_{AM} \cdot Q_{AM}[n, m] &= \sum_{y \in Y} R_{AM}[x_n, y] \cdot Q_{AM}[y, z_m] \\ &= \sum_{y \in Y} xRy \cdot yQz = \text{paths}_n(x, z) \end{aligned}$$

When $R = Q$, it follows easily that:

Theorem 26. $R_{AM}^2[n, m]$ is the number of paths that go from v_n to v_m in exactly 2 steps.

Theorem 27. R_{AM} explains whether or not two vertices v and u are connected. All possible paths are listed in $R^* = R_{AM} \cup R_{AM}^2 \cup R_{AM}^3 \cup \dots \cup R_{AM}^{|V_G|-1}$. We can find the length of the shortest path between two vertices u and v by

1.9.3 Planar graphs

The following theorems all deal with planar, connected graphs, and build up to Eulers theorem. Well use N for the number of nodes, E for the number of edges, and L for the number of loops.

Theorem 28. Adding a edge to a planar, connected graph means adding one loop: $\Delta N = 0 \wedge \Delta E = 1 \rightarrow \Delta L = 1$

Theorem 29. Adding x edges means adding x loops: $\Delta N = 0 \wedge \Delta E = x \rightarrow \Delta L = x$

Proof. Proof that $\Delta N = 0 \wedge \Delta E = x \rightarrow \Delta L = x$

By induction on ΔE .

Base case: $\Delta E = 1$ Proof that $\Delta L = 1$

| By theorem 11.

Thus $\Delta L = 1$

Let $\Delta N = 0 \wedge \Delta E = x \rightarrow \Delta L = x$.

Proof that $\Delta N = 0 \wedge \Delta E = x + 1 \rightarrow \Delta L = x + 1$

| Consider the graph from the induction hypothesis, then apply theorem 12.

Thus $\Delta N = 0 \wedge \Delta E = x + 1 \rightarrow \Delta L = x + 1$

Thus $\Delta N = 0 \wedge \Delta E = x \rightarrow \Delta L = x$

□

Theorem 30. When we add a new node to a graph and connect it to the graph using one or more edges, then the number of edges and loops behaves as such: $\Delta N = 1 \rightarrow \Delta L = \Delta E - 1$

Proof. Proof that $\Delta N = 1 \rightarrow \Delta L = \Delta E - 1$

Let $\Delta N = 1$.

By induction on ΔE .

Base case: $\Delta E = 1$ Proof that $\Delta L = 0$

| Graphically trivial.

Thus $\Delta L = 0$

Induction step: $\Delta L = \Delta E - 1$ Proof that $\Delta L' = \Delta E' - 1$

| $\Delta E' = \Delta E + 1$

Using theorem 12, that means: $\Delta L' = \Delta L + 1$

| Thus: $\Delta L + 1 = \Delta E + 1 - 1$

| Which reduces to: $\Delta L = \Delta E - 1$, which is true by the induction hypothesis.

Thus $\Delta L' = \Delta E' - 1$

Thus $\Delta N = 1 \rightarrow \Delta L = \Delta E - 1$

□

Theorem 31. *Eulers theorem:* $N + L = E + 1$

Proof. Proof that $N + L = E + 1$

By induction on N .

Base case: $N = 1$ or $N = 2$ Proof that $N + L = E + 1$

If $N = 1$: $1 + 0 = 0 + 1$
If $N = 2$: $2 + 0 = 1 + 1$

Thus $N + L = E + 1$

Induction step: Let $N + L = E + 1$.

Proof that $N + 1 + L' = E' + 1$

Using theorem 13: $\Delta L' = \Delta E - 1$
Thus: $N + 1 + L + \Delta L = E + \Delta E + 1$
 $N + L = E + 1$, which is true by the induction hypothesis.

Thus $N + 1 + L' = E' + 1$

Thus $N + L = E + 1$

□

1.9.3.1 Trees

Here is a fun little proof.

Proof. In a tree, the mean number of children (mcc) is always equal to $1 - 1/N$.

Proof that $mcc = 1 - \frac{1}{N}$

With what we have so far, this is almost trivial to prove. The mean child count can be written as

$$mcc = \frac{1}{N} \sum_N \text{outdeg } v_n$$

Using theorem 23 and lemma 1, this becomes

$$mcc = \frac{1}{N} E$$

Eulers formula in a tree becomes $N = E + 1$, since there are no loops in trees. Using this:

$$mcc = 1 - \frac{1}{N}$$

Thus $mcc = 1 - \frac{1}{N}$

□

When training a binary classifier, we usually need a measure of defining which of several trees is a good representation of the data. The number of nodes, however, cannot be used as a measure, because it is always going to be $2n - 1$.

Proof. In a binary tree, the number of leaves n equals $\frac{N+1}{2}$, regardless of the architecture of the tree.

Proof that $n = \frac{N+1}{2}$

We shall use the fact that in any tree, $\sum_{V_G} \text{indeg}(v_n) = \sum_{V_G} \text{outdeg}(v_n)$.

Consider the leaves of the tree. They all have $\text{indeg}(v_n) = 1$ and $\text{outdeg}(v_n) = 0$. Thus $\sum_{\text{leaves}} \text{indeg}(v_n) = n$

and $\sum_{\text{leaves}} \text{outdeg}(v_n) = 0$.

Now consider the nodes (there are $N - n$ of them). They all have $\text{outdeg}(v_n) = 2$ and $\text{indeg}(v_n) = 1$, except for the root-node. Thus $\sum_{\text{nodes}} \text{indeg}(v_n) = N - n - 1$ and $\sum_{\text{nodes}} \text{outdeg}(v_n) = 2(N - n)$.

Equating the in- and out-degrees, we obtain:

$$n + N - n - 1 = 0 + 2(N - n)$$

$$n = \frac{N + 1}{2}$$

Thus $n = \frac{N+1}{2}$

□

1.9.3.2 Triangulation

Proof. A n -gon, a polygon with n nodes, can be drawn with $n - 2$ triangles. The proof follows from induction on n , starting at $n = 3$. □

Proof. When drawing a n -gon ($n \geq 3$) using triangles, there are $n - 3$ shared edges.

- The n -gon has n nodes, n edges and one loop
- The $n - 2$ triangles have in total n nodes, $n - 2$ loops and (via Eulers formula) $2n - 3$ edges.
- The outer edges, that is, all the edges of the n -gon, are not shared. Thus, there are $e_{\text{triangles}} - e_{n-\text{gon}} = n - 3$ shared edges.

□

1.9.4 Stable marriage and Gayle-Shapely

In dating, the people doing the active flirting end up with partners higher up their preference-list than the people being flirted with.

```

class Person:
    def __init__(self, name):
        self.name = name
        self.prefs = []

    def setPreferences(self, prefs):
        self.prefs = prefs

    def __repr__(self):
        return self.name

class Woman(Person):
    def selectBest(self, suitors):
        for p in self.prefs:
            if p in suitors:
                print(f"{self} selected {p}")
                return p

class Man(Person):
    def getFavorite(self):
        print(f"{self} fancies {self.prefs[0]}")
        return self.prefs[0]

    def rejectedBy(self, w):
        print(f"{w} rejected {self}")
        self.prefs.remove(w)

rachel = Woman('Rachel')
phoebe = Woman('Phoebe')
monica = Woman('Monica')

ross = Man('Ross')
chandler = Man('Chandler')
joey = Man('Joey')

```

```

women = [rachel, phoebe, monica]
men = [ross, chandler, joey]

rachel.setPreferences([joey, ross, chandler])
phoebe.setPreferences([ross, chandler, joey])
monica.setPreferences([joey, chandler, ross])

ross.setPreferences([rachel, phoebe, monica])
chandler.setPreferences([rachel, monica, phoebe])
joey.setPreferences([phoebe, rachel, monica])

## Gayle - Shapely

flirting = {w: [] for w in women}

def isStable(flirting):
    for w in flirting:
        suitors = flirting[w]
        if len(suitors) != 1:
            return False
    return True

while not isStable(flirting):
    print("--- lets go, flirt! ---")

    # men move to their most desired woman
    for m in men:
        w = m.getFavorite()
        flirting[w].append(m)

    # women reject all but best suitor
    for w in women:
        suitors = flirting[w]
        fav = w.selectBest(suitors)
        for m in suitors:
            if m is not fav:
                m.rejectedBy(w)
        if fav:
            flirting[w] = [fav]

print(" ---- couples have formed: ---- ")
for w in flirting:
    m = flirting[w][0]
    print(f"{w} + {m}")

```

1.10 Calculus

1.10.1 Hyperreals

For the biggest part, we're going to deal with Nelson-style nonstandard-analysis.

List of external properties:

- std, nstd
- limt, nlimt
- inftsm, inft
- nearly cont

A statement using any external properties will be denoted as A_{ext} , one that *might* use external properties as $A_{(ext)}$.

We will use the following axioms:

1. $0 : std$
2. $\forall n \in \mathbb{N} : n : std \rightarrow (n + 1) : std$
3. $\exists n \in \mathbb{N} : n : nstd$
4. External induction: Induction over n_{std} about $A_{(ext)}$:
 $[A_{(ext)}(0) \wedge \forall n_{std} \in \mathbb{N} : A_{(ext)}(n) \rightarrow A_{(ext)}(n + 1)] \rightarrow \forall n_{std} \in \mathbb{N} : A_{(ext)}(n)$
5. Internal induction: Induction over $n_{(nstd)}$ about A:
 $[A(0) \wedge \forall n_{(nstd)} \in \mathbb{N} : A(n) \rightarrow A(n + 1)] \rightarrow \forall n \in \mathbb{N} : A(n)$

The rationale over the two induction-axioms is simple. Ordinary induction is about A over n_{std} . External induction is about $A_{(ext)}$ over n_{std} . This makes sure that statements about external stuff only apply to finite n , not to infinite ones. Internal induction is about A over $n_{(nstd)}$. This makes sure that when we talk about potentially infinite n 's, we only apply internal statements.

In other words: these two inductions ensure that we **never apply external statements to external numbers**.

Doing so would lead to logical inconsistencies. That's why there is no "fully external" induction.

However, note that axiom 2 is actually a case of "fully external" induction.

Theorem 32. $\forall n \in \mathbb{N} : n : nst \rightarrow (n + 1) : nst$

Proof. Suppose $n : nst$. Proof that $(n + 1) : nst$

By contradiction. Suppose $(n + 1) : std$. Proof that this leads to a contradiction.

$(n + 1) : std \rightarrow n : std$.

This contradicts the premise that $n : nst$.

Thus this leads to a contradiction.

Thus $(n + 1) : nst$

□

If you don't believe the argument in the previous proof, consider this:

Proof. Suppose all of the following:

$$[\forall n : Q(n) \rightarrow Q(n + 1)] \rightarrow \forall n : Q(n)$$

$$[\forall n : Q(n) \rightarrow Q(n + 1)]$$

This leads to $\forall n : Q(n)$.

Proof that $\forall n : Q(n + 1) \rightarrow Q(n)$

Let $n = n_0$ and suppose $Q(n_0 + 1)$. Proof that $Q(n_0)$

Since $\forall n : Q(n)$ holds, it must be true that $Q(n_0)$.

Thus $Q(n_0)$

Thus $\forall n : Q(n + 1) \rightarrow Q(n)$

□

We can use theorem 32 to prove the following:

Theorem 33. $\forall n, m \in \mathbb{N} : n : std \wedge m : nstd \rightarrow (n + m) : nst$

Proof. Let $m = m_0 : nst$. Proof that $\forall n \in \mathbb{N} : n : std \rightarrow (n + m_0) : nst$

By induction on n .

Base case. Let $n = 0$. Proof that $(0 + m_0) : nst$

$0 + m_0 = m_0$
 $m_0 : nst$

Thus $(0 + m_0) : nst$

Induction step. Proof that $[(n + m_0) : nst] \rightarrow [(n + 1 + m_0) : nst]$

Just apply theorem 32 to $n = (n + m_0)$.

Thus $[(n + m_0) : nst] \rightarrow [(n + 1 + m_0) : nst]$

Thus $\forall n \in \mathbb{N} : n : std \rightarrow (n + m_0) : nst$

□

It is notable that you can never reach a standard number when adding nonstandard numbers.

Theorem 34. $\forall n, m \in \mathbb{N} : n, m : nst \rightarrow (n + m) : nst$

Proof. We proceed by proving the equivalent $(n + m) : std \rightarrow (n : std \vee m : std)$ Suppose $(n + m) : std$ Proof that $(n : std \vee m : std)$

Without loss of generality, suppose $n : nst$ Proof that $m : std$

By contradiction. Suppose $m : nst$ Proof that this leads to a contradiction

We have already assumed that $(n + m) : std$.

Now, however, we also assume that $n, m : nst$.

Using theorem 32 however, we see that when $n, m : nst$, then it must be that $(n + m) : nst$.

Thus this leads to a contradiction

Thus $m : std$

Thus $(n : std \vee m : std)$

□

1.10.2 Limits

1.10.3 Sequences and series

1.10.3.1 Tailor

1.10.3.2 Fourier

1.10.3.3 Laplace

1.10.4 Eulers formula

Proof: Consider the function $f(t) = e^{-it}(cost + isint)$ for $t \in \mathbb{R}$. By the quotient rule

$$f'(t) = e^{-it}(i \cos(t) - \sin(t)) - ie^{-it}(\cos(t) + i \sin(t)) = 0$$

identically for all $t \in \mathbb{R}$. Hence, f is constant everywhere. Since $f(0) = 1$, it follows that $f(t) = 1$ identically. Therefore, $e^{it} = \cos t + i \sin t$ for all $t \in \mathbb{R}$, as claimed.

1.10.5 Integration

1.10.5.1 Integration strategies

u-substitution

Integration by parts is the last trick up our sleeve when all other strategies haven't helped. Consider the integral

$$\int xe^x dx$$

We can rewrite this integral as

$$\int u dv$$

where $u = x$ and $dv = e^x dx$. In general, you always want to pick u in such a way that u gets simpler after being differentiated. x does get a lot simpler after differentiation, whereas e^x doesn't, so the choice is clear.

We then use the following:

$$\int u dv = uv - \int v du$$

³

Since we chose $u = x$ we have $du = dx$, and from $dv = e^x dx$ we get $v = e^x$. This yields us:

$$\begin{aligned} xe^x - \int e^x dx \\ e^x(x - 1) \end{aligned}$$

as the solution.

1.10.6 Vector calculus

Integration over a vector, integration along a vector, integration along a surface.

You can find a nice introduction (mostly in the second part of) this pdf: http://www.maths.gla.ac.uk/~cc/2A/2A_notes/2A_chap4.pdf and here http://geocalc.clas.asu.edu/pdf-preAdobe8/SIMP_CAL.pdf

³The proof goes like this: ...

1.11 Number theory and cryptography

1.11.1 Divisibility

Definition 29. *Divisibility:* $a|b \leftrightarrow \exists k : ak = b$

Theorem 35. $a|b \wedge b|c \rightarrow a|c$

Proof. Proof that $a|b \wedge b|c \rightarrow a|c$

Suppose $\exists x_0 : ax_0 = b \wedge \exists x_1 : bx_1 = c$ Proof that $\exists x_2 : ax_2 = c$

| Try $x_2 = x_0x_1$

| Then $ax_2 = ax_0x_1 = bx_1 = c$

| Thus $\exists x_2 : ax_2 = c$

Thus $a|b \wedge b|c \rightarrow a|c$

□

Theorem 36. If a divides both b and c , then a divides any linear combination of b and c as well. $a|b \wedge a|c \rightarrow \forall v, w : a|(vb + wc)$

Proof. Proof that $a|b \wedge a|c \rightarrow \forall v, w : a|(vb + wc)$

Suppose $a|b \wedge a|c$ Proof that $\forall v, w : a|(vb + wc)$

| Let v_0, w_0 Proof that $\exists x : ax = (v_0b + w_0c)$

| Try $x = v_0x_0 + w_0x_1$

| Thus $\exists x : ax = (v_0b + w_0c)$

| Thus $\forall v, w : a|(vb + wc)$

Thus $a|b \wedge a|c \rightarrow \forall v, w : a|(vb + wc)$

□

It is neat how $a|b$ can be read in two ways: 'a divides b', but also 'b is a multiple of a'. This is very nice for interpreting the water-jug problem. Let $s = \alpha j_1 + \beta j_2$. If there is an a such that $a|j_1 \wedge a|j_2$ then, by the above proof, s is a multiple of a . Of course, every number can be divided at least by 1. So we can more generally write, for any linear combination: $s = \sum \alpha_i j_i = \beta \gcd(j_1, j_2, \dots)$. As a consequence: with two water-jugs we can only measure quantities that are multiples of $\gcd(j_1, j_2)$.

In more common terms, the above condition could be formulated as 'if the base-vectors have a common divisor'. As such, the above proof sheds some light on the relation between divisibility and linear combinations. If base-vectors have a common divisor, then they are linearly dependent (prove this for yourself, it's trivial). The inverse is not always true, however. Consider 2 and 3: they have no common divisor (other than 1), yet they are linearly dependent: $\frac{-3}{2}2 + 1 \cdot 3 = 0$

Theorem 37. $\forall a, b : \exists! q, r : a = qb + r \wedge 0 \leq r \leq q$

Theorem 38. Euclids algorithm: We can simplify the greatest common divisor like this: $\gcd(a, b) = \gcd(b, \text{rem}(a, b))$

Proof. Proof that $\gcd(a, b) = \gcd(b, \text{rem}(a, b))$

Proof that $\text{CD}_{(a,b)} = \text{CD}_{(b,\text{rem}(a,b))}$

| Let $d_0 \in \text{CD}_{(a,b)}$ Proof that Prove that $d_0 \in \text{CD}_{(b,\text{rem}(a,b))}$

| Since ...

| Thus Prove that $d_0 \in \text{CD}_{(b,\text{rem}(a,b))}$

| Let $d_0 \in \text{CD}_{(b,\text{rem}(a,b))}$ Proof that Prove that $d_0 \in \text{CD}_{(a,b)}$

| Since ...

| Thus Prove that $d_0 \in \text{CD}_{(a,b)}$

Thus $\text{CD}_{(a,b)} = \text{CD}_{(b,\text{rem}(a,b))}$

Proof that Since $A = B \rightarrow \max(A) = \max(B)$, it follows that $\gcd(a, b, \dots) = \gcd(b, \text{rem}(a, b, \dots))$

| It is easy to prove that $A \subseteq B \rightarrow \max(A) \leq \max(B)$. The reverse holds too, of course.

Thus Since $A = B \rightarrow \max(A) = \max(B)$, it follows that $\gcd(a, b, \dots) = \gcd(b, \text{rem}(a, b, \dots))$

Thus $\gcd(a, b) = \gcd(b, \text{rem}(a, b, \dots))$

□

Proof. Proof that $\forall a, b \exists \alpha, \beta : \gcd(a, b) = \alpha a + \beta b$

| The proof is not complicated but a little cumbersome. It is given in this stackexchange post

Thus $\forall a, b \exists \alpha, \beta : \gcd(a, b) = \alpha a + \beta b$

□

While we have now proven that the gcd can always be expressed as a linear combination of its arguments, we don't know yet just how to obtain the coefficients α and β . This is where the *extended euclidian algorithm* comes in.

1.11.2 Hashing

1.11.3 Encryption

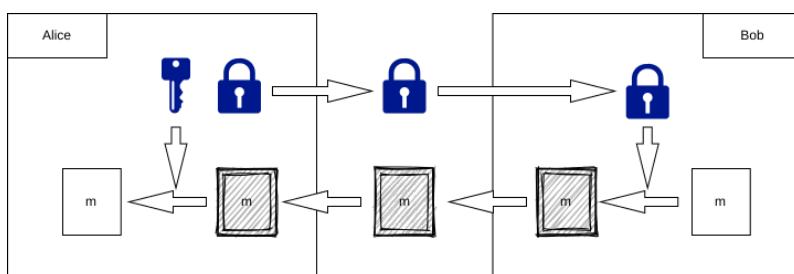
1.11.3.1 Symmetric encryption

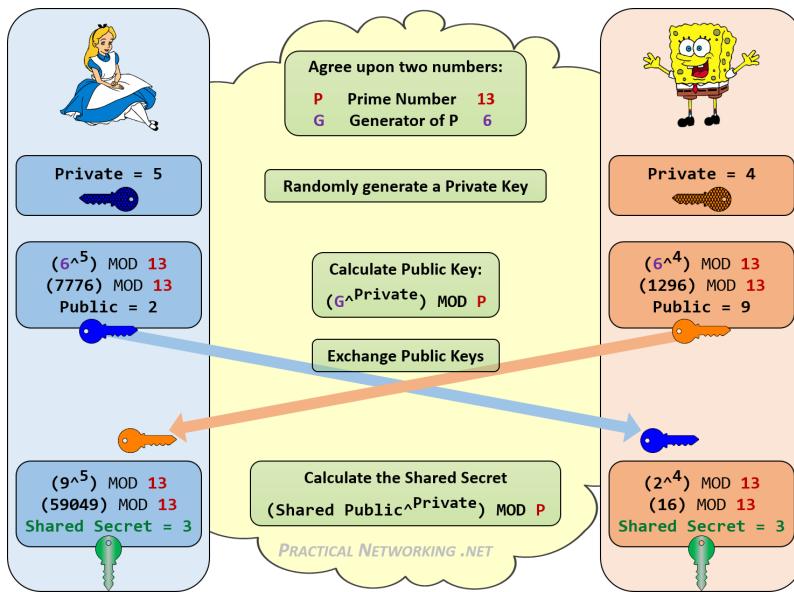
1.11.3.2 Asymmetric encryption

The basic idea of the Diffie-Hellman algorithm is as follows: instead of a secret key used for encryption and decryption, use a two-part-key. We can then find an algorithm where only the public part of the key is ever visible to Eve.

Here is an analogy:

- You have a lock (pubK) and a key (prvK). Send your lock publically to Bob, while you keep your key hidden in your vest.
- Bob writes his message to you and locks it with your lock in a box. Neither he nor Eve can now open the box again. Bob sends the box to you.
- You open the box with your key and read the message





This general concept is known as Diffie-Hellman asymmetric encryption. To implement it, one needs

- A function to generate a key-pair
- A function for encryption using the message and another persons public key
- A function for decryption using a cypher and your own private key

The first successful implementation of this principle is known as the WPA algorithm.

In practice WPA takes too long to be used very often, whereas symmetric encryption is much faster. For this reason, one usually follows a two step process:

- create a new symK
- exchange that symK via WPA
- use the symK to do symmetric encryption

1.12 Computation

1.12.1 Finite state machines

A FSM takes a series of inputs (a *sentence*), evaluates if the sentence is part of its language, and processes it by stepping through a few states to a final state.

The theory of FSM's does not really handle actions that are to be taken upon transitions. In some implementations, upon arriving at the final state, an (external) action is triggered based on that final state. In others, an action may be executed after any transition, not only after arriving at a final state.

It is important to note that the states of a FSM need not be one-dimensional. For example, in the die-hard problem, we model the system as a FDM where the state is a two-tuple; the first being the contents of a 3-liter jug, the second the content of a 5-liter jug. In such a situation it makes sense to draw the nodes of the transition function in a grid with the state-dimensions as the axes.

1.12.1.1 Definition

$$M = \{Q, q_0, F, \Sigma, \delta\} \quad (1.15)$$

where Q is the set of states, $F \subset Q$ is the set of final states, Σ is the alphabet of possible inputs, and δ is:

$$\delta(q, \sigma) = q' \quad (1.16)$$

$$\delta : Q \times \Sigma \rightarrow Q \quad (1.17)$$

$$\Delta(\vec{\sigma}) = \delta(\delta(\delta(q_0, \sigma_0), \sigma_1), \dots) \quad (1.18)$$

It should be noted that in reality state machines may differ from this definition in two ways:

- For one, usually every state may be considered a legitimate final state.
- Also usually a output function is added. There are two main designs for output-functions:
 - Moore machine: an output is generated after every state transition. This is the most common design in hardware applications, and also used in my NID-backend.
 - Mealy machine: every state has its own behavior, which reacts on inputs. One of the possible reactions to an input may be to trigger a state change, but not necessarily every input does cause a state change. This is the most common design in software-applications, especially where state-machines are used to model AI.

In both cases the output function is usually implemented as follows: A behavior is defined for each state. The output function takes the input and the current state as arguments, fetches the correct behaviour based on the current state, and lets the behavior handle the input. In code:

```
class FSM:
    State s

    Output handleInput(Input i):
        State ns = transitionFunction(s, i)
        s = ns # Might be just the same state
        Output o = outputFunction(s, i)

    State transitionFunction(State s, Input i):
        Behavior b = getTransBehavior(s)
        return b.handleInput(i)

    Output outputFunction(State s, Input i):
        Behavior b = getOutputBehavior(s)
        return b.handleInput(i)
```

1.12.1.2 Getting the language of a given machine M

$$L(M) = \{\vec{\sigma} | \Delta(\vec{\sigma}) \in F\} \quad (1.19)$$

Since δ can be expressed as a graph, a version of theorem 27 may be used to get $L(M)$.

1.12.1.3 Getting the simplest machine for a given language**1.12.1.4 Building machines from simpler machines****1.12.1.5 Induction using the invariant principle**

Often we want to prove that a certain desirable property $P(q)$ holds for any step along any sentence that goes into a FSM. This is something we use for Turing machines as well, and in fact in many other applications, like MCMC. We can prove that P holds at every step using the invariant principle.

1.12.2 Turing machines

A Turing machine is only moderately more complex than a FSM. Instead of stepping through each input in the given sentence from left to right, the machine can also choose (based on its state) to move back again, or to erase or change a letter⁴.

However, there seems to be no such thing as a Turing-Machine-Design-Pattern. Instead, the TM is used solely as a model for computation. Contrary to the FSM the TM has no relevant practical implementation.

⁴We use "letter" and "input" as synonyms.

1.13 Dynamic Programming

Dynamic programming tries to solve problems by expressing the problem recursively: The solution for some parameter-set $sol(para)$ equals a function of solutions of other parameter-sets, like for example:

$$sol(para) = sol(para') + x$$

or

$$sol(para) = sol(para') + sol(para'')$$

or most generally:

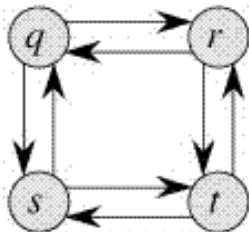
$$sol(para) = f(sol(para'), sol(para''), \dots)$$

For performance, subsolutions are stored in a lookup-table. Without the lookup-table, dynamic programming would be called "divide and conquer".

Dynamic programming is a technique that solves some particular type of problem in polynomial time. Also, dp-solutions can easily proven to be correct. But at first, we need to find out when dp-techniques apply.

1.13.1 Overlapping subproblems (always) and optimal substructure (most)

- Overlapping subproblems are those that make it worth to cache past results.
- Optimal substructure means this: if you have the solution to all parts of the problem, you get the solution to the whole problem. Consider a graph of nodes a, b, c, d . To get from a to d , consider the subproblems of getting from a to b and from b to d . The shortest path from a to d is the shortest path from a to b plus the shortest from b to d . Note that the opposite problem does not have optimal substructure: the longest path from a to d is in fact shorter than the longest paths from a to b plus the longest from b to d .



1.13.2 Worked example: possible ways to sum coins

Imagine there was a currency with coins of value 1, 5 and 7. The task is to list all possible combinations of these basic coins to obtain a total value of n .

Here is an idea how this could be implemented: Lets call our function `combs(n)`. Add 1 to each list in the results of `combs(n-1)`, add 5 to each list in the results of `combs(n-5)`, and add 7 to each list in the results of `combs(n-7)`.

```

def combs(n):
    if n < 1:
        return []
    if n == 1:
        return [[1]]
    if n == 5:
        return [[1,1,1,1,1], [5]]
    if n == 7:
        return [[1,1,1,1,1,1], [5,1,1], [1,5,1], [1,1,5], [7]]

    list1 = combs(n-1)
    list5 = combs(n-5)
    list7 = combs(n-7)

    fullList = []
    for c in list1:
        fullList.append(c + [1])
    for c in list5:
        fullList.append(c + [5])
    for c in list7:
        fullList.append(c + [7])

    return fullList
  
```

Now just add some caching, and you have an efficient algorithm as well.

Verify that this idea shows overlapping subproblems and that it shows optimal substructure

Proof that algorithm delivers correct results

counterexample: minimal number of coins

1.13.3 A strategy for finding the recursion

Dynamic programming tries to solve problems by breaking them apart into subproblems. But how do we find those subproblems?

The solution S we are looking for should be obtained by some function sol .

$$S = sol(para)$$

Try to find a way to split S into a large part S' and a small part s .

$$S = S' + \{s\}$$

The large part itself must be the result of sol for some other set of parameters $para'$.

$$S' = sol(para')$$

1.13.4 Markow chains

1.13.5 Markov decision process with known state-transition-function

MDPs are an extension of Markov chains, they add actions (allowing choice) and rewards (giving motivation). Conversely, if only one action exists for each state (e.g. "wait") and all rewards are the same (e.g. "zero"), a Markov decision process reduces to a Markov chain.

An MDP is a tuple (S, A, P_a, R_a) , where

- S : finite set of states
- A : finite set of actions
- $P_a(s, s') = P(s_{t+1} = s' | s_t = s, a_t = a)$
- $R_a(s, s')$ is the immediate reward received after transitioning from s to s' through action a .

The goal is to find a policy $\pi(s) = a$ that maximizes the cumulative reward $\sum_{t=0}^{\infty} \gamma_t R_{a_t}(s_t, s_{t+1})$, where γ_t is the discount-factor. The discount factor is used so that the decision maker favours taking actions early and doesn't postpone them indefinitely.

Because of the Markov property, the optimal policy for this particular problem can indeed be written as a function of s only, as assumed above.

1.13.6 Reinforcement learning: MDP without explicitly known state-transition-function

In more advanced cases we might instead partition S into sets S' and s .

1.14 Distributed systems

We implement distributed systems when our system becomes too big to be run from a single computer. Distribution always makes things more complex, so there is no reason to implement it when you don't have to. But in almost every case, you eventually will. Most of the following text is based on this lecture series.

1.14.1 Model

- **Communication graph G :** a directed, two-way graph
 - nodes: nodes_G
 - edges: edges_G
 - $\forall(u, v) \in \text{edges}_G : (v, u) \in \text{edges}_G$
- **Subgraph U_G :**
 - nodes: nodes_U
 - edges: $\text{edges}_U = \{(u, v) | u \in \text{nodes}_U \wedge v \in \text{nodes}_U\}$
- **Inedge border of U_G :** $\{(u, v) | u \notin \text{nodes}_U \wedge v \in \text{nodes}_U\}$
- **System \mathcal{G} :**
 - a communication graph G
 - a program A for each node n of G (we say: n runs A) (in the original paper, a program was called a *device*)
 - input to each node n of G
 - behaviour $\mathcal{B}_{\mathcal{G}}$

1.14.2 Axioms

The locality axiom states that a subsystems behaviour is only influenced by its inedge border.

The fault axiom states that there exists a (malicious) program F that can mimic other programs, so that it behaves towards node a like X , and towards b like Y .

1.14.3 Theorems

1.14.3.1 Coverings

Define the *neighborhood* of node u as $\text{nbh}_u = \{v \in \text{nodes}_G | (v, u) \in \text{edges}_G\}$. The neighborhood of u is simply the inedge border of a one-node-graph $U := \{u\}$. A graph S *covers* G if there is a mapping ϕ from any node $s \in S$ to $u \in G$ that preserves neighborhood, that is:

$$\phi : \text{nodes}_S \rightarrow \text{nodes}_G$$

$$\forall s \in \text{nodes}_S : \phi(s) = w \rightarrow \phi(\text{nbh}_s) = \text{nbh}_w$$

Under such a mapping, S looks locally like G . This is trivially proven. By the locality axiom, the behavior of a subsystem is only influenced by its inedge border. That means that the behaviour of the system on u is only influenced by its neighborhood. Since every node $s \in \text{nodes}_S$ has a neighborhood identical to that of some $u \in \text{nodes}_G$, the behaviour of every node in S is identical to that of some node in G .

In practice, we use this finding for the following strategy: Consider a big graph S . In S , we consider a subgraph S' . We might not know a lot about the behaviour of S' , but very likely, G' , which is covered by S' , is much simpler. Using the locality axiom we prove that S' has to behave like G' .

1.14.3.2 Byzantine agreement

Imagine three byzantine generals laying siege to an enemy city. They have to decide if they want to attack the city or not. An attack can only be successful if they all attack together, so they must agree on attack or abstain together. Each general has his own information on whether an attack, even with all three generals involved, will be successful or not. Finally, there is a traitor among them, who will do everything he can to make them decide on the wrong strategy. The two loyal (a.k.a. *correct*) generals follow some algorithm (like e.g. majority-vote, or seniority) for making their decision, while the traitor will do anything he wants. For byzantine agreement to be possible, the following must hold:

1. *Agreement*: The two loyal generals decide upon the same value, that is, whether to attack or not.
2. *Validity*: The decision is made soundly. If both generals think its a good idea to attack, the agreement must be on 'attack'; if both think that it isn't, then the agreement must be on 'abstain'. If one general thinks an attack would be smart while the other doesn't, this rule doesn't hold, and only item 1 holds.

In other words, if byzantine agreement were possible, it would mean that there is an one traitor cannot make the generals choose a bad strategy. However, we will proof that byzantine agreement is *not* possible.

We can easily proof this. We have three nodes u, v, w . In order to achieve validity, a correct node must decide on its own value if another node supports that value. The third node might disagree, but that node could be a traitor. If correct node u has input 0 and correct node v has input 1, the byzantine node w can fool them by telling u that its value is 0 and simultaneously telling v that its value is 1. This leads to u and v deciding on their own values, which results in violating the agreement condition. Even if u talks to v , and they figure out that they have different assumptions about w 's value, u cannot distinguish whether w or v is a traitor.

But there is another, more general proof, which we can use for the following theorems as well. The general idea of the proof is as follows.

1. Assume byzantine agreement (B.A.) is possible.
2. Consider S , a 6-n graph with inputs and programs, which we will construct in a cunning way. S does not need to exhibit B.A. It only needs to be free of contradictions.
3. Consider one flank S_1 of S . This flanks behaviour B_1 must be equal to that of a ceratain G , which exhibits B.A. and is covered by S_1 .
4. Consider another, adjacent flank S_2 of S . This flanks behaviour B_2 must be equal to that of another G , and also be in accordance with B_1 .
5. Consider a third, adjacent flank S_3 of S . This flanks behaviour B_3 must be equal to that of yet another G . However, B_3 will be a contradiction to B_2 .
6. Thus, if B.A. were possible, S could not possibly exist. But it is obvious that S can exist.

1.14.3.3 Week agreement

1.14.3.4 Byzantine fireing squad

1.14.3.5 Approximate agreement

1.14.3.6 Clock syncing

1.14.3.7 When timing does and does not matter

Operations on a single machine have a total order, on distributed machines we can at least hope to acchieve a partial order. Consider a transition function from one state to another. When is a transition function commutative? That would mean that the timing in which input arrives at a state machine does not matter. We know this is not always true. Just consider the state function $f(s, x) = sx + x$.

$$f : S \times X \rightarrow S$$

Theorem:

$$f : \text{some property} \rightarrow f(f(s_0, x_0), x_1) = f(f(s_0, x_1), x_0)$$

1.14.3.8 flp theorem

System model consists of:

In an asynchronous system, a server-failure can alter the outcome This is

When you are at a bivalent state, you can delay a message such that you end up at a bivalent state again This means that you can forever postpone the system from ever reaching a conclusion.

1.14.3.9 CAP-Theorem

1.14.4 Important algorithms

1.14.4.1 one-, two- and tree-phase commit

1.14.4.2 paxos algorithm

1.14.4.3 raft algorithm

1.14.4.4 map reduce algorithm

where in map reduce is the window of failure according to flp?

1.14.5 Databases

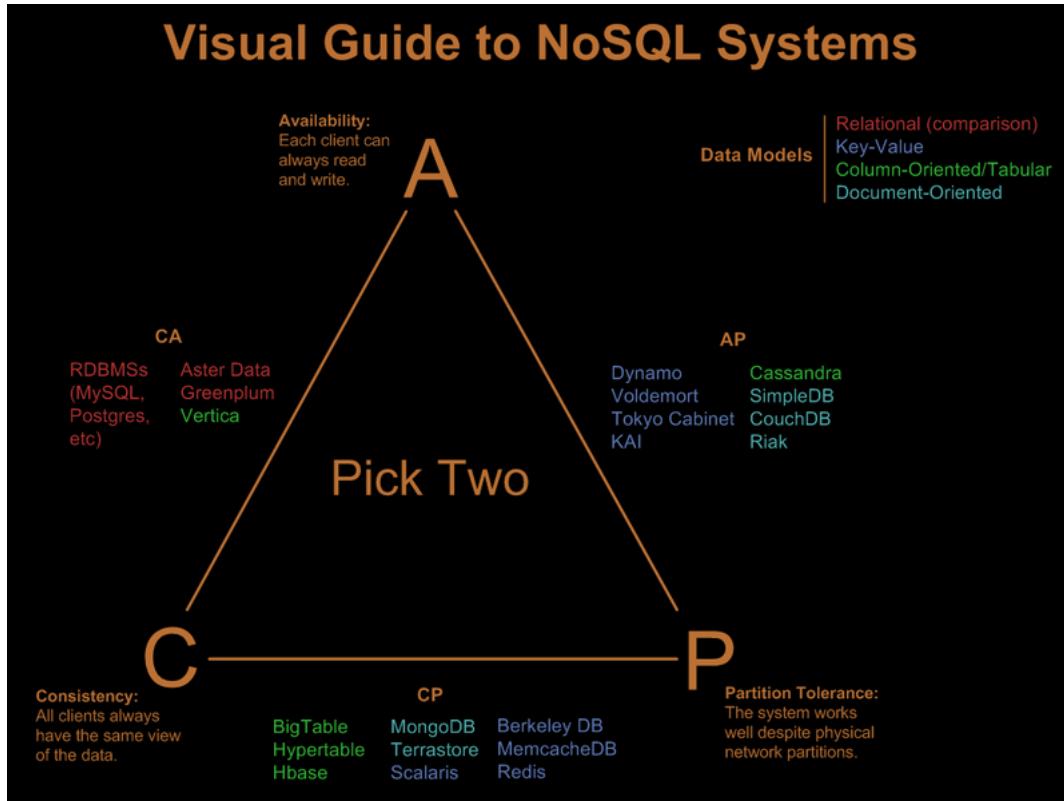
1.14.5.1 ACID

A transaction is a sequence of database operations that can be considered one independent, coherent unit of work. Database transaction should have the ACID-properties:

- Atomicity
- Consistency
- Isolation
- Durability

If a distributed database provides the acid-properties, then it must chose consistency over availability according to the cap theorem. A available database cannot provide acid-transactions.

1.14.5.2 NoSQL Database types



1.14.5.3 Normalisation

Normalisation is a way of structuring a relational database such that logical errors are minimized.

```

from DB import DB

class DictOfSets:
    data = {}

    def addData(self, d):
        for row in d:
            if row[0] in self.data:
                self.data[row[0]].add(row[1])
            else:
                self.data[row[0]] = set([row[1]])

    def addVals(self, v):
        for row in v:
            if row[0] in self.data:
                self.data[row[0]].add(row[1])

    def addKeys(self, k):
        for key in k:
            if not key in self.data:
                self.data[key] = set()

class DBAbstr:
    host = '10.173.202.110'
    user = 'lesen'
    password = ''
    database = 'gkd_chemie'

    def doQuery(self, query):
        result = []
        with DB(self.host, self.user, self.password, self.database) as db:
            result = db.queryToArray(query)
        return result

```

```

class SetTable(DBAbstr):
    vals = []

    def __init__(self, tableName, key):
        self.tableName = tableName
        self.key = key

    def getVals(self):
        if self.vals:
            return self.vals
        query = "select {0} from {1} group by {0}".format(self.key, self.tableName)
        result = self.doQuery(query)
        self.vals = result
        return result

class RelTable(DBAbstr):
    leftVals = DictOfSets()
    rightVals = DictOfSets()

    def __init__(self, tableName, setA, translKeyA, setB, translKeyB):
        self.tableName = tableName
        self.setA = setA
        self.translKeyA = translKeyA
        self.setB = setB
        self.translKeyB = translKeyB

    def getLeftVals(self):
        if self.leftVals.data:
            return self.leftVals.data
        query = "select a." + self.setA.key + ", c." + self.setB.key + " "
        query += "from " + self.setA.tableName + " as a"
        query += "left join " + self.tableName + " as b on b." + self.translKeyA + " = a." + self.setA.key + " "
        query += "left join " + self.setB.tableName + " as c on c." + self.setB.key + " = b." + self.translKeyB + " "
        query += "group by a." + self.setA.key + ", c." + self.setB.key
        result = self.doQuery(query)
        self.leftVals.addData(result)
        return result

    def getRightVals(self):
        if self.rightVals.data:
            return self.rightVals.data
        query = "select a." + self.setB.key + ", c." + self.setA.key + " "
        query += "from " + self.setB.tableName + " as a"
        query += "left join " + self.tableName + " as b on b." + self.translKeyB + " = a." + self.setB.key + " "
        query += "left join " + self.setA.tableName + " as c on c." + self.setA.key + " = b." + self.translKeyA + " "
        query += "group by a." + self.setB.key + ", c." + self.setA.key
        result = self.doQuery(query)
        self.rightVals.addData(result)
        return result

class Relation:
    vals = DictOfSets() # x -> f(x) = y
    invVals = DictOfSets() # y -> f^-1(y) = x

    def __init__(self, setTableX, setTableY, relTable):
        self.setTableX = setTableX
        self.setTableY = setTableY
        self.relTable = relTable

    def checkTransitive(self):
        pass

    def checkBijective(self):
        """ One-to-one and onto """
        return self.checkSurjective and self.checkInjective

    def checkSurjective(self):
        """ Onto: For any y in Y, there is a x in X: y = f(x) """
        invValsData = self.getInvValsData()
        for y in invValsData:
            if len(invValsData[y]) == 0:
                return False
        return True

    def checkInjective(self):
        """ One-to-One: For any a, b in X: f(a) = f(b) => a = b """
        invValsData = self.getInvValsData()
        for y in invValsData:
            if len(invValsData[y]) > 1:

```

	S2	H3
info	Quad-tree with indices (aka. geohash) along Hilbert-curve so that neighbors have similar index	Nested hexagons
computational geometry		All neighbors equidistant
shardable		
visual aspect	Squares distort a lot close to poles	Pretty!

```

        return False
    return True

def compose(self, rel2):
    pass

def getValsData(self):
    if self.vals.data:
        return self.vals.data
    xdata = self.setTableX.getVals()
    self.vals.addKeys(xdata)
    reldata = self.relTable.getLeftVals()
    self.vals.addVals(reldata)
    return self.vals.data

def getInvValsData(self):
    if self.vals.data:
        return self.vals.data
    ydata = self.setTableY.getVals()
    self.invVals.addKeys(ydata)
    reldata = self.relTable.getRightVals()
    self.invVals.addVals(reldata)
    return self.invVals.data


if __name__ == "__main__":
    pnstMappen = SetTable('gkd_chemie.LIMNO_PNST_MAPPE', 'MAPPE_NR')
    messnetze = SetTable('gkd_chemie.LIMNO_SL_PNST_MAPPE_TYP', 'TYP_NR')
    pnstMappePnst = RelTable('gkd_chemie.LIMNO_ZT_PNST_MAPPE_PNST', pnstMappen, 'MAPPE_NR', messnetze, 'MN_NR')
    mappenMessnetzRel = Relation(pnstMappen, messnetze, pnstMappePnst)
    isInj = mappenMessnetzRel.checkInjective()
    isSur = mappenMessnetzRel.checkSurjective()
    print isInj
    print isSur

```

1.14.5.4 Indices

In a database, you want to put indices on those columns that you are likely to read from / write to the most.

1.14.5.5 Spatial indices

For spatial data, the columns that you are likely to access the most are

- lon/lat for point-data
- bounding-boxes for polygons

You can select many types of indices for this kind of data. All of them are good for some operations and bad for others.

Common choices are S2 and H3. Both of these system provide a way to turn a coordinate on earth into a cell id that maps to a hexagonal or rectangular grid cell at a specific resolution. These cell ids have unique properties such as nearby cells have similar ids and you can find parent cells by truncating their length. These properties make operations such as aggregating data, finding nearby objects, measuring distances very fast.

1.14.6 Distributed cache: redis

1.15 Machine learning

1.15.1 Neural networks

1.15.1.1 Backpropagation

The analytical way of deriving the backpropagation algorithm consists of just a few steps. A few definitions:

- A layers outputvector \vec{y}^l is obtained by the activation function $\vec{y}^l = f(\vec{x}^l)$
- The layers inputvector is obtained as a weighted sum of previous outputs: $\vec{x}^l = \mathbf{W}^l \vec{y}^{l-1}$. We can express a single $x_t^l = \sum_f W_{t,f}^l y_f^{l-1}$
- We strive to minimize the error-function. Assuming only one single training item we get $e = \frac{1}{2} \sum_t (\vec{y}_t^* - \vec{y}_t^L)^2 = \frac{1}{2} (\vec{y}^* - \vec{y}^L) \odot (\vec{y}^* - \vec{y}^L)$

Let's first consider only the top layer.

$$\begin{aligned} \frac{de}{dx_{t_0}^L} &= \frac{1}{2} \sum_t \frac{d}{dx_{t_0}^L} (\vec{y}_t^* - \vec{y}_t^L)^2 \\ &= (\vec{y}_{t_0}^* - \vec{y}_{t_0}^L) f'(x_{t_0}) \end{aligned}$$

Or, in vector form:

$$\frac{de}{d\vec{x}^L} = (\vec{y}^* - \vec{y}^L)^T \odot f'(\vec{x}^L)$$

That part was easy. But how do we obtain the same differential for *any* layer l ?

$$\begin{aligned} \frac{de}{dx_{f_0}^l} &= \sum_t \frac{de}{dx_t^{l+1}} \frac{dx_t^{l+1}}{dx_{f_0}^l} \\ &= \sum_t \frac{de}{dx_t^{l+1}} \frac{d}{dx_{f_0}^l} \left(\sum_f W_{t,f}^{l+1} y_f^l \right) \\ &= \sum_t \frac{de}{dx_t^{l+1}} W_{t,f_0}^{l+1} f'(x_{f_0}^l) \end{aligned}$$

Or, in vector form:

$$\frac{de}{d\vec{x}^l} = \left(\frac{de}{d\vec{x}^{l+1}} \mathbf{W}^{l+1} \right) \odot f'(\vec{x}^l)$$

The smart part here was to not derive $\frac{de}{d\vec{x}^l}$ by going through $\vec{y}^L, \vec{x}^L, \mathbf{W}^L, \vec{y}^{L-1}, \vec{x}^{L-1}, \mathbf{W}^{L-1}, \dots$, but by instead creating a recurrence relation by differentiating by \vec{x}^{l+1} .

Finally, we can obtain the gradient at our weights as:

$$\begin{aligned} \frac{de}{dW_{t_0,f_0}^l} &= \frac{de}{dx_{t_0}^l} \frac{dx_{t_0}^l}{dW_{t_0,f_0}^l} \\ &= \frac{de}{dx_{t_0}^l} \frac{d}{dW_{t_0,f_0}^l} \left(\sum_f W_{t_0,f}^l y_f^{l-1} \right) \\ &= \frac{de}{dx_{t_0}^l} y_{f_0}^{l-1} \end{aligned}$$

Or, in vector form:

$$\frac{de}{d\mathbf{W}^l} = \left(\frac{de}{d\vec{x}^l} \right)^T \left(\vec{y}^{l-1} \right)^T$$

So we should change the weights by:

$$\begin{aligned}\Delta \mathbf{W}^l &= -\alpha \frac{de}{d\mathbf{W}^l} \\ &= -\alpha \frac{de}{d\vec{x}^l} \vec{y}^{l-1}\end{aligned}$$

It makes sense to reiterate the whole process in matrix-form.

First, we get δ^L :

$$\frac{de}{d\vec{y}^L} = (\vec{y}^* - \vec{y}^L)^T := \delta^L$$

Then we go through the highest layer:

$$\begin{aligned}\frac{de}{d\vec{x}^L} &= \delta^L \odot f'(\vec{x}^L) \\ \frac{de}{d\mathbf{W}^L} &= \left(\frac{de}{d\vec{x}^L} \right)^T (\vec{y}^{L-1})^T \\ \delta^{L-1} &= \frac{de}{d\vec{x}^L} \mathbf{W}^L\end{aligned}$$

Then we pass δ^{L-1} to the next layer.

$$\begin{aligned}\frac{de}{d\vec{x}^l} &= \delta^l \odot f'(\vec{x}^l) \\ \frac{de}{d\mathbf{W}^l} &= \left(\frac{de}{d\vec{x}^l} \right)^T (\vec{y}^{l-1})^T \\ \delta^{l-1} &= \frac{de}{d\vec{x}^l} \mathbf{W}^l\end{aligned}$$

1.15.1.2 Universal approximation

Let f be a function mapping images to labels. f stems from a vector space of functions \mathcal{F} . Let B be a basis for \mathcal{F} , meaning that

$$\forall f \in \mathcal{F} : \exists \vec{\alpha} : \sum \alpha_n b_n = f$$

So far, so simple. This holds for any basis of any vector space. Let's just propose that sigmoid functions do constitute a basis for these image-to-label functions. We cannot prove this, since we don't know what the image-to-label functions look like, but notice the potential: $\sum \alpha_n b_n$ is then just the output of one layer of a neural net!

It turns out that sigmoid functions do indeed form a basis for any continuous function on $[0, 1]^n$ ⁵.

There is an important point that the universal approximation theorem does not cover, however. The UAT only deals with a single layer net. We know from practice, however, that a multilayer net can approximate functions with far less nodes than what a single-layer net would need. There is some strength to having multiple layers.

1.15.1.3 Some data in n dimensions requires more than n neurons in a layer

The backpropagation algorithm requires a networks layers to transform its input data in a continuous fashion, i.e. distort the input surface without cutting it at any point. In topology, such a transformation is known as a homomorphism.

1.15.1.4 Some functions can be better approximated with a deep net than with a shallow one

Consider the case of a hierarchical function.

$$f(x_1, x_2) = h_2(h_{11}(x_1), h_{12}(x_2))$$

We will prove that a deep net needs less neurons than a shallow one to approximate this function.

⁵Note also that, while sigmoids do form a basis, they do not constitute an *orthogonal* basis, meaning that we cannot obtain weights with the inner-product-trick. We couldn't have obtained them anyway, because for that trick we need the analytical form of f , which is generally not known to us.

1.15.1.5 Convolutional networks

These are the networks most commonly found employed in image-classification. Really, they are just a simplified version of our backpropagation-networks (they are even trained using an only slightly altered algorithm). Instead of connecting every node from layer l to every node of layer $l+1$, they impose some restrictions on the connection-matrix:

- Nodes are connected in a pyramid scheme. A node on layer $l+1$ is connected to 9 nodes directly beneath it. Instead of a $n_l \times n_{l+1}$ connection matrix, we thus have several 9×1 matrices.
- The connection-strengths of these 9×1 matrices are all the same - so really there is only just one 9×1 matrix.

These restrictions are inspired by the physiology of the visual cortex. They have the nice effect that a network is trained much faster, since they massively reduce the amount of weights that need to be learned.

In practice, such networks have a few convolutional layers to reduce the dimension of the input-images, followed by a few conventional, fully connected layers that learn some logic based on the reduced images.

We give the backpropagation-steps for convolutional layers and pooling layers.

In convolutional layers, the forward step goes:

$$\begin{aligned}\vec{x}^l &= \vec{y}^{l-1} \circledast \vec{w}^l \\ \vec{y}^l &= f(\vec{x}^l)\end{aligned}$$

Where the convolution is defined (in our simplified case) as:

$$(\vec{y} \circledast \vec{w})_n = \sum_{m=-1}^1 \vec{y}_{n+m} \vec{w}_m$$

Differentiating a convolution is unfamiliar, but not too hard:

$$\begin{aligned}\frac{d(\vec{y} \circledast \vec{w})}{d\vec{w}} &= \begin{bmatrix} 0 & y_0 & y_1 \\ y_0 & y_1 & y_2 \\ y_1 & y_2 & y_3 \\ \dots & & \\ y_{l-1} & y_l & 0 \end{bmatrix} := tr(\vec{y}) \\ \frac{d(\vec{y} \circledast \vec{w})}{d\vec{y}} &= \begin{bmatrix} w_0 & w_1 & 0 & \dots \\ w_{-1} & w_0 & w_1 & \dots \\ 0 & w_{-1} & w_0 & \dots \\ \dots & & & \\ 0 & \dots & w_{-1} & w_0 \end{bmatrix} := br(\vec{w})\end{aligned}$$

Accordingly, the backwards step goes:

$$\begin{aligned}\frac{de}{d\vec{x}^l} &= \delta^l \odot f'(\vec{x}^l) \\ \frac{de}{d\vec{w}^l} &= \frac{de}{d\vec{x}^l} tr(\vec{y}) = \left(\sum_{n=1}^l e'_{x_n} y_{n+1}, \sum_{n=0}^l e'_{x_n} y_n, \sum_{n=0}^{l-1} e'_{x_n} y_{n+1} \right) \\ \delta^{l-1} &= \frac{de}{d\vec{x}^l} br(\vec{w}) = \frac{de}{d\vec{x}^l} \circledast \vec{w}\end{aligned}$$

In pooling layers, the forward step goes:

$$\begin{aligned}x_t^l &= \frac{1}{4} \sum_f y_{4t+f}^{l-1} \\ y_t^l &= x_t^l\end{aligned}$$

And the backwards step:

$$\begin{aligned}\frac{de}{d\vec{x}^l} &= \frac{de}{d\vec{y}^l} \frac{d\vec{y}^l}{d\vec{x}^l} = \delta^l \\ \frac{de}{d\vec{w}^l} &= 0 \\ \delta^{l-1} &= \frac{1}{4} \frac{de}{d\vec{x}^l}\end{aligned}$$

1.15.1.6 Self organizing maps

Self organizing maps are another, fundamentally different type of neural network. Where feedforward nets employ supervised learning with backpropagation, SOM's do unsupervised learning with a competitive algorithm.

1.15.2 Computer vision

Feature detection and pose estimation Say you want to locate the position of the nose in a portrait.

1.15.3 Feature extraction and dimensionality reduction

The art of preprocessing input has developed in a branch of machine learning itself. Classifiers like SVM's and nnets work better when they get cleaned up and encoded input instead of raw data.

1.15.4 Symbolic AI

Contrary to the before mentioned approaches, symbolic AI uses logical deduction instead of numerical processing to arrive at decisions. If neural nets and decision-trees learning from data can be called building *experience*, then an inference engine deducting from rules can be called building *expertise*. A good tutorial can be found here: codeproject.com. A inference engine can do the following:

- Learning:
 - Gather if-then statements (usually in the form of Horn-clauses).
 - Create a graph of dependencies between statements (a so called and-or-tree).
- Applying the learned things: Given a question:
 - Possible answers: find all potential answers to the problem
 - Backward pass: go through the graph to see what data is required
 - Data aquisition: ask user to provide the data
 - Forward pass: trace the graph forward again to arrive at a single one of the possible answers

Depending on how much effort you put into the expressions that the engine can understand, we differentiate between different levels of logic:

- 0th order logic: understands simple facts and chains them using modus ponens
- 1st order logic: understands variables: all of naive math can be written in 1st order logic.
- higher order logic: is better at inference; can create new if-then-statements as a result of inference or even explore

Popular expert-system-libraries are Prolog, CLIPS and Pyke.

There are many variants to how you can write an expert system. The most important variables are

- How are rules parsed? Do we allow for other logical connectives than 'AND'? See <https://medium.com/a-42-journey/expert-systems-how-to-implement-a-backward-chaining-resolver-in-python-bf7d8924f72f>.
- How is inference done? Forward pass on `addFact` or backward pass on

, or a mixture?

- How is pattern-matching done? The Rete-algorithm is very popular for this.
- Are metaheuristics used?
 - Does the engine try to infer more general rules while idle?
 - Does the engine keep track if a search down one branch takes very long (needs a watching strategy-module)?

```

class InferenceEngine:
    def __init__(self):
        self.facts = []
        self.rules = []

    def addFact(self, *fact):
        if fact not in self.facts:
            self.facts.append(fact) # add the fact
            print(f"Just learned that {fact}")
        for arg in fact[1:]:
            if not self._isVariable(arg):
                self.addFact(arg) # add all arguments

    def addRule(self, conditions, consequence):
        newRule = {'conditions': conditions, 'consequence': consequence}
        if not newRule in self.rules:
            self.rules.append(newRule)
            print(f"Just learned that {newRule}")

    def eval(self, *statement):
        print(f"Evaluating whether '{statement}' holds true")
        fact = self.__searchFacts(*statement)
        if fact:
            return fact
        fact = self.__tryToProve(*statement)
        if fact:
            return fact
        print(f"No, '{statement}' holds not true")
        return False

    def __searchFacts(self, *statement):
        for fact in self.facts:
            if self.__statementsMatch(fact, statement):
                return fact
        return False

    def __tryToProve(self, *statement):
        # a professional inference engine would probably use Rete here
        candidateRules = self.__findCandidateRules(*statement)
        for candidateRule in candidateRules:
            substitutedRule = self.__setRuleVariablesByStatement(candidateRule, statement)
            facts = self.__evalAll(substitutedRule['conditions'])
            if facts:
                substitutedRule = self.__setRuleVariablesByStatements(substitutedRule, facts)
                self.addFact(*substitutedRule['consequence'])
                return substitutedRule['consequence']
        return False

    def __evalAll(self, statements):
        facts = []
        for statement in statements:
            fact = self.eval(*statement)
            if not fact:
                return False
            else:
                facts.append(fact)
        return facts

    """
    ----- helpers -----
    """

    def __setRuleVariablesByStatements(self, rule, statements):
        newRule = dict(rule)
        for statement in statements:
            newRule = self.__setRuleVariablesByStatement(newRule, statement)
        return newRule

    def __setRuleVariablesByStatement(self, rule, statement):
        substitutionDict = {}
        for ruleArg, stateArg in zip(rule['consequence'][1:], statement[1:]):

```

```

        if self.__isVariable(ruleArg):
            substitutionDict[ruleArg] = stateArg
        return self.__setRuleVariablesByDict(rule, substitutionDict)

    def __setRuleVariablesByDict(self, rule, subsDict):
        newRule = {'conditions': [], 'consequence': None}
        for condition in rule['conditions']:
            newCondition = self.__setVariablesByDict(condition, subsDict)
            newRule['conditions'].append(newCondition)
        newConsequence = self.__setVariablesByDict(rule['consequence'], subsDict)
        newRule['consequence'] = newConsequence
        return newRule

    def __setVariablesByDict(self, statement, subsDict):
        newStatement = [statement[0]]
        for arg in statement[1:]:
            if self.__isVariable(arg) and arg in subsDict:
                newStatement.append(subsDict[arg])
            else:
                newStatement.append(arg)
        return newStatement

    def __findCandidateRules(self, *statement):
        candidates = []
        for rule in self.rules:
            if self.__statementsMatch(rule['consequence'], statement):
                candidates.append(rule)
        return candidates

    def __statementsMatch(self, statementOne, statementTwo):
        for wordOne, wordTwo in zip(statementOne, statementTwo):
            oneIsVal = not self.__isVariable(wordOne)
            twoIsVal = not self.__isVariable(wordTwo)
            if oneIsVal and twoIsVal: # ? (oneIsVar and twoIsVar) or (not oneIsVar and not twoIsVar):
                if wordOne != wordTwo:
                    return False
        return True

    def __isVariable(self, arg):
        return isinstance(arg, str) and arg[0].isupper()

if __name__ == '__main__':
    e = InferenceEngine()

    # Test 0: obligatory Socrates test
    e.addRule([
        ['man', 'X'],
        ['mortal', 'X']])
    e.addFact('man', 'socrates')
    print(e.eval('mortal', 'Y'))

```


Chapter 2

Algorithms and data-structures

2.1 General algorithm theory

2.1.1 Invariant principle

The invariant principle is not much more than using induction to prove that some statement holds over all iterations n from 0 to n_0 .

2.1.2 Well ordering

The well ordering principle is a fundamental theorem in discrete mathematics.

$$\forall S \subset \mathbb{N}^+ : S \neq \emptyset : \exists x_0 \in S : \forall x \in S : x_0 \leq x$$

We usually use it in existence proofs, by contradiction. However, there is a nice little template for proofs using the well ordering principle.

To prove that $P(n)$ holds true for $\forall n \in \mathbb{N}$:

- Define the set C of counterexamples to P being true. That is:

$$C = \{n \in \mathbb{N} | \neg P(n)\}$$

- For a proof by contradiction, assume C is nonempty.
- By the well ordering principle, there will be a smallest element $n_0 \in C$
- Reach a contradiction (often by showing how to use n_0 to find another member of $n_1 \in C$ that is even smaller than n_0)

2.1.3 Program verification according to Floyd

If the following properties are given, then an algorithm is be correct.

- Partial correctness: if an answer is returned, it will be correct.
- termination: The algorithm eventually reaches a halt - meaning that an answer is always returned eventually.

There is a simple method of proving an algorithm correct (by Floyd):

- Proving 'partial correctness' with invariant principle: there is some loop invariant that basically says: 'At this iteration the state is a little better than before'. For example, in a sorting problem, at the i 'th iteration the subarray `data[0:i]` might already be sorted.
- Proving 'termination' with well ordering principle

Note that there might be correct algorithms that don't have these properties. For example, some algorithms might eventually reach a desired state without ever having an ever-improving loop invariant. Also, there are other ways to prove an algorithm correct than the Floyd-method. However, in many real-world examples, this is the easiest approach to formally proving an algorithm correct.

2.1.3.1 Proving 'partial correctness' with invariant principle

2.1.3.2 Proving 'termination' with well ordering principle

2.1.4 Notation

An algorithm that takes $f(n)$ cycles is said to take $\Theta(g(n))$ time, iff:

$$f \in \Theta(g(n))$$

where

$$\Theta(g(n)) = \{f(n) | \exists c_1, c_2, n_0 : \forall n \geq n_0 : 0 \leq c_1 g(n) \leq f(n) \leq c_2 g(n)\}$$

Similarly, the set $O(g(n))$ is defined as:

$$O(g(n)) = \{f(n) | \exists c_2, n_0 : \forall n \geq n_0 : 0 \leq f(n) \leq c_2 g(n)\}$$

That is, $O()$ only has an upper bound, but no lower bound.

2.1.5 Order of basic operations

In the integer ram model of computation, addition, multiplication, remainder and bitwise operations all take $O(1)$.

2.1.6 Optimization

- *Linear programming*: Minimizing f (linear) subject to constraints (linear, inequalities)
- *Lagrange multipliers*: Minimizing f (non-linear) subject to constraints (non-linear, equalities)
- *Non-linear programming*: Minimizing f (non-linear) subject to constraints (non-linear, inequalities)
 - f and constraints are differentiable: use KKT (a generalization of Lagrange)
 - f and constraints are convex: use convex optimization
 - All other cases: use some form of gradient descent (simulated annealing, MCMC, ...)

2.1.7 Memoization: Remembering the state in loops

Sometimes, operations have to be repeatedly executed in a loop. These operations may take $O(g)$ time by themselves, but only $O(1)$ time if the result of the previous execution is already known.

In such a case it makes sense to keep track of the previous result. StateMachines are a natural vehicle to store such information.

Lets assume for now that modulo was not constant time. Then we could do the infamous FizzBuzz problem like this in linear time:

```
class SM:
    def __init__(self, nr):
        self.nr = nr
        self.nToGo = nr

    def feed(self):
        self.nToGo -= 1
        if self.nToGo < 0:
            self.nToGo = self.nr - 1

    def isMult(self):
        return self.nToGo == 0


threeM = SM(3)
fiveM = SM(5)
for i in range(1, 16):
    threeM.feed()
    fiveM.feed()
    if threeM.isMult() and fiveM.isMult():
        print("{} --> FizzBuzz".format(i))
    elif threeM.isMult():
        print("{} --> Fizz".format(i))
    elif fiveM.isMult():
        print("{} --> Buzz".format(i))
    else:
        print(i)
```

2.1.8 Using sub problems: Dynamic programming

Solve a multi-step-problem by solving its sub-problems¹. Such problems usually use recursion and memoization.

Example: Towers of Hanoi : You have three rods A , B and C . Move a tower of (say) 4 disks from A to C using the rules:

- move one disk at a time
- no larger disk may be placed on a smaller disk

Solution:

¹Note that the optimal solution to a multi-step-problems is *not always* equal to the sum of the optimal solutions of its sub-problems. Those weird problems are said to *not have optimal substructure*. Example: longest path problem.

```

||| def move([4, 3, 2, 1], from: 'A', to: 'C', via: 'B'):
|||     return
|||         move([3, 2, 1], from: 'A', to: 'B', via: 'C')
|||         + move([4],      from: 'A', to: 'C', via: 'B')
|||         + move([3, 2, 1], from: 'B', to: 'C', via: 'A')

```

2.1.9 Avoiding the need for dynamic programming: Solving recurrences to closed form

Sometimes we don't have to do a recursive solution at all. Many recursive problems can be reshaped into a explicit form (the so called closed form). Think of the Fibonacci-sequence:

$$fib(n) = fib(n - 1) + fib(n - 2) = \frac{(1 + \sqrt{5})^n - (1 - \sqrt{5})^n}{2^n \sqrt{5}}$$

This section will exemplify some methods for finding the closed form expression.

2.1.9.1 Solving linear recurrences

A homogeneous linear recurrence is one of the following form:

$$f(n) = a_1 f(n - 1) + a_2 f(n - 2) + \dots + a_d f(n - d)$$

We can solve it as follows:

1. Assume $f(n) = x^n$. We now try to find an expression for x .
2. Divide both sides in the hlr by x^{n-d} , leaving a (hyper-)quadratic equation.
3. Every root of the quadratic equation is a **homogeneous solution**. Also, if a root r occurs v times, $r^n, nr^{n-1}, n^2r^n, \dots, n^{v-1}r^n$ are also solutions.
4. Also, every linear combination of the above is also a solution. So, a solution might in general have a form like $ar_1^n + br_2^n + \dots$.
5. Finally, choose a, b, \dots such that they fulfill the boundary conditions (in Fibonacci those would be $f(0) = f(1) = 1$). This is the **concrete solution**.

We can extend the above schema to also solve (nonhomogeneous) linear recurrences:

$$f(n) = a_1 f(n - 1) + a_2 f(n - 2) + \dots + a_d f(n - d) + g(n)$$

1. Ignore $g(n)$, find the homogeneous solution from step 4 above.
2. Find a **particular solution** to the equation including $g(n)$.
3. homogeneous solution + particular solution = **general solution**
4. Now for the general solution, again plug in the boundary conditions as in step 5 above.

2.2 Important algorithms

2.2.1 Merge-sort

The function merge takes two already sorted lists and merges them into one sorted list.

```
def merge(a, b):
    c = []
    ia = ib = ic = 0
    while len(c) < len(a) + len(b):
        if a[ia] <= b[ib]:
            c[ic] = a[ia]
            ia++
        else:
            c[ic] = b[ib]
            ib++
        ic++
    return c
```

$Q(i)$: statement

Initial state $Q(0)$: sometext

Transition $Q(i) \rightarrow Q(i+1)$: sometext

Final state $Q(n)$: sometext

2.2.2 Matrix inverse

Gauss Jordan

2.2.3 Polynomials

In this section, we will deal with the polynomials $P_A(x) = 1x^3 + 2x^2 + 3x + 4$. Working with polynomials turns out to be quite important in everyday practice - especially when you're working with generating functions.

2.2.3.1 Coefficient-representation of polynomials

The most common representation of this polynomial is its coefficient representation [4, 3, 2, 1]. This representation is very convenient for fast evaluation at any x using Horner's method:

```
#include "stdio.h"
#include "stdlib.h"
#include "math.h"

/**
 * Coefficient representation of polynomials
 */

typedef struct PolynomialC {
    int length;
    int* coefficients;
} PolynomialC;

PolynomialC* polyc_create(int l, int* coeffs) {
    PolynomialC* p = malloc(sizeof(PolynomialC));
    p->length = l;
    p->coefficients = malloc(l*sizeof(int));
    for(int i = 0; i<l; i++){
        p->coefficients[i] = coeffs[i];
    }
    return p;
}

void polyc_delete(PolynomialC* p) {
    free(p->coefficients);
    free(p);
}

int polyc_evalAt(PolynomialC* p, int x) {
    int sum = 0;
    for(int i = p->length; i > 0; i--){
        sum = sum * x + p->coefficients[i-1];
    }
    return sum;
}
```

```

    }

PolynomialC* polyc_add(PolynomialC* p1, PolynomialC* p2) {
    int l1 = p1->length;
    int l2 = p2->length;
    int l = max(l1, l2);
    int coeffs[l];
    for(int i = 0; i < l; i++) {
        int v = 0;
        if(i < l1) v += p1->coefficients[i];
        if(i < l2) v += p2->coefficients[i];
        coeffs[i] = v;
    }
    PolynomialC* p3 = polyc_create(l, coeffs);
    return p3;
}

int main(void) {

    int coeffs1[4] = {4, 3, 2, 1};
    PolynomialC* p1 = polyc_create(4, coeffs1);

    int coeffs2[3] = {1, 2, 3};
    PolynomialC* p2 = polyc_create(3, coeffs2);

    PolynomialC* p3 = polyc_add(p1, p2);

    int r = polyc_evalAt(p3, 2);
    printf("Result: %d\n", r);

    polyc_delete(p1);
    polyc_delete(p2);
    polyc_delete(p3);

    return 0;
}
}

```

2.2.3.2 Point-value-representation of polynomials

Another way of representing this polynomial is the point-value representation $[(0, 4), (1, 10), (2, 26), (3, 58)]$. This representation is very convenient for fast evaluation of polynomial multiplication (assuming both polys are evaluated at exactly the same points):

```

#include "stdio.h"
#include "stdlib.h"
#include "math.h"

/** 
 * Point/Value representation of polynomials
 */

typedef struct PolynomialV {
    int length;
    int* points;
    int* values;
} PolynomialV;

PolynomialV* polyv_create(int l, int* pts, int* vls) {
    PolynomialV* p = malloc(sizeof(PolynomialV));
    p->length = l;
    p->points = malloc(sizeof(int)*l);
    p->values = malloc(sizeof(int)*l);
    for(int i = 0; i < l; i++) {
        p->points[i] = pts[i];
        p->values[i] = vls[i];
    }
    return p;
}

void polyv_delete(PolynomialV* p) {
    free(p->points);
    free(p->values);
    free(p);
}

PolynomialV* polyv_mult(PolynomialV* p1, PolynomialV* p2) {
    int l = p1->length;
    int vals[l];
    for(int i = 0; i < l; i++) {
        vals[i] = p1->values[i] * p2->values[i];
    }
}

```

```

    PolynomialV* p = polyv_create(l, p1->points, vals);
    return p;
}

int main(void) {
    int points[4] = {0, 1, 2, 3};

    int vals1[4] = {4, 10, 26, 58}; // [4, 3, 2, 1]
    PolynomialV* p1 = polyv_create(4, points, vals1);

    int vals2[4] = {1, 10, 49, 142}; // [1, 2, 3, 4]
    PolynomialV* p2 = polyv_create(3, points, vals2);

    PolynomialV* p3 = polyv_mult(p1, p2);

    polyv_delete(p1);
    polyv_delete(p2);
    polyv_delete(p3);

    return 0;
}

```

2.2.3.3 Naive transformation from coefficient to point-value and vice-versa

```

// O(n^2)
PolynomialV* poly_coeffToPv(PolynomialC* pc, int* points){
    int l = pc->length;
    int vals[l];
    for(int i = 0; i < l; i++) { // n operations
        vals[i] = polyc_evalAt(pc, points[i]); // O(n)
    }
    PolynomialV* pv = polyv_create(l, points, vals);
    return pv;
}

// O(n^3)
PolynomialC* poly_PvToCoeff(PolynomialV* p){
    int l = p->length;
    int points[l] = p->points;
    int values[l] = p->values;
    int** data = [[[1, x1, x1^2, ...],
                   [1, x2, x2^2, ...],
                   [1, x3, x3^2, ...]];
    Matrix* m = mtrix_create(l, l, data);
    Matrix* mi = mtrix_inverse(m); // Gauss Jordan: O(n^3)
    int coeffs[l] = mtrix_mult(mi, values);
    PolynomialC* pc = polyc_create(l, coeffs);
    return pc;
}

int main(void) {
    int coeffs[4] = {4, 3, 2, 1};
    PolynomialC* pc = polyc_create(4, coeffs);
    int points[4] = {0, 1, 2, 3};
    PolynomialV* pv = poly_coeffToPv(pc, points);
    PolynomialC* px = poly_PvToCoeff(pv);

    polyc_delete(pc);
    polyv_delete(pv);
    polyc_delete(px);

    return 0;
}

```

2.2.3.4 Fast transformation

Up to now we have used the points 0, 1, 2, 3 for our evaluation. However, there is no reason why we should have chosen these specific points - any other points of evaluation are just as good. It turns out that if we use the 4 roots of 1 as points of evaluation, we can cut down on computation costs.

In general, the k th root of 1 is $e^{2\pi i \frac{k}{n}}$. So let's do the evaluation at $(1^{1/4})_1, (1^{1/4})_2, (1^{1/4})_3, (1^{1/4})_4 = 1, i, -1, -i$.

2.2.4 Stream algorithms

Stream algorithms handle the case where a function is not applied once for a single input, but continuously on a neverending stream of inputs. This creates some constraints:

- The algorithm must not take longer than the feed-rate of the stream
- The algorithm must use only a constant amount of storage (because otherwise at some point the disc will be full)

Here is a good introduction to the topic. Side note: a popular software to provide a stream to (multiple) subscriber(s) would be Apache kafka.

2.3 Data-structures

2.3.1 Stack

Stacks perform push and pop in O(1). However, they do search in O(n).

```

||#include <stdio.h>
||#include <stdlib.h>
||#include <string.h>

typedef struct El {
    char* key;
    float val;
} El;

El* el_create(char* key, float val){
    char* keyk = strdup(key);
    El* el = malloc(sizeof(El));
    el->key = keyk;
    el->val = val;
    return el;
}

void el_destroy(El* el){
    free(el->key);
    free(el);
}

typedef struct Stack {
    El** data;
    int size;
    int top;
} Stack;

Stack* stack_create(int size){
    Stack* stack = malloc(sizeof(Stack));
    stack->data = malloc(sizeof(El*) * size);
    stack->size = size;
    stack->top = -1;
    return stack;
}

void stack_push_el(El* el, Stack* stack){
    stack->top += 1;
    stack->data[stack->top] = el;
}

void stack_push(char* key, float val, Stack* stack){
    El* el = el_create(key, val);
    stack_push_el(el, stack);
}

El* stack_pop(Stack* stack){
    stack->top -= 1;
    return stack->data[stack->top + 1];
}

El* stack_peek(Stack* stack){
    return stack->data[stack->top];
}

void stack_print(Stack* stack){
    int i;
    for(i = 0; i <= stack->top; i++){
        printf("%s -> %d", stack->data[i]->key, stack->data[i]->val);
    }
}

void stack_destroy(Stack* stack){
    while(stack->top >= 0){
        El* el = stack_pop(stack);
        el_destroy(el);
    }
    free(stack->data);
    free(stack);
}

```

```

int main(){

    Stack* s = stack_create(10);
    stack_push("eins", 1.1, s);
    stack_push("zwei", 1.2, s);
    stack_push("drei", 3.1, s);

    stack_print(s);

    stack_destroy(s);

    return 0;
}

```

2.3.2 Linked lists

Linked lists have the advantage that they have no prefixed size. Addition and insertion at any point is $O(1)$. On the other hand, searching takes $O(n)$ in the worst case.

2.3.3 Binary trees

Trees make sense when the keys have to be ordered in some sense.

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>

typedef struct El {
    int key;
    float val;
} El;

El* el_create(int key, float val){
    El* el = malloc(sizeof(El));
    el->key = key;
    el->val = val;
    return el;
}

void el_destroy(El* el){
    free(el);
}

typedef struct Node {
    struct El* data;
    struct Node* parent;
    struct Node* left;
    struct Node* right;
} Node;

Node* node_create(Node* parent, El* element){
    Node* leaf = malloc(sizeof(Node));
    leaf->parent = parent;
    leaf->data = element;
    return leaf;
}

void node_destroy(Node* node){
    if(node->left != NULL){ node_destroy(node->left); }
    if(node->right != NULL){ node_destroy(node->right); }
    el_destroy(node->data);
    free(node);
}

void node_elementInsert(Node* node, El* el){
    if(el->key < node->data->key){
        if(node->left){
            node_elementInsert(node->left, el);
        }else{
            node->left = node_create(node, el);
        }
    }else{
        if(node->right){

```

```

        node_elementInsert(node->right, el);
    }else{
        node->right = node_create(node, el);
    }
}

void node_dfsPrint(Node* node){
    printf("%d -> %d \n", node->data->key, node->data->val);
    if(node->left){
        node_dfsPrint(node->left);
    }
    if(node->right){
        node_dfsPrint(node->right);
    }
}

typedef struct Tree {
    Node* root;
} Tree;

Tree* tree_create(El* el){
    Tree* tree = malloc(sizeof(Tree));
    Node* root = node_create(NULL, el);
    tree->root = root;
    return tree;
}

void tree_destroy(Tree* tree){
    node_destroy(tree->root);
    free(tree);
}

void tree_elementInsert(Tree* tree, El* el){
    node_elementInsert(tree->root, el);
}

void tree_dfsPrint(Tree* tree){
    node_dfsPrint(tree->root);
}

int main(){
    El* baseEl = el_create(20, 14);
    El* el1 = el_create(1, 4);
    El* el2 = el_create(25, 9);
    El* el3 = el_create(2, 2);

    Tree* tree = tree_create(baseEl);

    tree_elementInsert(tree, el1);
    tree_elementInsert(tree, el2);
    tree_elementInsert(tree, el3);

    tree_dfsPrint(tree);

    tree_destroy(tree);
    return 0;
}
}

```

2.3.4 Hash-tables

Hash tables are basically a array with a hash-function. The hash function is there to associate string-keys to the integer-array-indices. The array needs as many elements as the hash-function can create integer-keys.

Through the hash-function the hash-table has insertion- and lookup-times of expected O(1). However, when the array gets many elements. it might happen that the hash-function gives the same integer-key to several string-keys. To accommodate such a case, the array-elements are implemented as linked lists. This makes the worst-case lookup-time O(n/k).

Why does one not simply use a perfect hash? Make a hash-function that generates a unique key for any string it's given. This would be known as a lookup-table. That works perfectly fine when we expect almost all possible string-keys to be really used. If however we expect only a few of the possible strings to be used, the array would contain a lot of unused space.

Chapter 3

Practical computer knowledge

3.1 Hardware

3.1.1 Memory

Data is moved from ram to the processor by the memory controller.

Each cell in memory holds 8 bits.

3.1.1.1 How data is stored

- `tiny int`: use just one bite - making 256 possible values.
- `int`: uses 4 bites (32 bits)
- `long int`: uses 8 bites (64 bits)
- `signed int`: reverse the leftmost bit
- `fractions`: store two numbers: the numerator and the denominator.
- `float`: also store two numbers: the number without the point and the position of the point

Fun fact: java tends to silently cause integer overflows. When it has to compute an addition of two big ints, say $255 + 1$ (`1111 1111 + 0000 0001`), it does the right computation (`1 0000 0000`) but throws away the leftmost binary because that won't fit into memory, causing the result to be stored as `0000 0000`.

```
||| System.out.println(Integer.MAX_VALUE);
Integer a = 2147483647;
Integer b = 1;
Integer c = a + b;
System.out.println(c);
```

Words and Pages:

3.1.2 Processors

Processors have a cache where they store copies of stuff they recently read from ram. This cache is even faster than reading from ram itself. Because programs tend to put their stuff in sequential order in the ram, the memory controller not only feeds the processors cache the contents of the requested addresses, but also a few of the nearby ones.

32 versus 64 bit.

3.1.3 Cables and Busses

A bus is the conductor that leads data from one hardware-component to another, like from the memory to the processor. Busses come in serial and parallel form.

- SATA: Designed for storage-devices. Serial. Carries no power. 1.5 - 6.0 Gbps (Gigabits per second).
- PCI: (or newer, PCIe): parallelized (that is, multiple data-links per lane), fast, specifically designed for graphics-cards and expansion-cards. Can be used bidirectionally. Also carries power. 7 Gbps per lane (usually 2, 4, or up to 32 lanes present.)
- USB: Used in all kinds of devices. Serial. Unidirectional. Also carries power. 5.0 (USB 3.1 Gen 1) - 20.0 (USB 3.2) Gbps.
- Firewire: Common in Apple and media-devices. Designed for large streams of data. Allows direct device-to-device-connection without a computer in between. Can be used bidirectionally. 3.2 - 6.4 Gbps.

3.1.4 Harddrives

Harddrives use the scasi-System to be made visible to the os.

3.2 Networking

A lot more information can be found here: <http://cnp3book.info.ucl.ac.be/>

3.2.1 Protocols

3.2.1.1 Layer 1: Physical

- Synchronous serial protocols: Master-slave relationship. Can only go one-way
 - I2C/TWI
 - SPI
- Asynchronous serial protocols:
 - TTL Serial
 - RS-232: your monitor cable
- Asynchronous serial bus protocols:
 - usb
 - RS-485

3.2.1.2 Layer 2 : Data link

Ethernet [Frames] Ethernet is barely a protocol. It works like morse: the sender writes the destination-mac on the frame and sends it out on the wire. On the wire, the frame is broadcasted to absolutely everyone, and every computer has to check for itself if the frame was meant for it or for someone else.

Because that is way too much work with several millions of computers in the world, we have switches. They connect two nets of computers, and if a destination-mac is on the other net, they allow the frame to be broadcast to all computers in both nets; otherwise they block the frame, which makes it only broadcast to its source-net. (Hubs are even simpler devices: they always broadcast to all nets, never even reading the destination-mac.)

Because in the 70's hubs and switches didn't have very much memory, the payload of a frame is limited to 1500 bytes. As a consequence, IP's payload is 1480 bytes, and TCP's payload is 140 bytes.

There are few problems that can arise with switches, since they are such simple components. The one you should know about is a broadcast-storm (aka switching loop). This is where multiple switches keep broadcasting each other in search of a particular node. This can really only happen when there is a circle between switches.

When traversing through the net, routers change the destination-mac to that of the next router on the way to the destination-ip.

Spanning tree protocol is how ...

3.2.1.3 Layer 3 : Network

IP IP is being sent in packages. Each contains the source- and destination-IP and the payload. Stateless.

Ip packages are routed by routers. A router looks at the destination-ip and finds the next router that is closer to the ip. Then it changes the destination-mac to that of the next router and sends the package on its way.

IP doesn't have a notion of a connection. You can blindly send thousands of IP-packages into the ether and never know if they arrived.

Discovering a router and obtaining an ip The computer uses DHCP to find a router: it requests an IP address by broadcasting a DHCPDiscover message to the local subnet. The router hosts a DHCP server that then sends the computer an answer containing the computers new ip-adress and the default-gateway that it should use.

```
michael@michael-ThinkPad-Edge-E540:~$ sudo tcpdump -vnes0 -i wlp4s0 port 67 or port 68
tcpdump: listening on wlp4s0, link-type EN10MB (Ethernet), capture size 262144 bytes
13:53:33.023390 0c:8b:fd:8f:8d:65 > ff:ff:ff:ff:ff:ff, ethertype IPv4 (0x0800), length 342: (tos 0x10, ttl 128, id 0, offf
0.0.0.0.68 > 255.255.255.255.67: BOOTP/DHCP, Request from 0c:8b:fd:8f:8d:65, length 300, xid 0x83aa8a15, Flags [none]
  Client-Ethernet-Address 0c:8b:fd:8f:8d:65
  Vendor-rfc1048 Extensions
    Magic Cookie 0x63825363
```

```

DHCP-Message Option 53, length 1: Request
Requested-IP Option 50, length 4: 192.168.1.178
Hostname Option 12, length 26: "michael-ThinkPad-Edge-E540"
Parameter-Request Option 55, length 18:
    Subnet-Mask, BR, Time-Zone, Default-Gateway
    Domain-Name, Domain-Name-Server, Option 119, Hostname
    Netbios-Name-Server, Netbios-Scope, MTU, Classless-Static-Route
    NTP, Classless-Static-Route, Classless-Static-Route-Microsoft, Static-Route
    Option 252, NTP
13:53:33.753685 e8:37:7a:39:ed:2e > 0c:8b:fd:8f:8d:65, ethertype IPv4 (0x0800), length 346: (tos 0x0, ttl 64, id 0, offset
192.168.1.1.67 > 192.168.1.178.68: BOOTP/DHCP, Reply, length 304, xid 0x83aa8a15, Flags [none]
Your-IP 192.168.1.178
Client-Ethernet-Address 0c:8b:fd:8f:8d:65
Vendor-rfc1048 Extensions
    Magic Cookie 0x63825363
    DHCP-Message Option 53, length 1: ACK
    Server-ID Option 54, length 4: 192.168.1.1
    RN Option 58, length 4: 302400
    RB Option 59, length 4: 529200
    Lease-Time Option 51, length 4: 604800
    Subnet-Mask Option 1, length 4: 255.255.255.0
    Default-Gateway Option 3, length 4: 192.168.1.1
    Domain-Name-Server Option 6, length 4: 192.168.1.1
    Domain-Name Option 15, length 9: "rheingold"
    BR Option 28, length 4: 192.168.1.255

```

The 4 packets to a successful DHCP:

- DISCOVER: Client connects to the network and sends out a broadcast discovery looking for its DHCP information.
- OFFER: The server offers the DHCP information to the client
- REQUEST: The client requests verification of the DHCP information
- ACK: The server acknowledges the DHCP request

Sometimes you will not see the DISCOVER / OFFER and just see the REQUEST / ACK. This happens when the client has already obtained a valid DHCP lease earlier and is just requesting to have it again before its lease time expires. Typically this is performed when half the lease has lapsed.

If the REQUEST is not valid anymore the server will send a NACK indicating to the client that it can no longer use this DHCP information. This should cause the client to start over with a DISCOVER.

Sometimes you will see repeated DISCOVER / OFFER but never a REQUEST from the client. This happens when the client either doesn't receive the OFFER or doesn't like it for some reason. Perhaps a firewall is blocking it, they have a poor connection, or simply they're using a Windows computer.

It's common for Windows Vista to never even start its DHCP process. It will just refuse to DISCOVER and complain that the connection is "limited or no connectivity". You can try to diagnose the problem and tell it to reset the network card and/or get new IP information. If this fails to start it then I find adding a static IP and then setting it back to DHCP will get it going. You may even need to restart the DHCPC service. Its Vista.. where you expecting it to work as advertised?

Routers are the hardware components that

3.2.1.4 Layer 4 : Transport

Transport-layer protocols contain source- and destination-port-number. This way, a package that has already arrived at a server (by its ip-address) can be sent to the right application to handle the request.

Common port numbers on the receiving site are:

- 21: ftp
- 22: ssh
- 25: smtp
- 53: dns
- 67/68: dhcp

- 110: pop
- 80: http
- 443: https (http in tls)
- 989/990: sftp (ftp in tls or ssh)

The port number only gets important once a package arrives at the destination-server. There, it must first pass the firewall. Then, the server decides what application should receive packages addressed to the given port.

TCP Continuous connection, stateful. The recipient regularly sends the sender an ACK with the last received package number. If the sender receives ACKs of the last sent package, it gradually increases the amount of packages it sends in a batch before waiting for another ACK. If the sender receives an ACK of the same package number multiple times, it assumes that all the following packages have been lost and retransmits them from there. Also, the batch-size is temporarily reduced.

UDP

3.2.1.5 Layer 5: Application layer

HTTP Http kind of throws away the continuous connection created using tcp by terminating said connection once a request and all its subsequent subrequests (like fetching the image on a site) has been served.

Http requests consist of

- a url,
- a method,
- and potentially parameters

Http-responses consist of

- a status-code,
- a mime-type (= Content-type)
- and the actual content.

It's good to know a few of these status codes:

- 200 : all ok
- 302 : item was already cached
- 404 : resource not found
- 500 : internal server error

Http is really a stupid protocol. But for static websites it's just about enough. Also, webservices send and receive their xml over http.

HTTPS This is http within ssl (aka tls). This is not the same as http within ssh - that would be called a tunnel. ssl (port 443) and ssh (port 21) both use Diffie-Hellmann to create a common key. But the difference is that ssl does not usually require authentication of the client. The server proves its identity by sending the client a certificate, but the client does not need to authenticate to the server like he does in ssh (by password or by public-key authentication).

In https, all the http-content, being url, headers and parameters, is encrypted. However, ip- and tcp-headers are not encrypted - otherwise the packages would not be routable. Also, for making the initial handshake with the server, the servers url has to pass through the net in public. But once https is established, clicks on further hyperlinks on the server will not be visible to the outside; only the associated ip and port.

WebSocket Contrary to http, which is just a request-response-close protocol, in websockets a server has the option to query the client actively at any time.

3.2.1.6 Layer 6 and higher

REST REST sits on top of HTTP. It defines a few methods:

- GET: Get a ressource. Get puts all its parameters right in the url and has no body
- POST: Put enclosed data in database, changing an existing resource. Post-Parameters are put in the request body.
- PUT: Put enclosed data in database, adding a new resource
- HEAD:

SOAP Simple Object Access Protocol

3.2.2 Email

Ok, so we've covered the ethernet/ip/tcp/etc stack. Email is an altogether different beast.

SMTP : per default on port 25 (or 465 for TLS). For sending email.

POP3 : per default on port 110 (or 995 for TLS). For pulling email from server.

IMAP : per default on port 143 (or 993 for TLS). For copying email from server and syncing read-status.

3.2.3 Security

3.2.3.1 Encryption

3.2.4 Making things talk: Remote data transmission

Before we go into details, we need to know some basics of communication.

How radio works

- Radio waves are the longest and contain the least energy of any electromagnetic wave. While lighwaves are the size of bacteria, radiowaves are the size of a car or even a mountain.
- AM-Radio (amplitude-modulation); the original standard.
 - Take a signal of constant amplitude and frequency - known as the carrier wave.
 - Add your music to that signal. Even though the music has changing freqs and amps, the carrier signal transports it at its carrying frequency.
 - Send the signal from a radio-station
 - Receive the signal with a radio, where you pick the channel by choosing the carrier-frequency. The radio then subtracts the carrier-signal, leaving only the original music intact.

This is how a signal can be transmitted on one single frequency.

- FM-Radio does the same, but instead of staying on one freq and listening to amplitude, it listens to a (small) spectrum of freqs and emits sound whenever the frequency changes.
 - better against interference, because intf often causes ampl spikes
 - smaller range: AM can transmit over long distance because it uses a smaller frequency (though not so good quality)
- in modems, isdn and dsl we use am, fm and phase-modulation together to allow for multiple channels.

Types of wires.

- copper: traditional telephone wires. ISDN runs on these; DSL does too, but somewhat slower (this is because, contrary to ISDN, DSL needs an IP-Station as soon as possible after leaving your house. In rural areas DSL becomes much slower this way). Can carry electricity, which is why in the past telephones still worked even when a houses electricity crashed.
- glass-fibre/voip: old copper wires are being replaced with cables that end in a station with IP-adress. These don't carry electricity anymore, however. ISDN does not work on these, either. Those new voip-cables are often glass-fibre cables¹. DSL over glass-fibres is called VDSL (whereas DSL on the old copper wires is simply called DSL), telephony over DSL/glass is called VOIP.

And here a recap on important hardware:

- Router: really consists of three devices:

- router: sends packets by ip either to your devices or the next router on the internet
- switch: broadcasts incoming packets to all your home-devices or moves outgoing packets to the router
- wireless access point

3.2.4.1 ISDN (wired)

Your computer is attached to a modem², which is attached to a telephone-cable or even a telephone-earpiece (!). The modem translates your TCP-requests to sound-signals which would be carried over the telephone-wire in the form of sound. With ISDN, you cannot use both your telephone and your internet at the same time (actually, with am, fm and pm you can; you'd obtain up to 3 channels). Actually, isdn is different from a modem. Modems came first. There you put your telephone directly on the earpiece. Isdn then skipped the earpiece and put the modem directly onto the copper wire - this way a higher sampling frequency was possible. Since everyone in a neighbourhood shares the same cable, your speed will drop when many people use netflix at the same time. On the other hand, ISDN uses a boosting technology that makes sure that your phone-conversation stays clear over large distances. This means that ISDN does not require a relay nearby - contrary to DSL. So in rural areas, ISDN may be actually faster. Speed: around 0.1 Mbit/s.

3.2.5 DSL (wired)

DSL (digital subscriber line) is using ordinary telephone wires or the more modern glass-wires (in that case it is called VDSL). Both a telephone-cable as well as a your routers output will be plugged into a DSL-splitter (shown



below).

With ISDN, you could not use your telephone and your computer at the same time. This is because your internet-traffic is going through the copper-wire on a different frequency-band than the one telephone-conversations use. ISDN was hampered by interference, because multiple wires would lie in the ground close to each other. While ISDN did already use am, fm and pm, DSL has a technique for interference-cancelling that allows it to increase its throughput even more. But this technology requires an ip-station at the

¹Instead of electricity, glass-fibre cables carry light-signals. This technology allows for far more frequency bands to be used.

²modem stands for modulator/demodulator: a device which receives electrical waves comming from the copper-wire and transforms them into analog (sound) or digital (bits) waves; and back.

end of the line. To further increase speed for the average user, there is *ADSL*, with A standing for asymmetric. Here, downstream traffic is much preferred to upstream, since the average user does not host a server. Contrary to ISDN, DSL connections are not shared between all users in a neighbourhood. Instead, a dedicated connection runs directly to a station with its own IP-address near your house. Speed: up to 16 Mbit/s, but only 1-2 Mbit/s in rural areas where IP-Stations are far from your house.

3.2.6 WiFi (really only those few meters up to your wired router)

WiFi really uses the same "everyone-shout-first-gets-served" protocol as MAC does. That means that every device has to wait its turn to use the router. There are two frequency-ranges at which devices may communicate. Within each range there are sub-areas known as *channels*. Parallelism is obtained by different devices using different channels.

- 2.4 GHz: slower, further reach. 14 channels, but only 1, 6 and 11 don't overlap. Interferes with microwaves and bluetooth.
- 5.0 Ghz: faster, but only at shorter distance. Many channels.

Factors to consider with wifi:

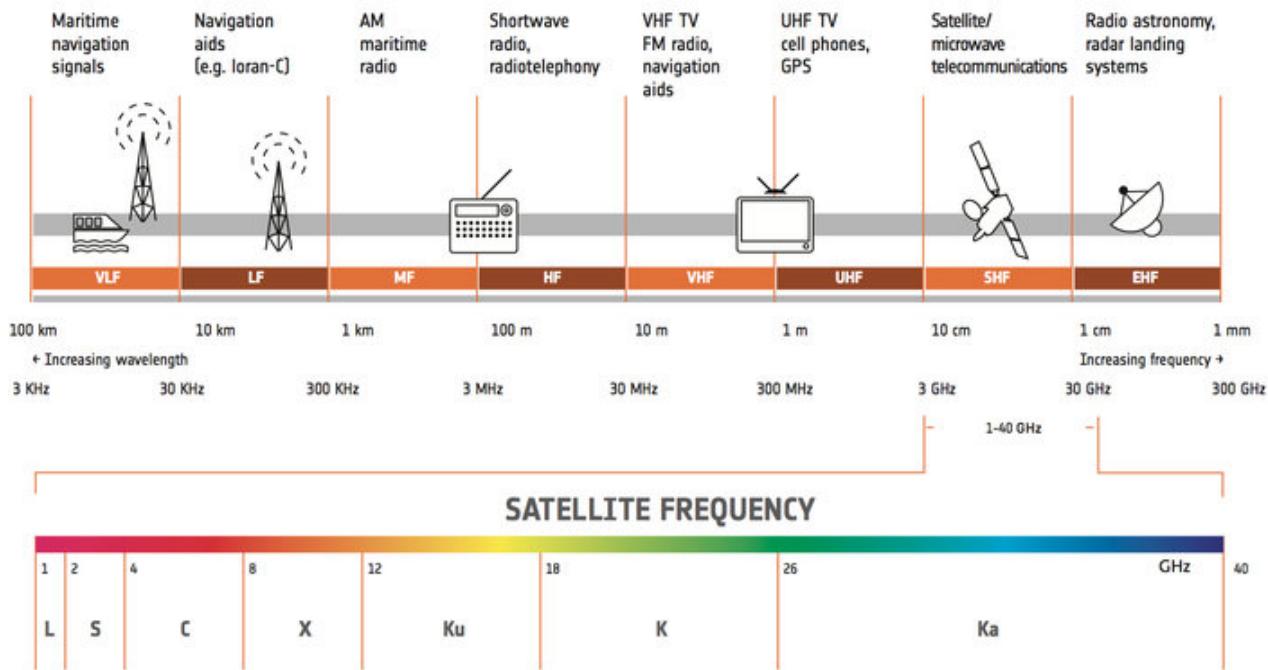
- Overloading: when all channels are busy, new devices will have to share channels that are already in use. Wifi allows connections under that regime, but speed goes down (wifi favors connectivity over speed).
- Channel width: may be broad (=40MHz) if there are no other routers around. Otherwise should be narrow (=20MHz), so as to avoid interference.
- Channel assignment: usually, routers chose a channel for you upon first handshake. Check if the router provides too few channels, or if the chosen channels are already in use by neighboring routers (good routers will automatically chose channels based on usage).
- Copper: wifi-connectivity is an important issue, but once at the router, all packages are funneled through one copper (or fibre) thread. So even with 5G, many devices may experience low performance because the *wired* connection is not fast enough.

3.2.7 GSM, EDGE, 3G=HSDPA, 4G=LTE=HSDPA(+) (radio)

When your device has no WLAN or LAN access to a router, you can only access the internet by going over radio-stations (or satellite, see later).

3.2.8 GPS (satellite)

Since with GPS a signal has to travel all the way to space and back, this is the slowest of the communication methods.



3.3 Internet - practical aspects

While the important theoretical aspects of the internet have been discussed in the section *networking*, there are so many important details that we devote another section purely to those.

3.3.1 Nomenclature

- url
- uri
- domain
-

3.3.2 Session basics

software :

- node: client-sessions (good crypto)
- bcryptjs for hashing
- csurf for CSRF

http session :

- http: Since http is a stateless protocol, we need some helper to store state over several page-visists
- session: webserver tells browser to remember the user
 - http-response-header: "Set-Cookie": "session=12345"
 - cookies: server stores information here
 - * that info is sent with every http-request's header
 - * "Cookie": "session=12345"
 - * Cookies are stored in the browser in a secret spot (not the usual browser-cache or db). They are only sent to a site if the sites url matches the cookie-domain. This means that sites cannot see cookies from other sites.
 - client-sessions has some very recomended headers for cookies:
 - * httpOnly: true // dont let js access cookies
 - * secure: true // only set cookies over https
 - * ephemeral: true // destroy cookies when browser closes

storing passwords : passwords must be hashed.

SSL/TLS

- Those cookies are just strings in the HTTP-header! They can be stolen! So they need to be encrypted.
- What is not encrypted by ssl?
 - SSL does encrypt everything, except ...
 - * the initial handshake target (ie your banks url)
 - * every request's url
 - but it **does** encrypt all headers (hence cookies), and post- **and** get parameters.

Tokens :(more specifically, JWT, OR, a bit more extensive, OAuth2. Not to be confused with USB-stick tokens)

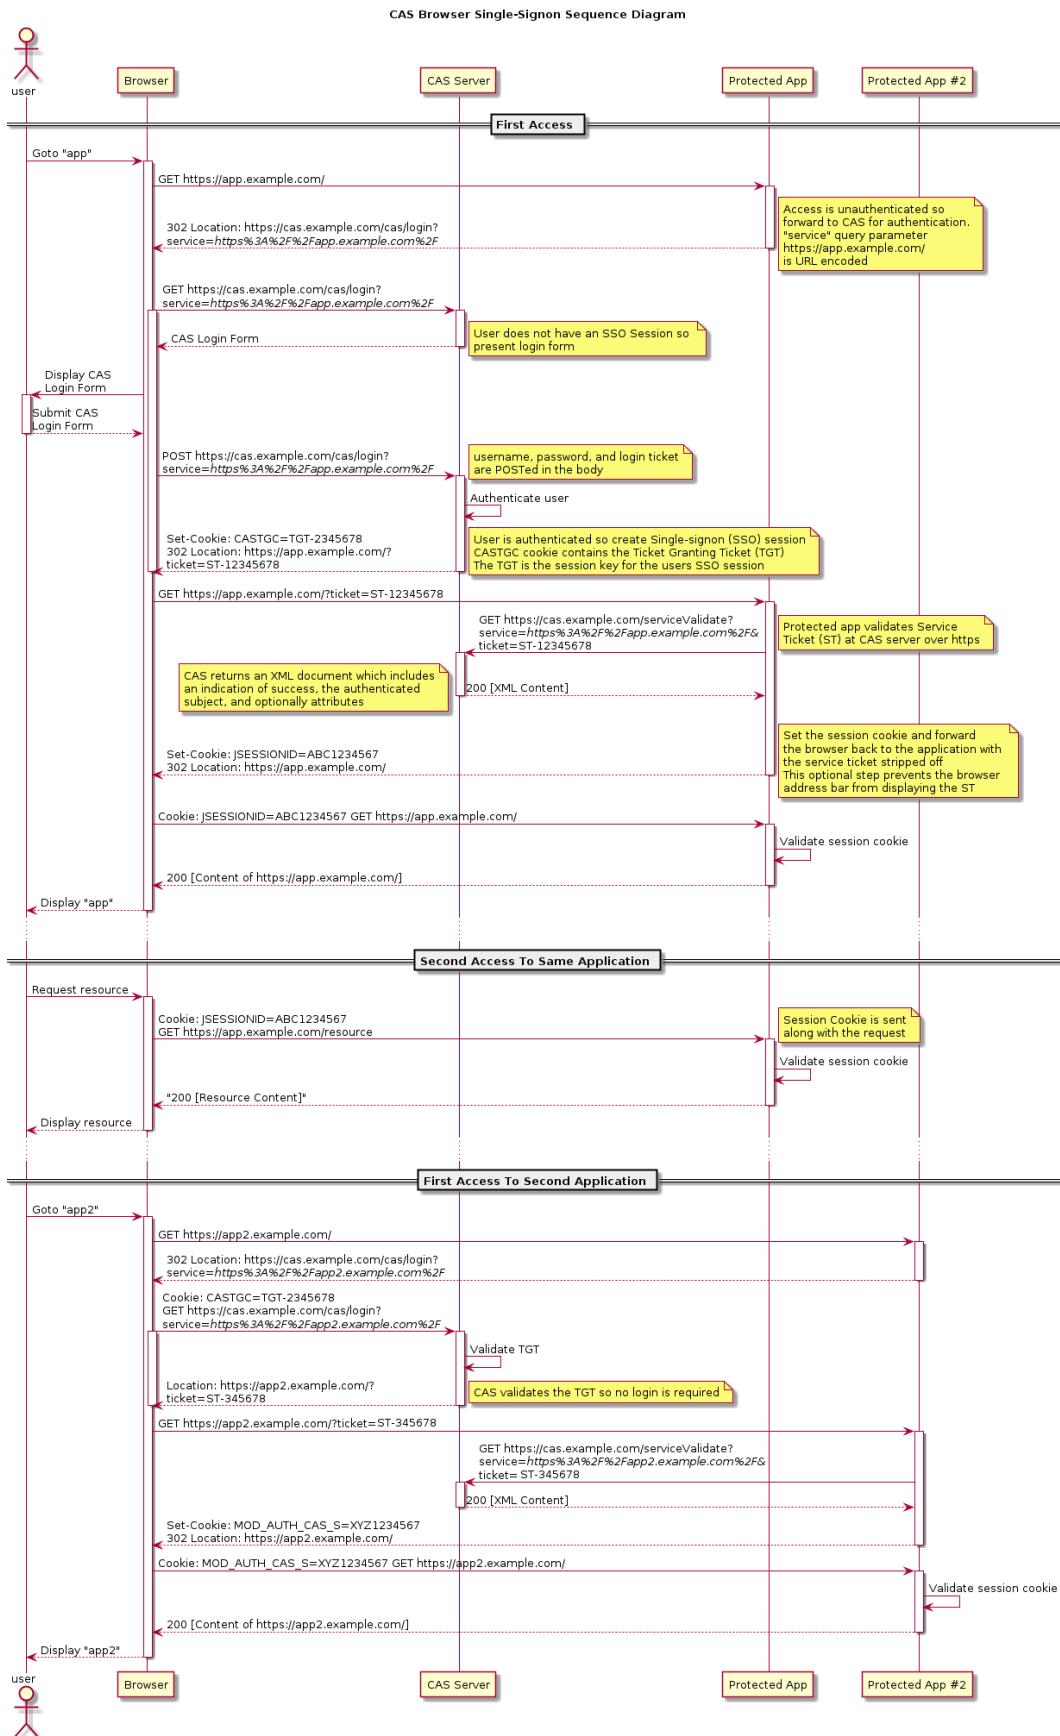
- Originally, tokens were designed as an alternative to cookies that would also work in apps.
- Based on this, however, two protocols (JWT and OAuth) grew, that go beyond what you could do with cookies.
- Here's what you cannot do with cookies. They are a two-party thing.
 - they are bound to a single domain. but your app might have multiple backend-services.
- OAuth/JWT is for multi-party things.
- JWTs are like a Vollmacht:
 - your app (tinder) wants to access your facebook photos
 - you login at facebook to tell it that tinder may use those photos
 - facebook gives you a token that reads: "photoPermission": "1234abcd"
 - you pass that token to tinder
 - tinder sends the token to facebook
 - facebook verifies that you have indeed given that permission
 - facebook gives tinder the photos
- Technically, ...
 - ...they are commonly implemented as a special type of cookie.
 - But they could be implemented as Post-Header + localStorage, too.
 - (more-often so, actually, because contrary to browsers, mobile-apps don't know about cookies)
- Some more stuff:
 - cookies usually contain any information that some programmer deemed necessary.
 - tokens follow strict standards (JWT or OAuth)
 - Tokens usually don't require tinder to maintain a session in its database (though they might)
 - * But does facebook still need to maintain a session?
 - Tokens don't require a remote-session because the token contains user-data.

3.3.3 Authentication, SSO, and authorization

Information from here. Assume that you want to check the identity of a user. That user is already member of some authentication-service like facebook. You could ask the user for his facebook-password. But authenticating at some place with your password is insecure, because your password needs to make it over the network. Instead, the user can ask facebook for a token, pass it to your site and you then check with facebook if this token is cool. Strictly speaking, OAuth is about authorisation (is the user allowed to use my service?), not about authentication (is the user really Michael?). The common analogy I've seen used while researching OAuth is the valet key to your car. The valet key allows the valet to start and move the car but doesn't give them access to the trunk or the glove box. An OAuth token is like that valet key. As a user, you get to tell the consumers what they can use and what they can't use from each service provider. You can give each consumer a different valet key. They never have the full key or any of the private data that gives them access to the full key.

	Centralized	Decentralized
Authorization (403)	CAS has an extension for authorization.	OAuth2: The user tells your app and a second app what your app may do on second app
Authentication(401)	CAS: your backend verifies that the user is really him (backend may be one of Kerberos, LDAP or ActiveDirectory)	OpenId: A usage of OAuth2 for authentication. The user tells your app on which second app that the user is really him Example: Sign-in-with-github

To confuse things a little, some CAS Implementations also support OAuth.



Both CAS and OAuth are 3rd party authentication-systems. In both cases, the guiding principle is the same; the situation is like a club-owner, a bouncer, and a guest.

- The guest (user, aka. client) doesn't like the bouncer (CAS-service, aka. classified app, aka. cas-client) and doesn't trust him with his real name. So instead, he gets a ticket (service-ticket) from the club-owner (CAS-server) to present to the bouncer.
- The bouncer doesn't trust the guest: he checks with the club-owner if that ticket is really valid.
- The club-owner
 - gives the guest a ticket, but doesn't write on it what the guest is entitled to, so the guest can't mess with his ticket.
 - verifies the ticket with the bouncer and tells the bouncer that the ticket is worth entry and a beer, but no VIP-lounge.

CAS and OAuth both ...

- keep guest and bouncer on a need-to-know basis
- it is the bouncer that needs to ensure the proper implementation of ...
 - directing the user to the boss (the CAS-server) for login
 - verify the ticket (service-ticket) presented by the user with the boss
- The login-form is provided by the boss, not the bouncer. The bouncer is oblivious to the user's creds or any other details except the privileges communicated by the boss.

That being said, there is a CAS-REST API that you could use to have the user pass his credentials to the app (in the app's own, custom form) so that the app can then post those creds to the CAS-server's REST-endpoint on the users behalf. However, this requires the user to trust the app a lot more: he now gave the app his credentials and must trust it to not abuse them (for example to log in at another service). Also note that there is no way to implement CAS in the frontend alone - you need something to protect your backend. Having CAS in the frontend but not protecting your fileserver in the backend is like having a bouncer control the parking lot: people can still sneak in by just walking directly to the club's door.

3.3.4 Security threats


CSRF :

- someone makes you click a link that is actually a form-post.
- Counter measures:
 - additionally to the session-cookie, we create a token for every page-load.
 - token provided by server
 - token added to form (as hidden field)
 - token sent back with form when user hits submit
 - server compares form-token with given token
 - if not identical, then
 - some bad guy sent us a link that we didn't

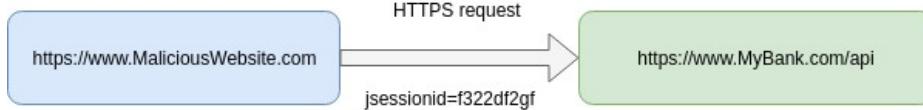
XSS : ...

	XSS: evil code being sneaked into frontend	XSRF: evil request being sneaked into backend
Roles	User trusts compromised site	Site trusts mislead user
Mechanism	Attacker injects script into site	Attacker tricks user into issuing request
Method	Attacker posted a comment that was not properly validated and now displays evil scripts to another user	Attacker made user click on something that caused a request to the backend
Counters	Form validation: don't allow script tags and similar in posts.	Form-tokens: a form-post must have been expected Cors: a form-post can only come from backend's domain

CORS (cross origin resource sharing) is the act of one site (www.mysite.com) requesting data from another, third-party site (www.puppyimagehost.org?img=3241). Doing so is forbidden by default, and browsers will block such requests. The reason is that a hacker could use his browser to do something malicious. This here is a perfectly fine request:



After that request, the `jsessionid` cookie will be stored on the browser. A hacker could now steal that cookie and do this:



This is called Cross-site-request-forgery. Honestly, it seems that this is something that should be fought on the *server*side, not from the client.

However, some sites depend on CORS.

- For one, images and scripts can be loaded from third-party sites (using the `src` attribute). As such, for historical reasons, images are exempt from CORS-blocking
- Maybe www.puppyimagehost.org *wants* to expose an API for other sites to call. In this case, www.puppyimagehost.org can set an additional header on its response: `Access-Control-Allow-Origin: www.mysite.com` or `Access-Control-Allow-Origin: *`

Note that there is no way that www.mysite.com can deactivate CORS-blocking; it *must* be allowed by www.puppyimagehost.org. An individual person might deactivate the browsers CORS-settings (if the browser allows you to do that), but the common user will never do that.

CORS when applied to an OL-canvas Really, using 3rd party data is a bit like a one-night-stand: you need consent from both sides. The server gives consent to the usage of its images by setting `Access-Control-Allow-Origin`, the client does so by setting `crossorigin="Anonymous"` (or `crossorigin="use-credentials"`). Now a canvas is a special case. For one, often we load images from more than one external source, so we need to get two-way-consent with a whole series of services. On the other hand: even if we don't actively set `crossorigin`, data is displayed - errors only start to occur once we want to call `canvas.toDataURL(...)`. Why is that? Well, canvases load `img`'s, and those are, as mentioned above, exempt from CORS-restrictions by default. However once we want to do any manipulations on that data, CORS kicks in again. That is why you can show any layer on a canvas, but cannot export it to pdf as long as it displays any layer that does not have two-way-consent.

CORB (cross origin read blocking) is when your browser deletes a request's response body before it can be read by your site. This is to prevent www.puppyimagehost.org from accidentally leaking sensitive information. Even with CORS, an attacker might use html tags like `img` to circumvent CORS-blocking. This is not a case of cross-site scripting, but rather of building a malicious website from scratch just to sneakily access data.

- First, load remote data to its site using `` or `<script src="https://your-bank.example/balance.json">`
- The browser would notice that the data in `src` is not an image (or javascript) and consequently would not display the data as image.
- But the browser *would* still have the data in it's memory. The attacker can then use a browser-memory vulnerability like spectre to access this data.

Let's break down how CORB works. A website can request two types of resources from a server:

- data resources such as HTML, XML, or JSON documents
- media resources such as images, JavaScript, CSS, or fonts

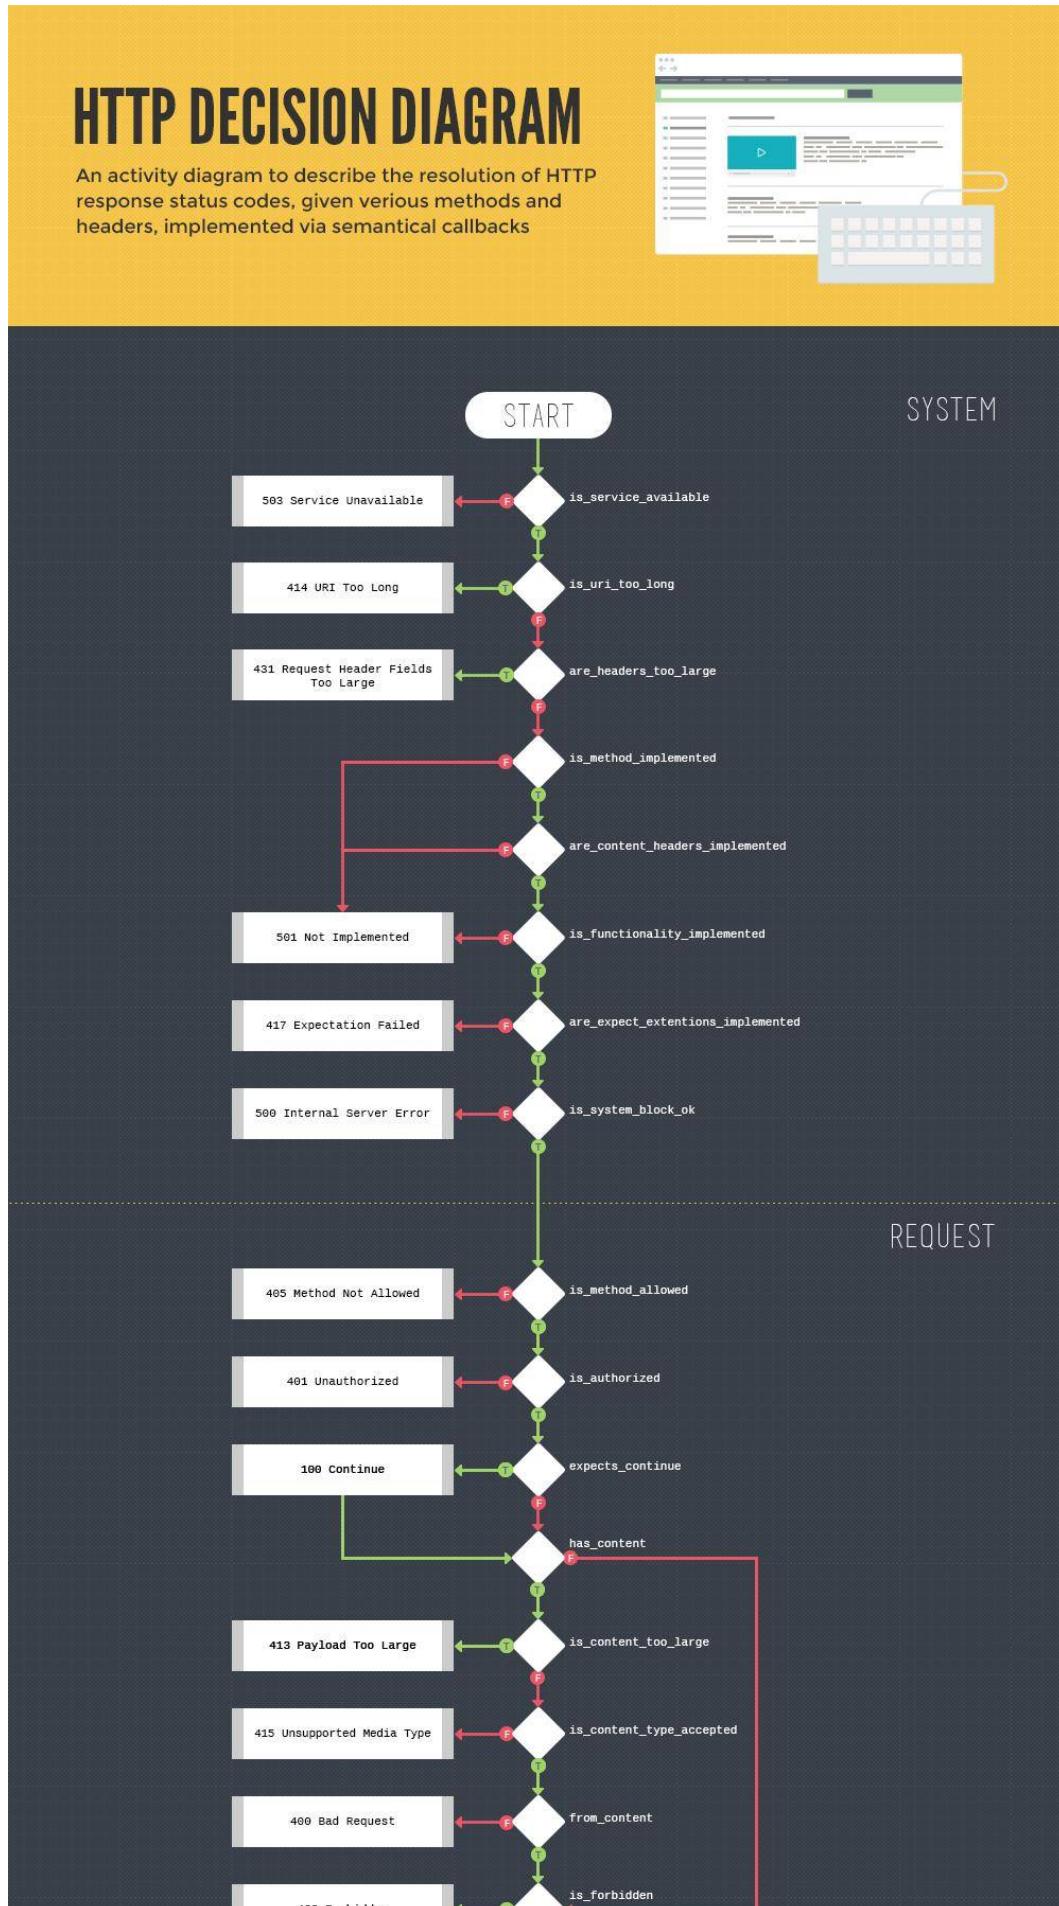
A website is able to receive data resources from its own origin or from other origins with permissive CORS headers such as `Access-Control-Allow-Origin: *`. On the other hand, media resources can be included from any origin, even without permissive CORS headers.

CORB will block the response of a request if all of the following are true:

- The resource is a "data resource". Specifically, the content type is HTML, XML, JSON
- The server responds with an X-Content-Type-Options: nosniff header, or if this header is omitted, Chrome detects the content type is one of HTML, XML, or JSON from inspecting the file
- CORS does not explicitly allow access to the resource

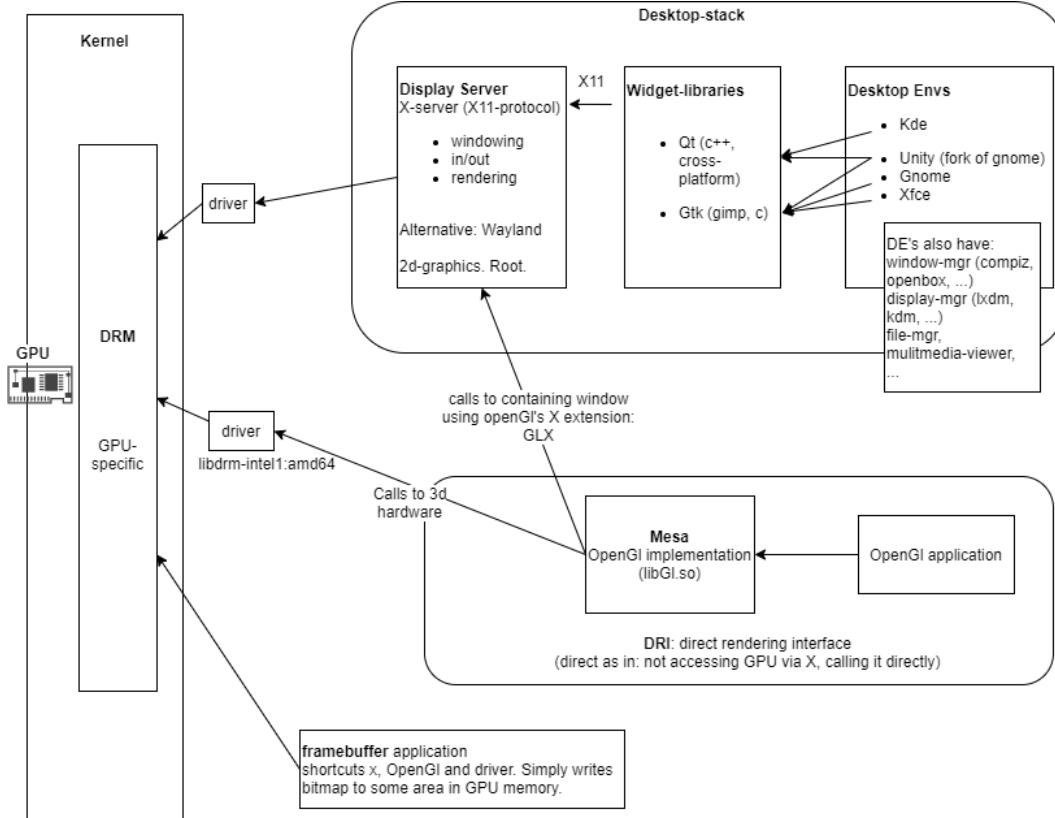
CORB protects neither against XSS nor XSRF. Instead, it prevents one mechanism of side-channel-attacks.

3.3.5 Status codes



3.4 Linux render-stack

Figure 3.1: Drivers: everything on the right of a driver is hardware-agnostic, everything on the left hardware-specific. Note that Xorg and mesa don't use the same drivers. Historically, linux used to only do 2d-graphics. For a brief time after 3d-cards arrived, OpenGL instructions were passed through Xorg. Only later was mesa allowed to bypass Xorg - hence the name *direct* rendering infrastructure.



3.5 Linux sound-stack

All sound-in (microphone, jack/cinch aka line-in) and sound-out (loudspeakers, headphones) cables lead to your soundcard, which has the sole task of transforming analog sound signals to digital `short[]` arrays. The whole routing is done via the ALSA System. Because ALSA is a little complicated, linux has a layer on top of it. Two applications can be used as this top-layer: pulseaudio and jackd.

3.5.1 Audio Standards

Raw audio consists of `short[]` arrays containing raw amplitude data.

Midi comes in two forms: JACK-Midi and ALSA-Midi.

3.5.2 Hardware

Use `aplay -l` to list all your sound cards.

PCH is your actual soundcard.

HDMI is combined audio/video transport. Mostly used by monitors, but also shows up in `aplay -l` because it can carry sound as well.

3.5.3 Software

There are different systems that control your soundcards.

ALSA is linux's standard sound routing mechanism.

Portaudio is an alternative to alsa that can also run on windows and mac.

Jack is a router that lets you plug one program's sound output into another program's sound input - or to the system's soundcard.

3.5.4 JACK

Jack deserves its own section.

Frames per period * periods per buffer yields your buffer size. Big buffer means low CPU load, but large latency. Small buffer means small latency, but high CPU load. If the CPU load approaches 100%, you'll hear *xruns* (which stands for buffer-Under- or Over-run), that is clicking and artifacts caused by your CPU not being able to process buffer-batches on time. Jack displays CPU usage under the term *DSP load* in qjackc.

Realtime (RT) mode requires your kernel to support realtime scheduling (called low-latency kernel). When that is given, jack will be given priority in memory and CPU over most other tasks.

	Public	Downstream license	Changes	Pat
GPL3	Sourcecode must be made public	All modifications must also use GPL3	All changes must be documented	If t
GPL2	"	"	"	-
Apache2	Needs not be made public	Modificators may use any license	"	San
BSD2	"	"	Changes need not be documented	-
BSD3	"	"	"	-
MIT	"	"	"	-

Table 3.1: Licenses from strict (top) to permissive (bottom)

	Static linking (into an Apache-licensed app)	Dynamic linking (into an Apache-licensed app)
Example	Webapp where all files are bundled and obfuscated.	Node-app: Libraries are included via require()
GPL3	X	
GPL2	X	
Apache2	✓, preserve the NOTICE file	✓
BSD2	✓, add license and copyright in documentation of compiled product	✓
BSD3	✓, add license and copyright in documentation of compiled product	✓
MIT	✓	✓

Table 3.2: Licenses kick in when your code is distributed. Putting files on a webserver counts as distribution, with maybe the exception of code that only serves as a service (depends on the case). Source: <https://medium.com/@vovabilonenko/licenses-of-npm-dependencies-bacaa00c8c65>

3.6 Legal and opensource

3.6.1 Popular licenses

All the below licenses are approved by the OSI.

3.6.2 Adding 3rd party licenses to your apache2-javascript product

3.6.3 Deutsches/Europäisches Recht

3.6.3.1 Haftung

Obwohl Lizenzen üblicherweise sagen, dass der Anbieter nicht haftbar gemacht werden kann für Schäden, gilt das im deutschen Recht nicht. Laut deutschem Recht ist ein Anbeiter nach dem 'Schenkungsrecht' für Schäden durch Fahrlässigkeit haftbar.

3.6.3.2 Exportkontrolle

Prinzipiell gilt der Grundsatz des freien Warenverkehrs. Aber Exportkontrolle kann greifen,

- Wenn Deutsche, Europäische oder Verbündete (sprich: USA) Sicherheit gefährdet ist,
- um Embargos zu erzwingen,
- Um in den einführenden Staaten nicht misbraucht werden zu können.

Insbesondere gilt hier die Europäische Dual-Use Verordnung.

Chapter 4

Programming languages

A good language

- has you thinking about the problem, not the language
- has enough libraries and community

4.1 C

4.1.1 General theory

4.1.1.1 Memory allocation

By default, memory is allocated statically to the stack. By using `malloc` and `free`, you can instead allocate memory dynamically on the heap.

4.1.1.2 Threading

Threads are created by making a copy of the original process. Threads share the same memory with their parents.

Table 4.1: Processes versus Threads

	Sub-Process	Thread
Creation	Copy of mother process	
Memory	Own memory	Shared memory
Communication	Communicates with mother through syscalls, pipes and files	Can directly call methods of mother process
Usecase	Ideal if mother and subprocesses must be separated for security reasons, like apache-server does.	Ideal if thread output to be processed by mother process, because no piping necessary

```
#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>

// A normal C function that is executed as a thread
// when its name is specified in pthread_create()
void *myThreadFun(void *vargp)
{
    sleep(1);
    printf("Printing GeeksQuiz from Thread \n");
    return NULL;
}

int main()
{
    pthread_t tid;
    printf("Before Thread\n");
    pthread_create(&tid, NULL, myThreadFun, NULL);
    pthread_join(tid, NULL);
    printf("After Thread\n");
    exit(0);
}
```

As mentioned above, all threads share data segment. Global and static variables are stored in data segment. Therefore, they are shared by all threads. The following example program demonstrates the same.

```
#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>

// Let us create a global variable to change it in threads
int g = 0;

// The function to be executed by all threads
void *myThreadFun(void *vargp)
{
    // Store the value argument passed to this thread
    int *myid = (int *)vargp;

    // Let us create a static variable to observe its changes
    static int s = 0;

    // Change static and global variables
    ++s; ++g;

    // Print the argument, static and global variables
}
```

```

    printf("Thread ID: %d, Static: %d, Global: %d\n", *myid, ++s, ++g);
}

int main()
{
    int i;
    pthread_t tid;

    // Let us create three threads
    for (i = 0; i < 3; i++)
        pthread_create(&tid, NULL, myThreadFun, (void *)i);

    pthread_exit(NULL);
    return 0;
}

gfg@ubuntu:~/ $ gcc multithread.c -lpthread
gfg@ubuntu:~/ $ ./a.out
Thread ID: 1, Static: 1, Global: 1
Thread ID: 0, Static: 2, Global: 2
Thread ID: 2, Static: 3, Global: 3
gfg@ubuntu:~/ $

```

Please note that above is simple example to show how threads work. Accessing a global variable in a thread is generally a bad idea. What if thread 2 has priority over thread 1 and thread 1 needs to change the variable. In practice, if it is required to access global variable by multiple threads, then they should be accessed using a mutex.

4.1.1.3 Getting data from C to another programm

Table 4.2: Pipes versus sockets

Pipe	Socket
Unidirectional	Bidirectional
Fifo	
Limited volume (~0.5 MB); can block!	
all in memory	all in memory
no protocol	packets (UDP or TCP), can go over network

Per file Simplest, but also slowest because it involves disk i/o.

Per pipe or socket Faster than file-writing, because no disc i/o is involved. But still slow because data is chunked into packages that are wrapped in a rather verbose protocol.

Per JNI This should allow you to convert C-datastructures to Java-datastructures.

4.1.2 Quirks and features you need to know

4.1.2.1 * syntax

```

int a = 1;
int * a_ptr = &a;
*a_ptr; // <--- yields 1

```

When assigning, * means: "make it a pointer".

When querying, * means: "get the value behind the pointer".

& always means: "get the address".

4.1.2.2 Pointer arithmetic

Array indexing is actually a rather complex issue with a lot of syntactical sugar.

```
|| int arr[3] = {10, 20, 30};
|| arr[2] = 15;
```

- *arr* gets you the pointer to the first element in the array. In other words, we have $arr = \&arr[0]$
- [2] adds to *arr* the size of 2 ints. Thus we have $arr[2] = \&arr[0] + 2 * sizeof(int)$.
- Finally, the actual value behind the expression is adjusted when we assign the value of 15. Thus the expression reduces to

```
|| *(arr[0] + 2 * sizeof(int)) = 15;
```

We can use array indexing to our advantage: with *buffer overflows* we can read arbitrary parts of the memory. Consider this simple example:

```
|| int main () {
||     int arr[3] = {1, 2, 3};
||     int i;
||     for(i=0; i < 100; i++){
||         printf("%d th entry: location: %p value: %d \n", i, &arr[i], arr[i]);
||     }
||     return 0;
|| }
```

Especially when creating arrays, we need to be aware of a few things:

- ... what is the current content of the accessed memory?
- ... is the accessed memory protected against overwriting by others?

```
#include<stdio.h>
#include <stdlib.h>

int main() {

    const int size = 1000;
    //double data[size] = {2.1,2.1,2.1, 2.1, 2.1};    //--- puts in vals; fills rest with zeros. This mem won't be overwritten
    //double data[size];      //--- does not put in zeros - keeps whatever was in there before. This mem won't be overwritten
    //double* data = (double*) malloc(size * sizeof(double));           //--- puts in zeros. This mem won't be overwritten
    //double* data;          //--- does not put in zeros - keeps whatever was in there before. This mem can be overwritten by others
    for(int i = 0; i < size; i++) {
        double val = data[i];
        printf("val at %i: %f \n", i, val);
    }

    return 0;
}
```

When we state that *this mem won't be overwritten*, of course we exclude overwriting by buffer overflow.

4.1.2.3 Array decay: Functions can't accept arrays

Functions never accept arrays, they only take pointers to the first element. This is known as array-decay.

```
|| int arr[3] = {1,2,3};
|| somefunc(arr); // ---- compiler turns this automatically into
|| somefunc(&arr[0]);
```

4.1.2.4 Array decay: Functions don't return arrays, either

Functions return single values just fine, but arrays only by pointer. *arrFunc* must save array on heap and return pointer to the heap. The calling function must know the arrays size in advance or be given a struct with metainfo about the size. The calling function must also later deallocate the array.

```
|| int * arrOnHeap(){
||     int * arr_ptr = malloc( sizeof(int) * 3 );
||     arr_ptr[0] = 10;
||     arr_ptr[1] = 20;
||     arr_ptr[2] = 30;
||     return arr_ptr;
|| }
|| int * arr_ptr = arrOnHeap();
```

4.1.2.5 Array syntax

Java and c have some weird differences in array initialisation. Consider array literals:

```
|| int coeffs[5] = {1, 2, 3, 4, 5}; // c
|| int[] coeffs = {1, 2, 3, 4, 5}; // java
```

And also standard initialisation:

```
|| int coeffs[5]; // c
|| int[] coeffs = new int[5]; // java
```

4.1.2.6 Struct syntax

There is a really important thing when creating struct-construction functions. Consider the following code.

```
|| typedef struct Island {
    char* name;
    char* opens;
    char* closes;
    struct Island* next;
} Island;

Island* island_create(char* name){
    Island* i = malloc(sizeof(Island));
    i->name = name;
    i->opens = "09:00";
    i->closes = "17:00";
    i->next = NULL;
    return i;
}

int main(){
    char* name;
    fgets(name, 80, stdin);
    Island* island1 = island_create(name);
    fgets(name, 80, stdin);
    Island* island2 = island_create(name);

    return 0;
}
```

You will find that the name of island1 equals that of island2! The reason is that their names are just a reference to `char*` `name` in the `main` method. We need a safety measure inside the constructor to allocate a copy of the input string so that we can sure a new call to the constructor does not give us the same pointer again.

This code fixes the problem:

```
|| Island* island_create(char* name){
    char* namec = strdup(name);
    Island* i = malloc(sizeof(Island));
    i->name = namec;
    i->opens = "09:00";
    i->closes = "17:00";
    i->next = NULL;
    return i;
}

void island_destroy(Island* i){
    free(i->name);
    free(i);
}
```

4.1.2.7 Passing functions as variables

This is our entry to functional programming in c. A functionname is really just a pointer to the memory location where the function code is stored. So we can use the function name as a pointer.

Of course, functions have different types. That's why we can't just write

```
|| function* f;
```

But instead must write:

```
|| char* (*reverseName)(char*);
```

Here, the first `char*` is the return type, whereas the second is the argument(-list) to the function. Let's see an example:

```

||#include <stdio.h>
||#include <stdlib.h>
||#include <string.h>

int num_ads = 7;
char* ads[] = {
    "William: SBM GSOH likes sports, TV, dining",
    "Matt: SWM NS likes art, movies, theater",
    "Luis: SLM ND likes books, theater, art",
    "Mike: DWM DS likes trucks, sports and bieber",
    "Peter: SAM likes chess, working out and art",
    "Josh: SJM likes sports, movies and theater",
    "Jed: DBM likes theater, books and dining"
};

int likes_sport_not_bieber(char* ad){
    int hit = (strstr(ad, "sports") && !strstr(ad, "bieber"));
    return hit;
}

int likes_sports_or_workout(char* ad){
    int hit = (strstr(ad, "sports") || strstr(ad, "workout"));
    return hit;
}

void find(int (*match_fn)(char*)){
    puts("Search results:");
    puts("-----");

    int i;
    for(i=0; i < num_ads; i++){
        if(match_fn(ads[i])){
            printf("%s\n", ads[i]);
        }
    }

    puts("-----");
}

int main(){
    find(likes_sport_not_bieber);
    find(likes_sports_or_workout);

    return 0;
}

```

We should note some syntactic sugar here. Calling functions is usually done like this:

```
|| int a = somFunc(b);
```

But that's just shorthand for the actual code:

```
|| int a = (*somFunc)(b);
```

Also, passing functions as arguments really compiles down to:

```
|| find(&likes_sport_not_bieber);
```

4.1.2.8 Strings end with a \0

That's why, when declaring an empty string, leave space for the \0 .

```

|| char word[3];
|| strcpy(word, "hi");
|| strlen(word); // <--- yields 2
|| sizeof(word); // <--- yields 3

```

4.1.2.9 char[] does not equal char*

```

|| char a[] = "hello";
|| char * b = "world";

```

a is equal to: `H E L L O \0`

b is equal to a `pointer` pointing to `W O R L D \0`

4.1.2.10 Header files

Header files are how we can expose only a subset of a file to main.c. Really they contain only the function signatures, not their implementation. In that way, c basically forces you to import *everything* as an interface.

Listing 4.1: main.c

```
#include <stdio.h>
#include "encrypt.h"

int main(){
    char msg[80];
    while(fgets(msg, 80, stdin)){
        encrypt(msg);
        printf("%s", msg);
    }
}
```

Listing 4.2: encrypt.h

```
|| void encrypt(char * msg);
```

Listing 4.3: encrypt.c

```
#include "encrypt.h"

void encrypt(char * msg) {
    char c;
    while(*msg){
        *msg = *msg ^ 31;
        msg++;
    }
}
```

4.1.2.11 Building and Makefiles

The building of an executable consists of compiling and linking.

- **Compiling your own code** The compiler generates an object file from your source-code. An object-file is a compiled piece of code together with some meta-information on what kind of functions and structures it contains.
- **Linking in libraries** After that, the include-directories are scanned for libraries that contain implementations of the header-files you included in your source-code. This phase is called linking. There are two ways you can link in external libraries.
 - Static linking means that a copy of the library is added to the final executable at compile time. The executable can then be moved to another machine without worries, because all required libraries are placed inside the executable. Static libraries are usually written like this: `libXXX.a` and created with the `ar` program.
 - Alternatively, libraries can be linked in a dynamic way: that means that the executable will search for an appropriate library once it needs to call its functions, that is, at run time. Shared/dynamic libraries are written like `libYYY.so` and created with `gcc -fPIC -c YYY.c && gcc -shared -o libYYY.so YYY.o`
 - Both `.a` and `.so` libraries are included with the `-l` and `-L` commands.

Makefiles are how we knot together different parts of a c program. It's really easiest to look at a specific example:

```
# Includes are header files without a concrete implementation.
INCLUDES = -I/usr/include/mysql
# Libraries are object files. -L adds a directory of libraries, -l adds a single library.
LIBS = -L/usr/lib/x86_64-linux-gnu -lmysqlclient -ljansson -pthread
WARNINGS = -Wall -Wextra
```

```

# Compiling
#   -c: compile (nur compile, nicht link!)
#   -g: fuer debugger

incl.o: incl.c incl.h
        gcc -g -c $(WARNINGS) $(INCLUDES) incl.c

jsn.o: jsn.c jsn.h incl.h
        gcc -g -c $(WARNINGS) $(INCLUDES) jsn.c

psql.o: psql.c jsn.h
        gcc -g -c $(WARNINGS) $(INCLUDES) psql.c

# Linking
#   -o: object-file: name der fertigen binary
#   -g: fuer debugger

psql: psql.o jsn.o incl.o
        gcc -g -o psql psql.o jsn.o incl.o $(LIBS)

clean:
        rm *.o

all: psql clean

```

Includes sind header files. Option `-I(dirname)`: hinzufügen dir zu Suchpfad für *.h files. Im Gegensatz zu **libs** kann man bei **incls** nur dirs spezifizieren, nicht einzelne files. Dafür gibt es ja schon die `#include` Direktive.

Libs sind die den header files zu Grunde liegenden so files. Option `-L(dirname)`: hinzufügen dir zu Suchpfad für *.so files. `-l(libname)`: mit Einbinden einer lib. `-pthread`: eine besondere Option; steht für -lpthread + definiere ein paar extra macros

Jansson ist ein externes Programm. Es muss erst ge-make-t werden (oder ge apt-get-tet), danach finden wir header per default in /usr/local/include und libs in /usr/local/lib

4.1.2.12 Valgrind

The `-g` flags from the above makefile were actually meant for use in valgrind. Valgrind analyses your memory allocation. It does so by creating its own, fake versions of `malloc` and `free`. Anytime these are called, valgrind does its bookkeeping to check if any allocated memory is left on the heap.

Analysing code with valgrind is easy. Just compile it and then run it like this:

```

|| make all
|| valgrind -v --leak-check=full --show-leak-kinds=all --log-file=val.log ./psql input.json

```

4.1.2.13 OpenMP

OpenMP is a set of macros used for control-flow for threads. Think of automatically distributing your for-loop over threads.

4.1.3 CMake and Autotools

While make does a good job at automating the build of c-programs, all the commands you type there are platform dependent. For that reason, larger projects use more sophisticated buildsystems. They don't require less configuration (in fact, they usually require *more* work) but they make it easier to create a build that will work on any platform. Usually they work by first scanning your environment and then creating makefiles for you.

Working with these is not exactly pleasant. This section only contains a *how-to* handle existing open-source projects that you have to get to compile somehow.

We begin with autotools.

- `./configure`
- `make`
- `make install`
- Troubleshooting: most likely, some shared libraries will be missing. Find out which by calling

4.1.4 Best practice

The above quirks gave us a lot of points where we need to be careful in c. For that reason, we should adhere to some best practices to make coding as save as possible.

Only expose top and lowest layer, not internals Datastructures in C tend to be actual data wrapped in structs wrapped in structs wrapped in structs. The highest level struct should only expose data, not any intermediate structs. This way, a user only needs to know the api for the highest level struct and the data itself. Also, every struct should manage its own and only its own memory.

Handling unknown datatype elements Your datastructures will usually contain elements of a type unknown to you. That is not really a problem, because you can reference them using a `void*`. But how about creating and destroying those unknown elements? Well, here we can use c's functional programming skills: just add to the datastructure a function for creating and one for destroying the element.

Never leave a pointer unassigned You can create a pointer without having it point anywhere in particular: `int* apt;.` But what if later in your code you want to check if that pointer has been pointed to an int yet? `apt` is initially just going to point to some random location. This means that you cannot check `if(apt == NULL)`, because it's never going to be 0! For this reason, even if you don't want `apt` to point to anything yet, at least make sure it points to `NULL`. So always create pointers with `int* apt = NULL;`

4.2 C++

C++ is the best compromise between raw power and versatility. Fortran is faster in large matrix-operations, but lacks the abstraction of OOP and any multi-media support. Java is even better abstracted and has an unbeatable ecosystem with modules for everything, but slower.

4.2.1 Some weird syntax and new concepts

Contrary to garbage-collected languages, where you only need to care about who knows about a given object, in cpp you also net to worry about where that object lives.

4.2.1.1 References

References are really just syntactic sugar around pointers. It's still useful to know a few of their basic properties, since they are intended to help you avoid errors in memory management.

Using pointers, to mutate an object you need to use pointer syntax for any operation on the object. This is known as dereferencing: using `*vpt` and `vpt->`.

```
|| void foo(Vec* vpt){
||   vpt->increment(1);
|| }
|| Vec vec = Vec();
|| foo(&vec);
```

Using references we can act as if the actual object and not just a pointer to it was in the function scope.

```
|| void foo(Vec& v){
||   v.increment(1);
|| }
|| Vec vec = Vec();
|| foo(vec);
```

Syntax:

```
|| int i = 3;
|| int* i_ptr = &i;
|| int& i_ref = i;
```

There are a few more noteworthy properties of references that are intended to make your life easier:

- A pointer may point to nothing¹ - you may write `int* nullpt;`. A reference however mustn't do that. This will not compile: `int& nullrf;`.
- You can point a pointer to a new object, but a reference allways points to the same object. This will have `ap` point to `b`:

```
|| int a = 1;
|| int b = 2;
|| int& ar = a;
|| ar = b;
```

This, however, will overwrite `a`:

```
|| int a = 1;
|| int b = 2;
|| int& ar = a;
|| ar = b;
|| cout << a << endl; // yields 2 (!)
```

4.2.1.2 Const keyword

Sometimes you want immutability; in cpp probably even more than in other languages, for security reasons. That is what the `const` keyword is there for. In most cases, you'll use this keyword when a function is passed some objects that it mustn't alter. Passing the objects as consts guarantees that the function has no sideeffects on the objects.

¹However, this is very bad practice. Always initialize your free pointers to null using `int* ap = nullptr;`

```

||| bool isEqual(const Vec & v1, const Vec & v2) {
    bool eq = false;
    if( v1.x == v2.x && v1.y == v2.y ){
        eq = true;
    }
    return eq;
}

```

Of course, we could also have avoided side-effects on v1 and v2 by just passing by value, because then the function would have only altered it's own copies of v1 and v2. But creating those copies is a costly operation, so we're cheeper of by passing by reference and using const to achieve immutability.

Unfortunately, `const` has some weird syntax rules. The trick is in figuring out whether a value or a pointer to a value is to be constant. The rule is that `const` always applies to the keyword to it's left. If there is nothing on the left, it applies to the keyword on the right.

```

||| const int a = 1;           // a constant integer
const int* a_ptr = nullptr;   // a variable pointer to a constant integer
int const* a_ptr;            // a variable pointer to a constant integer
int * const a_ptr;           // a constant pointer to a variable integer (you will rarely need this)

```

4.2.2 Object orientation

Here is a simple example of the class-syntax in cpp.

```

||| #include <iostream>
#include <string>

using std::string;

class Robot {

private:
    string name;
    int hitpoints;

public:
    Robot(string n, int hp) : name(n), hitpoints(hp){
        std::cout << "Created " << name << " with " << hitpoints << " health" << std::endl;
    }

    void hit(Robot* other){
        other->takeHit(10);
    }

    void takeHit(int strength){
        hitpoints -= strength;
        std::cout << name << " now has " << hitpoints << " health" << std::endl;
    }

    ~Robot(){
        std::cout << name << " destroyed." << std::endl;
    }
};

int main() {
    Robot b("Bender", 100);
    Robot v("Vender", 100);
    b.hit(&v);
    return 0;
}

```

From this basic example we can already note a few important points.

Constructors cpp knows *three* ways of creating objects:

- The default constructor `Robot b("Bender", 100);` is what you should usually make use of.
- Construction by copying an existing instance `Robot b("Bender", 100); Robot v = b;` requires a functioning implementation of the copy constructor `Robot(const Robot& otherRobot)`. If this isn't provided, the compiler generates one automatically. The copy constructor is called in three cases:
 - When an object is created by copying: `Robot b("Bender", 100); Robot v = b;` or `Robot b("Bender", 100); Robot v(b);`
 - When passing an object to a function by value: `void someFunc(Robot r);`. Note, however, that it with big objects it may be better to pass the object by (const) reference: `void someFunc(const Robot& r);`

- When returning an object from a function by value: `Robot someFunc(); Robot b = someFunc();`
- . Here, too, it is usually more efficient to use references instead of copying by value.
- Construction by assignment operator `Robot b("Bender", 100); Robot v("Vender", 90); v = b;` requires a functioning implementation of the assignment operator `Robot& operator=(const Robot& otherRobot){ ... return *this;}`. If this isn't provided, the compiler generates one automatically.

Here an example to clarify the difference between assignment and copy-constructor:

```
#include<iostream>
#include<stdio.h>

using namespace std;

class Test{
public:
    Test() {}
    Test(const Test &t) {
        cout<<"Copy constructor called "<<endl;
    }
    Test& operator = (const Test &t) {
        cout<<"Assignment operator called "<<endl;
        return *this;
    }
};

int main() {
    Test t1, t2;
    t2 = t1;
    Test t3 = t1;
    getchar();
    return 0;
}
```

Constructors as objects or as pointers Objects can be constructed as objects or as pointers to objects.

Construction as object:

```
Particle p = Particle(20, 30);
p.move();
```

Here the object is automatically destroyed once it goes out of scope. Note the missing `new` keyword.

Construction as pointer:

```
Particle* p = new Particle(20, 30);
p->move();
delete p;
```

Here, the object survives until you manually call `delete p`.

Whenever possible, you should prefer the first method. However, there are some situations when it makes sense to use pointers:

- when the object is returned from a function (or any other situation where the object needs to outlive its scope)
- When you're passing the object to a function and you want that function to change the object and not just a copy of the object.

Initialisation Note also this important difference. In java, when you write `Particle p;` this creates an uninitialized reference. In c++ however, this same command already does the initialisation by calling the default constructor.

4.2.2.1 Operator overloading

You can override any operator, as long as the operation makes use of at least one custom datatype.

```
#include <iostream>
#include <string>

using std::string;

class Vector {
public:
    int x;
```

```

int y;
int z;
Vector(int _x, int _y, int _z) : x(_x), y(_y), z(_z) {
    std::cout << "Created: [" << x << "," << y << "," << z << "]"
    << std::endl;
}
~Vector(){
    std::cout << "Going out of scope: [" << x << "," << y << "," << z << "]"
    << std::endl;
}

Vector operator+(Vector const & v1, Vector const & v2) {
    int x = v1.x + v2.x;
    int y = v1.y + v2.y;
    int z = v1.z + v2.z;
    Vector v3 = Vector(x, y, z);
    return v3;
}

int main() {

    Vector a = Vector(1, 2, 3);
    Vector b = Vector(2, 3, 4);
    Vector c = a + b;
    return 0;
}

```

A very special case: assignment operator The assignment operator is a very special case. It is only allowed to be defined inside a class. But it is very useful to get an insight into what happens when you do an assignment in C++. example with overwriting an object

From this we learn a very important lesson about the assignment operator. When one object is assigned to another, a copy of the actual values is made. In Java, copying an object variable merely establishes a second reference to the object. Copying a C++ object is just like calling clone in Java. Modifying the copy does not change the original.

```

Point q = p; /* copies p into q */
q.move(1, 1); /* moves q but not p */

```

Assignment by copying versus assignment by moving The last paragraph was not absolutely honest with you: you can force C++ to move an object instead of copying it when doing an assignment. ex with std::move

4.2.2.2 Rule of three

The rule of three states that if you write a custom constructor or destructor, you're also going to need a custom copy-constructor and a custom assignment-operator. The rule basically comes down to this: if you do any memory management during construction, you will also do memory management during copying and during destruction. The compiler cannot guess that memory management for you, so you need to specify it yourself.

```

class Ball {
private:
    int xpos;
    int ypos;
    int xvel;
    int yvel;

public:
    Ball(int x, int y) : xpos(x), ypos(y) { // Standard constructor
        std::cout << "ball created" << std::endl;
    }

    ~Ball() {
        std::cout << "ball deleted" << std::endl;
    }

    Ball(const Ball&) = delete; // Copy constructor (note that argument has no name when using delete)
    Ball& operator=(const Ball&) = delete; // Assignment operator (note that argument has no name when us)

    void move() {
        xpos += xvel;
        ypos += yvel;
    }

    void push(int dvx, int dvy) {
        xvel += dvx;
    }
}

```

```

    }           yvel += dvy;
};

}

```

4.2.2.3 Smart pointers

With a `unique_ptr` you can make sure that only one object, the owner of the pointer, can access the value behind the pointer. Consider this epic tale of a king and his magic sword:

```

#include <iostream>
#include <string>
#include <memory>

using std::string;

class Weapon {
private:
    string name;

public:
    Weapon(string n = "rusty sword") : name(n) {

    }

    ~Weapon(){
        std::cout << name << " destroyed" << std::endl;
    }

    string getName() {
        return name;
    }
};

class Hero {
private:
    string name;
    Weapon* weapon_ptr;

public:
    Hero(string n) : name(n), weapon_ptr(nullptr) {}

    ~Hero(){
        if(weapon_ptr != nullptr){
            std::cout << name << " now destroying " << weapon_ptr->getName() << std::endl;
            delete weapon_ptr;
        }
        std::cout << name << " destroyed." << std::endl;
    }

    void pickUpWeapon(Weapon* w) {
        if(weapon_ptr != nullptr) delete weapon_ptr;
        weapon_ptr = w;
    }

    string describe() {
        if(weapon_ptr != nullptr){
            return name + " now swings " + weapon_ptr->getName();
        } else {
            return name + " now swings his bare hands";
        }
    }
};

Weapon* blacksmithForge(string name){
    Weapon* wp = new Weapon(name);
    return wp;
}

int main() {
    Hero arthur = Hero("Arthur");
    Hero mordred = Hero("Mordred");
    Weapon* excalibur = blacksmithForge("Excalibur");

    std::cout << arthur.describe() << std::endl;
    arthur.pickUpWeapon(excalibur);
    std::cout << arthur.describe() << std::endl;
    mordred.pickUpWeapon(excalibur);
    std::cout << mordred.describe() << std::endl;
    std::cout << "Everything going out of scope." << std::endl;
}

```

```
|| return 0;
|| }
```

Note how Mordred has snatched away Arthurs sword. When mordred dies, he takes Excalibur with him into the grave. When Arthur dies, he tries to do the same, but has to find that the sword he thought he had was no longer there. We can fix this problem by using `unique_ptr`'s:

```
#include <iostream>
#include <string>
#include <memory>

using std::string;

class Weapon {
private:
    string name;

public:
    Weapon(string n = "rusty sword") : name(n) {}

    string getName() {
        return name;
    }
};

class Hero {
private:
    string name;
    std::unique_ptr<Weapon> weapon_ptr;

public:
    Hero(string n) : name(n), weapon_ptr(new Weapon(n + "'s bare hands")) {}

    void pickUpWeapon(std::unique_ptr<Weapon> w) {
        if(!w){
            std::cout << "Nothing to pick up." << std::endl;
            return;
        }
        weapon_ptr = std::move(w);
    }

    string describe() {
        return name + " now swings " + weapon_ptr->getName();
    }
};

std::unique_ptr<Weapon> blacksmithForge(string name){
    std::unique_ptr<Weapon> wp(new Weapon(name));
    return wp;
}

int main() {

    Hero arthur = Hero("Arthur");
    Hero mordred = Hero("Mordred");
    std::unique_ptr<Weapon> excalibur = blacksmithForge("Excalibur");

    std::cout << arthur.describe() << std::endl;
    arthur.pickUpWeapon(std::move(excalibur));
    std::cout << arthur.describe() << std::endl;

    std::cout << "Mordred now tries to snatch Excalibur." << std::endl;
    mordred.pickUpWeapon(std::move(excalibur));
    std::cout << mordred.describe() << std::endl;

    std::cout << "Everything going out of scope." << std::endl;
    return 0;
}
```

4.2.2.4 Ownership: unique, shared and weak pointers

Let us explain in a bit more detail what a smart pointer is. Really, it is just a class encapsulating a raw pointer. The idea is that by wrapping a pointer in a class, we can have the class object be allocated on the stack and have it free the memory behind the pointer when the object is destroyed. This way, you don't have to manually delete a pointer.

```
|| class SmartPointer{
```

```

private:
    T* ptr;
public:
    SmartPointer(T* p){
        ptr = p;
    }
    ~SmartPointer(){
        free(ptr);
    }
}

```

In this example, we allocate Excalibur on the heap. But because we're using smart pointers, we never have to clean up memory manually:

```

#include <iostream>
#include <memory>

using std::string;

class Weapon{
private:
    string name;
public:
    Weapon(string n) : name(n){}
    string getName(){
        return name;
    }
    ~Weapon(){
        std::cout << name << " is being destroyed" << std::endl;
    }
};

class Hero{
private:
    string name;
    std::unique_ptr<Weapon> weapon;
public:
    Hero(string n) : name(n), weapon(new Weapon("rusty sword")) {}
    void pickWeapon(std::unique_ptr<Weapon> w){
        std::cout << name << " picked up " << w->getName() << std::endl;
        weapon = std::move(w);
        std::cout << weapon->getName() << " is now on " << name << std::endl;
    }
    ~Hero(){
        std::cout << name << " is being destroyed" << std::endl;
    }
};

int main() {
    Hero arthur("Arthur");
    std::unique_ptr<Weapon> w(new Weapon("Excalibur"));
    arthur.pickWeapon(std::move(w));

    return 0;
}

```

Note how we used `move(w)` to overwrite `weapon`. What happens here is this: In the main-scope, the `w`'s internal pointer to Excalibur is replaced by a `nullptr`. Any further attempt of accessing `w` will cause a runtime-error. Inside the `pickWeapon` method, the old smartpointer to the rusty sword is destroyed, causing the sword itself to be freed. We couldn't have used `weapon = w;`; that would have caused an error (because `unique_ptr`'s copy-assignment method is deleted). This is deliberate: we don't want copies of a unique pointer to exist. We want a unique pointer to only ever exist in one place (in this case: first in the `main` function, later in the `arthur` instance), so that the memory behind the pointer can only get freed once (in this case: when `arthur` is destroyed).

Generally, we can proclaim the following rules:

- When you're given a reference, someone else will clean up the original.
- When you're given a unique pointer, you are now the sole owner of the smartpointer. If you move the pointer, it will go out of scope once you execute that move; if you don't, it will go out of scope once your method ends. Passing a unique pointer to a method that doesn't move the pointer means destroying the object! To pass a unique pointer to a method without it being destroyed or moved, don't pass the unique pointer, pass a `unique_ptr<T> const &`.
- When you're given a shared pointer, Shared pointers use reference counting just like the java garbage collector.

There are a lot more useful tips here and here.

Using ownership wisely Based on this, we can define different ownership strategies depending on whether an object is a value-object or an id-object.

4.2.3 Templates

Templates are like java generics.

4.2.3.1 Function template

For a function to accept a generic parameter, just prepend `template <class myType>` to the function definition.

```
#include <iostream>
#include <string>

using namespace std;

template <class myType>
myType GetMax (myType a, myType b) {
    return (a>b?a:b);
}

int main () {
    int i = 39;
    int j = 20;
    cout << "Max(i, j): " << GetMax(i, j) << endl;

    double f1 = 13.5;
    double f2 = 20.7;
    cout << "Max(f1, f2): " << GetMax(f1, f2) << endl;

    string s1 = "Hello";
    string s2 = "World";
    cout << "Max(s1, s2): " << GetMax(s1, s2) << endl;

    return 0;
}
```

4.2.3.2 Class template

Things work very similarly with classes:

```
#include <iostream>
using namespace std;

template <class T>
class MyPair {
private:
    T values [2];

public:
    MyPair (T first, T second) {
        values[0]=first;
        values[1]=second;
    }

    T getMax() {
        return (values[0] > values[1] ? values[0] : values[1]);
    }
};

int main() {
    MyPair<int> myobject(100, 75);
    cout << myobject.getMax();
    return 0;
}
```

One last note on notation: if we had defined class and implementation in separate files, we'd have to write:

```
// ---- class definition -----

template <class T>
class MyPair {
private:
    T values [2];
public:
    MyPair(T first, T second);
    T getMax();
```

```

};

// ---- class implementation ----

template <class T>
MyPair<T>::MyPair(T first, T second) {
    values[0] = first;
    values[1] = second;
}

template <class T>
MyPair<T>::getMax() {
    return (values[0] < values[1] ? values[1] : values[0]);
}

```

4.2.4 Threading

Compared to c's posix threads, c++ has somewhat simplified it's core threading library.

```

#include <iostream>
#include <thread>

void threadFunc(int i) {
    for(int j = 0; j < 100; j++){
        std::cout << "hi! this is operation " << j << " from thread nr " << i << std::endl;
    }
}

int main(){
    std::thread threads[10];

    for(int i = 0; i<10; i++){
        threads[i] = std::thread(threadFunc, i);
    }

    for(int i = 0; i<10; i++){
        threads[i].join();
    }

    return 0;
}

```

You can avoid some of the above problem using barriers in your code (`std::mutex`) which will let you synchronize the way a group of threads share a resource.

4.2.5 CMake

CMake input files are written in the “CMake Language” in source files named CMakeLists.txt or ending in a .cmake file name extension.

CMake Language source files in a project are organized into:

- *Directories* (CMakeLists.txt). When CMake processes a project source tree, the entry point is a source file called CMakeLists.txt in the top-level source directory. This file may contain the entire build specification or use the `add_subdirectory()` command to add subdirectories to the build. Each subdirectory added by the command must also contain a CMakeLists.txt file as the entry point to that directory. For each source directory whose CMakeLists.txt file is processed CMake generates a corresponding directory in the build tree to act as the default working and output directory.
- *Scripts* (script.cmake). An individual script.cmake source file may be processed in script mode by using the `cmake(1)` command-line tool with the -P option. Script mode simply runs the commands in the given CMake Language source file and does not generate a build system. It does not allow CMake commands that define build targets or actions.
- *Modules* (module.cmake). CMake Language code in either Directories or Scripts may use the `include()` command to load a module.cmake source file in the scope of the including context. See the `cmake-modules(7)` manual page for documentation of modules included with the CMake distribution. Project source trees may also provide their own modules and specify their location(s) in the `CMAKE_MODULE_PATH` variable.

4.2.5.1 Variables

Variables are the basic unit of storage in the CMake Language. Their values are always of string type, though some commands may interpret the strings as values of other types. The `set()` and `unset()` commands explicitly set or unset a variable, but other commands have semantics that modify variables as well. Variable names are case-sensitive and may consist of almost any text, but we recommend sticking to names consisting only of alphanumeric characters plus `_` and `-`.

Although all values in CMake are stored as strings, a string may be treated as a list in certain contexts, such as during evaluation of an Unquoted Argument. In such contexts, a string is divided into list elements by splitting on `,` characters not following an unequal number of `[` and `]` characters and not immediately preceded by a `\`. The sequence `\;` does not divide a value but is replaced by `,` in the resulting element.

A list of elements is represented as a string by concatenating the elements separated by `,`. For example, the `set()` command stores multiple values into the destination variable as a list:

```
|| set(srcs a.c b.c c.c) # sets "srcs" to "a.c;b.c;c.c"
```

Lists are meant for simple use cases such as a list of source files and should not be used for complex data processing tasks. Most commands that construct lists do not escape `,` characters in list elements, thus flattening nested lists:

```
|| set(x a "b;c") # sets "x" to "a;b;c", not "a;b\;c"
```

A variable reference has the form `${variable_name}` and is evaluated inside a Quoted Argument or an Unquoted Argument. A variable reference is replaced by the value of the variable, or by the empty string if the variable is not set. Variable references can nest and are evaluated from the inside out, e.g. `${outer_${inner_variable}_variable}`.

An environment variable reference has the form `$ENV{VAR}` and is evaluated in the same contexts as a normal variable reference.

4.2.5.2 Logic

- Conditional Blocks: The `if()`/`elseif()`/`else()`/`endif()` commands delimit code blocks to be executed conditionally.
- Loops: The `foreach()`/`endforeach()` and `while()`/`endwhile()` commands delimit code blocks to be executed in a loop. The `break()` command may be used inside such blocks to terminate the loop early.
- Command Definitions: The `macro()`/`endmacro()`, and `function()`/`endfunction()` commands delimit code blocks to be recorded for later invocation as commands.

4.2.5.3 Command arguments

- bracket argument:

```
|| message([=[
This is the first line in a bracket argument with bracket length 1.
No \-escape sequences or ${variable} references are evaluated.
This is always one argument even though it contains a ; character.
The text does not end on a closing bracket of length 0 like ]].
It does end in a closing bracket of length 1.
]=] )
```

- quoted argument:

```
|| message("This is a quoted argument containing multiple lines.
This is always one argument even though it contains a ; character.
Both \\-escape sequences and ${variable} references are evaluated.
The text does not end on an escaped double-quote like \".
It does end in an unescaped double quote.
")
```

- unquoted argument:

```
|| foreach(arg
  NoSpace
  Escaped\ Space
  This;Divides;Into;Five;Arguments
  Escaped\;Semicolon
)
  message("${arg}")
endforeach()
```

4.2.5.4 Important predefined variables

There is a list of CMake's global variables. You should know them.

- `CMAKE_BINARY_DIR`: if you are building in-source, this is the same as `CMAKE_SOURCE_DIR`, otherwise this is the top level directory of your build tree
- `CMAKE_SOURCE_DIR`: this is the directory, from which cmake was started, i.e. the top level source directory
- `EXECUTABLE_OUTPUT_PATH`: set this variable to specify a common place where CMake should put all executable files (instead of `CMAKE_CURRENT_BINARY_DIR`):
|| `SET(EXECUTABLE_OUTPUT_PATH ${PROJECT_BINARY_DIR}/bin)`
- `LIBRARY_OUTPUT_PATH`: set this variable to specify a common place where CMake should put all libraries (instead of `CMAKE_CURRENT_BINARY_DIR`):
|| `SET(LIBRARY_OUTPUT_PATH ${PROJECT_BINARY_DIR}/lib)`
- `PROJECT_NAME`: the name of the project set by `PROJECT()` command.
- `PROJECT_SOURCE_DIR`: contains the full path to the root of your project source directory, i.e. to the nearest directory where CMakeLists.txt contains the `PROJECT()` command. Now, you have to compile the test.cpp. The way to do this task is too simple. Add the following line into your CMakeLists.txt:
|| `add_executable(hello ${PROJECT_SOURCE_DIR}/test.cpp)`

4.3 Java

We tend to use java when we want to create enterprise level software. The whole development-experience is way different from c. While in c you control every detail of the program, a java-app has you reuse as many standardized components as possible. There are layers upon layers of abstraction. Usually, there is some way to build on top of a framework, leaving you to only fill out a few xml-configurations and implement a few skeleton-beans with loads of annotations.

4.3.1 General

4.3.1.1 Basic theory

4.3.1.2 Environment variables

First things first, to find out if you're using a 64 bit version of the jdk, just execute

```
|| java -d64 -version
```

This will throw an error if your version is not made for 64 bit.

There are a few environment variables that you should know about:

- PATH: Sys looks for exes here. Your JAVA_HOME/bin should be part of PATH
- CLASSPATH: Java looks for code here
- JAVA_HOME: location jdk (example: /usr/lib/jvm/java-7-openjdk-amd64/bin)
- JRE_HOME: location jre (example: /usr/lib/jvm/java-7-openjdk-amd64/jre/bin)

The most important setting for the JVM is the heap size: how much memory will we allocate to the java-process?

There are two settings:

- `-Xms<size>` - Set initial Java heap size
- `-Xmx<size>` - Set maximum Java heap size

These are either set in or used directly when invoking java with `java -Xms512m -Xmx1024m JavaApp`.

You can display the default settings like this:

```
$ java -XX:+PrintFlagsFinal -version | grep -iE 'HeapSize|PermSize|ThreadStackSize'
  uintx InitialHeapSize          := 64781184      {product}
  uintx MaxHeapSize             := 1038090240    {product}
  uintx PermSize                = 21757952      {pd product}
  uintx MaxPermSize             = 174063616     {pd product}
  intx ThreadStackSize          = 1024          {pd product}
java version "1.7.0_51"
OpenJDK Runtime Environment (IcedTea 2.4.4) (7u51-2.4.4-0ubuntu0.13.10.1)
OpenJDK 64-Bit Server VM (build 24.45-b08, mixed mode)
```

Here are some suggested values:

- Heap = `-Xms512m -Xmx1024m`
- PermGen = `-XX:PermSize=64m -XX:MaxPermSize=128m`
- Thread = `-Xss512k`

4.3.1.3 Versions

Unfortunately, with java11 and its module system and the deprecation of former JavaEE packages, java has introduced its own breaking change á lá python2/3. Often, you will find that older programs only run in java8 or lower. You can see which java versions you have installed and choose the appropriate one by

```
|| update-alternatives --config java
```

Usually, java programs will have their own, linux-specific startup-script, like `/etc/init.d/openfire`. Here, these scripts will pick java based on `$JAVA_HOME`. To set this variable for your own user only, append

```
|| export JAVA_HOME="/usr/lib/jvm/java-8-openjdk-amd64/jre"
|| export PATH=$JAVA_HOME/bin:$PATH
```

to `~/.bashrc`, to set it for every user, append to `/etc/profiles` for login-shells and `/etc/bash.bashrc` for non-login-shells. Even better, put it in `/etc/environment` to cover both cases at once.

4.3.2 Maven

Maven is a build- and dependency-resolution tool. Building eg. webprojects with maven is encouraged, also because maven takes away buildsteps from your IDE. Relying on maven instead of an IDE to do your building then makes your builds portable between different development environments.

Maven can initialise a project structure for you from the command-line:

```
mvn archetype:generate -DarchetypeArtifactId=maven-archetype-quickstart
```

It will then ask you all further necessary information.

cli syntax The syntax generally consists of mvn [lifecycle]:[phase]:[args]

```
|| mvn default:package
|| mvn [lifecycle]:[phase]
|| mvn default:package:help
```

build lifecycles Maven has a quite hierarchical structure to it.

- lifecycles: default (compiling et al), clean (remove artifacts), site (create documentation). Plugins might add further lifecycles.
 - phases : default build lifecycle consists of several phases: validate, compile, test, package, verify, install, deploy
 - * goals: each phase executes one or more goals. these are the actual code instructions.
 - profiles : here we can adjust some environment settings. should never really be necessary.
- plugin: a set of extra goals for maven. plugins might even define whole new lifecycles like 'mvn jetty:run', or hook into the build cycle to generate code from xml like cxf does.
 - mojo : maven pojo. a class representing an executable goal
 - configuration of plugins. executions : specifies how the mojo should be executed; ie in which phase of the maven build-lifecycle

Directory structure Everything that you put under *src/main/resources* during development will be put under the root-directory on compilation. Same thing holds for the contents of *src/test/resources*: they will be on the root path during execution of tests. So, when using relative paths, *src/*resources/* must be omitted.

Custom archetypes

Custom plugins

4.3.3 Threading

Java was designed to make threading easy (though, contrary to clojure, it was not designed to make threading *safe*).

Basic threading with shared data is the simplest, but somewhat dangerous method.

```
public class Main {
    public static void main(String[] args) throws InterruptedException {
        LinkedList<Integer> sharedData = new LinkedList<Integer>();
        Producer p = new Producer(sharedData);
        Consumer c = new Consumer(sharedData);
        p.start();
        Thread.sleep(1000);
        c.start();
    }
}
```

```

public class Producer extends Thread {
    private LinkedList<Integer> sharedData;

    public Producer(LinkedList<Integer> sharedData) {
        this.sharedData = sharedData;
    }

    @Override
    public void run() {
        Random r = new Random();
        while(true) {
            try {
                Integer i = r.nextInt(5);
                Thread.sleep(100 * i);
                sharedData.add(i);
                System.out.println("Producer just added " + i + " to the shared data.");
                System.out.println("Data now looks like this: " + sharedData.toString());
            } catch (InterruptedException e) {
                e.printStackTrace();
            }
        }
    }
}

public class Consumer extends Thread {
    private LinkedList<Integer> sharedData;

    public Consumer(LinkedList<Integer> sharedData) {
        this.sharedData = sharedData;
    }

    @Override
    public void run() {
        Random r = new Random();
        while(true) {
            try {
                Thread.sleep(100 * r.nextInt(5));
                i = sharedData.removeFirst();
                System.out.println("Consumer just took " + i + " from the shared data.");
                System.out.println("Data now looks like this: " + sharedData.toString());
            } catch (InterruptedException e) {
                e.printStackTrace();
            }
        }
    }
}

```

This results in an output like this:

```

Producer just added 3 to the shared data.
Data now looks like this: [0, 0, 3]
Consumer just took 0 from the shared data.
Data now looks like this: [0, 3]
Consumer just took 0 from the shared data.
Data now looks like this: [3]
Consumer just took 3 from the shared data.
Data now looks like this: []
Exception in thread "Thread-1" java.util.NoSuchElementException
    at java.util.LinkedList.removeFirst(LinkedList.java:270)
    at threadingStuff.Consumer.run(Consumer.java:20)

```

Notice how at some point the consumer overtook the producer, leading to a `NoSuchElementException`. We can use signaling to handle this kind of situation:

- `wait()`
- `notify()`

Alternatively, a blocking queue would automatically handle such a case by just blocking the consumer until data is available again.

Using a blocking queue makes sense when

- you need protection from `NoSuchElementException`.
- it is ok for the consumer to sometimes be blocked while waiting for data to arrive.

```

public class Main {
    public static void main(String[] args) throws InterruptedException {
        ArrayBlockingQueue<Integer> sharedData = new ArrayBlockingQueue<Integer>(100);
        Producer p = new Producer(sharedData);
        Consumer c = new Consumer(sharedData);
        p.start();
        Thread.sleep(1000);
        c.start();
    }
}

public class Producer extends Thread {
    private ArrayBlockingQueue<Integer> sharedData;
    public Producer(ArrayBlockingQueue<Integer> sharedData) {
        this.sharedData = sharedData;
    }

    @Override
    public void run() {
        Random r = new Random();
        while(true) {
            try {
                Integer i = r.nextInt(5);
                Thread.sleep(100 * i);
                sharedData.put(i);
                System.out.println("Producer just added " + " to the shared data.");
                System.out.println("Data now looks like this: " + sharedData.toString());
            } catch (InterruptedException e) {
                e.printStackTrace();
            }
        }
    }
}

public class Consumer extends Thread {
    private ArrayBlockingQueue<Integer> sharedData;
    public Consumer(ArrayBlockingQueue<Integer> sharedData) {
        this.sharedData = sharedData;
    }

    @Override
    public void run() {
        Random r = new Random();
        Integer i;
        while(true) {
            i = null;
            try {
                Thread.sleep(100 * r.nextInt(5));
                i = sharedData.take();
                System.out.println("Consumer just took " + i + " from the shared data.");
                System.out.println("Data now looks like this: " + sharedData.toString());
            } catch (InterruptedException e) {
                e.printStackTrace();
            }
        }
    }
}

```

This results in an output like this:

```

Producer just added 3 to the shared data.
Data now looks like this: [0, 0, 3]
Consumer just took 0 from the shared data.
Data now looks like this: [0, 3]
Consumer just took 0 from the shared data.
Data now looks like this: [3]
Consumer just took 3 from the shared data.
Data now looks like this: []      // ----- no exception thrown here! Consumer just has to wait.
Producer just added 3 to the shared data.
Data now looks like this: [3]
Producer just added 4 to the shared data.
Data now looks like this: [3, 4]

```

Using a pipe makes only sense under very specific conditions. A pipe is optimized for serialized in- and output, so you'd first have to serialize and then deserialize anything that has to pass through the pipe. Also, a pipe can always only exist between exactly two threads, no more.

Using an executor service makes sense when ...

Calling another thread's methods is actually possible without much ado:

```

class Main {
    public static void main() {
        RT rt = new RT();
        rt.start();
        rt.updateV(3);
    }
}

class RT extends Thread {
    private int v;

    public RT() {
        v = 1;
    }

    public void run() {
        while(true){
            Thread ct = Thread.currentThread();
            String name = ct.getName();
            System.out.println("Method run() being executed on " + name);
            System.out.println("v now is: " + v);
        }
    }

    public void updateV(int newV) {
        Thread ct = Thread.currentThread();
        String name = ct.getName();
        System.out.println("Method updateV() being executed on " + name);
        v = newV;
    }
}

```

A thread's public methods may be called from anywhere. So the main thread can call `accessibleMethod()` without problems. However, the execution of that public method will than take place on the calling thread - in that case, the main thread. You can verify this by sparkling your code with a few `Thread.currentThread().getName()` calls. But the called method may still manipulate a variable that is private to the called thread; so the main thread may manipulate `v`. This is possible because threads share the same memory base.

4.3.4 Functional programming

Java 8 exposes a few new Classes that can be used as functions.

- `Function<Integer, Integer> add3 = a -> a + 3;`
- `BiFunction<Integer, Integer, Integer> add = (a, b) -> a + b;`
- `Predicate<Integer> isEven = a -> a%2 == 0;` A predicate always returns a boolean.
- `Consumer<Book> storeAway = b -> b.save();` A consumer returns void.

```

public static void main() {
    BiFunction<Integer, Integer, Integer> add = (a, b) -> a + b;
    BiFunction<Integer, Integer, Integer> sub = (a, b) -> a - b;
    Integer result1 = compute(add, 3, 4);
    Integer result2 = compute(sub, 3, 4);
}

public static Integer compute(BiFunction<Integer, Integer, Integer> function, Integer a, Integer b) {
    return function.apply(a, b);
}

```

These functions could also have been passed anonymously:

```

public static void main() {
    Integer result1 = compute((a, b) -> a + b, 3, 4);
    Integer result2 = compute((a, b) -> a - b, 3, 4);
}

public static Integer compute(BiFunction<Integer, Integer, Integer> function, Integer a, Integer b) {
    return function.apply(a, b);
}

```

4.3.5 Annotations

We mostly use annotations to let others read out meta-information about our pojos. Consider the following example.

In this code, we define what kind of metadata we want to create.

```

import java.lang.annotation.ElementType;
import java.lang.annotation.Retention;
import java.lang.annotation.RetentionPolicy;
import java.lang.annotation.Target;

@Retention(RetentionPolicy.RUNTIME)
public @interface TableRecord {
    String table();
    String database();
}

@Target(ElementType.FIELD)
@Retention(RetentionPolicy.RUNTIME)
public @interface TableField {
    String name();
}

```

Then we decorate a pojo with the new annotations.

```

import org.apache.camel.dataformat.bindy.annotation.CsvRecord;
import org.apache.camel.dataformat.bindy.annotation.DataField;

@CsvRecord(separator = ";")
@TableRecord(table="iw_mn", database="gkddat")
public class Messnetz {

    @DataField(pos = 1)
    @TableField(name="Messnetznummer")
    private int Messnetznummer;

    @DataField(pos = 2)
    @TableField(name="Messnetzname")
    private String Messnetzname;

    public Messnetz() {}

    public Messnetz(int msnr, String mnname) {
        Messnetznummer = msnr;
        Messnetzname = mnname;
    }

    ... getters and setters ...
}

```

And finally we read out the annotations to create an sql-query:

```

public class SqlConverter implements DataFormat {

    /**
     * This is to convert any object into a "INSERT" sql statement.
     * Makes use of the @TableRecord and @TableField annotations to figure out
     * what parts of the pojo need to be persisted.
     */
    @Override
    public void marshal(Exchange ex, Object graph, OutputStream stream) throws Exception {

        Object pojo = ex.getIn().getBody();
        TableRecord tr = pojo.getClass().getAnnotation(TableRecord.class);
        Field[] fields = pojo.getClass().getDeclaredFields();

        String database = tr.database();
        String table = tr.table();

        ArrayList<String> fieldNames = new ArrayList<String>();
        ArrayList<String> fieldValues = new ArrayList<String>();
        for(Field f : fields) {
            TableField tf = f.getAnnotation(TableField.class);

```

```

    if(tf == null) continue;

    String fieldName = tf.name();
    fieldNames.add(fieldName);

    f.setAccessible(true);
    Object objFieldValue = f.get(pojo);
    String fieldValue = objFieldValue.toString();
    if(f.getType() == String.class) {
        fieldValue = "}" + fieldValue + "}";
    }
    fieldValues.add(fieldValue);
}
String fieldList = StringUtils.join(fieldNames);
String valueList = StringUtils.join(fieldValues);

String sql = "insert into " + database + "." + table + " (" + fieldList + ") values (" + valueList + ");";
stream.write(sql.getBytes());
}
}

```

4.3.6 Transforming textformats: marshaling

This is known as marshaling. The general idea is always a simple two-step proces. To convert one format into another,

- Unmasrshal the file (that uses format 1) into a pojo,
- then marshal the pojo into another file (that uses format 2).

4.3.6.1 CSV to POJO and back: Bindy

Important annotations in the pojo are:

4.3.6.2 JSON to POJO and back: Jackson

Important annotations in the pojo are:

4.3.6.3 XML to POJO and back: JAXB

Important annotations in the pojo are:

4.3.7 Reading from documents

You don't always want to load a whole file into a pojo. Sometimes it's enough to just extract one simple information out of the file.

4.3.7.1 Extracting from xml: XPath

4.3.7.2 Extracting from json: JsonPath

4.3.8 Databases

Java allows for simple query-based database access (JDBC) and for automated marshaling of pojos into tables (JPA and hibernate).

4.3.8.1 First level db-access: JDBC

Every kind of database (MySQL, JavaDB, filesystembased, ...) can be accessed by its own *driver*. Because those drivers are vendor-specific, they are abstracted away using a *driver manager*.

```

Connection conn = DriverManager.getConnection("jdbc:mysql://10.112.70.133", "michael", "meinpw");
Statement stmt = conn.createStatement();
ResultSet res = stmt.executeQuery("select * from amoado.meldung");
while(res.next()){
    ...
}

```

However, nowadays it is more common to set up a DataSource instead of using the driver manager. The datasource will take care of managing a pool of connections, whereas with a drivermanager you would have to handle a connection-pool manually. Note how the concrete implementation of the Datasource is vendor specific. It therefore makes sense to create a factory.

```
MysqlDataSource mysqlDS = new MysqlDataSource();
mysqlDS.setURL("jdbc:mysql://localhost:3306/UserDB");
mysqlDS.setUser("michael");
mysqlDS.setPassword("meinpw");

...
DataSource ds = myFactory.getDs("mysql");
Connection con = ds.getConnection();
Statement stmt = con.createStatement();
ResultSet rs = stmt.executeQuery("select empid, name from Employee");
while(rs.next()){
    ...
}
```

4.3.8.2 Second level db-access: JPA and Hibernate

The JPA is a specification implemented by Hibernate (say: hibernate is a "jpa provider"), among others. Its function is to persist pojos in a relational database - it's an ORM. Hibernate will accept pojos and scan them for annotations explaining how the pojo should be persisted (if you don't want to use annotations, you can instead create a mapping.xml).

Hibernate Hibernate will generate automatic sql for you. It will also create tables where needed. Therefore, it makes no sense to use it on an existing database where the structure mustn't change. Under such conditions, you're way better off using JDBC (or even better spring jdbc-template).

```
<dependency>
    <groupId>org.hibernate</groupId>
    <artifactId>hibernate-core</artifactId>
    <version>4.3.5.Final</version>
</dependency>
<!-- Hibernate 4 uses Jboss logging, but older versions slf4j for logging -->
<dependency>
    <groupId>org.slf4j</groupId>
    <artifactId>slf4j-simple</artifactId>
    <version>1.7.5</version>
</dependency>
<dependency>
    <groupId>mysql</groupId>
    <artifactId>mysql-connector-java</artifactId>
    <version>5.0.5</version>
</dependency>

package hibernate;

import java.util.Date;

import javax.persistence.Entity;
import javax.persistence.GeneratedValue;
import javax.persistence.GenerationType;
import javax.persistence.Id;
import javax.persistence.Table;
import javax.persistence.UniqueConstraint;

@Entity
@Table(name="Employees", uniqueConstraints= {@UniqueConstraint(columnNames = { "id" })})
public class Employee {

    @Id
    @GeneratedValue(strategy=GenerationType.IDENTITY)
    private int id;
    private String name;
    private String role;
    private Date insertTime;

    ... bunch of getters and setters ...
}

<?xml version="1.0" encoding="UTF-8"?>
<!DOCTYPE hibernate-configuration PUBLIC
        "-//Hibernate/Hibernate Configuration DTD 3.0//EN"

```

```

    "http://hibernate.org/dtd/hibernate-configuration-3.0.dtd">

<hibernate-configuration>
    <session-factory>
        <property name="hibernate.connection.driver_class">com.mysql.jdbc.Driver</property>
        <property name="hibernate.connection.url">jdbc:mysql://localhost/testdb</property>
        <property name="hibernate.current_session_context_class">thread</property>
        <mapping class="hibernate.Employee"/>
    </session-factory>
</hibernate-configuration>

public class AppMain {

    public static void main(String[] args) {
        Configuration conf = new Configuration();
        conf.configure("hibernate-annotation.cfg.xml");
        ServiceRegistry sr = (ServiceRegistry) new StandardServiceRegistryBuilder().applySettings(conf.getProperties());
        SessionFactory sf = conf.buildSessionFactory(sr);
        Session s = sf.getCurrentSession();

        Employee e = new Employee();
        e.setName("Michael");
        e.setRole("programmer");
        e.setInsertTime(new Date());

        s.beginTransaction();
        s.save(e);
        s.getTransaction().commit();

        System.out.println("Employee id: " + e.getId());

        HibernateUtil.getSessionFactory().close();
    }
}

```

JPA Hibernate has later been generalised to JPA. JPA can use hibernate as backend, but might also use eclipselink or any other ORM. Unfortunately, JPA comes with a few changes compared to the hibernate syntax, not all of them for the better.

See here for an example-application:

```

EntityManagerFactory emp = Persistence.createEntityManagerFactory("employeeDB");
// This is a factory-factory. Talk about over-engineering :)
// employeeDB is the name of a persistance unit in your persistence.xml.
// Note that the position of this xml is defined nowhere - it MUST be placed in classroot/META-INF/persistence.xml
// and have exactly this name.

EntityManager em = emf.createEntityManager();
// The entity-manager is the most important class and basically your CRUD-Interface to the database.

EntityTransaction tx = em.getTransaction();

tx.begin();
Employee empl = new Employee("Homer", "Simpson", "97G");
em.persist(empl);
tx.commit();
em.close();

```

This class makes use of a `META-INF/persistence.xml` file. If this file is not found under this exact name in that specific location, a "no persistence provider found" exception is thrown.

```

<persistence
    xmlns="http://java.sun.com/xml/ns/persistence"
    xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
    version="1.0">

    <!-- Value given to createEntityManagerFactory -->
    <persistence-unit name="employeeDB" transaction-type="RESOURCE_LOCAL">

        <!-- All beans that you want to persist must be mentioned here -->
        <class>org.langbein.michael.Employee</class>

        <!-- Database-driver, credentials, and JPA-settings (differ per vendor, but generally logging, create tables etc)-->
        <properties>
            <property name="openjpa.ConnectionDriverName" value="org.hsqldb.jdbcDriver" />
            <property name="openjpa.ConnectionURL" value="jdbc:hsqldb:mem:order" />
            <property name="openjpa.ConnectionUserName" value="sa" />
            <property name="openjpa.ConnectionPassword" value="" />
            <property name="openjpa.jdbc.SynchronizeMappings" value="buildSchema" />
        </properties>

```

```

||     </persistence-unit>
|| </persistence>
```

This is not a very convenient setup. In the spring appendix we'll describe using the spring data jpa starter as an alternative way of using jpa.

4.3.9 Swing

Swing is a event based GUI-Framework. Data is wrapped in components. Components mainly do two things: tell the swing environment how they are to be rendered, and babble with other components or the user through listening to and sending of events.

Most components like buttons, labels etc. are trivial. But the interesting ones are JList, JTable and JTree. These all fire custom events. You can set a custom event-listener for these events. You can also pass a model to them to create a mapping between your actual data and the JList, and you can pass a renderer to them to instruct them how your data is supposed to be displayed.

4.3.9.1 Lists

Lists will be the first component we look at in more detail to illustrate how Swing handles subcomponents and subcomponent-models. As stated before, there are three ways in which we can customise a default list: add custom event handlers, add a model to map our data to the list, and add a renderer to customise the display of the elements.

Adding a custom event-handler for JListEvents This is as simple as passing in a implementation of a single method interface.

Defining a model With almost any Swing component, you can separate the underlying data structure containing the data from the GUI control that displays the data. It is rare to want to do that with simple controls like buttons, but with lists, tables, and trees, it is inconvenient to have to copy data out of an existing hash table, array, linked list, or other structure into the format needed by the GUI control. So Swing lets you use your existing structure directly. You simply implement an interface that tells Swing how to access the data, and then use the data directly in the component.

So you could add simple strings to a list, that JList can handle without any kind of a model:

```

|| String[] options = {"Option1", "Option2", "Option3"};
|| JList optionList = new JList<String>(options);
```

Instead of predetermining the data structure that holds list elements, Swing lets you use your own data structure as long as you tell it how to get the data into or out of that structure. You could either implement the ListModel-interface or extend the AbstractListModel-class.

```

|| Location[] locations = {
||     new Location("Berlin", "Germany", "nice"),
||     new Location("Ghent", "Belgium", "awesome"),
||     new Location("Gera", "Germany", "sucking balls")
|| };
|| JList list = new JList<Location>(new LocationListModel(locations));
|
| public class LocationListModel extends AbstractListModel<Location>{
| private Location[] locations;
| public LocationListModel(Location[] locations) {
|     this.locations = locations;
| }
| public Location getElementAt(int index) {
|     return locations[index];
| }
| public int getSize() {
|     return locations.length;
| }
|
| public class Location {
|     private String name;
|     private String country;
|     private String description;
|     public Location(String name, String country, String description) {
|         this.name = name;
|         this.country = country;
|         this.description = description;
|     }
| }
```

```

        }
        public String toString() {
            return (name + " in " + country + " is " + description);
        }
    }
}

```

Defining a renderer Swing has a few simple defaults for displaying values in lists, trees, and tables. In a JList, values that are Icons are drawn directly, while other objects are converted to strings via their `toString` method, then displayed via a JLabel. However, Swing also lets you define arbitrary mappings between values and display components, yielding whatever visual representation you want for values of specified types. This is done by building a "cell renderer" that takes the containing component, the value of the entry, a few state parameters that say things like whether or not the value is currently selected, and then returns a Component that is used to draw the value.

Normally Swing uses a JLabel, either directly from the String representation of the entry or directly from the value if the value is an Icon. But in this example, which continues the previous one, I want to include an image of the flag of the country for each JavaLocation. So I define a class that implements the `ListCellRenderer` interface, with a method `getListCellRendererComponent` that constructs the Component of interest given the list entry. I then associate that renderer with the JList via the JList's `setCellRenderer` method.

```

public class AppMain {
    public static void main(String[] args) {
        Location[] locations = {
            new Location("Berlin", "Germany", "nice"),
            new Location("Ghent", "Belgium", "awesome"),
            new Location("Gera", "Germany", "sucking balls")
        };
        JList list = new JList<Location>(new LocationListModel(locations));
        list.setCellRenderer(new MyCellRenderer());
        ...
    }
}

public class MyCellRenderer extends DefaultListCellRenderer {
    public Component getListCellRendererComponent(JList list, Object value, int index, boolean isSelected, boolean cellHasFocus) {
        JLabel label = (JLabel) super.getListCellRendererComponent(list, value, index, isSelected, cellHasFocus);
        Component out = label;
        if(isSelected) {
            JButton but = new JButton();
            but.setText(label.getText());
            out = but;
        }
        return out;
    }
}

```

4.3.9.2 Trees

Trees are very similar to lists. Here, too, we can define custom handlers, add models and add renderers.

Adding custom handlers for JTreeEvents

Adding a model A JTree uses a TreeModel to get its data. As with JList, you can replace the model altogether, specifying how to extract data from the custom model. In the case of JTree, however, the default TreeModel is very useful as it is. It is then most often better to write a custom TreeNode (by extending `DefaultMutableTreeNode`) and notify your datastructure of changes through the user by creating a `TreeModelListener`.

```

class AppMain {
    public static void main(String[] args) {
        Place e = new Place(null, "Europe");
        Place g = new Place(e, "Germany");
        Place m = new Place(g, "Munich");
        Place i = new Place(e, "Italy");
        Place r = new Place(i, "Rome");
        Place b = new Place(g, "Berlin");

        MyTreeCell cont = new MyTreeCell(e);
        cont.add(new MyTreeCell(g));
        cont.add(new MyTreeCell(i));

        JTree tree = new JTree(cont);
    }
}

```

```

||| public class MyTreeCell extends DefaultMutableTreeNode {
|||     public MyTreeCell(Place p) {
|||         super(p.getName());
|||     }
|||

```

Adding a renderer

Renderers can be changed just like you can in JLists.

```

||| public abstract class MyCellRenderer extends DefaultTreeCellRenderer {
|||     public Component getTreeCellRendererComponent(JTree tree, Object value, boolean sel, boolean expanded, boolean
|||         return super.getTreeCellRendererComponent(tree, value, sel, expanded, leaf, row, hasFocus);
|||     }
|||

```

4.3.9.3 Writing custom components

Extending JComponent Everything in Swing is a component. A component does the following things:

- contain data
- listen to user events
- fire own events in response
- provide a way for API-users to define event-listeners

Concretely, a implementation might look like this:

- Extend JComponent
- In the constructor, pass your actual data.
- Also in the constructor, add a few of these:

```

|||     addMouseListener(new MouseAdapter() {
|||         public void mousePressed(MouseEvent e) { myMouseEventHandler(e); }
|||     });
|||

```

- implement your myXEventHandler(e) methods.
- write a method fireEvent(), that goes through all the Listeners in the listenerList and allows them to react to a MyCompEvent.
- write a method addListener(Listener l), that accepts a listenerfunction (single method interface).
- Overwrite the method paintComponent

4.3.10 Servlets and JSP

Contrary to CGI programs, a servlet only is started once and then sits there running waiting for requests to arrive via the servlet-container (such as tomcat, jetty or undertow). For every client request, the servlet spawns a new thread (not a whole process like in cgi).

For a servlet to run in a servlet-container, it must simply extend the class HttpServlet. Also, the servlet needs a web.xml file. That file tells the container which requests should go to this servlet, where the base dir is, which class extended HttpServlet.

Place the following in an eclipse "dynamic-web-project"'s src-folder:

```

package meinServlet;

import java.io.IOException;
import javax.servlet.ServletException;
import javax.servlet.annotation.WebServlet;
import javax.servlet.http.HttpServlet;
import javax.servlet.http.HttpServletRequest;
import javax.servlet.http.HttpServletResponse;

public class MainServlet extends HttpServlet {
    private static final long serialVersionUID = 1L;
}

```

```

public MainServlet() {
    super();
}

protected void doGet(HttpServletRequest request, HttpServletResponse response) throws ServletException, IOException
{
    response.getWriter().append("Served at: ").append(request.getContextPath());
}

protected void doPost(HttpServletRequest request, HttpServletResponse response) throws ServletException, IOException
{
    doGet(request, response);
}
}

```

And put this DD in the projects WebContent/WEB-INF folder:

```

<?xml version="1.0" encoding="UTF-8"?>
<web-app xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
    xmlns="http://xmlns.jcp.org/xml/ns/javaee"
    xsi:schemaLocation="http://xmlns.jcp.org/xml/ns/javaee http://xmlns.jcp.org/xml/ns/javaee/web-app_3_1.xsd"
    id="WebApp_ID" version="3.1">

    <servlet>
        <servlet-name>Main</servlet-name>
        <servlet-class>meinServlet.MainServlet</servlet-class>
    </servlet>

    <servlet-mapping>
        <servlet-name>Main</servlet-name>
        <url-pattern>/hallo</url-pattern>
    </servlet-mapping>
</web-app>

```

With these two files, tomcat will know how to handle the servlet. Now if you go to <http://localhost:8080/meinServlet/hallo> should see the site.

4.3.10.1 Building with maven

- `mvn archetype:generate` with maven-archetype-webapp
- add folder `src/main/java`
- add package `org.langbein.michael.qa`
- add dependencies to pom, especially `http://maven.apache.org/jar`
- add servlet in `org.langbein.michael.qa`
- add `web.xml` in `src/main/webapp/WEB-INF`
- `mvn compile war:war`
- copy new war into `<tomcat>/webapps`
- restart tomcat
- visit `localhost:8080/qa/`

This is a ridiculously complex setup, especially since you'll have to repeat the last four steps for every iteration. That's why during development you should use the jetty plugin:

- Add the jetty-plugin to your pom:

```

<build>
    <finalName>qa3</finalName>
    <plugins>
        <plugin>
            <groupId>org.eclipse.jetty</groupId>
            <artifactId>jetty-maven-plugin</artifactId>
            <version>9.4.7.v20170914</version>
            <configuration>
                <scanIntervalSeconds>10</scanIntervalSeconds>
                <connectors>
                    <connector implementation="org.mortbay.jetty.nio.SelectChannelConnector">
                        <port>8080</port>
                        <maxIdleTime>60000</maxIdleTime>

```

```

        </connector>
    </connectors>
</configuration>
</plugin>
</plugins>
</build>

```

- run jetty with `mvn jetty:run`

4.3.10.2 Jsp and forms

It's really easy to template html and forms. Consider this servlet:

```

package ct;

public class QAController extends HttpServlet{
    private static final long serialVersionUID = 1L;

    protected void doGet(HttpServletRequest request, HttpServletResponse response) throws ServletException, IOException {
        request.getRequestDispatcher("/WEB-INF/qa.jsp").forward(request, response);
    }

    protected void doPost(HttpServletRequest request, HttpServletResponse response) throws ServletException, IOException {
        String answer = request.getParameter("answer");
        System.out.println("Got the answer " + answer);
        request.setAttribute("lastAnswer", answer);
        doGet(request, response);
    }
}

```

When a GET-request arrives at the server, it forwards this request to a jsp-page located under `/WEB-INF/qa.jsp`. That page looks like this:

```

<%@ page
language="java"
contentType="text/html; charset=UTF-8"
pageEncoding="UTF-8" %>
<!DOCTYPE html PUBLIC "-//W3C//DTD HTML 4.01 Transitional//EN" "http://www.w3.org/TR/html4/loose.dtd">
<html>

    <head>
        <meta http-equiv="Content-Type" content="text/html; charset=UTF-8">
        <title>Welcome to Michael's Q & A site!</title>
    </head>

    <body>
        <form name="qaForm" method="post" action="qa">
            Your last answer was: ${lastAnswer} <br />
            Question: ... blah ... <br />
            Answer: <input type="text" name="answer" /> <br />
            <input type="submit" value="Submit your answer" />
        </form>
        This page was generated at <%= new java.util.Date() %>
    </body>
</html>

```

Note three important aspects about this form:

- It has a POST-method and an input-element of the "submit" type
- The post is mapped to a relative url named "qa". This is the url of the servlet (could also be another servlet)
- The file uses a variable named `lastAnswer`. This variable has been set during the last request by the `doPost`-method.

4.3.10.3 Servlet context listener

Especially when our application depends on a database to exist on the server, ... Integrating a database requires us to do three things:

- add the connector jar to the apps libraries
- add the db credentials to the web.xml using `<context-param>`
- add a context listener, that has the servlet read out the database credentials on initialisation.

4.3.10.4 Servlet lifecycle

- Container creates *one* instance of the servlet.
- Container calls the `init()` method (*once*).
- From now on, every request is passed to that one instance's `doGet()` and `doPost()` methods.

4.3.11 Webservices

There are basically two kinds of webservices. Both build up on top of http.

- REST: a simple service using nothing but http's get, post etc. The details are defined in the JAX-RS specification. A concrete implementation of JAX-RS is called Jersey². A rest-service is described to the consumer using WADL.
- SOAP: a more complex service. SOAP is a xml-protocol that defines what kind of xml-documents the client may send and what kind of xml-documents the server sends in response. The details are defined in the JAX-WS specification. A concrete implementation of JAX-WS is Apache CXF. A soap-service is described to the consumer using WSDL.

4.3.11.1 Rest-container: Jersey

Jersey is a server and container to be used as a REST-frontend to your application. You just feed it annotated beans like the following, and jersey takes care of receiving the http-request and mapping them to the right point in your bean.

MyController.java

```
package myrest;

import javax.ws.rs.GET;
import javax.ws.rs.Path;
import javax.ws.rs.Produces;

@Path("/hello")
public class MyController {

    @GET
    @Produces("text/plain")
    public String getClicchedMessage() {
        return "Hello World";
    }
}
```

web.xml

```
<?xml version="1.0" encoding="UTF-8"?>
<web-app xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
          xmlns="http://xmlns.jcp.org/xml/ns/javaee"
          xsi:schemaLocation="http://xmlns.jcp.org/xml/ns/javaee http://xmlns.jcp.org/xml/ns/javaee/web-app_3_1.xsd"
          id="WebApp_ID" version="3.1">

    <servlet>
        <servlet-name>MyRestApp</servlet-name>
        <servlet-class>org.glassfish.jersey.servlet.ServletContainer</servlet-class>
        <init-param>
            <param-name>jersey.config.server.provider.packages</param-name>
            <param-value>myrest.MyController</param-value>
        </init-param>
        <load-on-startup>1</load-on-startup>
    </servlet>

    <servlet-mapping>
        <servlet-name>MyRestApp</servlet-name>
        <url-pattern>/rest/*</url-pattern>
    </servlet-mapping>
</web-app>
```

²Its good to know that there is a popular alternative to jersey known as restlet.

pom.xml

```

<project xmlns="http://maven.apache.org/POM/4.0.0" xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
  xsi:schemaLocation="http://maven.apache.org/POM/4.0.0 http://maven.apache.org/xsd/maven-4.0.0.xsd">
  <modelVersion>4.0.0</modelVersion>
  <groupId>org.langbein.michael</groupId>
  <artifactId>myrest</artifactId>
  <version>1</version>

  <packaging>war</packaging>

  <dependencies>
    <dependency>
      <groupId>org.glassfish.jersey.core</groupId>
      <artifactId>jersey-server</artifactId>
      <version>2.17</version>
    </dependency>
    <dependency>
      <groupId>org.glassfish.jersey.containers</groupId>
      <artifactId>jersey-container-servlet-core</artifactId>
      <version>2.17</version>
    </dependency>
  </dependencies>

  <build>
    <sourceDirectory>src</sourceDirectory>
    <plugins>
      <plugin>
        <artifactId>maven-compiler-plugin</artifactId>
        <version>3.6.1</version>
        <configuration>
          <source>1.8</source>
          <target>1.8</target>
        </configuration>
      </plugin>
      <plugin>
        <artifactId>maven-war-plugin</artifactId>
        <version>3.0.0</version>
        <configuration>
          <warSourceDirectory>WebContent</warSourceDirectory>
        </configuration>
      </plugin>
    </plugins>
  </build>
</project>
```

4.3.11.2 Soap-container: Apache CXF

JXF is a server and container to be used as either a REST- or a SOAP-frontend to your application³. You configure JXF as one of the two by

Then you feed it your beans like this:

....

4.3.11.3 JSP and JSTL

JSTL is meant as an extension of html. It allows you to create script-like logic in html, but without inserting any java (or php) code in html. Here an example:

```

<%@ taglib uri="http://java.sun.com/jstl/core" prefix="c" %>
<html>
  <head>
    <title>Count to 10 Example (using JSTL)</title>
  </head>

  <body>
    <c:forEach var="i" begin="1" end="10" step="1">
      <c:out value="${i}" />

      <br />
    </c:forEach>
  </body>
</html>
```

³Actually, cxf can speak even more protocols than just REST and SOAP, like CORBA. Also, it can not only listen to HTTP, but also JMS or JBI

4.3.12 Unit testing

JUnit is the primary tool for testing in Java. It is a container like tomcat: it takes classes and runs them according to instructions that it finds in those testclasses in the form of annotations. Every such testclass extends `TestCase`, this primarily gives you access to different `assert` methods. It is important to know in what order junit executes tests that it finds in a testclass:

- Junit counts the number of `@Test` methods you have in your class and creates that many different instances of the class. So if you have five tests, five separate instances are created, by calling the constructor five times.
- Then, for each instance,
 - JUnit calls the method `setUp()` and any method annotated with `@Before`
 - Junit calls the one `@Test` method that the instance has been created for
 - Junit calls the method `tearDown()` and any method annotated with `@After`

4.3.13 Logging

That's right, in an enterprise-application, not even logging is done without a framework. This section will introduce logback.

To include logback, you should add the following dependency:

```
<dependency>
    <groupId>ch.qos.logback</groupId>
    <artifactId>logback-classic</artifactId>
    <version>1.2.3</version>
</dependency>
<dependency>
    <groupId>org.slf4j</groupId>
    <artifactId>slf4j-jdk14</artifactId>
    <version>1.7.25</version>
</dependency>
```

This single dependency is enough, as it will transitively pull in the logback-core and slf4j-api dependencies. If you use spring boot, there is no need for that dependency; logback is already contained within the starter.

The usage of logback is rather simple.

```
import org.slf4j.Logger;
import org.slf4j.LoggerFactory;

public class SomeApp {
    private final Logger logger = LoggerFactory.getLogger("MeinDefaultLogger");

    public void someMethod(){
        logger.debug("This is a debug message");
        logger.info("This is an info message. It can have {} of {}", "all kinds of", "parameters");
        logger.warn("This is a warn message");
        logger.error("This is an error message");
    }
}
```

configuration It is up to you to decide how the different log-statements should be handled. To do so, place a `logback.xml` in your `src/main/resources` folder. This configuration file contains the following elements:

- An **appender** is responsible for writing out the logstatements to a destination. Properties: `name` , `class` , `level` . There are several types (indicated by the `class` attribute):
 - RollingFileAppender
 - SMTPAppender (email)
 - ConsoleAppender
 - SiftingAppender (separates logs based on a runtime attribute)

An **appender** can also have a **filter**, that decides what kind of messages it would display.

- `filter` .

- `root` is the logger's software interface. Properties: `level` (the log-level threshold that should actually be used).
 - `appender-ref`'s list all the elements appenders that should actually be used.
- `logger` is just like the `root` logger, but with a different name. You can get the logger "meinBesondererLogger" with `Logger mbl = LoggerFactory.getLogger("meinBesondererLogger"); mbl.warn("Pass besser auf!");`. Loggers can have ancestors; to create a child to an ancestor, just give it a name that includes the ancestors name, like "meinBesondererLogger.sonderFall".

An example would be:

```
<configuration debug="true" scan="true">

    <property name="meineLogFile" value="log/meinlog.txt" />
    <property name="meinLogFileArchiv" value="log/meinlog" />

    <appender name="consolenOutput" class="ch.qos.logback.core.ConsoleAppender">
        <encoder class="ch.qos.logback.classic.encoder.PatternLayoutEncoder">
            <pattern>%d{HH:mm:ss.SSS} [%thread] %-5level %logger{50} - %msg%n</pattern>
            <target>System.err</target>
        </encoder>
    </appender>

    <appender name="fileOutput" class="ch.qos.logback.core.rolling.RollingFileAppender">
        <encoder class="ch.qos.logback.classic.encoder.PatternLayoutEncoder">
            <pattern>%d{HH:mm:ss.SSS} [%thread] %-5level %logger{50} - %msg%n</pattern>
        </encoder>
        <file>${meineLogFile}</file>
        <rollingPolicy class="ch.qos.logback.core.rolling.TimeBasedRollingPolicy">
            <fileNamePattern>${meinLogFileArchiv}.%d{yyyy-MM-dd}.%i.log</fileNamePattern>
            <timeBasedFileNamingAndTriggeringPolicy class="ch.qos.logback.core.rolling.SizeAndTimeBasedFNATP">
                <maxFileSize>5MB</maxFileSize>
            </timeBasedFileNamingAndTriggeringPolicy>
            <maxHistory>7</maxHistory>
        </rollingPolicy>
    </appender>

    <appender name="roleSiftingAppender" class="ch.qos.logback.classic.sift.SiftingAppender">
        <discriminator>
            <key>userRole</key>
            <defaultValue>ANONYMOUS</defaultValue>
        </discriminator>
        <sift>
            <appender name="fileAppender" class="ch.qos.logback.core.FileAppender">
                <file>${userRole}.log</file>
                <encoder>
                    <pattern>%d [%thread] %-5level %mdc %logger{50} - %msg%n</pattern>
                </encoder>
            </appender>
        </sift>
    </appender>

    <root level="debug">
        <appender-ref ref="consolenOutput" />
    </root>

    <logger level="info" name="meinBesondererLogger">
        <appender-ref ref="rollingFileAppender" />
    </logger>
</configuration>
```

For the discriminator to have access to the `userRole` key, you need to place it in the MDC (Mapped Diagnostic Context). Simply put, MDC allows you to set information to be later retrieved by other Logback components, using a simple, static API:

```
|| MDC.put("userRole", "ADMIN");
|| logger.info("Admin Action");
```

4.4 Python

Python is a wonderful calculator. It is not intended for large projects, but for one-off scripts, prototypes, data-analyses and the like.

4.4.1 List comprehensions

List comprehensions allow us to write set-notation like declarative definitions of lists.

```
|| S = [x**2 for x in range(10)]
|| M = [x for x in S if x%2 == 0]
```

We can also do dict-comprehensions:

```
|| dict2 = {k*2:v for (k,v) in dict1.items()}
```

List comprehensions can be conditional:

```
|| listA = ["a", "b", "c", "d"]
|| listB = ["c", "d", "e", "f"]
|| [entry for entry in listA if entry in listB] # <-- [c, d]
```

4.4.2 Functional programming

Anonymous functions are declared by using the `lambda` keyword.

```
|| ls = [0, 1, 2, 3, 4]
|| squares = map(lambda x: x * x, ls)
|| evens = filter(lambda x: x%2 == 0, ls)
```

There is also support for currying:

```
|| from functools import partial
```

4.4.3 Decorators

Python decorators are annotations that extend the behavior of the function over which they are written. That lends them to do, for example, aspect oriented programming.

In vanilla python, we could do this:

```
|| def myDecorator(someFunc):
||     def wrappedFunc():
||         print("This is printed before the function is called")
||         someFunc()
||         print("This is printed after the function is called")
||     return wrappedFunc
|
|| def sayWhee():
||     print("Wheee!")
|
|| wrappedWhee = myDecorator(sayWhee)
|| wrappedWhee()
|
|| # >> "This is printed before the function is called"
|| # >> "Whee!"
|| # >> "This is printed after the function is called"
```

But there is syntactic sugar for this.

```
|| @myDecorator
|| def sayWhee():
||     print("Wheee!")
|
|| sayWhee()
```

This also works with functions that take arguments:

```
|| def addOne(func):
||     def wrapped(*args):
||         out = func(*args)
||         return out + 1
||     return wrapped
|
|| @addOne
|| def timesTwo(n):
||     return 2*n
|
|| timesTwo(2) # <-- 5
```

As you can see, a decorator takes a function as an argument. We can also create decorators that depend on an argument themselves:

```
def addN(n):
    def decorator(func):
        def wrapped(*args):
            out = func(*args)
            return out + n
        return wrapped
    return decorator

@addN(2)
def timesTwo(n):
    return 2*n

timesTwo(2) # <-- 6
```

4.4.4 Metaprogramming

On one hand we can change the way an object handles the standard operators (like `+`, `-`, `==` etc); on the other we can define our own new operators (like `innerprod`, `crossprod` etc). And finally we can extend existing datastructures to change their behaviour to suite our needs.

4.4.4.1 Changing the way an object handles operations

Python has a lot of built-in methods that are added to every class by default⁴. One of them is `__add__`. Consider this example:

```
class MyVec:

    def __init__(self, x, y):
        self.x = x
        self.y = y

    def __add__(self, other):
        x = self.x + other.x
        y = self.y + other.y
        return MyVec(x, y)

    def __str__(self):
        return "[{}, {}]".format(self.x, self.y)

v1 = MyVec(1, 2)
v2 = MyVec(2, 3)
print v1 + v2 # >> [3, 5]
```

4.4.4.2 Creating new operators

Defining custom operators is not internally supported by python, but you can use this:

```
# From http://code.activestate.com/recipes/384122/

class Infix:
    def __init__(self, function):
        self.function = function
    def __ror__(self, other):
        return Infix(lambda x, self=self, other=other: self.function(other, x))
    def __or__(self, other):
        return self.function(other)
    def __rlshift__(self, other):
        return Infix(lambda x, self=self, other=other: self.function(other, x))
    def __rshift__(self, other):
        return self.function(other)
    def __call__(self, value1, value2):
        return self.function(value1, value2)

x=Infix(lambda x,y: x*y)
print 2|x|4
```

Update: this has now made it into its own package: <https://pypi.python.org/pypi/infix/>

⁴This is because internally, every class inherits from a default-class named 'type'

```

||| from infix import or_infix as infix
|||
||| @infix
||| def add(a, b):
|||     return a + b
|||
||| a | add| b

```

4.4.4.3 Extending existing datastructures

Say you want lists to be indexed at 1.

```

||| class MyList(list):
|||     def __getitem__(self, key):
|||         return list.__getitem__(self, key-1)

```

However, it is often problematic to change basic data-structures because other modules might depend on their predefined behaviour. For that reason python exposes a few abstract base types. Extending from these base types is like forking from abstract list instead of committing to list trunk. A classical case of an existing datastructure that we might want to extend from is dictionaries for use as timeseries:

```

import datetime as dt
from UserDict import *

class TimeSeries(IterableUserDict):
    """ self.data verhaelt sich wie ein ganz normales dict """
    def __init__(self, formDtTime='%d.%m.%Y %H:%M'):
        UserDict.__init__(self)
        self.formDtTime = formDtTime

    def __setitem__(self, time, value):
        if isinstance(time, basestring):
            time = self.stringToDtTime(time)
        if not isinstance(time, dt.datetime):
            raise ValueError('Schlüssel muss ein Zeitstring oder ein Datetime sein')
        UserDict.__setitem__(self, time, value)

    def stringToDt(self, timeString, form = None):
        dtTime = self.stringToDtTime(timeString, form)
        return dtTime.date()

    def stringToDtTime(self, timeString, form = None):
        if not form:
            form = self.formDtTime
        return dt.datetime.strptime(timeString, form)

    def getFirstTime(self):
        return min(self.data)

    def getLastTime(self):
        return max(self.data)

    def getFirstVal(self):
        return self.data[self.getFirstTime()]

    def getLastVal(self):
        return self.data[self.getLastTime()]

    def getFirstPoint(self):
        time = self.getFirstTime()
        return [time, self.data[time]]

    def getLastPoint(self):
        time = self.getLastTime()
        return [time, self.data[time]]

    def getSpan(self):
        return self.getLastTime() - self.getFirstTime()

    def getSubseries(self, von = None, bis = None):
        if not von:
            von = self.getMinTime()
        if not bis:
            bis = self.getMaxTime()
        ts = TimeSeries()
        for time in self.data:
            if time < von or time > bis:
                continue
            ts[time] = self.data[time]
        return ts

```

```

def getFilteredSeries(self, funct):
    ts = TimeSeries()
    for time in self.data:
        if funct(time, self.data[time]):
            ts[time] = self.data[time]
    return ts

if __name__ == '__main__':
    # TimeSeries ist ein dict mit zusätzlichen Methoden
    ts = TimeSeries('%d.%m.%Y')
    ts['14.10.1986'] = 0
    ts['15.10.1986'] = 1
    ts['16.10.1986'] = 2
    ts['17.10.1986'] = 3
    print ts.getSpan()
    lw = ts.getLastPoint()
    print "Letzter Wert ist {} am {}".format(lw[0], lw[1])
    ss = ts.getFilteredSeries(lambda t, v: v >= 1)
    print ss

    # Iterieren funktioniert wie gewohnt
    for el in ts:
        print el

    # Beim Einfügen von Daten wird auf das richtige Format geachtet
    try:
        ts[4] = 2
    except Exception as e:
        print repr(e)

    # Map und reduce funktioniert wie gewohnt
    print reduce(lambda summe, key: summe + ts[key], ts, 0)

```

4.4.4.4 Named tuples

Named tuples take the best of simple classes (accessing fields by names) and tuples (accessing fields by index, immutable, iterable). They are wonderful in combination with infix custom operators for use as geometric objects.

```

>>> Point = namedtuple('Point', ['x', 'y'])
>>> p = Point(11, y=22)      # instantiate with positional or keyword arguments
>>> p[0] + p[1]             # indexable like the plain tuple (11, 22)
33
>>> x, y                  # unpack like a regular tuple
>>> x, y
(11, 22)
>>> p.x + p.y              # fields also accessible by name
33
>>> p                      # readable __repr__ with a name=value style
Point(x=11, y=22)

```

4.4.5 Plotting

Matplotlib is a matlab inspired library for plotting.

```

import matplotlib.pyplot as plt

plt.plot(ns, NsShallow)
plt.plot(ns, NsDeep)
plt.xlabel("input size")
plt.ylabel("number of nodes")
plt.legend(["node in shallow net to achieve accuracy = {}".format(epsilon), "node in deep net to achieve accuracy = {}"])
plt.show()

```

In general, you want to avoid `plt`. `plt` is just the header of the current plotting state. Use axes instead - axes have all the same methods as `plt`. But contrary to `plt`, they can be passed around freely.

```

fig, axesArr = plt.subplots(2, 2)
doFirstPlot(axesArr[0])
doSecondPlot(axesArr[1])
...
plt.show()

```

There are two important plotting functions: `plot` and `scatter`. `Plot` expects you to pass in all values as an array, `scatter` allows you to add each point one by one.

```

import matplotlib.pyplot as plt
from mpl_toolkits.mplot3d import Axes3D

fig = plt.figure()
ax = fig.gca(projection='3d')
for point in line:
    ax.scatter(point.x, point.y, point.z)
plt.show()

import matplotlib.pyplot as plt
from mpl_toolkits.mplot3d import Axes3D

fig = plt.figure()
ax = fig.gca(projection='3d')
ax.plot(xarray, yarray)
plt.show()

```

Also, often we want to plot heatmaps:

```

data = np.random.random((100, 100))
plt.imshow(data, cmap='hot', interpolation='nearest')
plt.show()

```

We can add multiple images to one single figure.

```

fig = plt.figure()

ax1 = fig.add_subplot(231, projection='3d')
ax1.set_title("Original body")
for point in signalCart:
    ax1.scatter(point[0], point[1], point[2])

ax2 = fig.add_subplot(232)
ax2.set_title("Flattened")
ax2.imshow(signal)

ax3 = fig.add_subplot(233)
ax3.set_title("Fourier Transform")
ax3.imshow(log(abs(amps)))

ax4 = fig.add_subplot(234, projection='3d')
ax4.set_title("New body")
for point in signalNewCart:
    ax4.scatter(point[0], point[1], point[2])

ax5 = fig.add_subplot(235)
ax5.set_title("Flattened")
ax5.imshow(log(abs(signalNew)))

ax6 = fig.add_subplot(236)
ax6.set_title("Altered amplitudes")
ax6.imshow(log(abs(ampsNew)))

plt.show()

```

4.4.6 SymPy

Sympy is fairly good at doing symbolic calculations. You can put restrictions to the type of a variable in the Symbol class.

```

from sympy import Symbol, sin, cos

x = Symbol('x')
y = Symbol('y')

diff(sin(x), x) # --> cos(x)
cos(x).series(x, 0, 10) # --> 1 - x^2/2 + x^4/24 - ...
sin(x+y).expand(trig=True) # --> sin(x)cos(y) + sin(y)cos(x)

from sympy import Matrix
M = Matrix([[1,0],[0,1]])

from sympy.solvers import solve

solve(x**2 -1, x) # --> [-1, 1]
solve([x+y-5, x-y+1], [x,y]) # --> {x:2,y:3}

integrate(t * exp(t), t)
>>>
integrate(t * exp(t), (t, 0, oo))

```

```

>>> from sympy import Symbol, exp, I
>>> x = Symbol("x")
>>> exp(I*x).expand()
I*x
e
>>> exp(I*x).expand(complex=True)
-I*x
-I*x
I*e
    *sin(re(x)) + e
    *cos(re(x))
>>> x = Symbol("x", real=True)
>>> exp(I*x).expand(complex=True)
I*sin(x) + cos(x)

```

For many common integrals sympy already has a predefined function:

```

>>> from sympy import fourier_transform, exp
>>> from sympy.abc import t, k
>>> fourier_transform(exp(-t**2), t, k)
sqrt(pi)*exp(-pi**2*k**2)
>>> fourier_transform(exp(-t**2), t, k, noconds=False)
(sqrt(pi)*exp(-pi**2*k**2), True)

```

The same holds for statistics: many common distributions are already given:

```

>>> from sympy.stats import Poisson, density, E, variance, cdf
>>> from sympy import Symbol, simplify
>>> rate = Symbol("lambda", positive=True)
>>> z = Symbol("z")
>>> X = Poisson("x", rate)
>>> density(X)(z)
lambda**z*exp(-lambda)/factorial(z)
>>> E(X)
lambda
>>> simplify(variance(X))
lambda
>>> cdf(X)(2)
lambda**2*exp(-lambda)/2 + lambda*exp(-lambda) + exp(-lambda)

```

You can even use standard python function definitions in sympy:

```

def myf(x):
    return x**2 + 1

x = Symbol("x")

diff(myf(x), x)
>> 2*x

```

4.4.7 Numpy and Scipy

Scipy has, amongst others, a bunch of useful statistical methods:

```

from scipy.stats import poisson

def pois(lmbd, n):
    return poisson.pmf(n, lmbd)

def Pois(lmbd, n):
    return poisson.cdf(n, lmbd)

def poisCond(lmbd, n, n0):
    if n < n0:
        return 0
    else:
        return (pois(lmbd, n)) / (1 - Pois(lmbd, n0-1))

def expt(lmbd):
    ex = 0
    for i in range(samplesize):
        ex += i * pois(lmbd, i)
    return ex

def exptCond(lmbd, n0):
    ex = 0
    for i in range(samplesize):
        ex += i * poisCond(lmbd, i, n0)
    return ex

```

4.4.8 Gradual typing

Since python 3.6, gradual typing is possible with the mypy-checker. Here is some examplecode:

```
class Person:
    def __init__(self, name: str) -> None:
        self.name = name

    def sayHi(person: Person) -> str:
        return "Hello, {}".format(person.name)

m = Person("Michael")
print(sayHi(m))
```

However, the python-interpreter itself does, so far, not check types. Python does allow type-hints, but does not enforce them by itself. To check types ahead of type, use `mypy prog.py`.

4.4.9 Piping and currying

```
from infix import or_infix

@or_infix
def pipe(val, func):
    return func(val)

def add2(x):
    return x + 2

print(3 | pipe| add2 | pipe| (lambda x:x**2))

import inspect

def curry(func, *args, **kwargs):
    """
    This decorator make your functions and methods curried.
    Usage:
    >>> adder = curry(lambda (x, y): (x + y))
    >>> adder(2, 3)
    5
    >>> adder(2)(3)
    5
    >>> adder(y = 3)(2)
    5
    """

    assert inspect.getargspec(func)[1] == None, 'Currying can\'t work with *args syntax'
    assert inspect.getargspec(func)[2] == None, 'Currying can\'t work with *kwargs syntax'
    assert inspect.getargspec(func)[3] == None, 'Currying can\'t work with default arguments'

    if (len(args) + len(kwargs)) >= func.__code__.co_argcount:
        return func(*args, **kwargs)

    return (lambda *x, **y: curry(func, *(args + x), **dict(kwargs, **y)))
```

4.4.10 Generators

Generators create lists that are not kept in memory, but only evaluated on demand. This is useful for performance if we need to work with very large lists. We can potentially define even infinite lists.

```
# As function. Note the 'yield' instead of a 'return'
def evens(n):
    for i in range(n):
        if i%2 == 0:
            yield i

# As a list comprehension. Note the () instead of []
evens = (i for i in range(n) if i%2 == 0)
```

Both these methods don't return [1, 2, 4, ..., n**2], but <generator_n>, which will be evaluated when neccessary.

4.4.11 Oslash: advanced functional programming

Very often, we don't know if an actual value or null is passed into or returned from a function. Functional programming languages manage this problem by wrapping those values in a maybe-monad. More generally, very

often operations are nondeterministic: a function might return nothing or multiple results at once. Then, it makes sense to wrap these possible outcomes in a monad.

One important thing to note is this: monads are not the essence of functional programming, but a means of dealing with uncertain operation-results. It will make sense to unwrap monads as soon as you can.

Here we will show functors, applicative functors (which are just beefed up functors) and finally monads (which are just beefed up applicatives).

- A functor is a data type that implements the Functor abstract base class. You apply a function to a wrapped value using map or %
- An applicative is a data type that implements the Applicative abstract base class. You apply a wrapped function to a wrapped value using * or lift
- A monad is a data type that implements the Monad abstract base class. You apply a function that returns a wrapped value, to a wrapped value using — or bind

A Maybe implements all three, so it is a functor, an applicative, and a monad.

There are several common monads, like for example:

- maybe : if a value might be null
- list : if a value might be 0, 1 or more values
- io
- reader
- writer
- state
- identity : there is no reason to use this - only useful for theory.

Wrapped values

Functors Unfortunately, you cannot just stick a wrapped integer to a function that accepts an integer:

```
def add3(x):
    return x + 3

val = Just(2)

add3(val)
>> Error!
```

This is what functors are for.

```
def add3(x):
    return x + 3

val = Just(2)

val.map(add3)
>> Just 5
```

This works because the `Just` wrapper implements the abstract `Functor` baseclass, which forces it to define a `map(self, func)` method.

Applicatives It was a little weird that we had to write the value before the function in `val.map(add3)`. But we can also wrap functions in a wrapper. If the wrapper implements the `Applicative` class, we can instead write:

```
add3Wrapped = Just(lambda: x: x+3)
val = Just(2)

add3Wrapped.lift(val)
>> Just 5
```

4.4.12 Multimedia: Pygame

In general when we want to do a bigger multimedia-app, pygame is a save choice. It's python's equivalent of the Processing framework for Java.

```

import sys, pygame

pygame.init()

size = width, height = 320, 240
speed = [2, 2]
black = 0, 0, 0

screen = pygame.display.set_mode(size)

ball = pygame.image.load("ball.gif")
ballrect = ball.get_rect()

while True:
    for event in pygame.event.get():
        if event.type == pygame.QUIT:
            sys.exit()

    ballrect = ballrect.move(speed)
    if ballrect.left < 0 or ballrect.right > width:
        speed[0] = -speed[0]
    if ballrect.top < 0 or ballrect.bottom > height:
        speed[1] = -speed[1]

    screen.fill(black)
    screen.blit(ball, ballrect)
    pygame.display.flip()

```

screen aka. surface: `pygame.display.set_mode()`.

`display.flip()` will update the contents of the entire display. `display.update()` allows to update a portion of the screen, instead of the entire area of the screen. Passing no arguments, it updates the entire display. To tell PyGame which portions of the screen it should update (i.e. draw on your monitor) you can pass a single `pygame.Rect` object, or a sequence of them to the `display.update()` function. A `Rect` in PyGame stores a width and a height as well as a x- and y-coordinate for the position.

PyGame's built-in dawning functions and the `.blit()` method for instance return a `Rect`, so you can simply pass it to the `display.update()` function in order to update only the "new" drawn area.

Due to the fact that `display.update()` only updates certain portions of the whole screen in comparison to `display.flip()`, `display.update()` is faster in most cases.

4.4.13 Regex

4.4.14 Modules and Packages

In python, it is very important from which directory a script is called, because that determines how `import` statements are resolved. That means that when you write `import` statements, you are forcing the user to call your code from a certain directory.

Every python-file is a module. Every folder containing python-files is a package. Packages can additionally have a `__init__.py` file with additional information.

When the interpreter encounters the code ...

```

||   from ie.engine import InferenceEngine
... in a file main.py, it searches for a file engine.py in a folder ie inside

```

- the current directory or the location of the script `main.py`,

- The location of the `main.py` file will be added if you ran the script with `python somedir/main.py`
- The current directory will be added if you ran the script with `python -m somedir.main`

- the `PYTHONPATH`. You can append to this path with:

```

||       import sys
||       sys.path.append('path/to/ie')

```

If the directory `ie` contains an `__init__.py` file, this file is executed when a module in that dir is loaded for the first time. This way, the `__init__.py` can be used to initialize any global variables or such that the module needs.

4.4.15 Miniconda

Honestly, the python2/3 shism is going to kill you. And if that doesn't, then the multiple c-dependencies python has. There is really only one way around this, which is using miniconda as a venv manager. Miniconda first installs a root-environment for you under `/miniconda`. The packages in there will be included in any further venv you create. You create a new venv with `source activate meinTestEnv`. Note that this does not create a folder at your current location. It does, however, create one under `/miniconda/envs/meinTestEnv`. This is where all your env-specific packages will be stored. A venv may even use a different python-version than your root-env. You can deactivate the venv again with `source deactivate`.

- `conda env list`
- `conda create --name meinEnv`
- `source activate meinEnv`
- `source deactivate`

Conda gets its packages from so-called channels - just like ubuntu-repositories. You can change the priority at which conda searches through channels, and add new ones.

- `conda config --get channels`
- Search for packages here: <https://anaconda.org/anaconda/repo>
- `conda config --append channels newchannel`
- `conda list`
- `conda search tensorflow`
- `conda install tensorflow=1.0.1`
- If you use pip while in a conda venv, the pip-package will also be installed in the venv. However, it does *not* include things you install through `apt-get`.

You can export and import environment-specifications like this:

- `conda env export > environment.yml`
- `conda env create -f environment.yml`

4.4.16 IPythonNotebook aka. Jupyter

Jupyter is a client-server program. You start it on your server, where it runs in some conda-environment. The server exposes a web-interface, where the user can enter python-commands into the client that will then be executed on the server. This is advantageous for big-data environments, where you want to bring the code to the data, because getting the data to the user takes too much network traffic. Instructions can be found here.

4.4.17 Pandas

```
|| df.size
```

4.5 Racket

The hallmark of the Lisp family is that programs are defined in terms of data structures rather than in terms of a text-based syntax. The most visible consequence is a rather peculiar visual aspect, which is dominated by parentheses. The more profound implication, and in fact the motivation for this uncommon choice, is the equivalence of code and data. Program execution in Lisp is nothing but interpretation of a data structure. It is possible, and common practice, to construct data structures programmatically and then evaluate them. The most frequent use of this characteristic is writing macros (which can be seen as code preprocessors) to effectively extend the language with new features. In that sense, all members of the Lisp family are “programmable programming languages”. There is a nice set of tutorials here.

4.5.1 Basics

- *list*: '(1 2 3)
- *syntax*: #'(+ 1 2), or #'car

```
#lang racket

(provide make-complex get-real get-complex +c -c *c)

(define (make-complex r c)
  (cons r c))

(define (get-real a)
  (car a))

(define (get-complex a)
  (cdr a))

(define (piecewise op a b)
  (make-complex
    (op (get-real a) (get-real b))
    (op (get-complex a) (get-complex b)))))

(define (add-complex a b)
  (piecewise + a b))

(define (substr-complex a b)
  (piecewise - a b))

(define (mult-complex a b)
  (let ((ar (get-real a))
        (ac (get-complex a))
        (br (get-real b))
        (bc (get-real b)))
    (make-complex
      (- (* ar br) (* ac bc))
      (+ (* ar bc) (* ac br)))))

(define (reduce func arglist carry)
  (let* ((first-el (car arglist))
         (rest-els (cdr arglist))
         (carry2 (func carry first-el)))
    (if (<= (length rest-els) 0)
        carry2
        (reduce func rest-els carry2)))))

(define (+c firstarg . restargs)
  (reduce add-complex restargs firstarg))

(define (-c firstarg . restargs)
  (reduce substr-complex restargs firstarg))

(define (*c firstarg . restargs)
  (reduce mult-complex restargs firstarg))

;; (define a (make-complex 1 1))
;; (define b (make-complex 2 1))
;; (define c (make-complex 3 3))
;; (*c a b c)
;; (+c a a a b)
```

4.5.2 Macros

Is there a way we can do something like this?

```
|| (define (faculty n)
  (if (= n 1) n
      (* n (faculty (- n 1)))))

|| (hypothetical-macro-function (n !)
  (faculty n))

|| (4 !) ; ---> yields 24
```

Macros are functions that take one syntax object #'`syntx-obj` as input and return another syntax object. These syntax objects contain literal code, packaged with metadata like lexical context and source location. The syntax object passed to a macro contains the whole calling expression that invoked the macro. So if we invoke and like this:

```
|| (and (expr a) (expr b) (expr c))
```

and does not get three arguments as input, the way an ordinary function would. Rather, it gets a syntax object like this, which retains a reference to the lexical context of the calling site:

```
|| (and #'(and (expr a) (expr b) (expr c)))
```

4.5.3 Nomenclature

Lisps are often used for metaprogramming. For that reason, we need a firm grasp on the intricacies of the nomenclature of the language. We want to be able to give distinctive names to every part of a statement.

- statement: anything that is written
 - combination: aka. expression, aka. form
 1. eval all subexpressions
 2. apply leftmost value to rest
 - special forms are all statements that are not combinations.
 - * Definitions: `(define name "Michael")` does not evaluate all subexpressions, because `name` does not exist yet and can therefore not be evaluated. What `define` does is to save a new entry in the workspace-memory under the label `name`

4.5.4 Datastructures

On the most basic level, racket knows only one datastructure: the `pair`:

```
|| (cons "a" "b") ; ; => ("a" . "b")
```

A list is a pair whose tail is also a pair:

```
|| (cons "a" (cons "b" '())) ; ; => '("a" "b")
```

4.6 Javascript

At its heart, javascript is a simple language. Its has however two very unconventional features that are not used in any other languages: *this* and *prototypes*.

- Javascript does not have methods. It has properties that refer to functions.
- A function is an object. (That's why every function has a *this*). Note that this means that there are executable objects.

this : every function implicitly gets a parameters named *this*, containing the context.

```
||| function add (x, y) {
    console.log(this);
    return x + y;
}
```

The context *this* can also be explicitly assigned.

```
||| const add3 = add.bind("michael", 2);
```

Note that this does not work for arrow-functions, which keep their once assigned context!


prototype : Prototype-languages should really only achieve inheritance by copying objects. Javascript implements this badly, though, because it offers *two* means of inheritance: object-copying and constructor-functions.

- Instances
 - Every object has a *.__proto__*.
 - *.__proto__* contains the list of inherited props and methods and *should* not be modified.
 - *bender.prototype* is bender's own prototype, and *bender.__proto__* is a reference to bender's parent's prototype.
- Constructing new objects
 - With *Object.create*: passing *theObject.__proto__* and *theObject*'s instance-properties.
 - * Calling `const vender = Object.create(bender); vender.name = 'Vender';` will give *vender* these properties: *bender.__proto__* and *bender*'s instance-properties;
 - With constructor functions: passing *TheFunction.prototype*
 - * Functions are objects, so they, too, have a *.__proto__*
 - * Additionally, functions have a *.prototype*. This is because any function may serve as a constructor, so they need this *.prototype*.
 - * Even though not all functions are intended as constructors, js expects them to be that. So every function can be called with *new*. Doing this expects the function to do two things:
 - it expects to have the function assign some things to its internal *this*. These assignments will be instance-properties.
 - it expects *TheFunction.prototype* to have a few props that are to be shared between all instances. In other words: adding a function via *TheFunction.prototype.myMethod = () => ...* instead of *function TheFunction () this.myMethod = () => ...* makes the function static. There is only one instance of that function on the heap, not a separate one for each instance.
 - * Contrary to *.__proto__*, *.prototype* may and should be modified.
 - * So, calling `const bender = new Robot('Bender');` will give *bender* these properties: *Robot.prototype*;
 - Funfact: *Object* is a constructor-function, which is why there is a *Object.prototype*. Indeed, calling `ame: 'Michael'` is shorthand for `ame: 'Michael' new Object(n)`

```
||| const Robot = function(name) {
    this.name = name;
}
Robot.prototype.sayHi = function() {
    return 'Hi from ' + this.name;
};
const bender = new Robot('Bender');
bender.sayHi();
Robot.prototype.sayHi = function() {
    return 'Bugger off from ' + this.name;
}
bender.sayHi();
```

4.6.1 The js-runtime

Figure 4.1: *Stack*: Todo's in the current frame. Makes calls to the *Web-APIs*. *Queue*: From another thread the browser sends messages onto the queue. Some of these messages may have been triggered by the JS-stack calling for example a `setTimeout` Web-API call. They're handled FIFO.



4.6.2 Asynchronous programming

Rxjs is a popular, stream based implementation of the observer-pattern.

```

export interface Observer<T> {
  onNext: (val: T) => void,
  onCompleted: () => void,
  onError: (e: Error) => void
}

const nullObserver: Observer<any> = {
  onNext: (val: any) => {},
  onCompleted: () => {},
  onError: (e: Error) => {}
}

/**
 * Subscriptions are just handles to observers,
 * so that they can be unsubscribed if need be.
 */
export interface Subscription {
  unsubscribe: () => void;
}

export class SimpleSubscription<T> implements Subscription {
  constructor(
    private observer: Observer<T>
  ) {}

  unsubscribe () {
    // unsubscription == overwrite original observer
    this.observer = nullObserver;
  }
}

```

```

/**
 *
 */
export class Observable<T> {

    constructor(
        private _subscribe: (observer: Observer<T>) => Subscription
    ) {}

    subscribe(downstreamObserver: Observer<T>): Subscription {
        return this._subscribe(downstreamObserver);
    }

    /**
     * List of static creational methods.
     * They each return new observables.
     * Note that those new observables do not execute the _subscribe method yet.
     * that method is only executed when subscribe is executed.
     */
    static of<X>(args: X[]): Observable<X> {
        return new Observable<X>((downstreamObserver: Observer<X>) => {
            console.log('executing subscription-body (returned from \'of\' method)')

            args.forEach(val => downstreamObserver.onNext(val));
            downstreamObserver.onCompleted();

            return new SimpleSubscription(downstreamObserver);
        })
    }

    static fromEvent(source: Element, event: string): Observable<Event> {
        return new Observable<Event>((downstreamObserver: Observer<Event>) => {
            console.log('executing subscription-body (returned from \'fromEvent\' method)')

            const callback = (e: Event) => downstreamObserver.onNext(e);
            source.addEventListener(event, callback);
            return {
                unsubscribe: () => source.removeEventListener(event, callback)
            }
        });
    }

    /**
     * list of non-static creational methods.
     * They each return new observables.
     * Note that the subscribe methods are not called yet with this method.
     */
    map<Y>(mapFunc: (v: T) => Y): Observable<Y> {
        return new Observable<Y>((downstreamObserver: Observer<Y>) => {
            console.log('executing subscription-body (returned from \'map\' method)')

            const mappingObserver: Observer<T> = {
                onNext: (val: T) => {
                    const y: Y = mapFunc(val);
                    downstreamObserver.onNext(y);
                },
                onCompleted: () => downstreamObserver.onCompleted(),
                onError: (e: Error) => downstreamObserver.onError(e)
            };

            return this.subscribe(mappingObserver);
        });
    }

    filter(filterFunc: (v: T) => boolean): Observable<T> {
        return new Observable<T>((downstreamObserver: Observer<T>) => {
            console.log('executing subscription-body (returned from \'filter\' method)')

            const filteringObserver: Observer<T> = {
                onNext: (val: T) => {
                    if (filterFunc(val)) {
                        downstreamObserver.onNext(val);
                    }
                },
                onCompleted: () => downstreamObserver.onCompleted(),
                onError: (e: Error) => downstreamObserver.onError(e)
            };
        });
    }
}

```

```

        };

        return this.subscribe(filteringObserver);
    });

}

take(nr: number): Observable<T> {
    return new Observable<T>((downstreamObserver: Observer<T>) => {
        console.log(`executing subscription-body (returned from 'take' method)`)
        let i = 0;
        const takingObserver: Observer<T> = {
            onNext: (val: T) => {
                if (i < nr) {
                    downstreamObserver.onNext(val);
                    i += 1;
                } else {
                    downstreamObserver.onCompleted();
                }
            },
            onCompleted: () => downstreamObserver.onCompleted(),
            onError: (e: Error) => downstreamObserver.onError(e)
        };
        return this.subscribe(takingObserver);
    })
}

const list$ = Observable.of([1, 2, 3, 4, 5, 6, 7, 8]);
const list2$ = list$.map((v) => v+1);
const list3$ = list2$.filter((v) => v % 2 === 1);
const list4$ = list3$.take(3);

list4$.subscribe({
    onNext: (val: number) => console.log(val),
    onCompleted: () => {},
    onError: () => {}
});

```

Zones allow you to wrap multiple, potentially asynchronous call-frames in one common environment.

```

interface ZoneSpec {
    name: string;
    props?: object;
    onFork?;
    onIntercept?;
    onInvoke?;
    onScheduleTask?;
    onInvokeTask?;
    onCancelTask?;
    onHasTask?;
}

class Zone {
    private static _current: Zone = new Zone(null, {name: 'Base'});
    private _parent: Zone;
    private zoneSpec: ZoneSpec;

    constructor(parent: Zone, zoneSpec: ZoneSpec) {
        this._parent = parent;
        this.zoneSpec = zoneSpec;
    }

    static get current() {
        return Zone._current;
    }

    get name() {
        return this.zoneSpec.name;
    }

    get parent() {
        return this._parent;
    }

    get(key: string) {
        return this.zoneSpec.props ? this.zoneSpec.props[key] : null;
    }
}

```

```

    fork(newZoneSpec: ZoneSpec) {
      for (const key in this.zoneSpec) {
        if (key != 'name') {
          if (!newZoneSpec[key]) {
            newZoneSpec[key] = this.zoneSpec[key];
          }
        }
      }
      return new Zone(Zone.current, newZoneSpec);
    }

    run(callback) {
      Zone._current = this;
      callback();
      Zone._current = this.parent;
    }
  }

  const _setTimeout = (callback, time) => {
    const zoneOnCreateTime = Zone.current;
    const wrappedCallback = () => {
      zoneOnCreateTime.run(callback);
    };
    setTimeout(wrappedCallback, time);
  }

  const zoneBC = Zone.current.fork({
    name: 'BC',
    props: {
      message: "Hi! You can only see me inside BC!"
    }
  });

  function c() {
    console.log("executing c in zone ", Zone.current.name);
    console.log("Here's the message data: ", Zone.current.get('message'));
  }

  function b() {
    console.log("executing b in zone ", Zone.current.name);
    console.log("Here's the message data: ", Zone.current.get('message'));
    _setTimeout(c, 2000);
    // c();
  }

  function a() {
    console.log("executing a in zone ", Zone.current.name);
    console.log("Here's the message data: ", Zone.current.get('message'));
    zoneBC.run(b);
  }

  a();
  console.log("this is root, running in zone ", Zone.current.name);
}

```

4.6.3 Module system

4.6.3.1 Building with webpack

My builder of choice is webpack. Basically, webpack makes ES6-import and -export work. There is some nice documentation here.

paragraphGeneral concepts: Some concepts and nomenclature.

- Loaders vs Plugins: loaders work on the individual file-level, plugins work at the bundle- or chunk-level.
- compiling: compiling means transforming (eg.) ts-code to some (older version of) js.
- bundling: bundling takes compiled code and uses the import statements to create one large file.

Module-resolution : Webpack can work with most module-types that javascript has to offer.

- Javascript: webpack works fine with commonJs, nodeJs, and AMD, but will sometimes have problems with UMD.

	CommonJs formerly node	NodeJs node	AMD requirejs	UMD	ES2015 aka <code>export const s</code>
used by module-file		<code>exports.area = (r) = PI * r * r;</code>	<code>area: (r)defPhe* r * r; }</code>	First tries AMD, then commonJs, then exports as global.	
user-file loads modules implementations		const module = require('./module'); module.exports = (r) => PI * r * r;	synchronously node, webpack, browserify	synchronously webpack, requirejs browserify	<code>import { sqrt</code>
Notes	Used to be a candidate for node, but now abandoned.			Does not work well with node (there is amdefine to help, but not ideal, either).	Official JS standard but so far implemented in browser or node

- if a package.json has a *module* entry, this field indicates that there is a ES6 version of the code. This will be loaded then.
 - if it has a *main* entry, that one leads to the entry-file of a (usually UMD) module.
 - otherwise, you cannot import a whole module, but must load by filename.
- Typescript: `import b from "moduleB";`. If in tsconfig the property `compilerOptions.module` is not '`AMD`' '`System`' — '`ES2015`'—, then a variation of the node-resolution is used:
 - `'/node_modules/moduleB.ts'`
 - `'/node_modules/moduleB.tsx'`
 - `'/node_modules/moduleB.d.ts'`
 - `'/node_modules/moduleB/package.json'` (if it specifies a "types" property)
 - `'/node_modules/@types/moduleB.d.ts'`
 - `'/node_modules/moduleB/index.ts'`
 - `'/node_modules/moduleB/index.tsx'`
 - `'/node_modules/moduleB/index.d.ts'`
 - A note on Typescript's `import` and `require` statements: With TypeScript, `import` can be used if there is a declaration file for the module. If there isn't a declaration file, the TypeScript compiler doesn't know if the module exists, so you need to use `require` instead which lacks the compilation checking.
 - Other languages: how to resolve other languages (ts, coffee, sass, shader, ...) is left over to loaders.

Shimming : Webpack only loads a module when it sees it required at some point. This makes for smaller bundles, of course, but here is one problem this can cause. Say you have a dependency to jQuery and bootstrap. JQuery exposes `window.$` as a global object. Bootstrap requires there to be a global `window.$` object. However, since bootstrap just assumes that `window.$` is there without ever calling `import`, webpack won't resolve the dependency and never include jQuery in the package. To avoid such problems, there is something called *shimming*.

```
plugins: [
  new CopywebpackPlugin([
    new CopywebpackPlugin({ from: path.resolve(__dirname, 'node_modules/cesium/Build/Cesium/Workers'), to: 'Workers' }),
    new CopywebpackPlugin({ from: path.resolve(__dirname, 'node_modules/cesium/Source/Assets'), to: 'Assets' })
  ]),
  new CopywebpackPlugin([
    new CopywebpackPlugin({ from: path.resolve(__dirname, 'node_modules/cesium/Source/Widgets'), to: 'Widgets' })
  ]),
  new webpack.DefinePlugin({
    CESIUM_BASE_URL: JSON.stringify(''),
  }),
  new webpack.ProvidePlugin({
    Cesium: 'cesium/Build/Cesium/Cesium'
  })
],
```

The `ProvidePlugin` makes Cesium available to the global workspace. If you don't want this, you might want to instead use the `module.rules.imports-loader`.

Path-resolution : When you ask to import a string that doesn't have a leading `./`, it tells Webpack: "Look for the file everywhere on my hard drive except the current directory". I'm not kidding. You can see the output desperately climbing up your directory hierarchy looking for a directory called `node_modules` in which it hopes to find `index.js`. When you precede the module name with a `./`, it tells Webpack to look in the current directory for the file. So, the rule of thumb is: your code usually has a `./` and third-party libraries don't.

3rd party libs `d3` is a good example of a big library moving to the new ES6 module-syntax. `d3` has many 3rd party libs that depend on and extend `d3`. As of `D3 4.0`, none of the `D3` modules (formerly known as plugins) rely on a global `d3` object. Each module is available either as a UMD bundle for use in Node or browsers, and as ES modules for use in bundlers and other modern JavaScript environments. Note that `d3-plus` hasn't somehow still not managed to switch.

How webpack does it Webpack compiles this code:

```
// helpers.ts
export function sayHi(name: string): string {
  return `Hi, ${name}!`;
}

// index.ts
import { sayHi } from "./helpers";
console.log(sayHi('Michael'));
```

to this code:

```
// self calling function
(function(modules) {
  // The module cache
  var installedModules = {};

  // The require function
  function __webpack_require__(moduleId) {

    // Check if module is in cache
    if(installedModules[moduleId]) {
      return installedModules[moduleId].exports;
    }
    // Create a new module (and put it into the cache)
    var module = installedModules[moduleId] = {
      i: moduleId,
      l: false,
      exports: {}
    };
    // Execute the module function
    modules[moduleId].call(module.exports, module, module.exports, __webpack_require__);

    // Flag the module as loaded
    module.l = true;

    // Return the exports of the module
    return module.exports;
  }

  // expose the modules object (__webpack_modules__)
  __webpack_require__.m = modules;

  // expose the module cache
  __webpack_require__.c = installedModules;

  // define getter function for harmony exports
  __webpack_require__.d = function(exports, name, getter) {
    if(!__webpack_require__.o(exports, name)) {
      Object.defineProperty(exports, name, { enumerable: true, get: getter });
    }
  };

  // define __esModule on exports
  __webpack_require__.r = function(exports) {
    if(typeof Symbol !== 'undefined' && Symbol.toStringTag) {
      Object.defineProperty(exports, Symbol.toStringTag, { value: 'Module' });
    }
    Object.defineProperty(exports, '__esModule', { value: true });
  };

  // create a fake namespace object
  // mode 0: value is a module id, require it
  // mode 1: value is a module id, merge all properties of value into the ns
  __webpack_require__.t = function(value, mode) {
    if(mode & 1) {
      Object.assign(exports, value);
    }
    return value;
  };

  // expose the require function
  __webpack_require__.n = function(module) {
    return module && module.__esModule ? module.default : module;
  };

  // expose the module function
  __webpack_require__.o = function(object, property) {
    return Object.prototype.hasOwnProperty.call(object, property);
  };

  // expose the module function
  __webpack_require__.p = '';
  __webpack_require__.l = true;
})()
```

```

// mode & 4: return value when already ns object
// mode & 8/1: behave like require
__webpack_require__.t = function(value, mode) {
    if(mode & 1) value = __webpack_require__(value);
    if(mode & 8) return value;
    if((mode & 4) && typeof value === 'object' && value && value.__esModule) return value;
    var ns = Object.create(null);
    __webpack_require__.r(ns);
    Object.defineProperty(ns, 'default', { enumerable: true, value: value });
    if(mode & 2 && typeof value !== 'string') for(var key in value) __webpack_require__.d(ns, key, function(key) {
        return ns[key];
    });
}

// getDefaultExport function for compatibility with non-harmony modules
__webpack_require__.n = function(module) {
    var getter = module && module.__esModule ?
        function getDefault() { return module['default']; } :
        function getModuleExports() { return module; };
    __webpack_require__.d(getter, 'a', getter);
    return getter;
};

// Object.prototype.hasOwnProperty.call
__webpack_require__.o = function(object, property) { return Object.prototype.hasOwnProperty.call(object, property); }

// __webpack_public_path__
__webpack_require__.p = "";

// Load entry module and return exports
return __webpack_require__(__webpack_require__.s = 0);
})({


// dependency
"./src/dependency.ts": (function(module, __webpack_exports__, __webpack_require__) {
    function sayHi(name) {
        return "Hi, " + name + "!";
    }
    __webpack_require__.r(__webpack_exports__);
    __webpack_require__.d(__webpack_exports__, "sayHi", function() { return sayHi; })
}),
// script
"./src/index.ts": (function(module, __webpack_exports__, __webpack_require__) {
    __webpack_require__.r(__webpack_exports__);
    var _helpers__WEBPACK_IMPORTED_MODULE_0__ = __webpack_require__(./src/helpers.ts);
    console.log(Object(_helpers__WEBPACK_IMPORTED_MODULE_0__["sayHi"])(`Michael`));
})
});
}

```

4.6.4 Promises, async-await and rxjs

`async - await` lets you handle asynchronous code as if it was synchronous.

4.6.5 Webworkers: multithreading

4.6.6 WASM

- If wasm does not get a garbage-collector,
 - there will be no easy way for non-gc languages to run in wasm. They need to bring their own gc (or even runtime) with them (like c# does now).
 - In that case, a situation like python + numy will probably evolve.
- If wasm will not allow DOM-manipulation,
 - then there really is no point for non numeric-focussed languages to be implemented in wasm.
 - In that case, a situation like python + numy will probably evolve.
- If, however, wasm will get a gc and allow DOM-manipulation,
 - Then python, ruby etc. will all develop angular-clones to run on the web.
 - However, given the large js-ecosystem, those will not kick in too mightily.

4.7 CSS

4.7.1 Position

- **static**: the default
- **relative**: default + allows top, bottom, left and right
- **fixed**: relative to viewport
- **absolute**: relative to nearest positioned parent (positioned: anything but static)
- **sticky**: like fixed, but relative to container instead of viewport.

4.7.2 Display

Per default, divs are displayed in blockstyle, while spans are displayed inlinestyle.

- **block**
- **inline**
- **inline-block**: inline, but allows width, height, padding etc.
- **none**
- **table**
- **grid**
- **flex** see 4.7.2.1

4.7.2.1 Flexbox

4.8 xml

XML is the de-facto standard language for specifying capabilities, interfaces etc.

4.8.1 Structure

XML consist of elements with attributes. The general name for anything in xml is a *node* - a node is an element, an attribute or a text.

4.8.2 Namespaces: xmlns

Because the naming of xml-elements and attributes can lead to conflicts, they are usually prefixed with a namespace `xmlns`. The id of the namespace need not resolve to a URL of a xsd. If it doesn't, you can still provide a hint for where to search for a schema using the extra-element `<namespacename>:schemaLocation="<schemaurl>"`.

4.8.3 Schema definitions: xsd

With xsd we can consult an online-resource to make sure a given xml-file conforms to a specification.

Xsd's description of elements (aka. element type) is used to describe what elements look like. Xsd defines simple and complex types.

- `<simpleType>` simple types can only have content directly contained between the element's opening and closing tags. They cannot have attributes or child elements.
- `<complexType>` complex types can have attributes, can contain other elements, can contain a mixture of elements and text, etc etc.

Xsd also has its own, custom elements that are used to describe certain requirements.

- `restriction` restricts an elements' content to a certain type (string, numeric, ...)
- `sequence`

Here is an example from a simple xsd:

```

<element name="email">
  <simpleType>
    <restriction base="xsd:string" />
  </simpleType>
</element>

<element name="getPreferences">
  <complexType>
    <sequence>
      <element name="email" type="email" />
    </sequence>
  </complexType>
</element>

```

Xsd's description of element's contents (aka. content type) is used to describe what is expected to be inside an elements brackets.

- A simple type's content can be one of:
 - atomic types, which have indivisible values, such as #000 and #AACCDD
 - list types, which have whitespace-separated lists of indivisible values, such as blue green red
 - union types, which have either atomic or list values, but they can be the union of other types, such as blue #000 red for a set of colors
- Complex types have a “content model,” which refers to how the content (the data between the element's opening and closing tags) is arranged:
 - simple content is only character data, no child elements allowed

- element-only content is only children, no data allowed
- mixed content means character data and child elements can be intermingled
- empty content means the element is empty (`<foo/>`) and either conveys information by just existing, or has attributes but no content.

Just to clarify: elements have an *element type*, and their content has a *content type*. By the way, attributes can only have simple types, because they cannot themselves have attributes or children.

4.8.4 Marshalling

XML to Pojo : JAXB

```
<jaxb:bindings
  xmlns:jaxb="http://java.sun.com/xml/ns/jaxb"
  xmlns:xs="http://www.w3.org/2001/XMLSchema"
  xmlns:xjc="http://java.sun.com/xml/ns/jaxb/xjc"
  jaxb:version="2.0">

  <!-- <jaxb:bindings schemaLocation="http://www.w3.org/1999/xlink.xsd">
    <jaxb:bindings node="//xs:attributeGroup[@name='locatorAttrs']/xs:attribute[@ref=xlink:title]">
      <jaxb:property name="LinkTitle"></jaxb:property>
    </jaxb:bindings>
  </jaxb:bindings> -->

  <jaxb:bindings scd="x-schema::xlink" xmlns:xlink="http://www.w3.org/1999/xlink">
    <jaxb:bindings scd="/group::xlink:locatorModel/model::sequence/xlink:title">
      <jaxb:property name="LocatorTitle"/>
    </jaxb:bindings>
  </jaxb:bindings>
</jaxb:bindings>
```

XML to JSON : Jsonix here

4.9 Design patterns

4.9.1 UML

4.9.1.1 Classes

Figure 4.2: UML arrows





Figure 4.3: UML classes



4.9.1.2 Sequences

4.9.1.3 Usecases

4.9.2 Dependencies versus microservices

Dependencies can be included right in your sourcode. Microservices on the other hand are not included at all, but called remotely from your programm.

Table 4.3: Dependencies versus microservices

	Dependencies	Microservices
Changes	You can include the exact needed version of a dependency	Your program needs to adapt immediately if the service


4.9.3 Strategy pattern

The strategy pattern (also known as the policy pattern) is a behavioral software design pattern that enables selecting an algorithm at runtime. Instead of implementing a single algorithm directly, code receives run-time instructions as to which in a family of algorithms to use.[1]

Strategy lets the algorithm vary independently from clients that use it.[2] Strategy is one of the patterns included in the influential book *Design Patterns* by Gamma et al.[3] that popularized the concept of using design patterns to describe how to design flexible and reusable object-oriented software. Deferring the decision about which algorithm to use until runtime allows the calling code to be more flexible and reusable.

For instance, a class that performs validation on incoming data may use the strategy pattern to select a validation algorithm depending on the type of data, the source of the data, user choice, or other discriminating factors. These factors are not known until run-time and may require radically different validation to be performed. The validation algorithms (strategies), encapsulated separately from the validating object, may be used by other validating objects in different areas of the system (or even different systems) without code duplication.

Figure 4.4: Strategy design pattern



4.9.4 State pattern

The state pattern is a behavioral software design pattern that allows an object to alter its behavior when its internal state changes. This pattern is close to the concept of finite-state machines. The state pattern can be interpreted as a strategy pattern, which is able to switch a strategy through invocations of methods defined in the pattern's interface.

The state pattern is used in computer programming to encapsulate varying behavior for the same object, based on its internal state. This can be a cleaner way for an object to change its behavior at runtime without resorting to conditional statements and thus improve maintainability.

The state pattern is set to solve two main problems:

- An object should change its behavior when its internal state changes.
- State-specific behavior should be defined independently. That is, adding new states should not affect the behavior of existing states.

4.9.5 Observer pattern

The Observer pattern addresses the following problems:

- A one-to-many dependency between objects should be defined without making the objects tightly coupled.
- It should be ensured that when one object changes state an open-ended number of dependent objects are updated automatically.
- It should be possible that one object can notify an open-ended number of other objects.

Figure 4.5: State design pattern





Figure 4.6: Observer design pattern



4.9.6 Orchestration: Workflow pattern

4.9.7 Choreography: Actor pattern

4.9.8 Monads

4.9.8.1 Maybe-monad

Maybe-monads make sense when a function might return null, which we want to avoid. A user can only access the value after he has specified what should happen if the value is null. In its simplest form, this could be done like this: A nice implementation can be found here.

```

class Maybe<T> {
    private value: T | null = null;

    constructor(val: T | null) {
        this.value = val;
    }

    public getVal(ifNull: T): T {
        if (this.value) return this.value;
        else return ifNull;
    }
}

// a function that might return null
function divide(dividend: number, divisor: number): Maybe<number> {
    return new Maybe<number>(dividend / divisor);
}

```

```
}

// we can only get the result of the calculation after specifying what should happen if null is returned.
function divideAndResolve(dividend: number, divisor: number, deflt: number): number {
    let maybeResult = divide(dividend, divisor);
    return maybeResult.getVal(deflt);
}

console.log("10/2 = ...", divideAndResolve(10, 2, 0));
console.log("10/0 = ...", divideAndResolve(10, 0, 0));
```


Chapter 5

Stacks, platforms and libraries

5.1 Git

5.1.1 Lingo

- origin: default name for your remote repository
- head: *your last commit* on the current branch. Does not yet include non-committed changes.
 - When you switch branches with `git checkout`, the HEAD revision changes to point to the tip of the new branch.
 - When you do a `git pull`, your working directory may contain new stuff, which your head is lagging behind. To see the difference between the working dir and your last commit, do `git diff HEAD`
- index: The index is a single, large, binary file in `baseOfRepo/.git/index`, which lists all files in the current branch, their sha1 checksums, time stamps and the file name

Merging remote changes on master into local master It can happen that somebody else has been messing with the master before you had a chance to commit your own changes. You will notice this by seeing your `git push` fail. In that case, try to first merge the remote changes into your local copy before pushing again.

```
|| git commit // first saving our latest changes
|| git fetch // getting remote stuff in remote-tracking-branches. Doesn't change our files yet
|| git diff master origin/master // compares what you have in master vs. what happened to master remotely
|| git merge remotes/origin/master
```

This will try to merge the upstream changes into yours. If it fails, the conflicting file will contain annotations like these:

```
|| <<<<< HEAD
|| print("159 just edited this file")
|| =====
|| print("133 just edited this file")
|| >>>> master
```

As you can guess, `<<<<< HEAD` is your code, `==` is a separator, and `>>>>` `master` is the end of the other guys code. Use a texteditor to resolve all conflicts. Eclipse actually has a graphical interpreter for the conflict-tags. Once the resolving is completed, a `git push` and `git commit` will complete the merge.

```
|| git status // shows that there is still a conflict
|| git add .
|| git status // shows that conflict is resolved, but merging still in progress
|| git commit -m "phew! managed to resolve the conflict!"
|| git push origin master
```


5.1.2 Changing origin

```
|| git remote -v
|| git remote set-url origin https://github.com/USERNAME/REPOSITORY.git
```

5.1.3 Merge vs. rebase


(a) git merge b

- splices in commits from b
- creates merge commit



(b) git rebase a

- deletes own history
- replays own history on latest commit of a



Reverting and resetting is best explained in this post: <https://stackoverflow.com/questions/4114095/how-to-revert-git-repository-to-a-previous-state>. Basically, we have these options:

- `git revert <lastGoodCommitHash>..HEAD` will execute the inverse of all the operations and commit them as a new commit. This way, the old history stays intact.
- `git reset <lastGoodCommitHash>` will unstage all the commits from `<firstBadCommitHash>..HEAD`. This way, you're rewriting history, which is really only good if you're working alone on a repo.

5.1.4 Creating your own repository

Basically, any git folder can be used as a repository. But you still need to do some work to expose that folder over the web. The easiest way to work is certainly over gitlab. For alternative means if you don't have a hoster, this site lists a few of the possibilities: <http://www.jedi.be/blog/2009/05/06/8-ways-to-share-your-git-repository/>.

5.2 Docker

Containers and virtual machines.

- A container runs natively on Linux and shares the kernel of the host machine with other containers. It runs a discrete process, taking no more memory than any other executable, making it lightweight.
- By contrast, a virtual machine (VM) runs a full-blown “guest” operating system with virtual access to host resources through a hypervisor. In general, VMs provide an environment with more resources than most applications need.

Images are described by *Dockerfiles* that list all the contents of an image. These image files can be downloaded from a central repository called *docker hub*. Usually, you write your own dockerfile that specifies as a dependency a more general dockerfile as your base.

- an *image* is an executable package that includes everything needed to run an application—the code, a runtime, libraries, environment variables, and configuration files. Includes execution of all `RUN` statements
- a *container* is a runtime instance of an image - what the image becomes in memory when executed (that is, an image with state, or a user process). On startup, the `CMD` directive is executed.

5.2.1 Command overview

- `docker`
 - `image`
 - * `pull`
 - * `rm <imagename>`
 - * `build --tag=<nameofappalllowercase> .`: package your app and send it to `/var/lib/docker/images`
 - `container`
 - * `run -p 4000:80 -d --name=<nameofyourcontaineralllowercase> <nameofyourappalllowercase>`: execute that package (where `-p 4000:80` means map the containers port 80 to the systems port 4000 and `-d` stands for detached, i.e. get back control of your command-line after starting the container). `run` is actually a shorthand for `create` and `start`.
 - * `run -it ...` creates a new container from an image
 - * `start -i ...` starts up an existing container
 - * `logs --tail 100 --since <minutes> --timestamps --follow <containerid>`: All `stdout` and `stderr` (?) goes to this log
 - * `cp /path/to/local/file.html my-nginx:/var/www/html`
 - * `exec -it my-nginx /bin/bash`
 - * `inspect --format={{.LogPath}} <containerid>`: Hardware-info
 - * `rm <containername>`
 - `volume`
 - * `ls`
 - * `rm`
 - * `create --name <volumename>`
 - * To be used with `docker container run -v <volumename>:/<container>/fs/path <imagename>`
- `docker-compose`
 - `up -d`
 - `build`: like `up`, but rebuilds when a dockerfile changes.
 - `start`: starts a `stop`'ed set of containers.
 - `stop`: stops the containers.
 - `down`: stops the containers and deletes them (but leaves the images intact).

5.2.2 Dockerfile syntax

Dockerfile

```
# Parent image
FROM python:2.7-slim
# All subsequent commands are executed here. May be called multiple times.
WORKDIR /app
# Copy the current directory contents into the container at /app
COPY . /app
# Every run command creates another layer
RUN pip install --trusted-host pypi.python.org -r requirements.txt
# Make port 80 available to the world outside this container
EXPOSE 80
# Define environment variable
ENV NAME World
# Cmd is called when the container launches. This command's stout goes directly into the docker logs.
CMD ["python", "app.py"]
```

5.2.3 Composefile syntax

docker-compose.yml

```
version: '3'
services:
  mysql:
    build: ./mysql
    env_file:
      - ./mysql/resources/mysql.env
    ports:
      - "127.0.0.1:3306:3306"
    volumes:
      - ./mysql/resources/init:/docker-entrypoint-initdb.d
      - mysql-data:/var/lib/mysql
    # See https://www.drupal.org/project/drupal/issues/2966523
    command: --default-authentication-plugin=mysql_native_password

  drupal:
    build: ./drupal
    env_file:
      - ./mysql/resources/mysql.env
      - ./drupal/resources/drupal.env
    depends_on:
      - mysql
    volumes:
      - html:/var/www/html
      - ./dev:/var/www/html/sites
    expose:
      - 9000

  nginx:
    build: ./nginx
    env_file:
      - ./nginx/resources/nginx.env
    depends_on:
      - mysql
      - drupal
    ports:
      - "127.0.0.1:8080:80"
    volumes:
      - html:/var/www/html
      - ./dev:/var/www/html/sites
      - ./nginx/resources/default.conf:/etc/nginx/conf.d/default.conf:ro

volumes:
  html:
  mysql-data:
```

5.2.4 Layers and continuing failed builds

Each directive in a dockerfile creates a new *layer* (aka. *intermediate image*), which is cached in the engine to speed up the build of other dockerfiles. If your build breaks at a certain point in the execution, like this:

```
Step 16/25 : RUN cd geoserver/src/src
--> Running in dff9e492bfc7
Removing intermediate container dff9e492bfc7
--> 16c2728b1c46
Step 17/25 : RUN mvn clean install -Pwps,wps-remote,importer -DskipTests
--> Running in f102f7633b49
```

```
[INFO] Scanning for projects...
[INFO] BUILD FAILURE
[WARNING] The requested profile "wps" could not be activated because it does not exist.
[WARNING] The requested profile "wps-remote" could not be activated because it does not exist.
[WARNING] The requested profile "importer" could not be activated because it does not exist.
[ERROR] The goal you specified requires a project to execute but there is no POM in this directory (/). Please verify ...
```

You can start the container at the last successful intermediate image like this:

```
|| docker image ls
|| REPOSITORY          TAG      IMAGE ID      CREATED      SIZE
|| <none>              <none>   16c2728b1c46  About an hour ago  1.1GB
|| docker history 16c2728b1c46
|| IMAGE              CREATED      CREATED BY
|| COMMENT
|| 16c2728b1c46      About an hour ago  /bin/sh -c cd geoserver/src/src
|| 63bb6b9446f9      About an hour ago  /bin/sh -c source /etc/profile.d/maven.sh
|| e236f5e2c9bc      About an hour ago  /bin/sh -c ln -s /opt/apache-maven-3.6.1/ /o...  24B
|| docker container run --rm -it 16c2728b1c46 (bash)
```

Note the `--rm` argument. This removes a container again once it has stopped running. You can start a very small container for a single instruction and immediately remove it again like this:

```
|| docker container run --rm alpine:latest bin/sh -c "whoami"
```

5.3 Spring

Spring is such a vast framework that it is worth its own section. In general the principle is as follows:

- A main class starts the context. From the context, it starts the app.
- The context wires together all the loose beans of your model and your utils using a bean.xml

5.3.1 Dependency injection

A basic spring-app may be set up like this:

```

<dependencies>
    <dependency>
        <groupId>org.springframework</groupId>
        <artifactId>spring-context</artifactId>
        <version>4.3.10.RELEASE</version>
    </dependency>
</dependencies>

package model;

public interface Knight {
    public void embarkOnQuest();
}

package model;

public class BraveKnight implements Knight {
    private Quest quest;
    public BraveKnight(Quest quest) {
        this.quest = quest;
    }
    public void embarkOnQuest() {
        quest.embark();
    }
}

package model;

public interface Quest {
    public void embark();
}

package model;

public class DamselRescueQuest implements Quest {
    public void embark() {
        System.out.println("Knight will now rescue the damsel!");
    }
}

package model;

public class DragonSlayingQuest implements Quest {
    public void embark() {
        System.out.println("Knight will now slay the dragon!");
    }
}

package main;

import org.springframework.context.annotation.Bean;
import org.springframework.context.annotation.ComponentScan;
import org.springframework.context.annotation.Configuration;

import model.BraveKnight;
import model.DamselRescueQuest;
import model.Knight;
import model.Quest;
```

```

@Configuration
@ComponentScan("model")
public class KnightsConfig {

    @Bean
    public Knight knight() {
        return new BraveKnight(quest());
    }

    @Bean
    public Quest quest() {
        return new DamselRescueQuest();
    }
}

```

Annotations:

- Configuration: Use this class instead of a xml file for the wiring-instructions
- ComponentScan: look in these packages to scan for beans
- Bean: put this object in the spring context. By default, the bean will be given an ID that is the same as the @Bean-annotated method's name.

```

package main;

import org.springframework.context.annotation.AnnotationConfigApplicationContext;

import model.Knight;

public class AppMain {

    public static void main(String[] args) {
        AnnotationConfigApplicationContext acac = new AnnotationConfigApplicationContext(KnightsConfig.class);

        Knight k = acac.getBean(Knight.class);
        k.embarkOnQuest();

        acac.close();
    }
}

```

As you can see, the basic idea is this: Use a DI framework to ensure that all beans are maximally decoupled and easily interchangeable.

5.3.2 Aspects

Aspects allow you to keep utility objects, like logging, authentication, security, etc out of your model.

```

<dependencies>
    <dependency>
        <groupId>org.springframework</groupId>
        <artifactId>spring-context</artifactId>
        <version>4.3.10.RELEASE</version>
    </dependency>

    <dependency>
        <groupId>org.springframework</groupId>
        <artifactId>spring-aop</artifactId>
        <version>4.3.10.RELEASE</version>
    </dependency>

    <dependency>
        <groupId>org.springframework</groupId>
        <artifactId>spring-aspects</artifactId>
        <version>4.3.10.RELEASE</version>
    </dependency>
</dependencies>

package util;

import java.io.PrintStream;

public class Minstrel {

    private PrintStream stream;

    public Minstrel(PrintStream stream) {
        this.stream = stream;
    }
}

```

```

        }

        public void singBeforeQuest() {
            stream.println("Fa la la la, the knight is so brave!");
        }

        public void singAfterQuest() {
            stream.println("Tee hee hee, the brave knight did finish the quest!");
        }
    }

<?xml version="1.0" encoding="UTF-8"?>


<beans xmlns="http://www.springframework.org/schema/beans"
       xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
       xmlns:aop="http://www.springframework.org/schema/aop"
       xsi:schemaLocation="http://www.springframework.org/schema/aop
http://www.springframework.org/schema/aop/spring-aop-3.2.xsd
http://www.springframework.org/schema/beans
http://www.springframework.org/schema/beans/spring-beans.xsd">

    <bean id="knight" class="model.BraveKnight">
        <constructor-arg ref="quest"></constructor-arg>
    </bean>

    <bean id="quest" class="model.DamselRescueQuest">
    </bean>

    <bean id="minstrel" class="util.Minstrel">
        <constructor-arg value="#{T(System).out}"></constructor-arg>
    </bean>

    <aop:config>
        <aop:aspect ref="minstrel">
            <aop:pointcut id="embark" expression="execution(* *.embarkOnQuest(..))" />
            <aop:before pointcut-ref="embark" method="singBeforeQuest" />
            <aop:after pointcut-ref="embark" method="singAfterQuest" />
        </aop:aspect>
    </aop:config>
</beans>
```



We put the knights.xml in src/main/resources. This is because files are loaded from the classpath, that is, the root of the produced jar. And resources is always put right into the root of the jar.

```

package main;

import org.springframework.context.annotation.AnnotationConfigApplicationContext;
import org.springframework.context.support.ClassPathXmlApplicationContext;

import model.Knight;

public class AppMain {

    public static void main(String[] args) {
        ClassPathXmlApplicationContext ctx = new ClassPathXmlApplicationContext("knights.xml");

        Knight k = ctx.getBean(Knight.class);
        k.embarkOnQuest();

        ctx.close();
    }
}
```

Some jargon knowledge is in order:

- advice: the action taken by an aspect. In the above example, `println`. Spring knows five kinds of advice: before, after, after-returning, after-throwing, and around.
- join-points: moments in the spring workflow where a advice may be applied. Internally, spring here loops through all registered aspects and checks if any of them applies in the given situation. (Equivalent to places

where hooks are looped through in drupal.) In other words: any place in spring where *some* aspect *might* be applied.

- pointcuts: a subset of all join-points, where *an actual* aspect *is* applied.
- introduction: An introduction allows you to add new methods or attributes to existing classes. For example, you could create an Auditable advice class that keeps the state of when an object was last modified. This could be as simple as having one method, setLast-Modified(Date), and an instance variable to hold this state. The new method and instance variable can then be introduced to existing classes without having to change them, giving them new behavior and state.

Aspects are woven in sometime during the execution of the application. Typically, an AOP container dynamically generates a proxy object that delegates to the target object while weaving in the aspects. This is how Spring AOP aspects are woven. However, JAspectJ can do this during compile-time!

You really don't need to use xml-configuration to be able to use aspects. We can also use the annotation-api borrowed from AspectJ:

```
@Aspect
public class Audience {

    @Pointcut("execution(** concert.Performance.perform(..))")
    public void performance() {}

    @Before("performance()")
    public void silenceCellPhones() {
        System.out.println("Silencing cellphones");
    }

    @AfterThrowing("performance()")
    public void demandRefund() {
        System.out.println("Demanding a refund");
    }
}
```

Don't forget to tell spring that it should look for Aspects though!

```
@Configuration
@EnableAspectJAutoProxy
@ComponentScan
public class AppConfig {

    @Bean
    public Performance performance() {
        return new RollingStonesConcert();
    }

    @Bean
    public Audience audience() {
        return new Audience();
    }
}
```

5.3.3 Templates

5.3.4 Bean lifecycle

1. Spring instantiates the bean.
2. Spring injects values and bean references into the bean's properties.
3. If the bean implements BeanNameAware, Spring passes the bean's ID to the setBeanName() method.
4. If the bean implements BeanFactoryAware, Spring calls the setBeanFactory() method, passing in the bean factory itself.
5. If the bean implements ApplicationContextAware, Spring calls the setApplicationContext() method, passing in a reference to the enclosing application context.
6. If the bean implements the BeanPostProcessor interface, Spring calls its postProcessBeforeInitialization() method.

7. If the bean implements the InitializingBean interface, Spring calls its afterPropertiesSet() method. Similarly, if the bean was declared with an initmethod, then the specified initialization method is called.
8. If the bean implements BeanPostProcessor, Spring calls its postProcessAfterInitialization() method.
9. At this point, the bean is ready to be used by the application and remains in the application context until the application context is destroyed.
10. If the bean implements the DisposableBean interface, Spring calls its destroy() method. Likewise, if the bean was declared with a destroy-method, the specified method is called.

By default, all beans in Spring are singletons.

5.3.5 Wiring

There are three ways of letting spring wire your components.

1. Automatically using annotations. Pro: simple. Con: you cannot set @Component or @Autowire in external libraries.
2. With a Java @Configuration class. Pro: it's java. Con: You might be tempted to put business logic in the configuration.
3. With xml. Pro: xml has one and only one purpose: config. Con: You cannot use any logic in your wiring.

5.3.6 Mocking

5.3.7 Spring MVC

A spring-mvc app can be created in the same folder structure as any other eclipse java app.

The app consists really only of three parts:

- the servlet-container - usually a embeded tomcat.
- the servlet, known as the dispatcher servlet
 - configured with
 - redirects all requests to `@Controller`'s
- the ContextLoaderListener
 - reads the spring configuration-file
 - creates spring-context
 - creates web-application-context

There are a few basic elements to every mvc-app:

- `@Component`: generic stereotype for any Spring-managed component
- `@Repository`: stereotype for persistence layer (jdbc, jpa, ...)
- `@Service`: stereotype for service layer (anything between controller and persistance layer)
- `@Controller`

Don't be afraid to use these annotations. Contrary to the general case, those mark classes that can only be part of a spring mvc app. It doesn't make sense to share them with other apps, so we might as well litter them with annotations.

5.3.7.1 General structure

Generally, every `@Controller` gets requests based on the rules defined in the `@RequestMapping`. The result of a request to a controller is always a string of a view-name.

5.3.7.2 Configuration

- AppInitializer: extends `AbstractAnnotationConfigDispatcherServletInitializer`. An alternative to the traditional web.xml file. Provides two configuration-classes: The servlet-config and the root-config.
- Servlet-config: extends `WebMvcConfigurerAdapter`. The configuration for the servlet.
- Root-config: the configuration for the backend, consisting of beans that are fully unaware of being used by a web-app.

5.3.7.3 Application: a minimal setup

src/main/resources/appconfig

```
<?xml version="1.0" encoding="UTF-8"?>

<beans xmlns="http://www.springframework.org/schema/beans"
       xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
       xmlns:context="http://www.springframework.org/schema/context"
       xsi:schemaLocation="http://www.springframework.org/schema/beans
                           http://www.springframework.org/schema/beans/spring-beans.xsd
                           http://www.springframework.org/schema/context
                           http://www.springframework.org/schema/context/spring-context.xsd ">

    <context:component-scan base-package="sm" />

</beans>
```

src/main/resources/webconfig

```
<?xml version="1.0" encoding="UTF-8"?>

<beans xmlns="http://www.springframework.org/schema/beans"
       xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
       xmlns:context="http://www.springframework.org/schema/context"
       xmlns:mvc="http://www.springframework.org/schema/mvc"
       xsi:schemaLocation="http://www.springframework.org/schema/beans
                           http://www.springframework.org/schema/beans/spring-beans.xsd
                           http://www.springframework.org/schema/mvc
                           http://www.springframework.org/schema/mvc/spring-mvc.xsd
                           http://www.springframework.org/schema/context
                           http://www.springframework.org/schema/context/spring-context.xsd ">

    <context:component-scan base-package="sm" />

    <mvc:annotation-driven />
    <mvc:resources mapping="/resources/**" location="/resources/" />

    <bean class="org.springframework.web.servlet.view.InternalResourceViewResolver">
        <property name="viewClass" value="org.springframework.web.servlet.view.JstlView" />
        <property name="prefix" value="/WEB-INF/views/" />
        <property name="suffix" value=".jsp" />
    </bean>

</beans>
```

src/main/webapp/WEB-INF/web.xml

```
<?xml version="1.0" encoding="UTF-8"?>

<web-app xmlns="http://xmlns.jcp.org/xml/ns/javaee"
          xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
          xsi:schemaLocation="http://xmlns.jcp.org/xml/ns/javaee
                             http://xmlns.jcp.org/xml/ns/javaee/web-app_3_1.xsd" version="3.1">

    <context-param>
        <param-name>contextConfigLocation</param-name>
        <param-value>classpath:appconfig.xml</param-value>
    </context-param>

    <listener>
        <listener-class>org.springframework.web.context.ContextLoaderListener</listener-class>
    </listener>
```

```

<servlet>
    <servlet-name>MyApp</servlet-name>
    <servlet-class>org.springframework.web.servlet.DispatcherServlet</servlet-class>
    <init-param>
        <param-name>contextConfigLocation</param-name>
        <param-value>classpath:webconfig.xml</param-value>
    </init-param>
    <load-on-startup>1</load-on-startup>
</servlet>

<servlet-mapping>
    <servlet-name>MyApp</servlet-name>
    <url-pattern>/</url-pattern>
</servlet-mapping>

</web-app>

```

src/main/java/sm/controller/AppMain.java

```

package sm.controller;

import org.springframework.stereotype.Controller;
import org.springframework.ui.Model;
import org.springframework.web.bind.annotation.RequestMapping;
import org.springframework.web.bind.annotation.RequestMethod;

@Controller
@RequestMapping("/")
public class HomeController {

    @RequestMapping(method = RequestMethod.GET)
    public String index(Model model){
        model.addAttribute("message", "Spring MVC XML Config Example");
        return "index";
    }
}

```

src/main/webapp/WEB-INF/views/index.jsp

```

<%@ page contentType="text/html; charset=UTF-8" language="java" %>
<html>
<head>
    <title>Spring MVC XML Configuration Example</title>
</head>
<body>
    ${message}
</body>
</html>

```

5.3.7.4 Application: what goes into the root-config?

I find that in practice nearly all of your non-trivial Spring MVC applications will require an application context (as opposed to only the spring MVC dispatcher servlet context). It is in the application context that you should configure all non-web related concerns such as:

- Security
- Persistence
- Scheduled Tasks

To make this a bit more concrete, here's an example of the Spring configuration I've used when setting up a modern (Spring version 4.1.2) Spring MVC application. Personally, I prefer to still use a WEB-INF/web.xml file but that's really the only xml configuration in sight.

WEB-INF/web.xml

```

<?xml version="1.0" encoding="UTF-8"?>
<web-app xmlns="http://xmlns.jcp.org/xml/ns/javaee" xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance" xsi:schemaLocati
<filter>
    <filter-name>openEntityManagerInViewFilter</filter-name>
    <filter-class>org.springframework.orm.jpa.support.OpenEntityManagerInViewFilter</filter-class>

```

```

    </filter>

    <filter>
        <filter-name>springSecurityFilterChain</filter-name>
        <filter-class>org.springframework.web.filter.DelegatingFilterProxy
        </filter-class>
    </filter>

    <filter-mapping>
        <filter-name>springSecurityFilterChain</filter-name>
        <url-pattern>/*</url-pattern>
    </filter-mapping>

    <filter-mapping>
        <filter-name>openEntityManagerInViewFilter</filter-name>
        <url-pattern>/*</url-pattern>
    </filter-mapping>

    <servlet>
        <servlet-name>springMvc</servlet-name>
        <servlet-class>org.springframework.web.servlet.DispatcherServlet</servlet-class>
        <load-on-startup>1</load-on-startup>
        <init-param>
            <param-name>contextClass</param-name>
            <param-value>org.springframework.web.context.support.AnnotationConfigWebApplicationContext</param-value>
        </init-param>
        <init-param>
            <param-name>contextConfigLocation</param-name>
            <param-value>com.company.config.WebConfig</param-value>
        </init-param>
    </servlet>

    <context-param>
        <param-name>contextClass</param-name>
        <param-value>org.springframework.web.context.support.AnnotationConfigWebApplicationContext</param-value>
    </context-param>

    <context-param>
        <param-name>contextConfigLocation</param-name>
        <param-value>com.company.config.AppConfig</param-value>
    </context-param>

    <listener>
        <listener-class>org.springframework.web.context.ContextLoaderListener</listener-class>
    </listener>

    <servlet-mapping>
        <servlet-name>springMvc</servlet-name>
        <url-pattern>/</url-pattern>
    </servlet-mapping>

    <session-config>
        <session-timeout>30</session-timeout>
    </session-config>

    <jsp-config>
        <jsp-property-group>
            <url-pattern>*.jsp</url-pattern>
            <scripting-invalid>true</scripting-invalid>
        </jsp-property-group>
    </jsp-config>

```

</web-app>

WebConfig.java

```

@Configuration
@EnableWebMvc
@ComponentScan(basePackages = "com.company.controller")
public class WebConfig {

    @Bean
    public InternalResourceViewResolver getInternalResourceViewResolver() {
        InternalResourceViewResolver resolver = new InternalResourceViewResolver();
        resolver.setPrefix("/WEB-INF/views/");
        resolver.setSuffix(".jsp");
        return resolver;
    }
}

```

AppConfig.java

```

|| @Configuration
|| @ComponentScan(basePackages = "com.company")
|| @Import(value = {SecurityConfig.class, PersistenceConfig.class, ScheduleConfig.class})
public class AppConfig {
    // application domain @Beans here...
}

```

Security.java

```

|| @Configuration
|| @EnableWebSecurity
public class SecurityConfig extends WebSecurityConfigurerAdapter {
    @Autowired
    private LdapUserDetailsMapper ldapUserDetailsMapper;

    @Override
    protected void configure(HttpSecurity http) throws Exception {
        http.authorizeRequests()
            .antMatchers("/*").permitAll()
            .antMatchers("/**/js/**").permitAll()
            .antMatchers("/**/images/**").permitAll()
            .antMatchers("/**").access("hasRole('ROLE_ADMIN')")
            .and().formLogin();

        http.logout().logoutRequestMatcher(new AntPathRequestMatcher("/logout"));
    }

    @Autowired
    public void configureGlobal(AuthenticationManagerBuilder auth) throws Exception {
        auth.ldapAuthentication()
            .userSearchBase("OU=App_Users")
            .userSearchFilter("sAMAccountName={0}")
            .groupSearchBase("OU=Development")
            .groupSearchFilter("member={0}")
            .userDetailsContextMapper(ldapUserDetailsMapper)
            .contextSource(getLdapContextSource());
    }

    private LdapContextSource getLdapContextSource() {
        LdapContextSource cs = new LdapContextSource();
        cs.setUrl("ldaps://ldapServer:636");
        cs.setBase("DC=COMPANY,DC=COM");
        cs.setUserDn("CN=administrator,CN=Users,DC=COMPANY,DC=COM");
        cs.setPassword("password");
        cs.afterPropertiesSet();
        return cs;
    }
}

```

PersistenceConfig.java

```

|| @Configuration
|| @EnableTransactionManagement
|| @EnableJpaRepositories(transactionManagerRef = "getTransactionManager", entityManagerFactoryRef = "getEntityManagerFactory")
public class PersistenceConfig {

    @Bean
    public LocalContainerEntityManagerFactoryBean getEntityManagerFactory(DataSource dataSource) {
        LocalContainerEntityManagerFactoryBean lef = new LocalContainerEntityManagerFactoryBean();
        lef.setDataSource(dataSource);
        lef.setJpaVendorAdapter(getHibernateJpaVendorAdapter());
        lef.setPackagesToScan("com.company");
        return lef;
    }

    private HibernateJpaVendorAdapter getHibernateJpaVendorAdapter() {
        HibernateJpaVendorAdapter hibernateJpaVendorAdapter = new HibernateJpaVendorAdapter();
        hibernateJpaVendorAdapter.setDatabase(Database.ORACLE);
        hibernateJpaVendorAdapter.setDatabasePlatform("org.hibernate.dialect.Oracle10gDialect");
        hibernateJpaVendorAdapter.setShowSql(false);
        hibernateJpaVendorAdapter.setGenerateDdl(false);
        return hibernateJpaVendorAdapter;
    }

    @Bean
    public JndiObjectFactoryBean getDataSource() {
        JndiObjectFactoryBean jndiFactoryBean = new JndiObjectFactoryBean();
        jndiFactoryBean.setJndiName("java:comp/env/jdbc/AppDS");
    }
}

```

```
    return jndiFactoryBean;
}

@Bean
public JpaTransactionManager getTransactionManager(DataSource dataSource) {
    JpaTransactionManager jpaTransactionManager = new JpaTransactionManager();
    jpaTransactionManager.setEntityManagerFactory(getEntityManagerFactory(dataSource).getObject());
    jpaTransactionManager.setDataSource(dataSource);
    return jpaTransactionManager;
}
}
```

ScheduleConfig.java

```
@Configuration
@EnableScheduling
public class ScheduleConfig {
    @Autowired
    private EmployeeSynchronizer employeeSynchronizer;

    // cron pattern: sec, min, hr, day-of-month, month, day-of-week, year (optional)
    @Scheduled(cron="0 0 0 * * *")
    public void employeeSync() {
        employeeSynchronizer.syncEmployees();
    }
}
```

As you can see, the web configuration is only a small part of the overall spring web application configuration. Most web applications I've worked with have many concerns that lie outside of the dispatcher servlet configuration that require a full-blown application context bootstrapped via the org.springframework.web.context.ContextLoaderListener in the web.xml.

5.3.8 Adding further servlets to spring mvc

5.3.9 Spring REST services

Rest works just like SpringMvc. One single difference is that a controllers methods return types are not Strings with view names, but POJOs and annotated with `@ResponseBody`. The `@ResponseBody` annotation tells Spring MVC not to render a model into a view, but rather to write the returned object into the response body. It does this by using one of Spring's message converters. Because Jackson 2 is in the classpath, this means that `MappingJackson2HttpMessageConverter` will handle the conversion of Greeting to JSON if the request's Accept header specifies that JSON should be returned.

5.3.10 Spring Database access

```
@Bean
public JdbcTemplate jdbcTemplate(DataSource dataSource) {
    return new JdbcTemplate(dataSource);
}

@Bean
public DataSource dataSource() {
    return new EmbeddedDatabaseBuilder()
        .setType(EmbeddedDatabaseType.H2)
        .addScripts('schema.sql', 'data.sql')
        .build();
}
```

5.4 Spring Boot

Spring boot takes away a lot of the boilerplate from spring. It reads all your classes and wires them together mostly by itself.

5.4.1 Example

We can reimplement the knight example like this:

TtApplication

```
@SpringBootApplication
public class TtApplication {
    public static void main(String[] args) {
        ConfigurableApplicationContext context = SpringApplication.run(TtApplication.class, args);
        Knight k = context.getBean(Knight.class);
        String result = k.goOnQuest();
        System.out.println(result);
    }
}
```

TtConfig

```
@Configuration
public class TtConfig {
    @Bean
    public Knight knight(Quest quest) {
        return new NotSoBraveKnight(quest);
    }
    @Bean
    public Quest quest() {
        return new ScaryQuest();
    }
}
```

Running this results in "The not-so-brave knight reluctantly goes on a quest to slay a dangerous dragon".

5.4.2 Configuration

Spring boot will automatically detect any classes marked with `@Configuration` in your classpath.

5.4.3 Database access

You can create a Jpa instance by just creating an interface extending the interface `JpaRepository`. Spring boot will automatically check your classpath for any jpa-provider and use that to create the actual implementation of the interface. If you don't want to use jpa, you can instead ...

5.4.4 Websites

Spring is primarily a framework for web-enabled apps. Consequently, spring boot makes developing webapps quite easily. First of all, we add two dependencies to our pom: `starter-web`, `starter-thymeleaf` and `devtools`.

```
<parent>
    <groupId>org.springframework.boot</groupId>
    <artifactId>spring-boot-starter-parent</artifactId>
    <version>1.5.9.RELEASE</version>
</parent>

<!-- Add typical dependencies for a web application -->
<dependencies>
    <dependency>
        <groupId>org.springframework.boot</groupId>
        <artifactId>spring-boot-starter-web</artifactId>
    </dependency>
    <dependency>
        <groupId>org.springframework.boot</groupId>
        <artifactId>spring-boot-starter-thymeleaf</artifactId>
    </dependency>
    <dependency>
        <groupId>org.springframework.boot</groupId>
    </dependency>
</dependencies>
```

```

<artifactId>spring-boot-devtools</artifactId>
<optional>true</optional>
</dependency>
</dependencies>

<!-- Package as an executable jar -->
<build>
    <plugins>
        <plugin>
            <groupId>org.springframework.boot</groupId>
            <artifactId>spring-boot-maven-plugin</artifactId>
        </plugin>
    </plugins>
</build>

```

Now we need a main-class from where spring can run, a controller for the webrequests, and finally a file in the resources folder that contains the template.

```

@SpringBootApplication
public class Main {
    public static void main(String[] args) {
        SpringApplication.run(Main.class, args);
    }
}

@Controller
public class WebsiteMainController {

    @RequestMapping("/greeting")
    public String startSite(@RequestParam(value="name", required=false, defaultValue="World") String name, Model model) {
        model.addAttribute("name", name);
        return "greeting"; // name of the template html-file to be used
    }
}

```

Here is the template, situated in `src/main/resources/templates/greeting.html`.

```

<!DOCTYPE HTML>
<html xmlns:th="http://www.thymeleaf.org">
<head>
    <title>Getting Started: Serving Web Content</title>
    <meta http-equiv="Content-Type" content="text/html; charset=UTF-8" />
</head>
<body>
    <p th:text="Hello, ' + ${name} + '!" />
</body>
</html>

```

You can start the app with the command `mvn spring-boot:run`. Going to the url `localhost:8080/greeting?name=Michael` will show you the expected site. Spring will automatically detect all changes you make to the source, and then re-build and restart the app. You can also build the app with `mvn clean package` and then run it with `java -jar target/myWebApp-1.0.jar`

5.4.5 Testing

For all your model classes, you can use the usual procedure of JUnit tests. But you might also want to test that spring itself is working fine. This section describes how to do this.

First, add Testing-support as a dependency:

```

<dependency>
    <groupId>org.springframework.boot</groupId>
    <artifactId>spring-boot-starter-test</artifactId>
    <scope>test</scope>
</dependency>

```

Now you can create Tests. For spring stuff, we use two special annotations: `@RunWith` and `@SpringBootTest` (Tells the Testser to go look for a main configuration class (one with `@SpringBootApplication` for instance), and start spring from there).

```

@RunWith(SpringRunner.class)
@SpringBootTest
public class MainTest {
    @Test
    public void contextLoads() throws Exception {
    }
}

```

Tests are run by simply calling `mvn test`.

5.4.6 Camel

Camel is quite easily integrated into spring boot.

```
package org.langbein.michael.springcamel;

import org.springframework.boot.SpringApplication;
import org.springframework.boot.autoconfigure.SpringBootApplication;

@SpringBootApplication
public class SpringcamelApplication {

    public static void main(String[] args) {
        SpringApplication.run(SpringcamelApplication.class, args);
    }
}

package org.langbein.michael.springcamel.controllers;

import org.apache.camel.ProducerTemplate;
import org.springframework.beans.factory.annotation.Autowired;
import org.springframework.web.bind.annotation.RequestMapping;
import org.springframework.web.bind.annotation.RestController;

@RestController
public class CamelController {

    @Autowired
    ProducerTemplate producerTemplate;

    @RequestMapping(value="/")
    public void startCamel() {
        producerTemplate.sendBody("direct:firstRoute", "Calling via Spring Boot Rest Controller");
    }
}

package org.langbein.michael.springcamel.routes;

import org.apache.camel.builder.RouteBuilder;
import org.springframework.stereotype.Component;

@Component
public class CamelRoutes extends RouteBuilder {

    @Override
    public void configure() throws Exception {
        from("direct:firstRoute")
            .log("camel body: ${body}");
    }
}
```

5.5 Android

5.5.1 ADB

It really makes a lot of sense to use a physical phone instead of an emulated one. To be able to communicate with your phone, however, you need the adb-server to be aware of the phone. There is a commandline-tool, `adb`, that you can use to have the adb-server talk to your phone. However, to be able to do that, you must first set up your phone accordingly.

- on your linux-machine, install `mtp-tools` and `mtpfs`.
- mount your phone using a *data* usb-cable
- enable developer-mode
- have your phone use the MTP protocol

With that given, you should be able to see your phone with `lsusb` and `adb devices`.

5.5.2 Basic concepts

The structure of an app is as follows: the compiled java-bytecode¹, together with all libraries and resources, is packaged into an APK (android application package), analog to a jar or a war.

The manifest mostly tells the phone some meta-info about the APK; like what rights are required, what SDK version is required, what the main-activity is, and such.

A layout describes the app's appearance on one screen. Written in XML. The layout can be found in `res/layout`. Within a layout file, you can reference data from `res/values/strings.xml` by using the string "`@string/app_someStringName`".

An activity is a single, defined thing that your user can do. Activities are usually associated with one screen. In code, an activity is represented by a single class extending the class `Activity`. It largely has the role of *controller*, negotiating between the `view.xml` and a custom-class-backend. All activities need to be declared in the manifest.

An intent is a type of message that one activity can send to another. It can be used to start another activity within the same app or even in another app. It can even be broadcasted to all apps on the phone so that any app that might be able to handle the request may use it. Here is an example of how we might use an intent to switch to another activity:

```

letsGoButton = (Button) findViewById(R.id.startGameButton);
letsGoButton.setOnClickListener(new View.OnClickListener() {
    public void onClick(View view) {
        Intent goToGameActivity = new Intent(MainActivity.this, PlayActivity.class);
        startActivity(goToGameActivity);
    }
});
```

A view is a custom drawable thing, like buttons or menus. You create such a view when the predefined UI-elements don't suit your needs, like when you want to create an animation. You create a view by extending `View` and implementing `onDraw(Canvas canvas)`, `onTouchEvent(MotionEvent event)`, etc.

The file R.java is an automatically created file that android uses to keep track of the resources available to your app. For example, your activity will tell the phone to use the `activity_play.xml` layout file by going `setContentView(R.layout.activity_play)`; inside the `onCreate` method. We can access things in R very much like we can in the DOM.

¹actually, your code gets compiled to java-bytecode (i.e. several .class files), but from there is compiled on to a single .DEX file to be run on android's version of the JRE, the ART.

5.5.3 Userinput

Android apps behave very much like html forms.

The steps are:

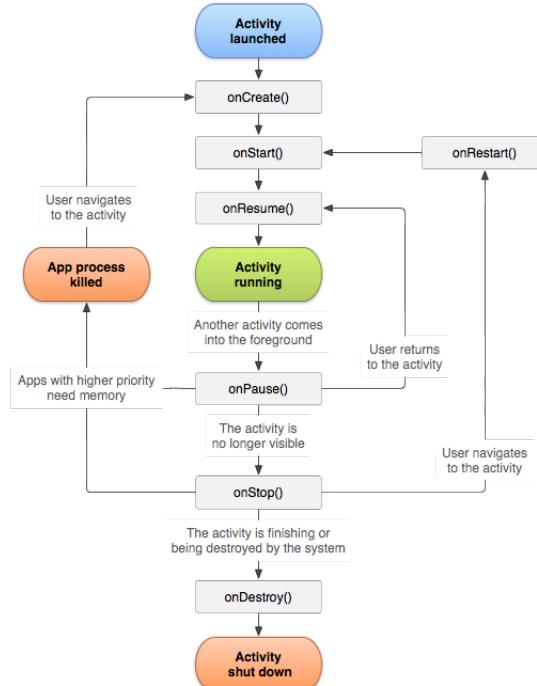
- The main-activity chooses a layout (in this case, `activity_find_beer.xml`) that is to be displayed by using

```
protected void onCreate(Bundle savedInstanceState) {
    super.onCreate(savedInstanceState);
    setContentView(R.layout.activity_find_beer);
}
```

- The layout specifies which method to call in the activity when a button is pressed. A button can have an attribute `android:onClick:"onClickFindBeer"`
- A click now is passed through to the activity, calling the method `public void onClickFindBeer(View view)`. In this case, `view` refers to the button that was pressed.
- Anywhere in the application logic, the activity can get a hold of any element of the layout by using `findViewById()` like this:

```
TextView brands = (TextView) findViewById(R.id.brands);
brands.setText("Gottle of geer");
```

5.5.4 Activities



This section handles activities in more detail.

5.5.5 Threading

Threads come in three different flavors:

- the main thread (aka. ui-thread)
 - listens for intents
 - listens for input (touch/speak) events
- the render thread

- You don't normally interact with that thread. The main-thread passes it views to render, and the render thread executes the views `onDraw` methods.
- However, you can steal work from it by creating a special kind of view: a `SurfaceView`. This kind of view exposes a canvas that is to be rendered in a custom thread.
- custom threads
 - simple custom threads extending `Thread`
 - android-specific custom threads extending `AsyncTask`

Handlers are a different thing than threads. They are not separate threads, they just postpone the execution of code. There basically promises: some code will be executed asynchronously, but still on the main thread. There are two usecases for handlers:

- You really want an action to happen on the main-thread, but at a later time. For example,
- Custom threads are not allowed to manipulate UI-elements on the main-thread. Instead, pass them a handler from the main-thread so they can call the handler. In other words: handlers are a (android-specific) method of inter-thread communication.

5.5.6 Drawing

There are three ways of drawing in android apps.

- When you only need a single animated item in an overall static UI, draw your animations into a view.
- When you want your drawing to be updated regularly, use a canvas.
- When you want to do 3d-rendering, use the OpenGL-API.

Using a custom view is certainly the most straightforward way of creating graphics. Consider this simple example:

```

public class PlayActivity extends AppCompatActivity {

    @Override
    protected void onCreate(Bundle savedInstanceState) {
        super.onCreate(savedInstanceState);
        setContentView(new FullScreenView(this));
    }
}

public class FullScreenView extends View {

    private Rect lilRect;
    private Paint painter;
    private static final int SQUARE_SIDE_LENGTH = 200;

    public FullScreenView(Context context) {
        super(context);
        lilRect = new Rect(30, 30, SQUARE_SIDE_LENGTH, SQUARE_SIDE_LENGTH);
        painter = new Paint();
        painter.setColor(Color.MAGENTA);
    }

    @Override
    protected void onDraw(Canvas canvas) {
        canvas.drawRGB(39, 111, 184);
        canvas.drawRect(lilRect, painter);
    }

    @Override
    public boolean onTouchEvent(MotionEvent event) {
        lilRect.left = (int) (event.getX() - (SQUARE_SIDE_LENGTH / 2));
        lilRect.right = (int) (event.getX() + (SQUARE_SIDE_LENGTH / 2));
        lilRect.top = (int) (event.getY() - (SQUARE_SIDE_LENGTH / 2));
        lilRect.bottom = (int) (event.getY() + (SQUARE_SIDE_LENGTH / 2));
        invalidate();
        return true;
    }
}

```

Using the canvas is a lot more direct than using a view-object. We still make use of a custom view, but instead of calling `invalidate()` when we're ready to draw, we host the game-loop inside the view and act upon the canvas directly.

```

public class MainActivity extends AppCompatActivity {

    @Override
    protected void onCreate(Bundle savedInstanceState) {
        super.onCreate(savedInstanceState);
        SceneLogic boardScene = new BoardScene();
        MainView mainView = new MainView(this, boardScene);
        setContentView(mainView);
        mainView.startRenderThread();
    }
}

/**
 * This class is there to present the SurfaceHolder as well as any touch- or system-events
 * to the RenderThread.
 */
public class MainView extends SurfaceView implements SurfaceHolder.Callback {

    private RenderThread renderThread;

    public MainView(Context context, SceneLogic sceneLogic) {
        super(context);
        SurfaceHolder surfaceHolder = getHolder();
        surfaceHolder.addCallback(this);
        renderThread = new RenderThread(context, surfaceHolder, sceneLogic);
    }

    public void startRenderThread() {
        renderThread.start();
    }

    public void stopRenderThread() {
        renderThread.setInactive();
    }

    @Override
    public boolean onTouchEvent(MotionEvent event) {
        return renderThread.getLogic().onTouch(event);
    }

    @Override
    public void surfaceCreated(SurfaceHolder surfaceHolder) {}

    @Override
    public void surfaceChanged(SurfaceHolder surfaceHolder, int i, int i1, int i2) {}

    @Override
    public void surfaceDestroyed(SurfaceHolder surfaceHolder) {}
}

public class RenderThread extends Thread {

    private final long frameTime = 17; // How long (in millisec) a frame may take
    private SceneLogic sceneLogic;
    private Context context;
    private SurfaceHolder surfaceHolder;
    private boolean running;

    public RenderThread(Context context, SurfaceHolder surfaceHolder, SceneLogic sceneLogic) {
        this.context = context;
        this.surfaceHolder = surfaceHolder;
        this.running = true;
        this.sceneLogic = sceneLogic;
    }

    @Override
    public void run() {
        long updateDuration = 0;
        long sleepDuration = 0;

        while(running) {

            // Step 0 : how long did the loop take?
            long beforeUpdateRender = System.nanoTime();
            long delta = sleepDuration + updateDuration;

            // Step 1 : scene logic
            sceneLogic.update(delta);
        }
    }
}

```

```
// Step 2: scene rendering
Canvas canvas = surfaceHolder.lockCanvas();
if(canvas != null) {
    sceneLogic.draw(canvas);
    surfaceHolder.unlockCanvasAndPost(canvas);
}

// Step 3: sleep for remainder of frameTime
updateDuration = (System.nanoTime() - beforeUpdateRender) / 1000000L;
sleepDuration = Math.max(2, frameTime - updateDuration);
try{
    Thread.sleep(sleepDuration);
} catch (Exception e) {
    e.printStackTrace();
}

}

public SceneLogic getLogic() {
    return sceneLogic;
}

public void setInactive() {
    running = false;
}
```

5.6 C++ on android

<https://www.youtube.com/watch?v=zOsUZjWbdI&list=PL0C9C46CAAB1CFB2B&t=3>

5.7 VR on android

5.7.1 Basic app

The manifest should basically be a variant of this:

```
<manifest ...>
    <uses-permission android:name="android.permission.READ_EXTERNAL_STORAGE" />
    ...
    <uses-sdk android:minSdkVersion="19" android:targetSdkVersion="24" />
    ...
    <uses-feature android:glEsVersion="0x00020000" android:required="true" />
    <uses-feature android:name="android.software.vr.mode" android:required="false" />
    <uses-feature android:name="android.hardware.vr.high_performance" android:required="false" />
    ...
    <application
        android:theme="@style/VrActivityTheme"
        ...
        <activity
            android:name=".TreasureHuntActivity"
            ...
            android:screenOrientation="landscape"
            android:configChanges="orientation|keyboardHidden|screenSize|uiMode|navigation"
            android:enableVrMode="@string/gvr_vr_mode_component"
            android:resizeableActivity="false">
            ...
            <intent-filter>
                <action android:name="android.intent.action.MAIN" />
                <category android:name="android.intent.category.LAUNCHER" />
                <category android:name="com.google.intent.category.CARDBOARD" />
                <category android:name="com.google.intent.category.DAYDREAM" />
            </intent-filter>
        </activity>
    </application>
</manifest>
```

Gradle handles the specific dependencies that are required to make a 3d-app. At least you're going to need:

```
allprojects {
    repositories {
        jcenter() // this is where we get the vr libraries
    }
}
model {
    android {
        ...
        buildTypes {
            release {
                minifyEnabled true
                // Proguard minifies the APK. This line makes sure it doesn't obfuscate vr-libs.
                proguardFiles.add(file('.../proguard-gvr.txt'))
            }
        }
    }
}
dependencies {
    classpath 'com.android.tools.build:gradle:2.3.3'
    compile 'com.google.vr:sdk-base:1.120.0'
    compile 'com.google.vr:sdk-audio:1.120.0'
}
```

The layout in a vr-app is very simple. There is only one element to use: the `com.google.vr.sdk.base.GvrView`.

The main task should in most case extend `GvrActivity`. This way, your activity gets access to a load of VR events.

5.8 Tensorflow

5.8.1 Low-level API

Working with tf always entails a two-step process. First, define a graph, and second, execute it.

```
import tensorflow as tf

# Step 1: define a graph
a = tf.Variable(initial_value=2) # Tensors are immutable by default. We must wrap them in a Variable to allow them to change
b = tf.Variable(initial_value=3)
c = a * b

# Step 2: execute graph
with tf.Session() as sess:
    sess.run(tf.global_variables_initializer())
    result = sess.run(c)
    print(result)
```

Running a tensor means evaluating it - in the case of `c`, that means calculating `a * b`. There are two things that can be run by a session: tensors (to be evaluated) and operations (to be executed).

Tf is declarative: `c = a*b` does not return the result of the multiplication of `a` and `b`, but a tensor that is still waiting to be evaluated. This can make debugging pretty hard, though. For debugging, we'd like tf to behave more imperatively, where every expression is immediately evaluated. We can achieve this by using an interactive session, where every tensor can be evaluated immediately by calling `eval()`:

```
sess = tf.InteractiveSession()
c.eval()
```

Optimizers are a very common example of operations, so we show a very simple example of their usage.

```
import tensorflow as tf
import numpy as np
import matplotlib.pyplot as plt

N = 100
a = 0.2
b = 5
xs_training = np.random.random(N) * 100
ys_training = a*xs_training + b + np.random.random(N)
xs_validation = np.random.random(N) * 100
ys_validation = a*xs_validation + b + np.random.random(N)

a_tf = tf.Variable(0.3421)
b_tf = tf.Variable(0.41)
xs_tf = tf.placeholder(tf.float32, shape=[N])
ys_tf = tf.placeholder(tf.float32, shape=[N])
model_tf = a_tf * xs_tf + b_tf
loss_tf = tf.reduce_sum((model_tf - ys_tf)**2)
optimizer = tf.train.AdamOptimizer(learning_rate=0.1)
training_operation = optimizer.minimize(loss_tf, var_list=[a_tf, b_tf])

with tf.Session() as sess:
    sess.run(tf.global_variables_initializer())

    for i in range(1000):
        sess.run(
            training_operation,
            {xs_tf: xs_training, ys_tf: ys_training}
        )

        prediction = sess.run(
            model_tf,
            {xs_tf: xs_validation}
        )

    plt.scatter(xs_validation, prediction)
    plt.scatter(xs_validation, ys_validation)
    plt.show()
```

5.8.2 Even lower level, and understanding TF internals

Tensorflow is an automatic differentiation engine. That is not quite the same as a symbolic differentiation engine. The difference becomes apparent when you create a custom operation to add to TF. To do this, you need to:

- register a new operation in c++
 - operation-name
 - inputs including their shapes
- implement the operation in c++
 - operation
 - gradient of operation for all inputs
- create a python wrapper for the public python-tf-api (not really necessary, because tf creates a default wrapper automatically)

So contrary to a symbolic engine, you have to explain to tf by hand how to calculate a gradient.

5.8.3 High-level API

When using keras, you usually don't see the low-level api at all - no tensors, no sessions, no `run`.

```

import numpy as np
import tensorflow as tf
import tensorflow.keras as k
import matplotlib.pyplot as pyplot

# 1. Data
N = 100
inputs = []
outputs = []
for n in range(N):
    i1 = np.random.randint(0,2)
    i2 = np.random.randint(0,2)
    o = 1 * ((i1 + i2) == 1)
    inputs.append([i1, i2])
    outputs.append([o])
inputs_np = np.array(inputs)
outputs_np = np.array(outputs)

# 2. Model
model = k.Sequential(layers=[
    k.layers.Dense(2, activation=k.activations.sigmoid),
    k.layers.Dense(1, activation=k.activations.sigmoid)
])

model.compile(
    optimizer=tf.train.AdamOptimizer(learning_rate=0.1),
    loss=tf.losses.mean_squared_error
)

# 3. Running

model.fit(x=inputs_np, y=outputs_np, epochs=1000)

result = model.predict(np.array([[0,0], [0,1], [1,0], [1,1]]))

print(result)

```

5.8.4 Topics

5.8.4.1 Convolutional nets

```

import tensorflow as tf
import tensorflow.keras as k
import matplotlib.pyplot as plt
import numpy as np
import dummydata as dd
import utils as u

# 1. Data
imageWidth = imageHeight = 80
batchSize = 10

```

```

timeSteps = 10
#dd.createDummyDataset(batchSize, timeSteps, imageWidth, imageHeight)
data_in, data_out = dd.loadDummyDataset()

# 2. Net
model = k.Sequential([
    k.layers.Conv3D(5, (2,2,2), name="conv3d_1", input_shape=(timeSteps, imageWidth, imageHeight, 1)),
    k.layers.MaxPool3D(name="maxpool_1"),
    k.layers.Conv3D(5, (2,2,2), name="conv3d_2"),
    k.layers.MaxPool3D(name="maxpool_2"),
    k.layers.Flatten(),
    k.layers.Dense(30, name="dense_1"),
    k.layers.Dense(10, name="dense_2"),
    k.layers.Dense(1, name="dense_3")
])

model.compile(
    optimizer=k.optimizers.Adam(),
    loss=k.losses.mean_squared_error
)

model.fit(x=data_in, y=data_out, batch_size=batchSize, epochs=10)

model.predict(np.array([data_in[0]]))

```

5.8.4.2 LSTM

Lstm's can be tricky in that they need input in a very specific format.

```

import numpy as np
import tensorflow as tf
import tensorflow.keras as k

# 1. Data: sine-curve
nr_samples = 10
nr_timesteps = 4
nr_pixels = 1
inpt = np.zeros([nr_samples, nr_timesteps, nr_pixels])
outpt = np.zeros([nr_samples, nr_pixels])
for sample in range(nr_samples):
    for timestep in range(nr_timesteps):
        for pixel in range(nr_pixels):
            inpt[sample, timestep, pixel] = (np.sin(timestep * 0.1 + sample) + 1.0) / 2.0
    for pixel in range(nr_pixels):
        outpt[sample, pixel] = (np.sin((nr_timesteps + 1) * 0.1 + sample) + 1.0) / 2.0

# 2. Network: simple lstm
model = k.Sequential([
    k.layers.LSTM(1) # 1 = output dimension. for every timeseries we pass in, we want the one next value
])

model.compile(
    optimizer=tf.train.AdamOptimizer(learning_rate=0.1),
    loss=tf.losses.mean_squared_error
)

model.fit(x=inpt, y=outpt, epochs=500)

model.predict(
    np.array([
        [0.5, 0.54991671, 0.59933467, 0.6477601]
    ])
)

```

5.8.4.3 Deep dream

Usually in tensorflow, we want to minimize the difference between a label and the network's output. Deep dreaming is somewhat different. Here we pick a layer that spots interesting shapes (in this case "layer7", let's say it spots cat-eyes) and maximize it's activation by repeatedly nudging the input image in direction of the gradient.

Listing 5.1: Pseudocode

```

# variables ending in "T" denote a tensor
# variables ending in "Val" denote a value of a tensor

inputImgVal = load_image("testImage.png")

model = VGG16("imagenet")

inputImgT = model.input
targetT = model.layers["layer7"].output
gradsT = K.gradients(targetT, inputImgT)[0]
pipe = K.function([inputImgT], [targetT, gradsT])

for i in range(10):
    targetVal, gradVal = pipe(inputImgVal)
    inputImgVal += alpha * gradVal

```

Really, the name deep dream is chosen somewhat badly. It should rather be "middle-deep interpretation" or "reading into things". Here is what an actual deep dream shouls look like in our opinion.

```

inputImgVal = load_image("randomNoise.png")
targetT = model.layers[-1].output

```

5.9 Angular

5.9.1 Elements

- *Element*: Wrapper around `nativeElement`. Wrapped so as to abstract away differences between native elements in browser and on mobile.
- *View*: Smallest group of elements that are created and destroyed together.
 - HostView: A view created by a component
 - EmbeddedView: A view created by a template
- *Component*
- *Directive*: An attribute that may be appended to a component. Like a component, but without an html-file.
- *Template*: Angular-html-files are usually static. The only dynamics that are allowed are via subcomponents or templates `<ng-template>`. Templates are really like simple, inline html-only subcomponents. They are even more like inline-functions in that they have access to their parent-component's template-variables.
- *ViewContainer*: removing html-elements from the DOM is dangerous, because angular will not be aware of that removal - it will keep the node in memory and continue running change detection. So how do you remove elements? That's what ViewContainers are for: wrapping a node in a ViewContainer provides us with an API to add or remove components or templates.

5.9.2 Life cycle

5.9.3 Change detection

Change detection goes top down: if a user-interaction happened on component X, then the whole tree starting at app-root is updated. If component X has `changeDetectionStrategy: OnPush`, then X and its children are shielded from updates from outside of X's subtree. But events originating from X or its children still traverse the *whole* tree from app-root down.

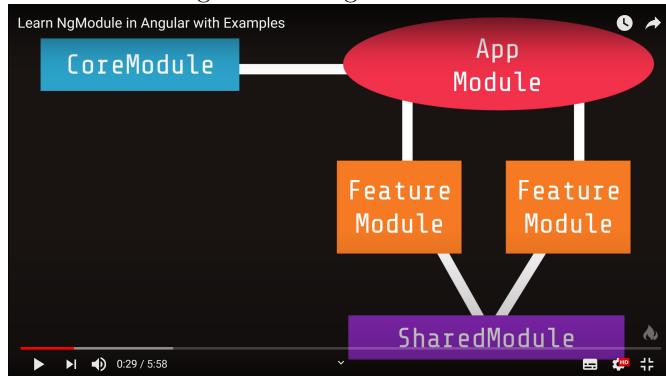
5.9.4 Service availability

Services are by default provided by app-root. The instance is used by every child component - except if that child, lets call it X, has its own provided instance of the service. Then all children of X use X's instance of the service.

5.9.5 Building blocks

- The *compiler*
 - Runs offline and is not included in the bundled application (except when in development mode).
 - Takes `*.component.ts` and `*.component.html` files and turns them into a `*.ngfactory.ts` file.
 - The `*.ngfactory.ts` file contains a `ComponentDef` class with a `factory` method and a `template` method.
- The *runtime*: `core.js/tick()`
 - Module bootstrap
 - View creation
 - * Instantiates a view by calling `const viewDef = compDef.factory(depInjector)`
 - A ViewDef has a list of nodes `viewDef.nodes[]`
 - Nodes can be accessed as `ElementRef`, `TemplateRef` or `ViewContainerRef`
 - A ViewDef has a `template()` method that calls render-functions like `text()`, `heading()`, `div`, ...
 - * Calls the ViewDef's template-function `viewDef.template()`
 - Change detection
 - *

Figure 5.1: Angular modules



The different parts of an angular application are namespaced in an interesting way.

- The *compiler* inspects html-files and writes components, directives and pipes into them
 - Components, directives and pipes are namespaced in *modules*
- The *injector* injects services and guards into components
 - Services and guards are not namespaced - they are always provided in one big pool in app-root.
 - If you really need to namespace services, there is the `forRoot()` hack.
- app module: root. `imports` other modules and `bootstrap`'s the root-component.
- core module: exports shared singleton services and guards in its `providers`. Should be imported only in app-module.
- feature module: exports thematic components, directives and pipes in its `declarations` and `exports`. Should be imported only in app-module. If necessary, feature modules may provide services, but then you need to make sure that the module is only imported in one single place (usually the app-module).
- shared module: exports globally used components, directives and pipes in its `declarations` and `exports`. Should be imported into any module that requires it. Shared modules *must not* export services, because they are imported in several other modules. The presence of the same service in multiple modules will have the injector commit harakiri.
- router module: imported into app module. Side-effect-configures the angular-core-router using `forRoot`.

Chapter 6

Sciences

Satellite	Bands	Data available at	Orbit	Data frequency	Description
Landsat					
Modis					
Copernicus					
Sentinel					
TerraSarX					
RapidEye					
Ikonos					

6.1 Remote sensing

6.1.1 Electromagnetic radiation

Satellites all work the same way: they are flying machines that carry a camera. The camera has one or more sensors, each of which can observe some part of the electromagnetic spectrum. Radar satellites (usually) additionally have their own 'light' source: they actively send out radio-radiation and collect the reflection.

6.1.2 Orbital periods and acquisition

Satellites usually have an orbital period of 1 - 2 hours. They revisit the same spot on earth every 2-3 weeks. (This example data is taken from landsat: an orbit takes 99 minutes, and it visits the same spot every 16 days. Note that landsat always stays on the sunny side of earth - most visible-light satellites do this.) Commonly, satellite series have offset periods so one machine of the same series visits the same spot every nth of a full orbital period. However, some satellites can narrow or broaden their sensors, so even if a satellite visits the same spot again, it might not make the same snapshot again. To handle the high demand and low availability of satellite images, acquisition plans are being made. Usually, these take into account seasonal effects, water vapor, and also a ranking of requests from different scientific institutions.

6.1.3 Important satellites and service providers

Landsat program is the longest-running enterprise for acquisition of satellite imagery of Earth. On July 23, 1972 the Earth Resources Technology Satellite was launched. This was eventually renamed to Landsat.[1] The most recent, Landsat 8, was launched on February 11, 2013. The instruments on the Landsat satellites have acquired millions of images. The images, archived in the United States and at Landsat receiving stations around the world, are a unique resource for global change research and applications in agriculture, cartography, geology, forestry, regional planning, surveillance and education, and can be viewed through the U.S. Geological Survey (USGS) 'EarthExplorer' website. Landsat 7 data has eight spectral bands with spatial resolutions ranging from 15 to 60 meters; the temporal resolution is 16 days.[2] Landsat images are usually divided into scenes for easy downloading. Each Landsat scene is about 115 miles long and 115 miles wide (or 100 nautical miles long and 100 nautical miles wide, or 185 kilometers long and 185 kilometers wide).

Modis ...

Copernicus ...

Sentinel ESA is currently developing seven missions under the Sentinel programme. The Sentinel missions include radar and super-spectral imaging for land, ocean and atmospheric monitoring. Each Sentinel mission is based on a constellation of two satellites to fulfill and revisit the coverage requirements for each mission, providing robust datasets for all Copernicus services.

CHIRPS is ...

EODC : Earth Observation Data Center for Water Resources Monitoring

Eurac : Eurac is a private research company ...

Google Earth Engine provides ...

NASA's ECS (Earth observation center Core System) is a vast catalogue of ...

6.1.4 Datatypes, protocols and libraries

Remote sensing data is usually massive in size and requires some special software to parse. This section is intended to give an overview of the different datatypes out there and the tools for reading them.

Datacubes are tensors of (usually very large) (geo-)data.

OGC-WCPS: OGC's Web Coverage Processing Service is a language for filtering and processing multi-dimensional raster coverages.

Open-EO ia an API to speak to different geo-processing services like GEE. Unfortunately, open-EO stands in concurrence to WCPS.

geoTiff files are geo-referenced images. They contain different bands. Every band is a 2d-array of 16-bit numbers, representing the image at one frequency-band. Usually, you'll have one band around the green color, one around the red, and one around the blue (plus many more bands at non-visible frequencies). When you color these three bands and put them together as an rgb-image, you can usually obtain a quite decent color-photo. Every geoTiff has an associated

attribute. This attribute is an affine transformation matrix and maps pixel-locations (row,col) to (x,y) spatial positions.

```

import rasterio

with rasterio.open("example.tif") as dataset:
    print(dataset.crs)
    print(dataset.transform)
    print(dataset.transform * (0, 0))
    print(dataset.transform * (dataset.width, dataset.height))

```

GeoPackage (GPKG) is an open, non-proprietary, platform-independent and standards-based data format for geographic information system implemented as a SQLite database container. Defined by the Open Geospatial Consortium (OGC)[1] with the backing of the US military[2] and published in 2014, GeoPackage has seen wide widespread support from various government, commercial, and open source organizations[3].

6.1.5 Projections

Wgs84 aka. Epsg4326 is a global spherical (aka non-projected, using lat/lon) CRS and the stuff that comes from GPS-satellites. GeoJson should by now only use Wgs84 in its data (it used to be allowed to add a field named "crs").

UTM aka. Epsg20** is a projected CRS using a two-dimensional cartesian grid in meters. Projections needs to be done on individual areas, which is why there several Epsg-codes for this projection.

Web Merkator aka. EPSG:900913 aka. Epsg3857 is the google projection and a variant of UTM.

6.1.6 NDVI

The normalized difference vegetation index (NDVI) is used to assess the state of vegetation. In living plants chlorophyll-A, from the photosynthetic machinery, strongly absorbs red color; on the other hand, near-infrared light is strongly reflected. Live, healthy vegetation reflects around 8The NDVI is calculated as:

$$NDVI = \frac{\lambda_{NIR} - \lambda_{red}}{\lambda_{NIR} + \lambda_{red}}$$

where:

- λ_{NIR} is near-infrared Band 5
- λ_{red} is color red Band 4

By its formulation, NDVI ranges from -1 to +1. In practice, an area of an image containing living vegetation will have NDVI in the range 0.3 to 0.8. High water content clouds and snow will have negative values of the index. Bodies of water, having low reflectance in both Band 4 and 5, exhibit very low positive or negative index. Soil, having slightly higher reflectance in near-infrared than in red, will produce low positive values of the index.

Calculation of the NDVI is sensitive to to a few factors:

- Atmospheric composition and appropriate modeling influence the calculation, especially if the correct water and aerosol content are initially incorrectly estimated.
- Thin, hard to spot clouds like cirrus can significantly affect the calculation.
- Sensor effects such as Sun angle not being calculated on per-pixel basis can influence the index estimation.

6.1.7 Image (co-)registration

Image registration is the process of transforming different sets of data into one coordinate system.

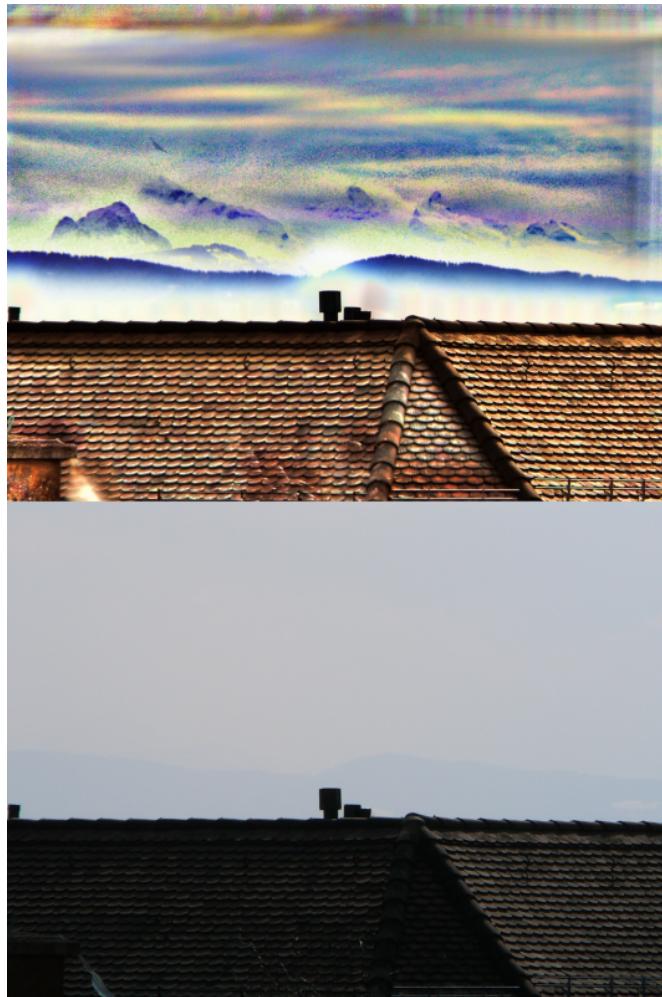


Figure 6.1: Registering and summing multiple exposures of the same scene improve signal to noise ratio, allowing one to see things previously impossible to see. In this picture, the distant Alps are made visible, although they are tens of kilometers into the haze.

6.2 Economics

A *preference relation* \succsim is a bivariate relation on a set X . It is *rational* if it is

- complete
- transitive

A *utility function* u is a function $u : X \rightarrow \mathbb{R}$ that has the property $\forall x, y \in X : x \succsim y \rightarrow u(x) \geq u(y)$. We can prove that

$$\succsim: \text{rational} \leftrightarrow \exists u : \text{utility function of } \succsim$$

6.3 Projektwesen

Glossar : einige Worte sind nicht ganz selbsterklärend

- Funding: Finanzierung
- Tender: Ausschreibung

Geldgeber

- EU
 - <https://ec.europa.eu/info/funding-tenders/opportunities/portal/screen/home>
 - <https://emergency.copernicus.eu/>

Chapter 7

Art

7.1 Design

Aesthetics is not just about making things stand out. It's about encoding information.

Empathy (UXD) means understanding whom you are working for - how they interpret your choice of color.

7.1.1 Layout

When designing something, just make a note for each of the following points, justifying your choice. This way you can be sure that you did not overlook a useful design-decision. Common layout-frameworks are bootstrap, clarity, 960, foundation, and neat.io.

Grids consist of columns, gutter, modules, negative space, margins, and baseline. Some common print-layouts are manuscript, columns and blocks.

Rhythm states that, no matter the width/font-size/color of any element, its width/text-height/hue is always a multiple of some base-unit. Pick elements/font-sizes/colors according to some scheme, like fibonacci, magic square, etc. Note, however: just like in music, people only recognize patterns when they are repeated. When implementing a pattern, it should appear more than just once on the screen.

Hierarchy : important elements are bigger and/or have negative space around them. Important items should be close to each other - don't force the users eyes to jump.

Colors change with time. Currently, the following colorschemes are popular: Material Design, pastels with crisp photography, greyscale with one accent color. Use color only when there is little text on the page. The more text, the more you should cut back on color.

7.1.2 UI

First commandment of UI : Per default, show as much as necessary. But allow a user to see his specific points of interest when he wants it.

7.2 Math

7.2.1 Points

Distance in 1-D:

$$d = \Delta x$$

Distance between two points in 2-D:

$$d = \sqrt{\Delta x^2 + \Delta y^2}$$

Distance between two points in 3-D:

$$d = \sqrt{\sqrt{\Delta x^2 + \Delta y^2}^2 + \Delta z^2} = \sqrt{\Delta x^2 + \Delta y^2 + \Delta z^2}$$

7.2.2 Geometric objects in matrix notation

$x = 1$ defines a plane, as does $x = 2y$. $x = 1 \wedge y = 2$ defines a line. More generally, the dimension of a geometrical object equals 3 minus the number of equations needed to describe the object. Informally written: $\dim(obj) = 3 - \#\{=\}$.

In more rigorous terms, the "number of equations needed to describe the object" is called the rank of the matrix. 3 in our case is the dimension of the space. The geometrical object is really the column-space of the matrix.

7.2.2.1 Plane

Defined by one point \vec{b} and two lines \vec{p}_1, \vec{p}_2 :

$$P = \{\vec{x} | \exists a, b : \vec{x} = \vec{b} + a\vec{p}_1 + b\vec{p}_2\}$$

Or defined by one point \vec{b} and one normal \vec{n} :

$$P = \{\vec{x} | (\vec{x} - \vec{b}) \cdot \vec{n} = 0\}$$

7.2.2.2 Line

Defined by two points:

$$L = \{\vec{x} | \exists \alpha : \vec{x} = \vec{b} + \alpha \cdot \vec{\Delta}\}$$

Or transformed in matrix-notation:

$$\vec{x} = [\vec{b} \quad \vec{\Delta}] \cdot \begin{bmatrix} 1 \\ \alpha \end{bmatrix}$$

$$[\vec{b} \quad \vec{\Delta}]^{-1} \cdot \vec{x} = \begin{bmatrix} 1 \\ \alpha \end{bmatrix}$$

7.2.2.3 Intersection Line/Plane

Substituting the definition of a line into the definition of a plane we get:

$$(\vec{b}_l + \alpha \cdot \vec{\Delta} - \vec{b}_p) \cdot \vec{n} = 0$$

This reduces to:

$$\alpha = \frac{(\vec{b}_l - \vec{b}_p) \cdot \vec{n}}{\vec{\Delta} \cdot \vec{n}}$$

Alternatively, we can make use of the two-line-definition of the plane and obtain:

$$\vec{b}_p + a\vec{p}_1 + b\vec{p}_2 = \vec{b}_l + \alpha \cdot \vec{\Delta}$$

$$\vec{b}_l - \vec{b}_p = [\vec{p}_1 \quad \vec{p}_2 \quad \vec{\Delta}] \begin{bmatrix} a \\ b \\ \alpha \end{bmatrix}$$

7.2.3 Projections

7.2.3.1 Finding a perpendicular

Having a vector \vec{a} , how do we find a vector \vec{a}^T that has the following properties:

- $\vec{a} \cdot \vec{a}^T = 0$
- $|\vec{a}| = |\vec{a}^T|$

These two requirements can be re-expressed as

- $ax + by = 0$
- $x^2 + y^2 = a^2 + b^2$


This is a nonlinear system, but still solveable. Indeed, we quickly find:

$$\vec{a}^T = \begin{bmatrix} -y \\ x \end{bmatrix}$$

How about the two perpendiculars to a 3d-vector? You can just use $\begin{bmatrix} -y & x & 0 \end{bmatrix}$ and $\begin{bmatrix} 0 & -z & y \end{bmatrix}$. All possible perpendiculars will be linear combinations of these two.

7.2.3.2 Projecting one vector onto another

Imagine we wanted to know α in the following graphic:



Here, α is the length of vector \vec{b} projected onto \vec{a} . How can we find α ?

The solution lies in imagingen the vector \vec{b} being expressed inside a coordinate-system consisting of \vec{a} and \vec{a}^T , like so:

$$[\vec{a}, \vec{a}^T] \begin{bmatrix} \alpha \\ \beta \end{bmatrix} = \vec{b}$$

7.2.3.3 Flattening: expressing the points on a 3d-plane in a 2d-system

From a screen placed in 3d-space we want to obtain a 2d-image that can be displayed on a computer. Consider the plane

$$P = \vec{x} | \vec{x} = \vec{b} + \alpha \vec{p}_1 + \beta \vec{p}_2$$

A point \vec{x}_0 on the screen can be expressed as

$$\vec{x}_0 = \vec{b} + \alpha \vec{p}_1 + \beta \vec{p}_2$$

In that case, we just use α and β as coordinates. If instead we were using

$$P = \vec{x} | (\vec{x} - \vec{b}) \vec{n} = 0$$

we can still obtain \vec{p}_1 and \vec{p}_2 as

$$\vec{p}_1 = \vec{z} \tag{7.1}$$

$$\vec{p}_2 = \vec{n} \times \vec{z} \tag{7.2}$$

7.2.4 Rotation

7.2.4.1 In 2d

To find the rotation matrix \mathbf{R}_2 in 2d, you could take the naive approach by taking a generic vector $\begin{bmatrix} x \\ y \end{bmatrix}$ and try to solve this equation:

$$r = \sqrt{x^2 + y^2} \quad (7.3)$$

$$\theta_0 = \cos^{-1}(x/r) \quad (7.4)$$

$$\mathbf{R}_2 \begin{bmatrix} x \\ y \end{bmatrix} = \begin{bmatrix} r \sin(\theta + \theta_0) \\ r \cos(\theta + \theta_0) \end{bmatrix} \quad (7.5)$$

While this is a feasible approach, we can have it much easier. Remember from 1.5.2 that

$$\vec{v} = \mathbf{T}\vec{v}_T$$

If \mathbf{T} is invertible:

$$\vec{v}_T = \mathbf{T}^{-1}\vec{v}$$

Where:

$$\mathbf{T}^{-1} = [\vec{x}_T, \vec{y}_T, \dots]$$

Specifically, in our case this results in

$$\mathbf{R}_2 = \begin{bmatrix} \cos \theta & -\sin \theta \\ \sin \theta & \cos \theta \end{bmatrix}$$

7.2.4.2 In 3d

While in 2d we only had one axis to rotate around, in 3d we can chose if we want to rotate around x, y, or z. In the general case:

$$\mathbf{R}_3 = \mathbf{R}_x \mathbf{R}_y \mathbf{R}_z$$

. It is easy to obtain \mathbf{R}_z , the rotation around the z-axis. Just imagine you viewed everything from above:

$$\mathbf{R}_x = \begin{bmatrix} \cos \theta & -\sin \theta & 0 \\ \sin \theta & \cos \theta & 0 \\ 0 & 0 & 1 \end{bmatrix}$$

This way we obtain

$$\begin{bmatrix} \cos(\beta) \cos(\gamma) & -\sin(\gamma) \cos(\beta) & \sin(\beta) \\ \sin(\alpha) \sin(\beta) \cos(\gamma) + \sin(\gamma) \cos(\alpha) & -\sin(\alpha) \sin(\beta) \sin(\gamma) + \cos(\alpha) \cos(\gamma) & -\sin(\alpha) \cos(\beta) \\ \sin(\alpha) \sin(\gamma) - \sin(\beta) \cos(\alpha) \cos(\gamma) & \sin(\alpha) \cos(\gamma) + \sin(\beta) \sin(\gamma) \cos(\alpha) & \cos(\alpha) \cos(\beta) \end{bmatrix}$$

7.2.4.3 Why is there one rotation axis in 2d, but 3 in 3d?

The rotation axis for any pair of axii that define a plane in space is a third axis perpendicular to that plane. In conventional 2D (XY plane), all rotation is done in the Z axis. But what about spherical coordinates? There we need only two angles! ...

7.2.5 Implicit versus parameterized representation of bodies

In geometry, we often represent bodies in set-notation, aka. implicit notation. Consider for example the ellipsoid:

$$\left\{ \vec{v} \mid \frac{x^2}{r_1^2} + \frac{y^2}{r_2^2} + \frac{z^2}{r_3^2} = 1 \right\}$$

This body can be parameterized as:

$$\begin{bmatrix} x \\ y \\ z \end{bmatrix} = \begin{bmatrix} r_1 \cos \theta \cos \phi \\ r_2 \cos \theta \sin \phi \\ r_3 \sin \theta \end{bmatrix}$$

In the set notation, the ellipsoid had no parameters (it had r_1, r_2, r_3 , but these are *constants*). With θ and ϕ we now found two parameters which, when varied over their whole domain $[0^\circ, 360^\circ]^2$, yield every point in the set.

Parameterisation means *finding a function of one or more parameters whose range equals the set*. That means nothing other than *putting the vector \vec{x} on the left side*. Conversely, you have an *implicit* equation when you can write the equation such that the left side is 0 (for clarity: we mean the number 0, not a zero-vector).

As a sidenote, when you find a set-expression of a body that contains a *thereis* statement, chances are that this statement already *is* parameterized.

Table 7.1: Explicit and parametric set of plane and ellipsoid

	implicit	parametric
linear	$\{\vec{x} (\vec{x} - \vec{b}) \cdot \vec{n} = 0\}$	$\left\{ \vec{x} \exists \alpha, \beta : \vec{x} = \begin{bmatrix} \vec{b} & \vec{p}_1 & \vec{p}_2 \end{bmatrix} \begin{bmatrix} 1 \\ \alpha \\ \beta \end{bmatrix} \right\}$
non-linear	$\left\{ \vec{x} \frac{x^2}{r_x} + \frac{y^2}{r_y} + \frac{z^2}{r_z} = 1 \right\}$	$\left\{ \vec{x} \exists \theta, \phi : \vec{x} = \begin{bmatrix} r_x \cos \theta \cos \phi \\ r_y \cos \theta \sin \phi \\ r_z \sin \theta \end{bmatrix} \right\}$

We can find more about this topic here: https://en.wikipedia.org/wiki/Implicit_surface. In this section, we have been careful to never mention the term *explicit*. In this section, we dealt with bodies. But if we only wanted surfaces, that is, bodies where there are never two z values at on x/y spot, there is also the explicit representation $z = f(x, y)$.

Table 7.2: Overview of implicit, explicit and parameterized functions

	in	out	notes
implicit	x, y, z		
explicit	x, y	z	doesn't allow more than one z per x/y
parameterized	θ, ϕ	x, y, z	

7.2.6 Maps

7.2.6.1 Projections

7.2.6.2 Zoom to width/height

Zoom is defined as follows: assume the whole world is displayed in a tile of $res * 360$ by $res * 180$ pixels. This tile has zoomlevel 0.

$$h = 1800.5^z$$

$$w = 3600.5^z$$

7.2.6.3 Apparent height

Imagine you were a bird flying over a flat map, your head pointing north and facing down. Your apparent height based on the maps width would be:

$$h = \frac{w}{2} \tan(\theta/2)$$

technology	use-case	programming style	
canvas/2d			single-buffering to bitmap. paintbrush-statemachine.
canvas/webgl	3d, interactive, many items	procedural	double-buffering to bitmap
svg + css	simple 2d graphics	declarative	no bitmap, but \inlinetoken{<svg>}. has concept of layers

7.3 Graphics Programming

7.3.1 Rendering in the browser

7.3.2 SVG

SVG consists of

- Objects
- Groups. Groups have ...
 - a transform
 - a style
 - Notably, they don't have a x, y, width or height attribute! If you need those, use SVG's instead
- SVG's.
 - x, y width, height


Objects have a nice, simple hierarchy.

- Objects: each object is a ...
 - Path: the most basic object: a Bezier-curve. Each of the following subtypes can be downgraded to a path again in inkscape.
 - * rect, star, ellipse, text
- Each object has a ...
 - fill, stroke, opacity

A bunch of software helps with SVG's:

- Inkscape: ideal for creating SVG logos
- d3: Data-driven, animated SVG's
- Raphael: animated SVG's; imperative, *not* data-driven

7.3.3 D3



```

<html>
<head>
    <script src="https://d3js.org/d3.v6.js"></script>
</head>
<body>
    <h1>Demonstrating the selection, enter & exit flow for animations</h1>

    <ul id="list">
        <li>one</li>
        <li>two</li>
    </ul>

    <p>
        To explore this behavior, call `displayData` a few times from the web-console.
    </p>
</body>
<script>
    function displayData(data) {
        // step 1: create element- & data-selections (= the blue and the red circles)
        const elsNdata = d3.select('#list') // selection of parent element
            .selectAll('li') // select all child elements: [li, li]
            .data(data);

        // step 2: handle li's that already have (possibly old) data on them
        elsNdata
            .text(d => d);
        // usually, there are initially no pre-defined li's, such that this step is not required.
        // this is good for performance: not handling old items saves on dom-operations.

        // step 3: handle new data points coming in
        elsNdata.enter() // all data points that have not yet got an associated li
            .append('li')
            .text(d => d);

        // step 4: handle old data points going out
        elsNdata.exit() // all li's that no longer have data on them
            .remove();
    }

    displayData([1, 2, 3])
</script>
</html>

```

7.3.4 WASM

- Module: compiled WASM. Can be passed around. Stateless.
- Instance: a module + it's state
- State:
 - Memory: resizable ArrayBuffer
 - Table: array of references to functions
- Calls from js to wasm are synchronous.
- LLVM: compiles C to intermediate (via the Clang-Frontend) and intermediate to WASM (via the WASM-Backend)
- Emscripten: includes LLVM. But also creates custom JS to be used by WASM to emulate `stdio` et.al.
 - Really, this includes all of `libc`, including `malloc()` and `free()`.

```

int add (int a, int b) {
    return a*a + b;
}

float int_arraySum (int* array, int length) {
    float total = 0;
    for (int i = 0; i < length; i++) {
        total += array[i];
    }
    return total;
}

void int_arrayAddition (int* arr1, int* arr2, int* arrSum, int arrLength) {

```

```

    for (int i = 0; i < arrLength; i++) {
        arrSum[i] = arr1[i] + arr2[i];
    }

float float_arraySum (float* array, int length) {
    float total = 0;
    for (int i = 0; i < length; i++) {
        total += array[i];
    }
    return total;
}

void float_arrayAddition (float* arr1, float* arr2, float* arrSum, int arrLength) {
    for (int i = 0; i < arrLength; i++) {
        arrSum[i] = arr1[i] + arr2[i];
    }
}

|| clang --target=wasm32 -O3 -flto -nostdlib -Wl,--no-entry -Wl,--export-all -Wl,--lto-O3 -o prog.wasm prog.c

const memory = new WebAssembly.Memory({
    initial: 100, // in pages (64KiB / Page)
    maximum: 1000
});

WebAssembly.instantiateStreaming(fetch('prog.wasm'), {
    js: {
        mem: memory
    }
}).then((response) => {
    console.log(response);

    // simple function
    const result1 = response.instance.exports.add(4, 1);
    console.log('4**2 + 1', result1);

    // pointer to top of memory array
    let offsetBytes = 0;
    const memBuf = response.instance.exports.memory.buffer;

    // wasm accesses js data
    const data1 = [1, 2, 3, 4.4, 5];
    const arr1 = new Float32Array(memBuf, offsetBytes, data1.length);
    arr1.set(data1);
    offsetBytes += arr1.length * arr1.BYTES_PER_ELEMENT;

    const result2 = response.instance.exports.float_arraySum(arr1.byteOffset, arr1.length);
    console.log('sum of ${arr1}', result2);

    // js accesses wasm data
    const l = 3;

    const data2 = [1, 2, 3.1];
    const arr2 = new Float32Array(memBuf, offsetBytes, 1);
    arr2.set(data2);
    offsetBytes += arr2.length * arr2.BYTES_PER_ELEMENT;

    const data3 = [4, 5.1, 6];
    const arr3 = new Float32Array(memBuf, offsetBytes, 1);
    arr3.set(data3);
    offsetBytes += arr3.length * arr3.BYTES_PER_ELEMENT;

    const arr4 = new Float32Array(memBuf, offsetBytes, 1);
    offsetBytes += arr4.length * arr4.BYTES_PER_ELEMENT;

    response.instance.exports.float_arrayAddition(arr2.byteOffset, arr3.byteOffset, arr4.byteOffset, 1);
    console.log('addition of ${arr2} and ${arr3} = ', arr4)

});
}

```

7.3.5 WebGL

```

import { flattenMatrix } from './engine.shapes';

/**
 * WEBGL
 *
 * A rasterization engine that allows to draw points, line segments, or triangles.
 *
 * Vertex shaders take whatever coordinates you use and return a 3-d array with elements between -1 and 1.
 * Basically, this is a 3d-array, but WebGL does not use the z-axis for real perspective, but only to differentiate
 * what pixel lies in front of another.
 * This is not like looking in a 3d-box, but rather like looking on multiple stacked sheets on a projector.
 * Actually, this is a lie. WebGL uses 4 coordinates: x, y, z and w. The above only holds if you keep w at 1.
 * After applying the vertex shader, WebGL divides all coordinates by w, yielding (x/w, y/w, z/w, 1).
 * This can be used to calculate projections - google for 'homogeneous coordinates' to find out more.
 * Compare this [site](https://www.todalling.com/blog/modern-opengl/explaining-homogenous-coordinates-and-projective)
 * and the shader 'basic3d.vert.glsl'.
 *
 * WebGL knows two data structures:
 * - buffers (generic byte arrays): usually positions, normals, texture-coordinates, vertex-colors etc.
 *   buffers are accessed in shaders as 'attributes'.
 *   note that buffers contain one entry for each vertex.
 * - textures (bitmap images).
 *
 * Shaders use these data structures in two different ways:
 * - Attributes are values, one per vertex.
 *   For the shader, attributes are read-only.
 *   Attributes default to [0, 0, 0, 1]
 * - Uniforms are values, one per shader.
 *   For the shader, uniforms are read-only.
 *
 * Apart from this, shaders know about two more types of data:
 * - Varyings are values that are passed from vertex-shader to fragment-shader.
 *   They are read-only only for the fragment-shader.
 * - Const: a compile-time constant.
 *
 * A program is just a list of compiled and linked vertex- and fragment-shaders.
 *
 *
 * Drawing: there's drawArrays and drawElements.
 * - drawArrays is the robust all-rounder.
 * - drawElements can be more performant if you share vertices between objects.
 *
 *
 * Rendering data is fast, but uploading it into GPU memory is slow.
 * Optimizing WebGL performance mostly means: Avoiding having GPU and CPU wait for each other.
 * The more the GPU can do in bulk, the better. The more often you have to upload data from CPU to GPU, the worse.
 * - So avoid switching programs, buffers and uniforms if you can.
 *   (You won't be able to avoid switching buffers, because every object is likely different. But sort your objects)
 * - Try to do translations, rotations and shears inside the vertex-shader instead of altering the object's buffer.
 * - If appropriate, create ueber-shaders and ueber-buffers, that contain information for more than just one object.
 *
 * There is another thing that affects performance:
 * WebGL will only run fragment-shaders when the object's pixels aren't already obscured by a larger object in front
 * That means it makes sense to first draw large objects that are close to the camera - all objects behind them won't
 *
 * All 'create*' functions unbind variables after setting their values. This is to avoid unwanted side-effects.
 */

const shaderInputTextureBindPoint = 0;
const textureConstructionBindPoint = 7;




/***
 * Compile shader.
 */
export const compileShader = (gl: WebGLRenderingContext, typeBit: number, shaderSource: string): WebGLShader => {
    const shader = gl.createShader(typeBit);
    if (!shader) {
        throw new Error('No shader was created');
    }
    gl.shaderSource(shader, shaderSource);
    gl.compileShader(shader);
    if (!gl.getShaderParameter(shader, gl.COMPILE_STATUS)) {
        gl.deleteShader(shader);
        throw new Error(`An error occurred compiling the shader: ${gl.getShaderInfoLog(shader)}. \n\n Shader code: ${shaderSource}`);
    }
    return shader;
};

```

```

    /**
     * Note that every program *must* have one and only one vertex-shader
     * and one and only one fragment shader.
     * That means you cannot add multiple fragment-shaders in one program. Instead, either load them in consecutively as
     * or generate an ueber-shader that contains both codes.
     */
    export const createShaderProgram = (gl: WebGLRenderingContext, vertexShaderSource: string, fragmentShaderSource: string): WebGLProgram => {
        const program = gl.createProgram();
        if (!program) {
            throw new Error('No program was created');
        }

        const vertexShader = compileShader(gl, gl.VERTEX_SHADER, vertexShaderSource);
        const fragmentShader = compileShader(gl, gl.FRAGMENT_SHADER, fragmentShaderSource);
        gl.attachShader(program, vertexShader);
        gl.attachShader(program, fragmentShader);

        gl.linkProgram(program);

        gl.detachShader(program, vertexShader);
        gl.detachShader(program, fragmentShader);
        gl.deleteShader(vertexShader);
        gl.deleteShader(fragmentShader);

        if (!gl.getProgramParameter(program, gl.LINK_STATUS)) {
            gl.deleteProgram(program);
            throw new Error('Unable to initialize the shader program: ' + gl.getProgramInfoLog(program));
        }

        return program;
    };

    export const setup3dScene = (gl: WebGLRenderingContext): void => {
        gl.viewport(0, 0, gl.canvas.width, gl.canvas.height); // making sure that shader-coordinate-system goes from 0 to width/height

        gl.enable(gl.DEPTH_TEST);
        gl.depthFunc(gl.LEQUAL);
        gl.cullFace(gl.BACK);

        clearBackground(gl, [0, 0, 0, 1]);
    };

    export const bindProgram = (gl: WebGLRenderingContext, program: WebGLProgram): void => {
        gl.useProgram(program);
    };

    export const clearBackground = (gl: WebGLRenderingContext, color: number[]): void => {
        gl.clearColor(color[0], color[1], color[2], color[3]);
        gl.clearDepth(1.0);
        gl.clear(gl.COLOR_BUFFER_BIT | gl.DEPTH_BUFFER_BIT);
    };

    /**
     * A generic buffer, together with it's metadata.
     */
    export interface BufferObject {
        buffer: WebGLBuffer;
        vectorSize: number;
        vectorCount: number;
        type: number;
        normalize: boolean;
        stride: number;
        offset: number;
        drawingMode: number; // gl.TRIANGLES, gl.POINTS, or gl.LINES
    }
}

/**
 * Create buffer. Creation is slow! Do *before* render loop.
 */
export const createFloatBuffer = (gl: WebGLRenderingContext, data: number[][][], drawingMode: number = gl.TRIANGLES): WebGLBuffer => {
    const dataFlattened = new Float32Array(flattenMatrix(data));

    const buffer = gl.createBuffer();
    if (!buffer) {
        throw new Error('No buffer was created');
    }
    gl.bindBuffer(gl.ARRAY_BUFFER, buffer);
}

```

```

gl.bufferData(gl.ARRAY_BUFFER, dataFlattened, gl.STATIC_DRAW);
// STATIC_DRAW: tells WebGL that we are not likely to change this data much.
gl.bindBuffer(gl.ARRAY_BUFFER, null); // unbinding

const bufferObject: BufferObject = {
  buffer,
  vectorSize: data[0].length,
  vectorCount: data.length,
  type: gl.FLOAT, // the data is 32bit floats
  normalize: false, // don't normalize the data
  stride: 0, // 0 = move forward size * sizeof(type) each iteration to get the next position. Only change
  offset: 0, // start at the beginning of the buffer. Only change this in very-high-performance jobs.
  drawingMode: drawingMode
};

return bufferObject;
};

export const drawArray = (gl: WebGLRenderingContext, bo: BufferObject): void => {
  gl.drawArrays(bo.drawingMode, bo.offset, bo.vectorCount);
};

export const updateBufferData = (gl: WebGLRenderingContext, bo: BufferObject, newData: number[][]): BufferObject => {
  const dataFlattened = new Float32Array(flattenMatrix(newData));

  gl.bindBuffer(gl.ARRAY_BUFFER, bo.buffer);
  gl.bufferData(gl.ARRAY_BUFFER, dataFlattened, gl.STATIC_DRAW);
  gl.bindBuffer(gl.ARRAY_BUFFER, null); // unbinding

  const newBufferObject: BufferObject = {
    buffer: bo.buffer,
    vectorSize: newData[0].length,
    vectorCount: newData.length,
    type: gl.FLOAT, // the data is 32bit floats
    normalize: false, // don't normalize the data
    stride: 0, // 0 = move forward size * sizeof(type) each iteration to get the next position. Only change
    offset: 0, // start at the beginning of the buffer. Only change this in very-high-performance jobs.
    drawingMode: bo.drawingMode,
  };

  return newBufferObject;
};

/**
 * Fetch attribute's location (attribute declared in some shader). Slow! Do *before* render loop.
 */
export const getAttributeLocation = (gl: WebGLRenderingContext, program: WebGLProgram, attributeName: string): number => {
  const loc = gl.getAttributeLocation(program, attributeName);
  if (loc === -1) {
    throw new Error(`Couldn't find attribute ${attributeName} in program.`);
  }
  return loc;
};

/**
 * Attributes vary from vertex to vertex - that means that there are *many* of them.
 * So it makes sense for WebGL to store attribute values in a dedicated data structure - the buffer.
 */
export const bindBufferToAttribute = (gl: WebGLRenderingContext, attributeLocation: number, bufferObject: BufferObject): void => {
  // Enable editing
  gl.enableVertexAttribArray(attributeLocation);
  // Bind buffer to ARRAY_BUFFER
  gl.bindBuffer(gl.ARRAY_BUFFER, bufferObject.buffer);
  // Bind the buffer currently at ARRAY_BUFFER to a vertex-buffer-location.
  gl.vertexAttribPointer(
    attributeLocation,
    bufferObject.vectorSize, bufferObject.type, bufferObject.normalize, bufferObject.stride, bufferObject.offset),
  // gl.disableVertexAttribArray(attributeLocation); <-- must not do this!
};

export interface IndexBufferObject {
  buffer: WebGLBuffer;
}

```

```

        count: number;
        type: number; // must be gl.UNSIGNED_SHORT
        offset: number;
        drawingMode: number; // gl.TRIANGLES, gl.POINTS, or gl.LINES
    }

    export const createIndexBuffer = (gl: WebGLRenderingContext, indices: number[][][], drawingMode: number = gl.TRIANGLES) {
        const indicesFlattened = new Uint16Array(flattenMatrix(indices));

        const buffer = gl.createBuffer();
        if (!buffer) {
            throw new Error('No buffer was created');
        }
        gl.bindBuffer(gl.ELEMENT_ARRAY_BUFFER, buffer);
        gl.bufferData(gl.ELEMENT_ARRAY_BUFFER, indicesFlattened, gl.STATIC_DRAW);
        gl.bindBuffer(gl.ELEMENT_ARRAY_BUFFER, null);

        const bufferObject: IndexBufferObject = {
            buffer,
            count: indicesFlattened.length,
            type: gl.UNSIGNED_SHORT,
            offset: 0,
            drawingMode: drawingMode
        };

        return bufferObject;
    };

    export const drawElements = (gl: WebGLRenderingContext, ibo: IndexBufferObject): void => {
        gl.bindBuffer(gl.ELEMENT_ARRAY_BUFFER, ibo.buffer); // @TODO: optimize this to somewhere else?
        gl.drawElements(ibo.drawingMode, ibo.count, ibo.type, ibo.offset);
    };

    export const draw = (gl: WebGLRenderingContext, b: BufferObject | IndexBufferObject): void => {
        if (b.type === gl.UNSIGNED_SHORT) {
            drawElements(gl, b as IndexBufferObject);
        } else {
            drawArray(gl, b as BufferObject);
        }
    };
}

export interface TextureObject {
    texture: WebGLTexture;
    width: number;
    height: number;
    level: number;
    internalformat: number;
    format: number;
    type: number;
}

/**
 * A shader's attributes get their buffer-values from the VERTEX_ARRAY, but they are constructed in the ARRAY_BUFFER.
 * Textures analogously are served from the TEXTURE_UNITS, while for construction they are bound to ACTIVE_TEXTURE.
 *
 * There is a big difference, however. Contrary to buffers which receive their initial value while still outside the array,
 * a texture does already have to be bound into the TEXTURE_UNITS when it's being created.
 * Since it'll always be bound into the slot that ACTIVE_TEXTURE points to, you can inadvertently overwrite another texture
 * currently in this place. To avoid this, we provide a dedicated 'textureConstructionBindPoint'.
 *
 * Buffers are easier in this, since with vertexAttribPointer we are guaranteed to get a slot in the VERTEX_ARRAY that
 * is not already occupied by another buffer.
 */
export const createTexture = (gl: WebGLRenderingContext, image: HTMLImageElement | HTMLCanvasElement): TextureObject => {
    const texture = gl.createTexture(); // analog to createBuffer
    if (!texture) {
        throw new Error('No texture was created');
    }
    gl.activeTexture(gl.TEXTURE0 + textureConstructionBindPoint); // so that we don't overwrite another texture in the array
    gl.bindTexture(gl.TEXTURE_2D, texture); // analog to bindBuffer. Binds texture to currently active texture-bindpoint
    gl.texImage2D(gl.TEXTURE_2D, 0, gl.RGBA, gl.RGBA, gl.UNSIGNED_BYTE, image); // analog to bufferData
    gl.generateMipmap(gl.TEXTURE_2D); // mipmaps are mini-versions of the texture.
    gl.bindTexture(gl.TEXTURE_2D, null); // unbinding
}

let w, h: number;
if (image instanceof HTMLImageElement) {
    w = image.naturalWidth;
}

```

```

        h = image.naturalHeight;
    } else {
        w = image.width;
        h = image.height;
    }

    const textureObj: TextureObject = {
        texture: texture,
        level: 0,
        internalformat: gl.RGBA,
        format: gl.RGBA,
        type: gl.UNSIGNED_BYTE,
        width: w,
        height: h
    };

    return textureObj;
};

export const createEmptyTexture = (gl: WebGLRenderingContext, width: number, height: number): TextureObject => {
    if (width <= 0 || height <= 0) {
        throw new Error('Width and height must be positive.');
    }
    const texture = gl.createTexture();
    if (!texture) {
        throw new Error('No texture was created');
    }
    gl.activeTexture(gl.TEXTURE0 + textureConstructionBindPoint); // so that we don't overwrite another texture in the
    gl.bindTexture(gl.TEXTURE_2D, texture);
    gl.texImage2D(gl.TEXTURE_2D, 0, gl.RGBA, width, height, 0, gl.RGBA, gl.UNSIGNED_BYTE, null);
    gl.texParameteri(gl.TEXTURE_2D, gl.TEXTURE_WRAP_S, gl.CLAMP_TO_EDGE);
    gl.texParameteri(gl.TEXTURE_2D, gl.TEXTURE_WRAP_T, gl.CLAMP_TO_EDGE);
    gl.texParameteri(gl.TEXTURE_2D, gl.TEXTURE_MIN_FILTER, gl.NEAREST);
    gl.texParameteri(gl.TEXTURE_2D, gl.TEXTURE_MAG_FILTER, gl.NEAREST);
    gl.bindTexture(gl.TEXTURE_2D, null);

    const textureObj: TextureObject = {
        texture: texture,
        level: 0,
        internalformat: gl.RGBA,
        format: gl.RGBA,
        type: gl.UNSIGNED_BYTE,
        width: width,
        height: height
    };

    return textureObj;
};

/**
 * Even though we reference textures as uniforms in a fragment shader, assigning an actual texture-value to that uniform
 * Normal uniforms have a concrete value.
 * Texture uniforms, on the other hand, are just an integer-index that points to a special slot in the GPU memory (the
 */
export const bindTextureToUniform = (gl: WebGLRenderingContext, texture: WebGLTexture, bindPoint: number, uniformLocation: number): void {
    if (bindPoint > gl.getParameter(gl.MAX_COMBINED_TEXTURE_IMAGE_UNITS)) {
        throw new Error(`There are only ${gl.getParameter(gl.MAX_COMBINED_TEXTURE_IMAGE_UNITS)} texture bind points, but you
    }
    if (bindPoint === textureConstructionBindPoint) {
        console.error(`You are about to bind to the dedicated texture-construction bind point (nr. ${bindPoint}). If after
            Consider using another bind point.`);
    }
    gl.activeTexture(gl.TEXTURE0 + bindPoint); // analog to enableVertexAttribArray
    gl.bindTexture(gl.TEXTURE_2D, texture); // analog to bindBuffer. Binds texture to currently active texture-bindpoint
    gl.uniform1i(uniformLocation, bindPoint); // analog to vertexAttribPointer
};

export const updateTexture = (gl: WebGLRenderingContext, to: TextureObject, newImage: HTMLImageElement | HTMLCanvasElement): void {
    gl.activeTexture(gl.TEXTURE0 + textureConstructionBindPoint); // so that we don't overwrite another texture in the
    gl.bindTexture(gl.TEXTURE_2D, to.texture); // analog to bindBuffer. Binds texture to currently active texture-bindpoint
    gl.texImage2D(gl.TEXTURE_2D, 0, gl.RGBA, gl.RGBA, gl.UNSIGNED_BYTE, newImage); // analog to bufferData
    gl.generateMipmap(gl.TEXTURE_2D); // mipmaps are mini-versions of the texture.
    gl.bindTexture(gl.TEXTURE_2D, null); // unbinding
}

let w, h: number;
if (newImage instanceof HTMLImageElement) {

```

```

        w = newImage.naturalWidth;
        h = newImage.naturalHeight;
    } else {
        w = newImage.width;
        h = newImage.height;
    }
    to.width = w;
    to.height = h;

    return to;
};

export interface FramebufferObject {
    framebuffer: WebGLFramebuffer;
    texture: TextureObject;
    width: number;
    height: number;
}

export const createFramebuffer = (gl: WebGLRenderingContext): WebGLFramebuffer => {
    const fb = gl.createFramebuffer(); // analog to createBuffer
    if (!fb) {
        throw new Error('Error creating framebuffer');
    }
    return fb;
};

/**
 * The operations 'clear', 'drawArrays' and 'drawElements' only affect the currently bound framebuffer.
 */
export const bindFramebuffer = (gl: WebGLRenderingContext, fbo: FramebufferObject) => {
    gl.bindFramebuffer(gl.FRAMEBUFFER, fbo.framebuffer);
    gl.viewport(0, 0, fbo.width, fbo.height);
    // Note that binding the framebuffer does *not* mean binding its texture. In fact, if there is a bound texture,
    // Therefore, a framebuffer's texture must not be bound when the framebuffer is.
};

export const bindOutputCanvasToFramebuffer = (gl: WebGLRenderingContext) => {
    gl.bindFramebuffer(gl.FRAMEBUFFER, null);
    gl.viewport(0, 0, gl.canvas.width, gl.canvas.height);
};

/**
 * A framebuffer can have a texture - that is the bitmap that the shader-*out*put is drawn on.
 * Shaders may also have one or more *in*put texture(s), which must be provided to the shader as a uniform sampler2D.
 * Only the shader needs to know about any potential input texture, the framebuffer will always only know about it's
 */
export const bindTextureToFramebuffer = (gl: WebGLRenderingContext, texture: TextureObject, fb: WebGLFramebuffer): FramebufferObject {
    gl.bindFramebuffer(gl.FRAMEBUFFER, fb);
    gl.framebufferTexture2D(gl.FRAMEBUFFER, gl.COLOR_ATTACHMENT0, gl.TEXTURE_2D, texture.texture, 0); // analog to bu

    if (gl.checkFramebufferStatus(gl.FRAMEBUFFER) !== gl.FRAMEBUFFER_COMPLETE) {
        throw new Error('Error creating framebuffer: framebuffer-status: ${gl.checkFramebufferStatus(gl.FRAMEBUFFER)}');
    }

    gl.bindFramebuffer(gl.FRAMEBUFFER, null);

    const fbo: FramebufferObject = {
        framebuffer: fb,
        texture: texture,
        width: texture.width,
        height: texture.height
    };
    return fbo;
};

/**
 * Fetch uniform's location (uniform declared in some shader). Slow! Do *before* render loop.
 */

```

```

export const getUniformLocation = (gl: WebGLRenderingContext, program: WebGLProgram, uniformName: string): WebGLUniform {
    const loc = gl.getUniformLocation(program, uniformName);
    if (loc === null) {
        throw new Error(`Couldn't find uniform ${uniformName} in program.`);
    }
    return loc;
};

export type UniformType = '1i' | '2i' | '3i' | '4i' | '1f' | '2f' | '3f' | '4f' | '1fv' | '2fv' | '3fv' | '4fv' | 'mat2fv' | 'mat3fv' | 'mat4fv';

/**
 * Contrary to attributes, uniforms don't need to be stored in a buffer.
 *
 * 'v' is not about the shader, but how you provide data from the js-side.
 * uniformfv(loc, [3.19]) === uniformf(loc, 3.19)
 *
 * |js|-----| shader |
 * |-----|
 * |uniformf(loc, 3.19)| uniform float u_pi; |
 * |uniform2f(loc, 3.19, 2.72)| uniform vec2 u_constants; |
 * |uniform2fv(loc, [3.19, 2.72])| uniform vec2 u_constants; |
 * |uniform1fv(loc, [1, 2, 3, 4, 5, 6])| uniform float u_kernel[6]; |
 * |uniform2fv(loc, [1, 2, 3, 4, 5, 6])| uniform vec2 u_observations[3]; |
 * |uniformMatrix3fv(loc, [...], [...])| uniform mat3 u_matrix; |
 *
 */
export const bindValueToUniform = (gl: WebGLRenderingContext, uniformLocation: WebGLUniformLocation, type: UniformType) {
    switch (type) {
        case '1i':
            gl.uniform1i(uniformLocation, values[0]);
            break;

        case '1f':
            gl.uniform1f(uniformLocation, values[0]);
            break;
        case '2f':
            gl.uniform2f(uniformLocation, values[0], values[1]);
            break;
        case '3f':
            gl.uniform3f(uniformLocation, values[0], values[1], values[2]);
            break;
        case '4f':
            gl.uniform4f(uniformLocation, values[0], values[1], values[2], values[3]);
            break;

        case '1fv':
            gl.uniform1fv(uniformLocation, values);
            break;
        case '2fv':
            gl.uniform2fv(uniformLocation, values);
            break;
        case '3fv':
            gl.uniform3fv(uniformLocation, values);
            break;
        case '4fv':
            gl.uniform4fv(uniformLocation, values);
            break;

        // In the following *matrix* calls, the 'transpose' parameter must always be false.
        // Quoting the OpenGL ES 2.0 spec:
        // If the transpose parameter to any of the UniformMatrix* commands is
        // not FALSE, an INVALID_VALUE error is generated, and no uniform values are
        // changed.
        case 'matrix2fv':
            gl.uniformMatrix2fv(uniformLocation, false, values);
            break;

        case 'matrix3fv':
            gl.uniformMatrix3fv(uniformLocation, false, values);
            break;

        case 'matrix4fv':
            gl.uniformMatrix4fv(uniformLocation, false, values);
            break;

        default:
            throw Error(`Type ${type} not yet implemented.`);
    }
};

```

```
/*
 * (From https://hacks.mozilla.org/2013/04/the-concepts-of-webgl/ and https://stackoverflow.com/questions/56303648/we
 * Ignoring handmade framebuffers, WebGL has two framebuffers that are always in use: the 'frontbuffer/displaybuffer'
 * WebGL per default renders to the 'drawingbuffer', aka. the 'backbuffer'.
 * There is also the currently displayed buffer, named the 'frontbuffer' aka. the 'displaybuffer'.
 * the WebGL programmer has no explicit access to the frontbuffer whatsoever.
 *
 * Once you called 'clear', 'drawElements' or 'drawArrays', the browser marks the canvas as 'needs to be composited'.
 * (Assuming 'preserveDrawingBuffer == false':) Immediately before compositing, the browser
 * - swaps the back- and frontbuffer
 * - clears the new backbuffer.
 * (If 'preserveDrawingBuffer === true':) Immediately before compositing, the browser
 * - copies the drawingbuffer to the frontbuffer.
 *
 * As a consequence, if you're going to use canvas.toDataURL or canvas.toBlob or gl.readPixels or any other way of ge
 * unless you read it in the same event then it will likely be clear when you try to read it.
 *
 * In the past, old games always preserved the drawing buffer, so they'd only have to change those pixels that have a
 *
 * A (almost brutal) workaround to get the canvas to preserve the drawingBuffer can be found here: https://stackoverf
 */
export const getCurrentFramebuffersPixels = (canvas: HTMLCanvasElement): ArrayBuffer => {
    const gl = canvas.getContext('webgl');
    if (!gl) {
        throw new Error('no context');
    }

    const format = gl.getParameter(glIMPLEMENTATION_COLOR_READ_FORMAT);
    const type = gl.getParameter(glIMPLEMENTATION_COLOR_READ_TYPE);

    let pixels;
    if (type === gl.UNSIGNED_BYTE) {
        pixels = new Uint8Array(gl.drawingBufferWidth * gl.drawingBufferHeight * 4);
    } else if (type === gl.UNSIGNED_SHORT_5_6_5 || type === gl.UNSIGNED_SHORT_4_4_4_4 || type === gl.UNSIGNED_SHORT_5_
    pixels = new Uint16Array(gl.drawingBufferWidth * gl.drawingBufferHeight * 4);
    } else if (type === gl.FLOAT) {
        pixels = new Float32Array(gl.drawingBufferWidth * gl.drawingBufferHeight * 4);
    } else {
        throw new Error(`Did not understand pixel data type ${type} for format ${format}`);
    }

    // Just like 'toDataURL' or 'toBlob', 'readPixels' does not access the frontbuffer.
    // It accesses the backbuffer or any other currently active framebuffer.
    gl.readPixels(0, 0, canvas.width, canvas.height, format, type, pixels);

    return pixels;
};

import { createShaderProgram, setup3dScene, createFloatBuffer, getAttributeLocation, bindBufferToAttribute, getUniform

// dead-simple hash function - not intended to be secure in any way.
const hash = function(s: string): string {
    let h = 0;
    for (const c of s) {
        h += c.charCodeAt(0);
    }
    return `${h}`;
};

export interface IProgram {
    program: WebGLProgram;
    id: string;
    vertexShaderSource: string;
    fragmentShaderSource: string;
}

export class Program implements IProgram {

    readonly program: WebGLProgram;
    readonly id: string;

    constructor(gl: WebGLRenderingContext,
               readonly vertexShaderSource: string,
               readonly fragmentShaderSource: string) {
        this.program = createShaderProgram(gl, vertexShaderSource, fragmentShaderSource);
        this.id = hash(vertexShaderSource + fragmentShaderSource);
    }
}

export interface IUniform {
```

```

    location: WebGLUniformLocation;
    type: UniformType;
    value: number[];
    variableName: string;
}

export class Uniform implements IUniform {
    readonly location: WebGLUniformLocation;
    readonly type: UniformType;
    readonly value: number[];
    readonly variableName: string;

    constructor(gl: WebGLRenderingContext, program: IProgram, variableName: string, type: UniformType, data: number[]) {
        this.location = getUniformLocation(gl, program.program, variableName);
        this.type = type;
        this.value = data;
        this.variableName = variableName;
    }
}

export interface ITexture {
    location: WebGLUniformLocation;
    bindPoint: number;
    texture: TextureObject;
    variableName: string;
}

export class Texture implements ITexture {
    readonly location: WebGLUniformLocation;
    readonly bindPoint: number;
    readonly texture: TextureObject;
    readonly variableName: string;

    constructor(gl: WebGLRenderingContext, program: IProgram, variableName: string, im: HTMLImageElement | HTMLCanvasElement) {
        this.location = getUniformLocation(gl, program.program, variableName);
        if (im instanceof HTMLImageElement || im instanceof HTMLCanvasElement) {
            this.texture = createTexture(gl, im);
        } else {
            this.texture = im;
        }
        this.bindPoint = bindPoint;
        this.variableName = variableName;
    }
}

export interface IAttribute {
    location: number;
    value: BufferObject;
    variableName: string;
}

export class Attribute implements IAttribute {
    readonly location: number;
    readonly value: BufferObject;
    readonly variableName: string;
    readonly drawingMode: number;

    constructor(gl: WebGLRenderingContext, program: IProgram, variableName: string, data: number[][][], drawingMode: number) {
        this.location = getAttributeLocation(gl, program.program, variableName);
        this.value = createFloatBuffer(gl, data, drawingMode);
        this.variableName = variableName;
        this.drawingMode = drawingMode;
    }
}

/**
 * @TODO: this is not yet used anywhere.
 */
export class ElementAttribute implements IAttribute {
    readonly location: number;
    readonly value: BufferObject;
    readonly indices: IndexBufferObject;
    readonly variableName: string;

    constructor(gl: WebGLRenderingContext, program: IProgram, variableName: string, data: number[][][], indices: number[][][]) {
        this.location = getAttributeLocation(gl, program.program, variableName);
    }
}

```

```

        this.value = createFloatBuffer(gl, data);
        this.indices = createIndexBuffer(gl, indices);
        this.variableName = variableName;
    }
}

function first<T>(arr: T[], condition: (el: T) => boolean): T | null {
    for (const el of arr) {
        if (condition(el)) {
            return el;
        }
    }
    return null;
}

function parseProgram(program: IProgram): [string[], string[], string[]] {
    const attributeRegex = /^s*attribute (int|float|vec2|vec3|vec4|mat2|mat3|mat4) (\w*);/gm;
    const uniformRegex = /^s*uniform (int|float|vec2|vec3|vec4|mat2|mat3|mat4) (\w*)(\[\d\])*;/gm;
    const textureRegex = /^s*uniform sampler2D (\w*);/gm;

    const shaderCode = program.fragmentShaderSource + '\n\n' + program.vertexShaderSource;

    const attributeNames = [];
    let attributeMatches;
    while ((attributeMatches = attributeRegex.exec(shaderCode)) !== null) {
        attributeNames.push(attributeMatches[2]);
    }
    const uniformNames = [];
    let uniformMatches;
    while ((uniformMatches = uniformRegex.exec(shaderCode)) !== null) {
        uniformNames.push(uniformMatches[2]);
    }
    const textureNames = [];
    let textureMatches;
    while ((textureMatches = textureRegex.exec(shaderCode)) !== null) {
        textureNames.push(textureMatches[1]);
    }

    return [attributeNames, uniformNames, textureNames];
}

export type RenderMode = 'points' | 'lines' | 'triangles';

interface IShader {
    program: IProgram;
    attributes: IAttribute[];
    uniforms: IUniform[];
    textures: ITexture[];
    bind: (gl: WebGLRenderingContext) => void;
    render: (gl: WebGLRenderingContext, background?: number[], frameBuffer?: FramebufferObject) => void;
    updateAttributeData: (gl: WebGLRenderingContext, variableName: string, newData: number[][][]) => void;
    updateUniformData: (gl: WebGLRenderingContext, variableName: string, newData: number[]) => void;
    updateTextureData: (gl: WebGLRenderingContext, variableName: string, newImage: HTMLImageElement | HTMLCanvasElement) => void;
}

export class Shader implements IShader {
    constructor(
        readonly program: IProgram,
        readonly attributes: IAttribute[],
        readonly uniforms: IUniform[],
        readonly textures: ITexture[]
    ) {
        const [attributeNames, uniformNames, textureNames] = parseProgram(program);
        for (const attrName of attributeNames) {
            const found = attributes.filter(a => a.variableName === attrName);
            if (found.length !== 1) {
                throw new Error(`Provided ${found.length} values for shader's attribute ${attrName}.`);
            }
        }
        for (const uniformName of uniformNames) {
            const found = uniforms.filter(a => a.variableName === uniformName);
            if (found.length !== 1) {
                throw new Error(`Provided ${found.length} values for shader's uniform ${uniformName}.`);
            }
        }
        for (const texName of textureNames) {
            const found = textures.filter(a => a.variableName === texName);
            if (found.length !== 1) {
                throw new Error(`Provided ${found.length} values for shader's texture ${texName}.`);
            }
        }
    }
}

```

```

        }

    public bind(gl: WebGLRenderingContext): void {
        bindProgram(gl, this.program.program);
        for (const a of this.attributes) {
            bindBufferToAttribute(gl, a.location, a.value);
        }
        for (const u of this.uniforms) {
            bindValueToUniform(gl, u.location, u.type, u.value);
        }
        for (const t of this.textures) {
            bindTextureToUniform(gl, t.texture.texture, t.bindPoint, t.location);
        }
    }

    public render(gl: WebGLRenderingContext, background?: number[], frameBuffer?: FramebufferObject): void {
        if (!frameBuffer) {
            bindOutputCanvasToFramebuffer(gl);
        } else {
            bindFramebuffer(gl, frameBuffer);
        }
        if (background) {
            clearBackground(gl, background);
        }

        const firstAttribute = this.attributes[0].value;
        drawArray(gl, firstAttribute);
    }

    public updateAttributeData(gl: WebGLRenderingContext, variableName: string, newData: number[][]): void {
        const attribute = first<IAtribute>(this.attributes, el => el.variableName === variableName);
        if (!attribute) {
            throw new Error(`No such attribute ${variableName} to be updated.`);
        }
        updateBufferData(gl, attribute.value, newData);
    }

    public updateUniformData(gl: WebGLRenderingContext, variableName: string, newData: number[]): void {
        const uniform = first<IUniform>(this.uniforms, el => el.variableName === variableName);
        if (!uniform) {
            throw new Error(`No such uniform ${variableName} to be updated.`);
        }
        uniform.value = newData;
    }

    public updateTextureData(gl: WebGLRenderingContext, variableName: string, newImage: HTMLImageElement | HTMLCanvasElement): void {
        const original = first<ITexture>(this.textures, t => t.variableName === variableName);
        if (!original) {
            throw new Error(`No such original ${variableName} to be updated.`);
        }
        const newTextureObject = updateTexture(gl, original.texture, newImage);
        original.texture = newTextureObject;
    }
}

export class Framebuffer {

    readonly fbo: FramebufferObject;

    constructor(gl: WebGLRenderingContext, width: number, height: number) {
        const fb = createFramebuffer(gl);
        const fbTexture = createEmptyTexture(gl, width, height);
        const fbo = bindTextureToFramebuffer(gl, fbTexture, fb);
        this.fbo = fbo;
    }
}

export function renderLoop(fps: number, renderFunction: (tDelta: number) => void): void {
    const tDeltaTarget = 1000 * 1.0 / fps;
    let tDelta = tDeltaTarget;
    let tStart, tNow, tSleep: number;

    const render = () => {
        tStart = window.performance.now();
        renderFunction(tDelta);
    }
}

```

```

        tNow = window.performance.now();
        tDelta = tNow - tStart;
        tSleep = Math.max(tDeltaTarget - tDelta, 0);
        setTimeout(() => {
            requestAnimationFrame(render);
        }, tSleep);
    };

    render();
}

interface IEntity {
    program: IProgram;
    attributes: IAttribute[]; // note that attributes must all have the same number of entries!
    uniforms: IUniform[];
    textures: ITexture[];
    update: (tDelta: number) => void;
}

export class Entity implements IEntity {

    constructor(
        readonly program: IProgram,
        readonly attributes: IAttribute[],
        readonly uniforms: IUniform[],
        readonly textures: ITexture[],
        readonly updateFunction: (tDelta: number, attrs: IAttribute[], unis: IUniform[]) => void) {}

    update(tDelta: number): void {
        this.updateFunction(tDelta, this.attributes, this.uniforms);
    }
}

export class Engine {

    readonly entities: IEntity[] = [];

    constructor() {}

    public renderLoop(gl: WebGLRenderingContext, fps: number): void {
        setup3dScene(gl);

        const tDeltaTarget = 1000 * 1.0 / fps;
        let tStart, tNow: number, tDelta: number, tSleep;
        let currentShader = '';
        const render = () => {
            tStart = window.performance.now();

            // Part 1: allow objects to update their state
            for (const e of this.entities) {
                e.update(tDeltaTarget);
            }

            // Part 2: do the actual rendering work here
            clearBackground(gl, [.7, .7, .7, 1]);
            for (const e of this.entities) {
                if (e.program.id !== currentShader) {
                    bindProgram(gl, e.program.program);
                    currentShader = e.program.id;
                }
                for (const a of e.attributes) {
                    bindBufferToAttribute(gl, a.location, a.value);
                }
                for (const u of e.uniforms) {
                    bindValueToUniform(gl, u.location, u.type, u.value);
                }
                for (const t of e.textures) {
                    bindTextureToUniform(gl, t.texture.texture, t.bindPoint, t.location);
                }
                gl.drawArrays(gl.TRIANGLES, 0, e.attributes[0].value.vectorCount);
            }
        };
    }
}

```

```
// Part 3: time-management
    tNow = window.performance.now();
    tDelta = tNow - tStart;
    tSleep = Math.max(tDeltaTarget - tDelta, 0);
    setTimeout(() => {
        requestAnimationFrame(render);
    }, tSleep);

};

render();

public addEntity(entity: IEntity): void {
    this.entities.push(entity);
    this.sortEntities();
}

private sortEntities(): void {
    this.entities.sort((a: IEntity, b: IEntity) => {
        return (a.program.id > b.program.id) ? 1 : -1;
    });
}

}
```

7.3.5.1 Shader gallery

7.3.6 CSS

7.3.6.1 Positioning

7.3.6.2 Events

7.3.6.3 Animations

Chapter 8

Applications

8.1 Where should we meet?

n people want to decide at who's house to meet up. Does it matter? Can the sum of the paths that each person needs to travel vary? In fact, it can! Imagine that $n - 1$ of them live in the same town, while the last of them lives far away on the other side of the country. If they all meet in the town, only one person has to travel very far, whereas if they meet on the other side of the country, $n - 1$ people need to travel a large distance.

The distance to travel can be represented by a weighted connection matrix. Consider this example:

$$\begin{bmatrix} 0 & 1 & 1 & 1 & 11 \\ 1 & 0 & 1 & 1 & 11 \\ 1 & 1 & 0 & 1 & 11 \\ 1 & 1 & 1 & 0 & 10 \\ 11 & 11 & 11 & 10 & 0 \end{bmatrix}$$

It is now trivial to see that we can find the optimal place to meet up as the column with the smallest sum - in this case this would be at the place of person d .

8.2 Effective wind on a ship

A ship can only take the force of the wind that is aligned with its own angle. To calculate the vector of the wind for the ship, we must project the actual wind-vector onto the ships direction-vector.

8.3 Making a touch-instrument

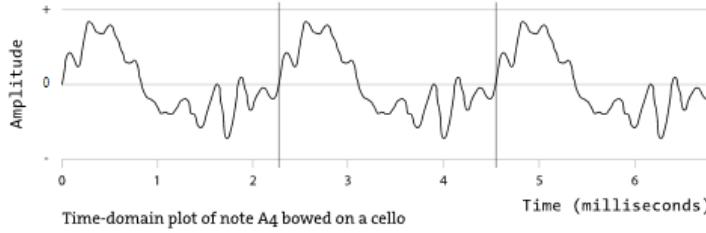
8.3.1 Music

8.3.1.1 Notes

A single tone x has only two attributes: frequency f_x and amplitude a_x a.k.a. pitch and loudness . The wave of a note is a function of these two: $w_x(t) = a_x \cdot \sin(f_x \cdot t)$

In practice, a note will stop being played at some time. Therefore, the real function for a note would be $w_x(t) \cdot s(t_1, t_2)$, where s is a step function. In FFT, however, we mostly treat a note as if it were a real periodic function - that is, never ending. We'll see details of this later.

By the way: The same note played on a piano still sounds different from a violin because physical instruments never play one pure note. Instead, depending on the material, different overtones are played along with the basenote.



8.3.1.2 Harmony

Instruments nowadays are tuned based on a , at 440 Hz. An octave higher would be a' at 880 Hz. There are twelve steps from a to a' , but their frequencies are not linearly, but exponentially spaced:

- linear: $f_x = \frac{12+x}{12} f_a$
- exponential: $f_x = 2^{\frac{x}{12}} f_a$

Some combinations sound nice together:

- Dur: 0,4,7; $w(t) = w_a(t) + w_{cis}(t) + w_e(t)$
- Dur-Sept: 0,4,7,11; $w(t) = w_a(t) + w_{cis}(t) + w_e(t) + w_{gis}(t)$
- Moll: 0,3,7; $w(t) = w_a(t) + w_c(t) + w_e(t)$
- Moll-Sept:

They do sound nice together when the wave of their sum does display a nice pattern of peaks - that is, when the sum, too, displays a repeating pattern.

8.3.1.3 Songs

Several tones can be played together, thus adding up their waves: $w(t) = \sum_x w_x(t)$ Also, with every beat, new notes may be played. Thus for every beat a new analysis is required.


8.3.2 Deconstructing $w(t)$ into f_x and a_x : the FFT

The sound we hear in a song is really just a timeseries. Its values are the sum of the current amplitudes of all notes, i.o.w. it's just one single, complicated wave $w(t)$. The FFT takes that wave, chops it into windows of, say, a second, and tries to deconstruct the signal in each window into the original, simple waves $w_x(t)$ of the single notes x .

So there are four steps to displaying a songs tones:

1. Chop the signal into parts using the window-function (usually just a step-function). The window-function's length should be equal to the length of one beat.
2. Take the first part and artificially elongate it by duplicating it over and over, thus creating a periodic signal.

3. Do the FFT: represent this periodic signal as a sum of sines of different frequencies and amplitudes.
4. Plot this.
5. Repeat from step two with the next part of the chopped-up signal.



Frequency-domain plot of a sustained note (220 Hz) bowed on a cello

8.4 Drawing a room onto a cylinder

The idea here is to imagine a scene taking place in a room. That scene is projected and drawn onto a paper cylinder. This way, a visitor can step inside the cylinder and see the scene around him as if he was actually there.

This requires two projections to happen to the objects in the scene. First we need to project them onto the cylinder. Second we need to roll out the cylinder so that we can actually print the scene on it with a printer or by hand.

8.4.1 Projecting a line in space onto a cylinder

We want to project the line from $\vec{a} = (10, 10, 1)$ to $\vec{b} = (0, 10, 1)$ onto a cylinder of radius $r = 1$. We assume that the visitors eyes will be at the coordinates $\vec{c} = (0, 0, 1.6)$.

8.4.1.1 Projecting a point onto the cylinder

8.4.1.2 Projecting a plane onto the cylinder

So we found out where the two edge points of the line should go onto the cylinder. But what would the line connecting them look like? What we need here is the intersection of the cylinder with the plane defined by the points $\vec{a}, \vec{b}, \vec{c}$. That plane is expressed as $P = \{\vec{x} | 0.07y + z = 1.7\}$.

To find the line we need an expression for $P \cap C$, that is, a solution to the system:

$$0.07y + z = 1.7$$

$$x^2 + y^2 = 1$$

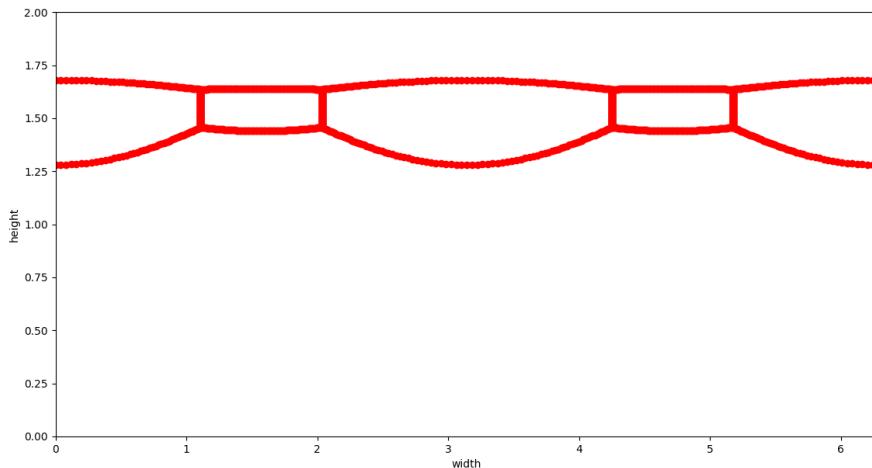
We can factor out y , leaving us with $x^2 + (\frac{1.7-z}{0.07})^2 = 1$

8.4.2 Rolling out the cylinder

We cut the cylinder open along the z axis. That means that for every z , we will apply a mapping $(x, y, z) \rightarrow (\theta, z)$. One particular mapping really stands out for that purpose - the polar representation of coordinates. Remember that

$$\cos\theta = \frac{\vec{v}\vec{u}}{|\vec{v}||\vec{u}|}$$

This means that we can achieve our mapping by using $(x, y, z) \rightarrow (\cos^{-1}x, z)$.



```
|| import matplotlib.pyplot as plt
from math import sqrt, pi, acos, atan
```

```

def distance(point1, point2):
    return sqrt((point2[0] - point1[0])**2 +
                (point2[1] - point1[1])**2 +
                (point2[2] - point1[2])**2)

def makeLine(point1, point2, num=100):
    line = []
    deltax = float(point2[0] - point1[0])/num
    deltay = float(point2[1] - point1[1])/num
    deltz = float(point2[2] - point1[2])/num
    for i in range(num):
        x = point1[0] + i*deltax
        y = point1[1] + i*deltay
        z = point1[2] + i*deltaz
        line.append([x, y, z])
    return line

def intersection(point, head, cylinder):
    rc = cylinder.radius
    xp = point[0]
    yp = point[1]
    xh = head[0]
    yh = head[1]
    alpha1 = (xh**2 - xh*xp + yh**2 - yh*yp - sqrt(rc**2*xh**2 - 2*rc**2*xh*xp + rc**2*xp**2 + rc**2*yh**2 - 2*rc**2*yh*yp))
    alpha2 = (xh**2 - xh*xp + yh**2 - yh*yp + sqrt(rc**2*xh**2 - 2*rc**2*xh*xp + rc**2*xp**2 + rc**2*yh**2 - 2*rc**2*yh*yp))
    ints1 = [0, 0, 0]
    ints1[0] = head[0] + alpha1*(point[0] - head[0])
    ints1[1] = head[1] + alpha1*(point[1] - head[1])
    ints1[2] = head[2] + alpha1*(point[2] - head[2])
    dist1 = distance(ints1, head)
    ints2 = [0, 0, 0]
    ints2[0] = head[0] + alpha2*(point[0] - head[0])
    ints2[1] = head[1] + alpha2*(point[1] - head[1])
    ints2[2] = head[2] + alpha2*(point[2] - head[2])
    dist2 = distance(ints2, head)
    if dist2 > dist1:
        return ints1
    else:
        return ints2

def toPolar(point):
    pointPolar = [0, 0, 0]
    radius = sqrt(point[0]**2 + point[1]**2)
    #theta = acos(point[0]/radius) * 360.0 / (2*pi)
    theta = atan(point[1]/(point[0] + 0.000001)) * 360.0 / (2*pi)
    if point[0] >= 0 and point[1] >= 0: # first quad
        theta = theta
    elif point[0] < 0 and point[1] >= 0: # second quad
        theta = theta + 180
    elif point[0] < 0 and point[1] < 0: # third quad
        theta = theta + 180
    else:
        theta = theta + 360
    pointPolar[0] = theta
    pointPolar[1] = radius
    pointPolar[2] = point[2]
    return pointPolar

def rollout(point, cylinder):
    pointPolar = toPolar(point)
    newpoint = [0, 0]
    newpoint[0] = 2*pi*cylinder.radius*pointPolar[0]/360
    newpoint[1] = point[2]
    return newpoint

class Cylinder:
    def __init__(self, radius):
        self.radius = radius

cylinder = Cylinder(1)
circumference = 2*pi*cylinder.radius
head = [0, 0, 1.6]
n = 100

# floor
line1 = makeLine([5, 10, 0], [-5, 10, 0], n)
line2 = makeLine([-5, 10, 0], [-5, -10, 0], n)
line3 = makeLine([-5, -10, 0], [5, -10, 0], n)

```

```

line4 = makeLine([5, -10, 0], [5, 10, 0], n)

# walls
line5 = makeLine([5, 10, 0], [5, 10, 2], n)
line6 = makeLine([-5, 10, 0], [-5, 10, 2], n)
line7 = makeLine([-5, -10, 0], [-5, -10, 2], n)
line8 = makeLine([5, -10, 0], [5, -10, 2], n)

# ceiling
line9 = makeLine([5, 10, 2], [-5, 10, 2], n)
line10 = makeLine([-5, 10, 2], [-5, -10, 2], n)
line11 = makeLine([-5, -10, 2], [5, -10, 2], n)
line12 = makeLine([5, -10, 2], [5, 10, 2], n)

line = line1 + line2 + line3 + line4 + line5 + line6 + line7 + line8 + line9 + line10 + line11 + line12

out = []
for i in range(len(line)):
    point = line[i]
    ints = intersection(point, head, cylinder)
    projp = rollout(ints, cylinder)
    out.append(projp)
    pointPolar = toPolar(ints)
    #print "mapping [{},{}],{} via ints [{},{}],{} and via polar [{},{}],{} to flat [{},{}],{}".format(point[0], point[1], p

xflat = []
yflat = []
for i in range(len(line)):
    xflat.append(out[i][0])
    yflat.append(out[i][1])

plt.plot(xflat, yflat, 'ro')
plt.xlabel('width')
plt.ylabel('height')
plt.axis([0, circumference, 0, 2])
plt.show()

```

8.5 Playing with projections and transformations

I like to start with a simple basefunction, and then distort it.

Nice basefunctions:

- simple geometric
- recursive

Nice distortions:

- sinusoidal (Fourier: add overtones)
- imaginary (Laplace: add more dimensions)
- stepfunction (Mandelbrot: cut off some of them)
- flat onto spherical (Mapprojection)
- circle rolled around circle (Spirograph)
- Copy and rotate
- shadows

Finally, we can refine the result.

- change paras over time
- heatmap
- path

We will try an example now. Take a simple geometric shape. Create it's fourier representation. Add some overtones to it. Convert it back into geometric space. What do you get?

There are basically two ways we can achieve this. The first is using ordinary (that is, one-dimensional) Fourier and parametrised shapes:

- a1: Take some parameterised Shape in \mathbb{R}^3 : $S(f(t))$
- a2: Extract $f(t)$
- a3: Express $f(t)$ in its Fourier representation: $f(t) = \sum \alpha_i b_i$, where b_i is one-dimensional
- a4: add some overtones to obtain $f'(t)$
- a5: Obtain S' as $S(f'(t))$

The second approach seems simpler, but requires tree dimensional Fourier:

- b1 (= a2): Take a shape $S(\vec{v})$
- b2 (= a3): $S(\vec{v}) = \sum \alpha_i \vec{b}_i$
- b3 (= a4, a5): Obtain S' by adding overtones to S

8.6 Fourier on sound and geometric objects

8.6.1 Basics: Fourier in one dimension

8.6.1.1 Representing a function under a Fourier base

The Fourier functions $\{\sin(nt), \cos(nt) | n \in \mathbb{N}\}$ form an orthogonal base for all functions in the space of square-integrable functions that are periodic between $-\pi$ and π . Lets first get some familiarity with Fourier analysis.

At first we'll try to express a simple function in a fourier base:

$$f(t) = \sin(t) = \alpha_0 + \sum_n^{\infty} \alpha_n \cos(nt) + \sum_n^{\infty} \beta_n \sin(nt)$$

It's comparatively easy to calculate the coefficients of this series¹:

$$\begin{aligned} \alpha_0 &= \frac{1}{\pi} \int_{-\pi}^{\pi} \sin(t) dt = 0 \\ \alpha_1 &= \frac{1}{\pi} \int_{-\pi}^{\pi} \sin(t) \cos(t) dt = 0 \\ \alpha_2 &= \dots = 0 \\ \alpha_3 &= \dots = 0 \\ \beta_1 &= \frac{1}{\pi} \int_{-\pi}^{\pi} \sin(t) \sin(t) dt = 1 \\ \beta_2 &= \frac{1}{\pi} \int_{-\pi}^{\pi} \sin(t) \sin(2t) dt = 0 \\ \beta_3 &= \dots = 0 \end{aligned}$$

In other words: the only amplitude different from 0 is $\beta_1 = 1$

We'll be using python for our Fourier-analysis, where it is more common to work with the exponential/imaginary representation of the Fourier functions. This representation is based on Eulers formula:

$$e^{it} = \cos(t) + i \sin(t)$$

Using this equation we can rewrite the Fourier base as $\{e^{int} | n \in \mathbb{N}\}$. Using this base we get as the only nonzero coefficient:

$$\gamma_1 = \frac{1}{\pi} \int_{-\pi}^{\pi} \sin(t) e^{it} dt = i$$

Having imaginary amplitudes may be intimidating at first, but really it's very simple. The real part of such an amplitude equals the amplitude of the equivalent cos term, whereas the imaginary part equals the amplitude of the sin term.

8.6.1.2 Bringing in arbitrary intervals

In the previous section, when we wrote the index n , we meant that the function would repeat itself n times within the interval $[-\pi, \pi]$ - it has a frequency of $f_n = \frac{n}{2\pi}$.

However, in reality we will deal with functions that are periodic over an unspecified interval $[-T, T]$ ². To accomodate this, we can expand our notion of Fourier basis to the space of all square integrable functions that are periodic between $[-T, T]$. Our base now consists of $\{\sin(\frac{2\pi n}{T}t), \cos(\frac{2\pi n}{T}t) | n \in \mathbb{N}\}$ or $\{e^{i\frac{2\pi n}{T}t} | n \in \mathbb{N}\}$.

Lets consider the function $f(t) = \sin(\frac{2\pi}{T}t)$. Just like before, it's easy to calculate this functions Fourier coefficients, and just like before, only one of them is nonzero:

$$\gamma_1 = \frac{1}{T} \int_{-T}^T \sin(\frac{2\pi}{T}t) e^{i\frac{2\pi}{T}t} dt = i$$

The frequency that goes associated with the index $n = 1$ is $f_n = \frac{n}{T} = 0.1$.

¹The term $\frac{1}{\pi}$ is not strictly neccessary for an orthogonal base, but it is very convenient, since it makes our base orthonormal.

²In *real* reality, we will not deal with periodic functions at all. But more on that later

8.6.1.3 Implementation and verification in python

It's time to get our hands dirty and learn about python's implementation of Fourier transforms. We'll try to use pythons `fft` library to recreate the previous section:

```
import numpy as np
import matplotlib.pyplot as plt

delta = 0.01
T = 10
t = np.arange(-T, T, delta)
data = np.sin(2 * np.pi * t / T)
amps = np.fft.fft(data)
frqs = np.fft.fftfreq(data.size, delta)

f, (ax1, ax2) = plt.subplots(1, 2)
ax1.plot(frqs, np.real(amps))
ax2.plot(frqs, np.imag(amps))
plt.show()
```

As expected, the plot shows a high amplitude at the frequency 0.1. Not quite as expected is the other peak at frequency -0.1. What's up with that? In fact, there are a few phenomena that warrant further explanation:

- Amplitudes are not really 0 where they should be, they only are very close to 0. This is merely an artifact of numerical computation that we won't bother with any further.
- The amplitudes are not normalized to T . This, too, should not pose any difficulties for us, as we will not be using the concrete values of the amplitudes in this project.
- How does fft chose which frequencies to consider? Here, for some reason, it's between -50 and 50.
- Why are there positive and negative frequencies?
- For now, we have chosen `t` to cover the interval $[-T, T]$ perfectly. But what if we have `t` too long or too short? Surprisingly, this doesn't seem to change the estimated frequencies very much.

8.6.1.4 Our first frequency domain operations: adding overtones

We went through all the trouble of representing $f(t)$ over a Fourier base for a reason: there are operations that are natural in the frequency domain that don't exactly come easily in time domain. One of those operations would be adding overtones.

```
import numpy as np
import matplotlib.pyplot as plt
import simpleaudio as sa

sampleRate = 44100

def addXHalfTonesTo(basefreq, steps):
    return basefreq * (2 ** (steps / 12.0))

def play(data, sampleRate):
    dataNrm = data * 32767 / np.max(np.abs(data))
    data16 = dataNrm.astype(np.int16)
    return sa.play_buffer(data16, 1, 2, sampleRate)

# creating the data
delta = 1.0 / sampleRate
T = 1
t = np.arange(0, T, delta)
data = np.sin(t * 440 * 2 * np.pi)
#play(data, sampleRate)

# transform to frequency domain
amps = np.fft.fft(data)
frqs = np.fft.fftfreq(data.size, delta)

# filter out the less important frqs
# add chord on top of basetone
thrsh = 0.1 * np.max(np.abs(amps))
ampsNew = np.zeros(np.shape(amps), dtype=np.complex128)
```

```

for i, a in enumerate(amps):
    if np.abs(a) > thrsh:
        f1 = frqs[i]
        f2 = np.round(addXHalfTonesTo(f1, 4))
        f3 = np.round(addXHalfTonesTo(f1, 7))
        i2 = np.where(frqs == f2)[0]
        i3 = np.where(frqs == f3)[0]
        ampsNew[i] += a
        ampsNew[i2] += a
        ampsNew[i3] += a

# convert back to timedomain
dataNew = np.fft.ifft(ampsNew)
play(dataNew, sampleRate)

# plot
f, ((ax1, ax2), (ax3, ax4)) = plt.subplots(2, 2)
ax1.plot(t, data)
ax2.plot(frqs, np.abs(amps))
ax3.plot(t, dataNew)
ax4.plot(frqs, np.abs(ampsNew))
plt.show()

```

8.6.1.5 Windowing

Now, in reality we will not be dealing with a allways constant sinewave. A piece of music is composed of different tones at different times. If we don't account for that, we get something like this:

```

import numpy as np
import matplotlib.pyplot as plt
import simpleaudio as sa

sampleRate = 44100

def addXHalfTonesTo(basefreq, steps):
    return basefreq * (2 ** (steps / 12.0))

def play(data, sampleRate):
    dataNrm = data * 32767 / np.max(np.abs(data))
    data16 = dataNrm.astype(np.int16)
    return sa.play_buffer(data16, 1, 2, sampleRate)

# creating the data
delta = 1.0 / sampleRate
T = 1
t = np.arange(0, T, delta)
fa = 440
fc = addXHalfTonesTo(fa, 4)
fd = addXHalfTonesTo(fc, 3)
a = np.sin(t * fa * 2 * np.pi)
c = np.sin(t * fc * 2 * np.pi)
d = np.sin(t * fd * 2 * np.pi)
t = np.concatenate((t, t, t, t))
data = np.concatenate((a, c, d, c))
#play(data, sampleRate)

# transform to frequency domain
amps = np.fft.fft(data)
frqs = np.fft.fftfreq(data.size, delta)

# filter out the less important frqs
# add chord on top of basetone
thrsh = 0.1 * np.max(np.abs(amps))
ampsNew = np.zeros(np.shape(amps), dtype=np.complex128)
for i, a in enumerate(amps):
    if np.abs(a) > thrsh:
        f1 = frqs[i]
        f2 = np.round(addXHalfTonesTo(f1, 4))
        f3 = np.round(addXHalfTonesTo(f1, 7))
        i2 = np.where(frqs == f2)[0]
        i3 = np.where(frqs == f3)[0]
        ampsNew[i] += a
        ampsNew[i2] += a
        ampsNew[i3] += a

```

```

# convert back to timedomain
dataNew = np.fft.ifft(ampsNew)
play(dataNew, sampleRate)

# plot
f, ((ax1, ax2), (ax3, ax4)) = plt.subplots(2, 2)
ax1.plot(t, data)
ax2.plot(frq, np.abs(amps))
ax3.plot(t, dataNew)
ax4.plot(frq, np.abs(ampsNew))
plt.show()

```

Thats why we use windowing. Windowing is the process of cutting an incomming signal in small chunks, for which we create an individual Fourier analysis. This way, at tone c, we don't have to drag along tone a any more; a has long been forgotten by the analysis.

```

import numpy as np
import matplotlib.pyplot as plt
import simpleaudio as sa

sampleRate = 44100

def first(data, cond):
    for i, d in enumerate(data):
        if cond(d):
            return i, d

def addXHalfTonesTo(basefreq, steps):
    return basefreq * (2 ** (steps / 12.0))

def overtone(steps, frq, amps):
    ampsNew = np.zeros(np.shape(amps), dtype=np.complex128)
    for i, a in enumerate(amps):
        if np.abs(a) > 0:
            f1 = frq[i]
            f2 = np.round(addXHalfTonesTo(f1, steps))
            i2 = np.where(frq == f2)
            #ampsNew[i] += a
            ampsNew[i2] += a
    return ampsNew

def play(data, sampleRate):
    dataNrm = data * 32767 / np.max(np.abs(data))
    data16 = dataNrm.astype(np.int16)
    po = sa.play_buffer(data16, 1, 2, sampleRate)
    po.wait_done()
    return po

def filterLower(amps, perc):
    thrsh = perc * np.max(np.abs(amps))
    ampsNew = np.zeros(np.shape(amps), dtype=np.complex128)
    for i, a in enumerate(amps):
        if np.abs(a) > thrsh:
            ampsNew[i] = a
    return ampsNew

def distort(frq, amps):
    ampsFiltered = filterLower(amps, 0.5)
    terz = overtone(4, frq, ampsFiltered)
    quint = overtone(7, frq, ampsFiltered)
    return ampsFiltered + terz + quint

def autotune(frq, amps):
    ampsFiltered = filterLower(amps, 0.9999)
    ampsTuned = np.zeros(np.shape(ampsFiltered), dtype=np.complex128)
    halfToneFreqs = [addXHalfTonesTo(220, x) for x in range(23)]
    for i, a in enumerate(ampsFiltered):
        if np.abs(a) > 0:
            f = frq[i]
            if f > 0:
                ihh, fhh = first(halfToneFreqs, lambda fc: fc >= f)
                fhl = halfToneFreqs[ihh-1]
                disth = np.abs(fhh - f)
                distl = np.abs(fhl - f)
                fh = fhl if distl < disth else fhh

```

```

        inew, fnew = first(frq, lambda fc: fc >= fh)
        ampsTuned[inew] = a
    return ampsTuned

def chunks(data, length):
    C = int(np.ceil(len(data) / length))
    chunks = []
    start = 0
    end = length
    for c in range(C):
        chunks.append(data[start:end])
        start += length
        end += length
    return chunks


delta = 1.0 / sampleRate
t = np.arange(0, 2, delta)
freq = lambda t: 220 + 220 * t / 2
data = np.sin(t * freq(t) * 2 * np.pi)

origAmps = []
tunedAmps = []
tunedChunks = []
chunksize = int(data.size / 4)
for chunk in chunks(data, chunksize):
    frqs = np.fft.fftfreq(chunk.size, delta)
    amps = np.fft.fft(chunk)
    ampsNew = autotune(frqs, amps)
    dataNew = np.fft.ifft(ampsNew)
    origAmps.append(amps)
    tunedAmps.append(ampsNew)
    tunedChunks.append(dataNew)

dataNew = np.concatenate(tunedChunks)
play(data, sampleRate)
play(dataNew, sampleRate)

start = 0
stop = 700
f, axes = plt.subplots(4)
for i in range(4):
    axes[i].plot(frqs[start:stop], np.abs(origAmps[i][start:stop]))
    axes[i].plot(frqs[start:stop], np.abs(tunedAmps[i][start:stop]), 'r')
plt.show()

```



8.6.2 Fourier in more than one dimensions

8.6.2.1 Multiple input parameters

Consider a simple black and white image. We'll describe the value at any pixel as $v(x, y)$, that is, $v : \mathbb{R}^2 \rightarrow \mathbb{R}$. Again, the base of the space of $(\mathbb{R}^2 \rightarrow \mathbb{R})$ functions must consist of functions that are themselves $(\mathbb{R}^2 \rightarrow \mathbb{R})$. The Fourier-representation of an image can be quite naturally obtained by:

$$v(x, y) = \sum_n^{\infty} \sum_m^{\infty} \alpha_{n,m} b_{n,m}(x, y)$$

where $b_{n,m}(x, y) = e^{i(\frac{n2\pi}{b_x-a_x}x + \frac{m2\pi}{b_y-a_y}y)}$ and

$$\alpha_{n,m} = \int_{a_x}^{b_x} \int_{a_y}^{b_y} v(x, y) e^{i(\frac{n2\pi}{b_x-a_x}x + \frac{m2\pi}{b_y-a_y}y)} dy dx$$

Generally, when a function takes two parameters, here x, y , then we also need two frequency parameters, here n, m . Analogously, for functions of three parameters we'd need three frequency parameters.

8.6.2.2 Multiple output parameters: Vector valued functions

In the previous section, we dealt with the space of functions mapping $\mathbb{R}^2 \rightarrow \mathbb{R}$. How about vector valued functions? Finally, we have arrived at the top level of abstraction when it comes to Fourier representations: the space of functions mapping $\mathbb{R}^a \rightarrow \mathbb{R}^b$.

For example, consider the surface $\{\vec{v} | v_z = v_x^2 + 3v_y\}$. The set-expression can be rewritten into a function-expression:

$$\vec{v}(x, y) = \begin{bmatrix} x \\ y \\ x^2 + 3y \end{bmatrix}$$

Here, we have an equation with two parameters (v_x and v_y) mapping onto a 3d-vector \vec{v} . Consequently, the Fourier base of the space of surfaces, too, must consist of functions mapping $\mathbb{R}^2 \rightarrow \mathbb{R}^3$.

As another example, consider an ellipsoid $\{\vec{v} | \frac{x^2}{r_1^2} + \frac{y^2}{r_2^2} + \frac{z^2}{r_3^2} = 1\}$. This body can be parameterized as:

$$\begin{bmatrix} x \\ y \\ z \end{bmatrix} = \begin{bmatrix} r_1 \cos \theta \cos \phi \\ r_2 \cos \theta \sin \phi \\ r_3 \sin \theta \end{bmatrix}$$

This means we can represent the ellipsoid function as a sum of Fourier-base functions that map $[0^\circ, 360^\circ]^2 \rightarrow \mathbb{R}^3$.

In general, we can represent any function mapping $\mathbb{R}^a \rightarrow \mathbb{R}^b$ by a Fourier base of functions from the same space. As for geometric bodies: as long as we can find a parameterisation of the body (that is continuous and square-integrable), we can also find a Fourier-representation³.

Surprisingly, this works just the same. Consider the ellipsoid in its parameterized representation:

$$\vec{v}(\phi, \theta) = \begin{bmatrix} r_x \cos \phi \cos \theta \\ r_y \cos \phi \sin \theta \\ r_z \sin \phi \end{bmatrix}$$

$\vec{v}(\phi, \theta)$ is a member of the space $S = \{f : \mathbb{R}^2 \rightarrow \mathbb{R}^3\}$, so the basis must consist of functions of the same shape. Let's call one of the members of the basis $\vec{b}_{n,m}$. It must have the following structure:

$$\vec{b}_{n,m}(\phi, \theta) = \begin{bmatrix} b_{n,m}^x(\phi, \theta) \\ b_{n,m}^y(\phi, \theta) \\ b_{n,m}^z(\phi, \theta) \end{bmatrix}$$

Then the coefficient $\alpha_{n,m}$ is calculated like this:

³We cannot find a Fourier-representation of non-parameterisable sets. The reason for this is that sets don't form an inner-product space - or at least not without some whacky definitions.

$$\begin{aligned}
\alpha_{n,m} &= \int_{a_\theta}^{b_\theta} \int_{a_\phi}^{b_\phi} \vec{v}(\phi, \theta) \vec{b}_{n,m}(\phi, \theta) d\phi d\theta \\
&= \int_{a_\theta}^{b_\theta} \int_{a_\phi}^{b_\phi} (\vec{i}v^x(\phi, \theta) + \vec{j}v^y(\phi, \theta) + \vec{k}v^z(\phi, \theta)) (\vec{i}b_{n,m}^x(\phi, \theta) + \vec{j}b_{n,m}^y(\phi, \theta) + \vec{k}b_{n,m}^z(\phi, \theta)) d\phi d\theta \\
&= \int_{a_\theta}^{b_\theta} \int_{a_\phi}^{b_\phi} 1v^x(\phi, \theta) b_{n,m}^x(\phi, \theta) + 0 + 0 + 0 + 1v^y(\phi, \theta) b_{n,m}^y(\phi, \theta) + 0 + 0 + 0 + v^z(\phi, \theta) b_{n,m}^z(\phi, \theta) d\phi d\theta \\
&= \int_{a_\theta}^{b_\theta} \int_{a_\phi}^{b_\phi} v^x(\phi, \theta) b_{n,m}^x(\phi, \theta) d\phi d\theta + \int_{a_\theta}^{b_\theta} \int_{a_\phi}^{b_\phi} v^y(\phi, \theta) b_{n,m}^y(\phi, \theta) d\phi d\theta + \int_{a_\theta}^{b_\theta} \int_{a_\phi}^{b_\phi} v^z(\phi, \theta) b_{n,m}^z(\phi, \theta) d\phi d\theta
\end{aligned}$$

Let's apply the Fourier base $\begin{bmatrix} b_{n,m}^x(\phi, \theta) \\ b_{n,m}^y(\phi, \theta) \\ b_{n,m}^z(\phi, \theta) \end{bmatrix} = \begin{bmatrix} e^{i(\frac{n2\pi}{b_\phi-a_\phi}\phi + \frac{m2\pi}{b_\theta-a_\theta}\theta)} \\ e^{i(\frac{n2\pi}{b_\phi-a_\phi}\phi + \frac{m2\pi}{b_\theta-a_\theta}\theta)} \\ e^{i(\frac{n2\pi}{b_\phi-a_\phi}\phi + \frac{m2\pi}{b_\theta-a_\theta}\theta)} \end{bmatrix}$ and the ellipsoid formula to this general expression. Then we get:

$$\begin{aligned}
\alpha_{n,m} &= \int_{a_\theta}^{b_\theta} \int_{a_\phi}^{b_\phi} r_x \cos \phi \cos \theta e^{i(\frac{n2\pi}{b_\phi-a_\phi}\phi + \frac{m2\pi}{b_\theta-a_\theta}\theta)} d\phi d\theta + \int_{a_\theta}^{b_\theta} \int_{a_\phi}^{b_\phi} r_y \cos \phi \sin \theta e^{i(\frac{n2\pi}{b_\phi-a_\phi}\phi + \frac{m2\pi}{b_\theta-a_\theta}\theta)} d\phi d\theta + \int_{a_\theta}^{b_\theta} \int_{a_\phi}^{b_\phi} r_z \sin \phi e^{i(\frac{n2\pi}{b_\phi-a_\phi}\phi + \frac{m2\pi}{b_\theta-a_\theta}\theta)} d\phi d\theta \\
\alpha_{n,m} &= r_x \int_0^{2\pi} \int_0^{2\pi} \cos \phi \cos \theta e^{i(n\phi+m\theta)} d\phi d\theta + r_y \int_0^{2\pi} \int_0^{2\pi} \cos \phi \sin \theta e^{i(n\phi+m\theta)} d\phi d\theta + r_z \int_0^{2\pi} \int_0^{2\pi} \sin \phi e^{i(n\phi+m\theta)} d\phi d\theta
\end{aligned}$$

Integrating yields:

...

Of course, the whole point is to graph the thing:

```

import numpy as np

def ellipsoid(theta, phi, rx, ry, rz):
    x = rx * np.cos(theta) * np.cos(phi)
    y = ry * np.cos(theta) * np.sin(phi)
    z = rz * np.sin(theta)
    return x, y, z

def filterAmps(data, perc=0.8):
    dataAbs = np.abs(data)
    x, y = np.shape(data)
    thrsh = perc * np.max(dataAbs)
    filtered = np.zeros((x,y), dtype=np.complex128)
    for c in range(x):
        for r in range(y):
            if dataAbs[c,r] >= thrsh:
                filtered[c,r] = data[c,r]
    return filtered

def addOvertones(amps):
    x, y = np.shape(amps)
    ampsNew = np.zeros((x, y), dtype=np.complex128)
    for c in range(x):
        for r in range(y):
            if np.abs(amps[c,r]) > 0:
                c2 = 2*c if 2*c < x else int(c/2)
                r2 = 2*r if 2*r < y else int(r/2)
                ampsNew[c,r] += amps[c,r]
                ampsNew[c,r2] += amps[c,r]
                ampsNew[c2,r] += amps[c,r]
                ampsNew[c2,r2] += amps[c,r]
    return ampsNew

rx = ry = rz = 1

```

```

data = np.zeros((360, 360, 3), dtype=np.float)
for t in np.arange(0,360):
    for p in np.arange(0, 360):
        data[t, p] = ellipsoid(t, p, rx, ry, rz)

ampsx = np.fft.rfft2(data[:, :, 0])
ampsy = np.fft.rfft2(data[:, :, 1])
ampsz = np.fft.rfft2(data[:, :, 2])

ampsxF = filterAmps(ampsx, 0.001)
ampsyF = filterAmps(ampsy, 0.001)
ampszF = filterAmps(ampsz, 0.001)
ampsxF = addOvertones(ampsxF)
ampsyF = addOvertones(ampsyF)
ampszF = addOvertones(ampszF)

dataNew = np.zeros((360, 360, 3), dtype=np.float)
dataNew[:, :, 0] = np.fft.irfft2(ampsxF)
dataNew[:, :, 1] = np.fft.irfft2(ampsyF)
dataNew[:, :, 2] = np.fft.irfft2(ampszF)

```

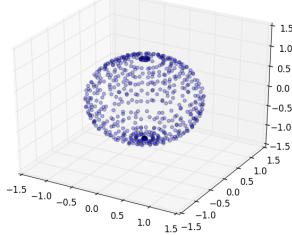


Figure 8.1: Original ellipsoid




Figure 8.2: Left: original spectra, right: adjusted spectra

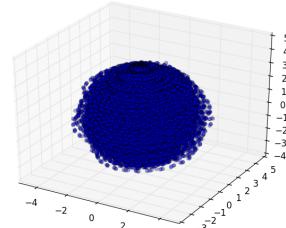


Figure 8.3: Ellipsoid backtransformed

8.7 Estimating the time needed for a task

Lets model the growth of a tree like this: we begin with the trunk. From the trunk, branches may grow. How many branches is decided by a Poisson-distribution, whose expected number of branches is dependent on the height from the ground.

```
from scipy.stats import poisson
from collections import namedtuple

Node = namedtuple('Node', 'parent, children')

def lamPerHeight(height):
    return 4 - height * 4 // 7

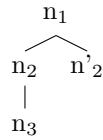
def randomTree(height=0):
    tree = Node(parent=None, children=[])
    lam = lamPerHeight(height)
    nrChildren = poisson.rvs(lam, size=1)
    for i in range(nrChildren):
        child = randomTree(height + 1)
        tree.children.append(child)
    return tree

print randomTree()
```

Using a Poisson-distribution, the probability of there being k branches on height h would be

$$pois_h(k) = \frac{e^{-\lambda_h} \lambda_h^k}{k!}$$

From the probability of a single branch we can now model the probabiltiy of a *tree*. Consider the following *tree*:




For this *treethe* probability is $P(tree) = pois_1(2) pois_2(1) pois_3(0) pois_2(0)$, or more generally:

$$P(tree) = \prod_{children} P(childtree)$$

```
def prob(tree, height=0):
    p = 1
    nrChildren = len(tree.children)
    lam = lamPerHeight(height)
    p *= poisson.pmf(nrChildren, lam)
    for child in tree.children:
        p *= prob(child, height + 1)
    return p

t = randomTree()
pt = prob(t)
```

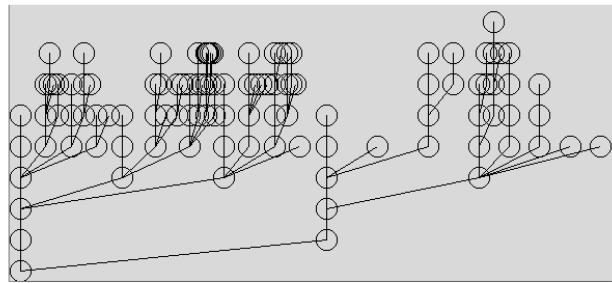
Now what about this tree:



This is an unfinished tree. It's probability is $P(\{tree \mid tree_0 \subseteq tree\}) = pois_1(\geq 2) pois_2(\geq 1) pois_3(\geq 0) pois_2(\geq 0) = pois_1(\geq 2) pois_2(\geq 1)$, or more generally:

$$P(\{tree \mid tree_0 \subseteq tree\}) = \prod_{children} P(\{childtree \mid childtree_0 \subseteq childtree\})$$

Figure 8.4: Using this model, we can spawn random trees



```

def probUnfinished(tree, height=0):
    p = 1
    nrChildren = len(tree.children)
    lam = lamPerHeight(height)
    p *= (1 - poisson.cdf(nrChildren, lam))
    for child in tree.children:
        p *= probUnfinished(child, height + 1)
    return p

```

Time as random variable Let time T be a random variable of a tree $tree$ defined as:

$$T_b(tree) = T_n(base) + \sum T_b(child)$$

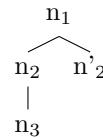
The expected time for a not yet grown tree would then be:

$$E[T_b] = \sum_{tree \in \Omega} T_b(tree) P(tree)$$

And the expected time given that we already have the first few nodes of a tree:

$$E[T_b]|_{tree_0} = \sum_{tree \in \Omega} T_b(tree) P(tree|tree_0)$$

This sum is pretty hard to compute, since Ω is an infinite set. But there is a way around this problem. Consider this tree:

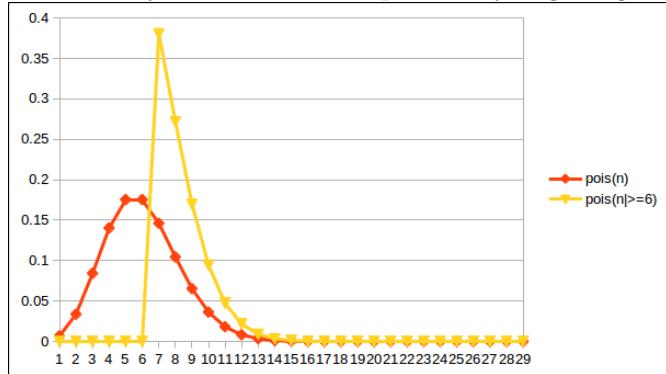


What is the expected number k of nodes on level 2, given that there are already two offspring?

$$\begin{aligned} E[k]|k \geq 2 &= \sum_{k=0}^{\infty} k P(k|2) = \sum_{k=2}^{\infty} k \frac{pois_{\lambda_2}(k)}{pois_{\lambda_2}(\geq 2)} \\ &= \frac{1}{pois_{\lambda_2}(\geq 2)} \left[\sum_{k=0}^{\infty} k pois_{\lambda_2}(k) - \sum_{k=0}^1 k pois_{\lambda_2}(k) \right] \\ &= \frac{1}{pois_{\lambda_2}(\geq 2)} [E[k] - \sum_{k=0}^1 k pois_{\lambda_2}(k)] \\ &= \frac{\lambda - \sum_{k=0}^{k=1} k pois_{\lambda_2}(k)}{1 - \sum_{k=0}^{k=1} pois_{\lambda_2}(k)} \\ &= \frac{\lambda - e^{-\lambda} \sum_{k=0}^{k=1} k \lambda^k / k!}{1 - e^{-\lambda} \sum_{k=0}^{k=1} \lambda^k / k!} \end{aligned}$$

We can check our predictions numerically:

Figure 8.5: When we already have 6 nodes, the probability of getting new nodes changes



```

from math import factorial, exp, floor

def pois(lmbd, k):
    return exp(-lmbd) * lmbd**k / factorial(k)

def poisCuml(lmbd, k):
    p = 0
    for j in range(k+1):
        p += pois(lmbd, j)
    return p

def poisCond(lmbd, k, k0):
    pb = 1 - poisCuml(lmbd, k0 - 1)
    if k < k0:
        pa = 0
    else:
        pa = pois(lmbd, k)
    return pa / pb

def expPois(lmbd):
    s = 0
    for i in range(int(50*lmbd)):
        s += i * pois(lmbd, i)
    return s

def expPoisCond(lmbd, k0):
    e = 0
    for i in range(k0, int(50*lmbd)):
        e += i * poisCond(lmbd, i, k0)
    return e

def expPoisCondAnal(lmbd, k0):
    sum1 = sum2 = 0
    for k in range(k0 + 1):
        sum1 += k * lmbd**k / float(factorial(k))
    for k in range(k0 + 1):
        sum2 += lmbd**k / float(factorial(k))
    sum1 *= exp(-lmbd)
    sum2 *= exp(-lmbd)
    a = float(lmbd - sum1)
    b = float(1 - sum2)
    return a / b

```

We will make use of this result shortly.

We can use this to find an intuitive expression for $E[T_b(tree)]$, which would be:

$$\begin{aligned}
E[T_b(tree)]|_{tree_0} &= E[T_n(base) + \sum T_b(child)]|_{tree_0} \\
&= E[T_n(base)]|_{tree_0} + E[\sum T_b(child)]|_{tree_0}
\end{aligned}$$

The term $E[\sum T_b(child)]|_{tree_0}$ is somewhat special: we cannot just equal it to $\sum E[T_b(child)]|_{tree_0}$, because we don't yet know how many children the tree will have. But it is reasonable to assume that:

$$E[\sum T_b(child)]_{|tree_0} = \sum_{children_0} E[T_b(child)] + \sum_{E[|children|]_{children_0} - children_0} E[T_b(child)]$$

With this, the previous equation becomes:

$$E[T_b(tree)]_{|tree_0} = E[T_n(base)]_{|tree_0} + \sum_{children_0} E[T_b(child)] + \sum_{E[|children|]_{children_0} - children_0} E[T_b(child)]$$

Now there are a bunch of terms in this equation that we can approximate statistically:

- $E[|children|]_{|children_0} = \frac{\lambda - e^{-\lambda} \sum_{k=0}^{k=children_0-1} k \lambda^k / k!}{1 - e^{-\lambda} \sum_{k=0}^{k=children_0-1} \lambda^k / k!}$. Here we can approximate λ by the average number of children on this level.
- $E[T_n(base)]_{|tree_0}$ can be approximated by the average net-time on this level
- $E[T_b(child)]$ knows two cases:
 - if we're dealing with a real child, the value is approximated with the same formula applied recursively
 - if we're dealing with a not-yet-spawned child, we have: $E[T_b] = E[T_n(1)] + E[\#c1]E[T_n(2)] + E[\#c1]E[\#c2]E[T_n(3)] + E[\#c1]E[\#c2]E[\#n3]E[T_n(4)] + \dots$

Time as random variable: revisited We can enhance our model quite a bit by recognizing that $T_b(base)$ is not a constant, but a random variable in its own right.

Then this random variable would still be defined as:

$$T_b(tree) = T_n(base) + \sum T_b(child)$$

But $T_n(base)$ itself would be random as well.

Under these assumptions, we get:

$$\begin{aligned} P(t|t_0) &= \begin{cases} \frac{P(t)}{1-P(t|t_0)} & \text{for } t \geq t_0 \\ 0 & \text{for } t < t_0 \end{cases} \\ E[T_n(base)]_{|t_0} &= \int_0^\infty t P(t|t \geq t_0) dt \\ &= \frac{1}{1 - P(t|t < t_0)} \int_{t_0}^\infty t P(t) dt \\ &= \frac{1}{1 - P(t|t < t_0)} [E[t] - \int_0^{t_0} t P(t) dt] \end{aligned}$$

Assuming that $P(t) = Gamma_{\alpha,\beta}(t) = \frac{\beta^\alpha e^{-\beta} t^{\alpha-1}}{\Gamma(\alpha)}$, we obtain:

$$\begin{aligned} \int_0^{t_0} t P(t) dt &= \frac{\beta^\alpha e^{-\beta}}{\Gamma(\alpha)} \int_0^{t_0} t^{\alpha-1} t dt \\ &= \frac{\beta^\alpha e^{-\beta} t_0^{\alpha-1}}{\Gamma(\alpha)} \frac{t_0^2}{\alpha+1} = Gamma(t_0) \frac{t_0^2}{\alpha+1} \end{aligned}$$

Here, the parameters α and β can be estimated by:

$$\begin{aligned} \alpha &\approx \frac{mean(t) mean(t)}{var(t)} \\ \beta &\approx \frac{mean(t)}{var(t)} \end{aligned}$$

We can once again verify this model numerically:

```

import scipy as sp
from scipy.stats import gamma
import matplotlib.pyplot as plt

def rateToScale(beta):
    return 1/beta

def scaleToRate(theta):
    return 1/theta

def probGamma(t, alpha, beta):
    theta = rateToScale(beta)
    return gamma.pdf(t, a=alpha, scale=theta)

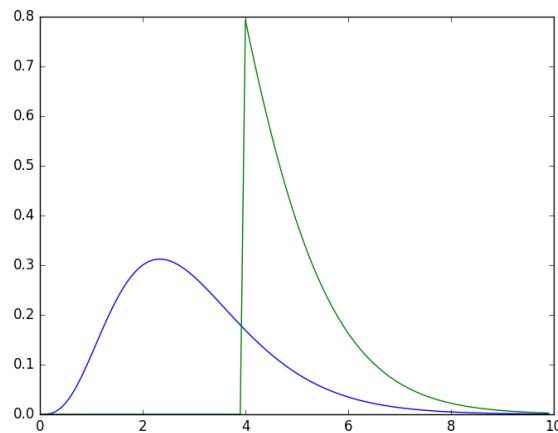
def probGammaCond(t, t0, alpha, beta):
    theta = rateToScale(beta)
    gam = gamma(a=alpha, scale=theta)
    tpart1 = t[sp.where(t<t0)]
    outpart1 = sp.zeros(sp.shape(tpart1))
    tpart2 = t[sp.where(t>=t0)]
    pt = gam.pdf(tpart2)
    Pt0 = gam.cdf(t0)
    outpart2 = pt / (1 - Pt0 )
    return sp.concatenate([outpart1, outpart2])

def expGammaCond(t0, alpha, beta):
    theta = rateToScale(beta)
    gam = gamma(a=alpha, scale=theta)
    expv = gam.mean()
    gt0 = gam.pdf(t0)
    gfac = t0 * to / (alpha + 1)
    gCumt0 = gam.cdf(t0)
    return (expv - gt0 * gfac) / (1 - gCumt0)

mean = 3.0
var = 2.0
alpha = mean * mean / var
beta = mean / var
t0 = 4
t = sp.arange(0, 10, 0.1)
pn = probGamma(t, alpha, beta)
pc = probGammaCond(t, t0, alpha, beta)

plt.plot(t, pn)
plt.plot(t, pc)
plt.show()

```



This model justifies another look at our samplespace. Note that

8.8 Scheduling

The proposed algorithm is simple:

1. Take the most urgent remaining task T_i .
2. Put it in the slot $[t_0, t_0 + \text{expt}(T_i)]$
3. If $t_0 + \text{expt}(T_i) > \text{deadline}(T_i)$, throw an exception
4. Set $t_0 = t_0 + \text{expt}(T_i)$ and $i = i + 1$ and loop to step 1.

Doublecheck: What if the task has scheduled subtasks?

8.8.1 Running time

8.8.2 Partial correctness

8.8.3 Termination

8.9 Webcamgiant

8.9.1 Webcam to RPy

8.9.2 RPy to computer

We need some way to transfer the images from the raspberry to our computer. Obviously, my favourite way would be over the internet - but to use the internet, you need a cable (ISDN or DSL) that is connected to some public router. Since we don't have that out in the field, we need to fall back onto a wireless transfer - GSM.

8.9.3 Computer to Oculus

8.10 Labeling radar-movies with likely summer-thunderstorms

Summer-thunderstorms are notoriously hard to predict. But we have an abundance of radar-data, together with decent records of actual precipitation.

- Create labeled data
 - When there is a thunderstorm about to come in the next three hours, label as "upcomming"
 - When there is a storm happening, label as "active"
 - When the storm is receding, label as "receding"
 - Otherwise, label as "clear"
- Train a net to recognize the labels given a 10-hour radar-movie
 - Lower layers: 3d-conv
 - Higher layers: dense
- FInd out what the model has learned
 - Are the conv-patterns interesting? Can they be used to find the outlines of a storm-cell?

8.11 Backpropagation for hydrological parameter estimation

Hydrological models consist of a network - a DAG - of landsurfaces, connected with flow paths. Each landsurface has one parameter c (the flow-constant). It seems like to find optimal values for the c 's, we could exploit the network-structure of a catchment and apply the backpropagation algorithm.

- Just like neural nets, we want to minimize an error function.
- Just like neural nets, we can do this by backpropagating the differential of a parameter.
- Unlike neural nets, we don't consider all possible connections between nodes - the elevation-map can take that burden from us.
- Unlike neural nets, each node has state. This might complicate things - but in reality, when the net runs long enough, different initial states converge.

8.11.1 Hydrological model

```

class Node:
    def __init__(self, id, V0, c, a):
        self.id = id
        self.V = V0
        self.c = c
        self.a = a

    def eval(N, qIns):
        inpt = N * self.a + sum(qIns)
        outpt = self.c * self.V
        self.V += (inpt - outpt)
        return outpt

class Model:
    def __init__(self, nodes, connections):
        self.nodes = nodes           # {id: node}
        self.connections = connections # [(idFrom, idTo)]

    def eval(Ns):      # Ns: {id: n}
        lastNode = self.nodes[0]
        qOut = evalNode(lastNode, Ns)
        return qOut

    def evalNode(self, node, Ns):
        upstreamNeighbours = self.getUpstreamNeighbours(node)
        qs = []
        for neighbour in upstreamNeighbours:
            qUpstream = self.evalNode(neighbour, Ns)
            qs.append(qUpstream)
        qOut = node.eval(Ns[node.id], qs)
        return qOut

    def getUpstreamNeighbours(self, node):
        neighbours = []
        for connection in self.connections:
            if connection[1] == node.id:
                neighbourId = connection[0]
                neighbour = self.nodes[neighbourId]
                neighbours.append(neighbour)
        return neighbours

    def gradients(self, sse):
        pass

    def updateParas(self, gradients):
        pass

def modelFactory(terrainMap):
    nodes = None
    connections = None
    model = Model(nodes, connections)
    return model

def optimize(model, N_timeSeries, q_timeSeries):
    errors = []
    for t, Ns in enumerate(N_timeSeries):

```

```
    qPred = model.eval(Ns)
    qObs = q_timeSeries[t]
    errors.append([qPred - qObs]**2)
    sse = sum(errors)
    grads = model.gradients(sse)
    model.update(grads)

terrainMap = None
N_timeSeries = None
q_timeSeries = None
model = modelFactory(terrainMap)
optimize(model, N_timeSeries, q_timeSeries)
```

8.11.2 Optimisation using tensorflow

8.12 Hydrological models as Hidden Markov Models

Whereas the above section dealt primarily with parameter estimation, here we are going to handle state estimation.